

PennOS

Generated by Doxygen 1.9.1



<b>1 Class Index</b>	<b>1</b>
1.1 Class List	1
<b>2 File Index</b>	<b>3</b>
2.1 File List	3
<b>3 Class Documentation</b>	<b>5</b>
3.1 <code>dir_entry_t</code> Struct Reference	5
3.1.1 Detailed Description	5
3.1.2 Member Data Documentation	5
3.1.2.1 <code>firstBlock</code>	5
3.1.2.2 <code>mtime</code>	6
3.1.2.3 <code>name</code>	6
3.1.2.4 <code>perm</code>	6
3.1.2.5 <code>reserved</code>	6
3.1.2.6 <code>size</code>	6
3.1.2.7 <code>type</code>	6
3.2 <code>fd_entry_t</code> Struct Reference	7
3.2.1 Detailed Description	7
3.2.2 Member Data Documentation	7
3.2.2.1 <code>filename</code>	7
3.2.2.2 <code>first_block</code>	7
3.2.2.3 <code>in_use</code>	7
3.2.2.4 <code>mode</code>	8
3.2.2.5 <code>position</code>	8
3.2.2.6 <code>ref_count</code>	8
3.2.2.7 <code>size</code>	8
3.3 <code>job_st</code> Struct Reference	8
3.3.1 Detailed Description	9
3.3.2 Member Data Documentation	9
3.3.2.1 <code>cmd</code>	9
3.3.2.2 <code>finished_count</code>	9
3.3.2.3 <code>id</code>	9
3.3.2.4 <code>num_pids</code>	9
3.3.2.5 <code>pgid</code>	10
3.3.2.6 <code>pids</code>	10
3.3.2.7 <code>state</code>	10
3.4 <code>parsed_command</code> Struct Reference	10
3.4.1 Detailed Description	10
3.4.2 Member Data Documentation	11
3.4.2.1 <code>commands</code>	11
3.4.2.2 <code>is_background</code>	11
3.4.2.3 <code>is_file_append</code>	11

3.4.2.4 num_commands	11
3.4.2.5 stdin_file	11
3.4.2.6 stdout_file	12
3.5 pcb_st Struct Reference	12
3.5.1 Detailed Description	13
3.5.2 Member Data Documentation	13
3.5.2.1 child_pcb	13
3.5.2.2 cmd_str	13
3.5.2.3 fd_table	13
3.5.2.4 input_fd	13
3.5.2.5 is_sleeping	14
3.5.2.6 output_fd	14
3.5.2.7 par_pid	14
3.5.2.8 pid	14
3.5.2.9 priority	14
3.5.2.10 process_state	14
3.5.2.11 process_status	15
3.5.2.12 signals	15
3.5.2.13 thread_handle	15
3.5.2.14 time_to_wake	15
3.6 spthread_fwd_args_st Struct Reference	15
3.6.1 Detailed Description	16
3.6.2 Member Data Documentation	16
3.6.2.1 actual_arg	16
3.6.2.2 actual_routine	16
3.6.2.3 child_meta	16
3.6.2.4 setup_cond	16
3.6.2.5 setup_done	17
3.6.2.6 setup_mutex	17
3.7 spthread_meta_st Struct Reference	17
3.7.1 Detailed Description	17
3.7.2 Member Data Documentation	17
3.7.2.1 meta_mutex	17
3.7.2.2 state	18
3.7.2.3 suspend_set	18
3.8 spthread_signal_args_st Struct Reference	18
3.8.1 Detailed Description	18
3.8.2 Member Data Documentation	18
3.8.2.1 ack	18
3.8.2.2 shutup_mutex	19
3.8.2.3 signal	19
3.9 spthread_st Struct Reference	19

3.9.1 Detailed Description	19
3.9.2 Member Data Documentation	20
3.9.2.1 meta	20
3.9.2.2 thread	20
3.10 vec_st Struct Reference	20
3.10.1 Detailed Description	20
3.10.2 Member Data Documentation	20
3.10.2.1 capacity	21
3.10.2.2 data	21
3.10.2.3 ele_dtor_fn	21
3.10.2.4 length	21
<b>4 File Documentation</b>	<b>23</b>
4.1 SRC/fs/fat_routines.c File Reference	23
4.1.1 Function Documentation	24
4.1.1.1 cat()	24
4.1.1.2 chmod()	27
4.1.1.3 cmpctdir()	28
4.1.1.4 cp()	28
4.1.1.5 ls()	29
4.1.1.6 mkfs()	30
4.1.1.7 mount()	32
4.1.1.8 mv()	32
4.1.1.9 rm()	34
4.1.1.10 touch()	35
4.1.1.11 unmount()	36
4.2 SRC/fs/fat_routines.h File Reference	36
4.2.1 Macro Definition Documentation	38
4.2.1.1 F_APPEND	38
4.2.1.2 F_READ	38
4.2.1.3 F_WRITE	39
4.2.1.4 FAT_EOF	39
4.2.1.5 FAT_FREE	39
4.2.1.6 PERM_EXEC	39
4.2.1.7 PERM_NONE	39
4.2.1.8 PERM_READ	39
4.2.1.9 PERM_READ_EXEC	40
4.2.1.10 PERM_READ_WRITE	40
4.2.1.11 PERM_READ_WRITE_EXEC	40
4.2.1.12 PERM_WRITE	40
4.2.1.13 TYPE_DIRECTORY	40
4.2.1.14 TYPE_REGULAR	40

4.2.1.15 TYPE_SYMLINK . . . . .	41
4.2.1.16 TYPE_UNKNOWN . . . . .	41
4.2.2 Function Documentation . . . . .	41
4.2.2.1 cat() . . . . .	41
4.2.2.2 chmod() . . . . .	44
4.2.2.3 cmpctdir() . . . . .	45
4.2.2.4 cp() . . . . .	46
4.2.2.5 ls() . . . . .	47
4.2.2.6 mkfs() . . . . .	48
4.2.2.7 mount() . . . . .	50
4.2.2.8 mv() . . . . .	51
4.2.2.9 rm() . . . . .	52
4.2.2.10 touch() . . . . .	53
4.2.2.11 unmount() . . . . .	55
4.3 SRC/fs/fs_helpers.c File Reference . . . . .	55
4.3.1 Function Documentation . . . . .	57
4.3.1.1 add_file_entry() . . . . .	57
4.3.1.2 allocate_block() . . . . .	58
4.3.1.3 compact_directory() . . . . .	59
4.3.1.4 copy_host_to_pennfat() . . . . .	61
4.3.1.5 copy_pennfat_to_host() . . . . .	63
4.3.1.6 copy_source_to_dest() . . . . .	64
4.3.1.7 decrement_fd_ref_count() . . . . .	65
4.3.1.8 find_file() . . . . .	66
4.3.1.9 get_free_fd() . . . . .	67
4.3.1.10 has_executable_permission() . . . . .	67
4.3.1.11 increment_fd_ref_count() . . . . .	67
4.3.1.12 init_fd_table() . . . . .	68
4.3.1.13 mark_entry_as_deleted() . . . . .	69
4.3.2 Variable Documentation . . . . .	69
4.3.2.1 block_size . . . . .	69
4.3.2.2 fat . . . . .	69
4.3.2.3 fat_size . . . . .	70
4.3.2.4 fd_table . . . . .	70
4.3.2.5 fs_fd . . . . .	70
4.3.2.6 is_mounted . . . . .	70
4.3.2.7 MAX_FDS . . . . .	70
4.3.2.8 num_fat_blocks . . . . .	70
4.4 SRC/fs/fs_helpers.h File Reference . . . . .	71
4.4.1 Function Documentation . . . . .	72
4.4.1.1 add_file_entry() . . . . .	72
4.4.1.2 allocate_block() . . . . .	74

4.4.1.3 compact_directory()	75
4.4.1.4 copy_host_to_pennfat()	77
4.4.1.5 copy_pennfat_to_host()	79
4.4.1.6 copy_source_to_dest()	80
4.4.1.7 decrement_fd_ref_count()	81
4.4.1.8 find_file()	82
4.4.1.9 get_free_fd()	83
4.4.1.10 has_executable_permission()	84
4.4.1.11 increment_fd_ref_count()	85
4.4.1.12 init_fd_table()	85
4.4.1.13 mark_entry_as_deleted()	86
4.4.2 Variable Documentation	87
4.4.2.1 block_size	87
4.4.2.2 fat	87
4.4.2.3 fat_size	87
4.4.2.4 fd_table	87
4.4.2.5 fs_fd	88
4.4.2.6 is_mounted	88
4.4.2.7 MAX_FDS	88
4.4.2.8 num_fat_blocks	88
4.5 SRC/fs/fs_kfuncs.c File Reference	88
4.5.1 Function Documentation	89
4.5.1.1 format_file_info()	89
4.5.1.2 k_close()	90
4.5.1.3 k_ls()	90
4.5.1.4 k_lseek()	91
4.5.1.5 k_open()	92
4.5.1.6 k_read()	94
4.5.1.7 k_unlink()	95
4.5.1.8 k_write()	96
4.5.2 Variable Documentation	99
4.5.2.1 current_fg_pid	99
4.5.2.2 current_running_pcb	100
4.6 SRC/fs/fs_kfuncs.h File Reference	100
4.6.1 Macro Definition Documentation	101
4.6.1.1 SEEK_CUR	101
4.6.1.2 SEEK_END	101
4.6.1.3 SEEK_SET	101
4.6.2 Function Documentation	101
4.6.2.1 k_close()	101
4.6.2.2 k_ls()	102
4.6.2.3 k_lseek()	104

4.6.2.4 k_open()	105
4.6.2.5 k_read()	107
4.6.2.6 k_unlink()	109
4.6.2.7 k_write()	110
4.7 SRC/fs/fs_syscalls.c File Reference	113
4.7.1 Function Documentation	114
4.7.1.1 s_close()	115
4.7.1.2 s_ls()	115
4.7.1.3 s_lseek()	115
4.7.1.4 s_open()	116
4.7.1.5 s_read()	116
4.7.1.6 s_unlink()	116
4.7.1.7 s_write()	117
4.8 SRC/fs/fs_syscalls.h File Reference	117
4.8.1 Macro Definition Documentation	118
4.8.1.1 STDERR_FILENO	118
4.8.1.2 STDIN_FILENO	118
4.8.1.3 STDOUT_FILENO	118
4.8.2 Function Documentation	118
4.8.2.1 s_close()	118
4.8.2.2 s_ls()	119
4.8.2.3 s_lseek()	120
4.8.2.4 s_open()	120
4.8.2.5 s_read()	121
4.8.2.6 s_unlink()	122
4.8.2.7 s_write()	122
4.9 SRC/kernel/kern_pcb.c File Reference	123
4.9.1 Function Documentation	124
4.9.1.1 create_pcb()	124
4.9.1.2 free_pcb()	125
4.9.1.3 k_proc_cleanup()	125
4.9.1.4 k_proc_create()	126
4.9.1.5 remove_child_in_parent()	127
4.9.2 Variable Documentation	127
4.9.2.1 current_pcb	127
4.9.2.2 current_running_pcb	127
4.9.2.3 next_pid	127
4.10 SRC/kernel/kern_pcb.h File Reference	128
4.10.1 Macro Definition Documentation	129
4.10.1.1 FILE_DESCRIPTOR_TABLE_SIZE	129
4.10.2 Typedef Documentation	129
4.10.2.1 pcb_t	129



4.10.3 Function Documentation . . . . .	129
4.10.3.1 create_pcb() . . . . .	129
4.10.3.2 free_pcb() . . . . .	130
4.10.3.3 k_proc_cleanup() . . . . .	131
4.10.3.4 k_proc_create() . . . . .	132
4.10.3.5 remove_child_in_parent() . . . . .	133
4.11 SRC/kernel/kern_sys_calls.c File Reference . . . . .	133
4.11.1 Function Documentation . . . . .	135
4.11.1.1 delete_from_explicit_queue() . . . . .	135
4.11.1.2 delete_from_queue() . . . . .	135
4.11.1.3 determine_index_in_queue() . . . . .	136
4.11.1.4 init_func() . . . . .	136
4.11.1.5 move_pcb_correct_queue() . . . . .	136
4.11.1.6 s_cleanup_init_process() . . . . .	137
4.11.1.7 s_echo() . . . . .	137
4.11.1.8 s_exit() . . . . .	138
4.11.1.9 s_kill() . . . . .	138
4.11.1.10 s_nice() . . . . .	139
4.11.1.11 s_ps() . . . . .	139
4.11.1.12 s_sleep() . . . . .	140
4.11.1.13 s_spawn() . . . . .	140
4.11.1.14 s_spawn_init() . . . . .	141
4.11.1.15 s_waitpid() . . . . .	141
4.11.2 Variable Documentation . . . . .	142
4.11.2.1 current_fg_pid . . . . .	142
4.11.2.2 current_pcb . . . . .	142
4.11.2.3 current_running_pcb . . . . .	143
4.11.2.4 one_priority_queue . . . . .	143
4.11.2.5 sleep_blocked_queue . . . . .	143
4.11.2.6 tick_counter . . . . .	143
4.11.2.7 two_priority_queue . . . . .	143
4.11.2.8 zero_priority_queue . . . . .	143
4.11.2.9 zombie_queue . . . . .	144
4.12 SRC/kernel/kern_sys_calls.h File Reference . . . . .	144
4.12.1 Function Documentation . . . . .	145
4.12.1.1 delete_from_explicit_queue() . . . . .	145
4.12.1.2 delete_from_queue() . . . . .	146
4.12.1.3 determine_index_in_queue() . . . . .	146
4.12.1.4 init_func() . . . . .	147
4.12.1.5 move_pcb_correct_queue() . . . . .	148
4.12.1.6 s_cleanup_init_process() . . . . .	148
4.12.1.7 s_echo() . . . . .	149

4.12.1.8 s_exit()	149
4.12.1.9 s_kill()	150
4.12.1.10 s_nice()	151
4.12.1.11 s_ps()	152
4.12.1.12 s_sleep()	152
4.12.1.13 s_spawn()	153
4.12.1.14 s_spawn_init()	154
4.12.1.15 s_waitpid()	154
4.13 SRC/kernel/logger.c File Reference	155
4.13.1 Function Documentation	156
4.13.1.1 log_generic_event()	156
4.13.1.2 log_nice_event()	157
4.13.1.3 log_scheduling_event()	158
4.14 SRC/kernel/logger.h File Reference	158
4.14.1 Function Documentation	159
4.14.1.1 log_generic_event()	159
4.14.1.2 log_nice_event()	160
4.14.1.3 log_scheduling_event()	161
4.14.2 Variable Documentation	161
4.14.2.1 log_fd	161
4.14.2.2 tick_counter	162
4.15 SRC/kernel/scheduler.c File Reference	162
4.15.1 Function Documentation	163
4.15.1.1 alarm_handler()	163
4.15.1.2 child_in_zombie_queue()	164
4.15.1.3 child_with_changed_process_status()	164
4.15.1.4 delete_process_from_all_queues()	164
4.15.1.5 delete_process_from_all_queues_except_current()	165
4.15.1.6 delete_process_from_particular_queue()	165
4.15.1.7 free_scheduler_queues()	165
4.15.1.8 generate_next_priority()	166
4.15.1.9 get_next_pcb()	166
4.15.1.10 get_pcb_in_queue()	167
4.15.1.11 handle_signal()	167
4.15.1.12 initialize_scheduler_queues()	168
4.15.1.13 put_pcb_into_correct_queue()	168
4.15.1.14 s_shutdown_pennos()	168
4.15.1.15 scheduler()	169
4.15.2 Variable Documentation	170
4.15.2.1 curr_priority_arr_index	170
4.15.2.2 current_pcb	170
4.15.2.3 current_running_pcb	171

4.15.2.4 det_priorities_arr . . . . .	171
4.15.2.5 log_fd . . . . .	171
4.15.2.6 one_priority_queue . . . . .	171
4.15.2.7 sleep_blocked_queue . . . . .	171
4.15.2.8 tick_counter . . . . .	171
4.15.2.9 two_priority_queue . . . . .	172
4.15.2.10 zero_priority_queue . . . . .	172
4.15.2.11 zombie_queue . . . . .	172
4.16 SRC/kernel/scheduler.h File Reference . . . . .	172
4.16.1 Function Documentation . . . . .	173
4.16.1.1 alarm_handler() . . . . .	173
4.16.1.2 child_in_zombie_queue() . . . . .	174
4.16.1.3 child_with_changed_process_status() . . . . .	174
4.16.1.4 delete_process_from_all_queues() . . . . .	175
4.16.1.5 delete_process_from_all_queues_except_current() . . . . .	175
4.16.1.6 delete_process_from_particular_queue() . . . . .	176
4.16.1.7 free_scheduler_queues() . . . . .	176
4.16.1.8 generate_next_priority() . . . . .	177
4.16.1.9 get_next_pcb() . . . . .	177
4.16.1.10 get_pcb_in_queue() . . . . .	178
4.16.1.11 handle_signal() . . . . .	179
4.16.1.12 initialize_scheduler_queues() . . . . .	179
4.16.1.13 put_pcb_into_correct_queue() . . . . .	180
4.16.1.14 s_shutdown_pennos() . . . . .	180
4.16.1.15 scheduler() . . . . .	181
4.17 SRC/kernel/signal.c File Reference . . . . .	182
4.18 SRC/kernel/signal.h File Reference . . . . .	183
4.18.1 Macro Definition Documentation . . . . .	183
4.18.1.1 CONT_BY_SIG . . . . .	183
4.18.1.2 EXITED_NORMALLY . . . . .	183
4.18.1.3 P_SIGCONT . . . . .	184
4.18.1.4 P_SIGSTOP . . . . .	184
4.18.1.5 P_SIGTERM . . . . .	184
4.18.1.6 P_WIFEXITED . . . . .	184
4.18.1.7 P_WIFSIGNALED . . . . .	184
4.18.1.8 P_WIFSTOPPED . . . . .	185
4.18.1.9 STOPPED_BY_SIG . . . . .	185
4.18.1.10 TERM_BY_SIG . . . . .	185
4.19 SRC/kernel/stress.c File Reference . . . . .	185
4.19.1 Function Documentation . . . . .	186
4.19.1.1 crash() . . . . .	186
4.19.1.2 hang() . . . . .	186

4.19.1.3 nohang()	186
4.19.1.4 recur()	187
4.20 SRC/kernel/stress.h File Reference	187
4.20.1 Function Documentation	187
4.20.1.1 crash()	187
4.20.1.2 hang()	188
4.20.1.3 nohang()	188
4.20.1.4 recur()	188
4.21 SRC/lib/pennos-errno.c File Reference	188
4.21.1 Variable Documentation	189
4.21.1.1 P_ERRNO	189
4.22 SRC/lib/pennos-errno.h File Reference	189
4.22.1 Macro Definition Documentation	190
4.22.1.1 P_EBADF	190
4.22.1.2 P_EBUSY	190
4.22.1.3 P_ECLOSE	190
4.22.1.4 P_ECOMMAND	190
4.22.1.5 P_EEXIST	190
4.22.1.6 P_EFS_NOT_MOUNTED	191
4.22.1.7 P_EFULL	191
4.22.1.8 P_EFUNC	191
4.22.1.9 P_EINTR	191
4.22.1.10 P_EINVAL	191
4.22.1.11 P_ELSEEK	191
4.22.1.12 P_EMALLOC	192
4.22.1.13 P_EMAP	192
4.22.1.14 P_ENOENT	192
4.22.1.15 P_ENULL	192
4.22.1.16 P_EOPEN	192
4.22.1.17 P_EPARSE	192
4.22.1.18 P_EPERM	193
4.22.1.19 P_EREAD	193
4.22.1.20 P_ESIGNAL	193
4.22.1.21 P_EUNKNOWN	193
4.22.1.22 P_EWRITE	193
4.22.1.23 P_INITFAIL	193
4.22.1.24 P_NEEDF	194
4.22.2 Variable Documentation	194
4.22.2.1 P_ERRNO	194
4.23 SRC/lib/spthread.c File Reference	194
4.23.1 Macro Definition Documentation	195
4.23.1.1 _GNU_SOURCE	195

4.23.1.2	MILISEC_IN_NANO . . . . .	195
4.23.1.3	SPTHREAD_RUNNING_STATE . . . . .	196
4.23.1.4	SPTHREAD_SIG_CONTINUE . . . . .	196
4.23.1.5	SPTHREAD_SIG_SUSPEND . . . . .	196
4.23.1.6	SPTHREAD_SUSPENDED_STATE . . . . .	196
4.23.1.7	SPTHREAD_TERMINATED_STATE . . . . .	196
4.23.2	Typedef Documentation . . . . .	196
4.23.2.1	pthread_fn . . . . .	196
4.23.2.2	spthread_fwd_args . . . . .	197
4.23.2.3	spthread_meta_t . . . . .	197
4.23.2.4	spthread_signal_args . . . . .	197
4.23.3	Function Documentation . . . . .	197
4.23.3.1	pthread_cancel() . . . . .	197
4.23.3.2	pthread_continue() . . . . .	197
4.23.3.3	pthread_create() . . . . .	198
4.23.3.4	pthread_disable_interrupts_self() . . . . .	199
4.23.3.5	pthread_enable_interrupts_self() . . . . .	199
4.23.3.6	pthread_equal() . . . . .	200
4.23.3.7	pthread_exit() . . . . .	200
4.23.3.8	pthread_join() . . . . .	200
4.23.3.9	pthread_self() . . . . .	200
4.23.3.10	pthread_suspend() . . . . .	201
4.23.3.11	pthread_suspend_self() . . . . .	201
4.24	SRC/lib/spthread.h File Reference . . . . .	202
4.24.1	Macro Definition Documentation . . . . .	203
4.24.1.1	SIGPTHD . . . . .	203
4.24.2	Typedef Documentation . . . . .	203
4.24.2.1	spthread_meta_t . . . . .	203
4.24.2.2	spthread_t . . . . .	203
4.24.3	Function Documentation . . . . .	203
4.24.3.1	pthread_cancel() . . . . .	204
4.24.3.2	pthread_continue() . . . . .	204
4.24.3.3	pthread_create() . . . . .	205
4.24.3.4	pthread_disable_interrupts_self() . . . . .	205
4.24.3.5	pthread_enable_interrupts_self() . . . . .	206
4.24.3.6	pthread_equal() . . . . .	206
4.24.3.7	pthread_exit() . . . . .	206
4.24.3.8	pthread_join() . . . . .	207
4.24.3.9	pthread_self() . . . . .	207
4.24.3.10	pthread_suspend() . . . . .	207
4.24.3.11	pthread_suspend_self() . . . . .	208
4.25	SRC/lib/Vec.c File Reference . . . . .	208

4.25.1 Function Documentation	209
4.25.1.1 <code>vec_clear()</code>	209
4.25.1.2 <code>vec_destroy()</code>	210
4.25.1.3 <code>vec_erase()</code>	210
4.25.1.4 <code>vec_erase_no_deletor()</code>	211
4.25.1.5 <code>vec_get()</code>	211
4.25.1.6 <code>vec_insert()</code>	212
4.25.1.7 <code>vec_new()</code>	213
4.25.1.8 <code>vec_pop_back()</code>	214
4.25.1.9 <code>vec_push_back()</code>	214
4.25.1.10 <code>vec_resize()</code>	215
4.25.1.11 <code>vec_set()</code>	216
4.26 SRC/lib/Vec.h File Reference	216
4.26.1 Macro Definition Documentation	218
4.26.1.1 <code>vec_capacity</code>	218
4.26.1.2 <code>vec_is_empty</code>	218
4.26.1.3 <code>vec_len</code>	218
4.26.2 Typedef Documentation	219
4.26.2.1 <code>ptr_dtor_fn</code>	219
4.26.2.2 <code>ptr_t</code>	219
4.26.2.3 <code>Vec</code>	219
4.26.3 Function Documentation	219
4.26.3.1 <code>vec_clear()</code>	219
4.26.3.2 <code>vec_destroy()</code>	220
4.26.3.3 <code>vec_erase()</code>	220
4.26.3.4 <code>vec_erase_no_deletor()</code>	221
4.26.3.5 <code>vec_get()</code>	221
4.26.3.6 <code>vec_insert()</code>	222
4.26.3.7 <code>vec_new()</code>	223
4.26.3.8 <code>vec_pop_back()</code>	224
4.26.3.9 <code>vec_push_back()</code>	224
4.26.3.10 <code>vec_resize()</code>	225
4.26.3.11 <code>vec_set()</code>	226
4.27 SRC/pennfat.c File Reference	226
4.27.1 Macro Definition Documentation	227
4.27.1.1 <code>PROMPT</code>	227
4.27.2 Function Documentation	227
4.27.2.1 <code>main()</code>	228
4.28 SRC/pennos.c File Reference	229
4.28.1 Function Documentation	230
4.28.1.1 <code>main()</code>	230
4.28.2 Variable Documentation	230

4.28.2.1 log_fd . . . . .	230
4.28.2.2 tick_counter . . . . .	231
4.29 SRC/shell/builtins.c File Reference . . . . .	231
4.29.1 Function Documentation . . . . .	231
4.29.1.1 u_perror() . . . . .	231
4.30 SRC/shell/builtins.h File Reference . . . . .	233
4.30.1 Function Documentation . . . . .	233
4.30.1.1 u_perror() . . . . .	233
4.31 SRC/shell/Job.h File Reference . . . . .	234
4.31.1 Typedef Documentation . . . . .	235
4.31.1.1 jid_t . . . . .	236
4.31.1.2 job . . . . .	236
4.31.2 Enumeration Type Documentation . . . . .	236
4.31.2.1 job_state_t . . . . .	236
4.32 SRC/shell/parser.c File Reference . . . . .	236
4.32.1 Macro Definition Documentation . . . . .	237
4.32.1.1 JUMP_OUT . . . . .	237
4.32.2 Function Documentation . . . . .	237
4.32.2.1 parse_command() . . . . .	238
4.32.2.2 print_parsed_command() . . . . .	240
4.32.2.3 print_parser_errcode() . . . . .	241
4.33 SRC/shell/parser.h File Reference . . . . .	241
4.33.1 Macro Definition Documentation . . . . .	242
4.33.1.1 EXPECT_COMMANDS . . . . .	242
4.33.1.2 EXPECT_INPUT_FILENAME . . . . .	243
4.33.1.3 EXPECT_OUTPUT_FILENAME . . . . .	243
4.33.1.4 UNEXPECTED_AMPERSAND . . . . .	243
4.33.1.5 UNEXPECTED_FILE_INPUT . . . . .	243
4.33.1.6 UNEXPECTED_FILE_OUTPUT . . . . .	243
4.33.1.7 UNEXPECTED_PIPELINE . . . . .	243
4.33.2 Function Documentation . . . . .	244
4.33.2.1 parse_command() . . . . .	244
4.33.2.2 print_parsed_command() . . . . .	246
4.33.2.3 print_parser_errcode() . . . . .	247
4.34 SRC/shell/shell.c File Reference . . . . .	247
4.34.1 Macro Definition Documentation . . . . .	248
4.34.1.1 MAX_BUFFER_SIZE . . . . .	249
4.34.1.2 MAX_LINE_BUFFER_SIZE . . . . .	249
4.34.1.3 PROMPT . . . . .	249
4.34.2 Function Documentation . . . . .	249
4.34.2.1 execute_command() . . . . .	249
4.34.2.2 fill_buffer_until_full_or_newline() . . . . .	251

4.34.2.3	free_job_ptr()	251
4.34.2.4	setup_terminal_signal_handlers()	252
4.34.2.5	shell()	252
4.34.2.6	shell_sigint_handler()	253
4.34.2.7	shell_sigstp_handler()	254
4.34.2.8	u_execute_command()	254
4.34.2.9	u_read_and_execute_script()	255
4.34.3	Variable Documentation	256
4.34.3.1	current_fg_pid	256
4.34.3.2	input_fd_script	256
4.34.3.3	job_list	256
4.34.3.4	next_job_id	256
4.34.3.5	output_fd_script	257
4.34.3.6	script_fd	257
4.35	SRC/shell/shell.h File Reference	257
4.35.1	Function Documentation	257
4.35.1.1	shell()	257
4.36	SRC/shell/shell_built_ins.c File Reference	259
4.36.1	Function Documentation	261
4.36.1.1	orphan_child()	261
4.36.1.2	u_bg()	261
4.36.1.3	u_busy()	261
4.36.1.4	u_cat()	262
4.36.1.5	u_chmod()	262
4.36.1.6	u_cp()	263
4.36.1.7	u_echo()	263
4.36.1.8	u_fg()	263
4.36.1.9	u_jobs()	264
4.36.1.10	u_kill()	264
4.36.1.11	u_logout()	265
4.36.1.12	u_ls()	265
4.36.1.13	u_man()	265
4.36.1.14	u_mv()	266
4.36.1.15	u_nice()	266
4.36.1.16	u_nice_pid()	267
4.36.1.17	u_orphanify()	267
4.36.1.18	u_ps()	268
4.36.1.19	u_rm()	268
4.36.1.20	u_sleep()	268
4.36.1.21	u_touch()	269
4.36.1.22	u_zombify()	269
4.36.1.23	zombie_child()	269



4.36.2 Variable Documentation . . . . .	269
4.36.2.1 get_associated_ufunc . . . . .	269
4.37 SRC/shell/shell_built_ins.h File Reference . . . . .	270
4.37.1 Function Documentation . . . . .	271
4.37.1.1 u_bg() . . . . .	272
4.37.1.2 u_busy() . . . . .	272
4.37.1.3 u_cat() . . . . .	272
4.37.1.4 u_chmod() . . . . .	273
4.37.1.5 u_cp() . . . . .	273
4.37.1.6 u_echo() . . . . .	274
4.37.1.7 u_fg() . . . . .	274
4.37.1.8 u_jobs() . . . . .	274
4.37.1.9 u_kill() . . . . .	275
4.37.1.10 u_logout() . . . . .	275
4.37.1.11 u_ls() . . . . .	276
4.37.1.12 u_man() . . . . .	276
4.37.1.13 u_mv() . . . . .	277
4.37.1.14 u_nice() . . . . .	277
4.37.1.15 u_nice_pid() . . . . .	278
4.37.1.16 u_orphanify() . . . . .	278
4.37.1.17 u_ps() . . . . .	278
4.37.1.18 u_rm() . . . . .	279
4.37.1.19 u_sleep() . . . . .	279
4.37.1.20 u_touch() . . . . .	280
4.37.1.21 u_zombify() . . . . .	280
<b>Index</b>	<b>281</b>



# Chapter 1

## Class Index

### 1.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

<a href="#">dir_entry_t</a>	Directory entry structure for files in the filesystem . . . . .	5
<a href="#">fd_entry_t</a>	File descriptor entry structure for open files . . . . .	7
<a href="#">job_st</a>	. . . . .	8
<a href="#">parsed_command</a>	. . . . .	10
<a href="#">pcb_st</a>	The PCB structure, which contains all the information about a process. Notably, it contains the thread handle, pid, parent pid, child pcbs, priority level, process state, command string, signals to be sent, input and output file descriptors, process status, sleeping status, and time to wake .	12
<a href="#">spthread_fwd_args_st</a>	. . . . .	15
<a href="#">spthread_meta_st</a>	. . . . .	17
<a href="#">spthread_signal_args_st</a>	. . . . .	18
<a href="#">spthread_st</a>	. . . . .	19
<a href="#">vec_st</a>	. . . . .	20



## Chapter 2

# File Index

### 2.1 File List

Here is a list of all files with brief descriptions:

SRC/ <a href="#">pennfat.c</a>	226
SRC/ <a href="#">pennos.c</a>	229
SRC/fs/ <a href="#">fat_routines.c</a>	23
SRC/fs/ <a href="#">fat_routines.h</a>	36
SRC/fs/ <a href="#">fs_helpers.c</a>	55
SRC/fs/ <a href="#">fs_helpers.h</a>	71
SRC/fs/ <a href="#">fs_kfuncs.c</a>	88
SRC/fs/ <a href="#">fs_kfuncs.h</a>	100
SRC/fs/ <a href="#">fs_syscalls.c</a>	113
SRC/fs/ <a href="#">fs_syscalls.h</a>	117
SRC/kernel/ <a href="#">kern_pcb.c</a>	123
SRC/kernel/ <a href="#">kern_pcb.h</a>	128
SRC/kernel/ <a href="#">kern_sys_calls.c</a>	133
SRC/kernel/ <a href="#">kern_sys_calls.h</a>	144
SRC/kernel/ <a href="#">logger.c</a>	155
SRC/kernel/ <a href="#">logger.h</a>	158
SRC/kernel/ <a href="#">scheduler.c</a>	162
SRC/kernel/ <a href="#">scheduler.h</a>	172
SRC/kernel/ <a href="#">signal.c</a>	182
SRC/kernel/ <a href="#">signal.h</a>	183
SRC/kernel/ <a href="#">stress.c</a>	185
SRC/kernel/ <a href="#">stress.h</a>	187
SRC/lib/ <a href="#">pennos-errno.c</a>	188
SRC/lib/ <a href="#">pennos-errno.h</a>	189
SRC/lib/ <a href="#">spthread.c</a>	194
SRC/lib/ <a href="#">spthread.h</a>	202
SRC/lib/ <a href="#">Vec.c</a>	208
SRC/lib/ <a href="#">Vec.h</a>	216
SRC/shell/ <a href="#">builtins.c</a>	231
SRC/shell/ <a href="#">builtins.h</a>	233
SRC/shell/ <a href="#">Job.h</a>	234
SRC/shell/ <a href="#">parser.c</a>	236
SRC/shell/ <a href="#">parser.h</a>	241
SRC/shell/ <a href="#">shell.c</a>	247
SRC/shell/ <a href="#">shell.h</a>	257
SRC/shell/ <a href="#">shell_built_ins.c</a>	259
SRC/shell/ <a href="#">shell_built_ins.h</a>	270



## Chapter 3

# Class Documentation

### 3.1 `dir_entry_t` Struct Reference

Directory entry structure for files in the filesystem.

```
#include <fat_routines.h>
```

#### Public Attributes

- char `name` [32]
- uint32\_t `size`
- uint16\_t `firstBlock`
- uint8\_t `type`
- uint8\_t `perm`
- time\_t `mtime`
- char `reserved` [16]

#### 3.1.1 Detailed Description

Directory entry structure for files in the filesystem.

Definition at line 46 of file `fat_routines.h`.

#### 3.1.2 Member Data Documentation

##### 3.1.2.1 `firstBlock`

```
uint16_t dir_entry_t::firstBlock
```

Definition at line 49 of file `fat_routines.h`.

### 3.1.2.2 mtime

```
time_t dir_entry_t::mtime
```

Definition at line 52 of file fat\_routines.h.

### 3.1.2.3 name

```
char dir_entry_t::name[32]
```

Definition at line 47 of file fat\_routines.h.

### 3.1.2.4 perm

```
uint8_t dir_entry_t::perm
```

Definition at line 51 of file fat\_routines.h.

### 3.1.2.5 reserved

```
char dir_entry_t::reserved[16]
```

Definition at line 53 of file fat\_routines.h.

### 3.1.2.6 size

```
uint32_t dir_entry_t::size
```

Definition at line 48 of file fat\_routines.h.

### 3.1.2.7 type

```
uint8_t dir_entry_t::type
```

Definition at line 50 of file fat\_routines.h.

The documentation for this struct was generated from the following file:

- SRC/fs/[fat\\_routines.h](#)



## 3.2 fd\_entry\_t Struct Reference

File descriptor entry structure for open files.

```
#include <fat_routines.h>
```

### Public Attributes

- int [in\\_use](#)
- int [ref\\_count](#)
- char [filename](#) [32]
- uint32\_t [size](#)
- uint16\_t [first\\_block](#)
- uint32\_t [position](#)
- uint8\_t [mode](#)

### 3.2.1 Detailed Description

File descriptor entry structure for open files.

Definition at line 59 of file fat\_routines.h.

### 3.2.2 Member Data Documentation

#### 3.2.2.1 filename

```
char fd_entry_t::filename[32]
```

Definition at line 62 of file fat\_routines.h.

#### 3.2.2.2 first\_block

```
uint16_t fd_entry_t::first_block
```

Definition at line 64 of file fat\_routines.h.

#### 3.2.2.3 in\_use

```
int fd_entry_t::in_use
```

Definition at line 60 of file fat\_routines.h.

#### 3.2.2.4 mode

```
uint8_t fd_entry_t::mode
```

Definition at line 66 of file `fat_routines.h`.

#### 3.2.2.5 position

```
uint32_t fd_entry_t::position
```

Definition at line 65 of file `fat_routines.h`.

#### 3.2.2.6 ref\_count

```
int fd_entry_t::ref_count
```

Definition at line 61 of file `fat_routines.h`.

#### 3.2.2.7 size

```
uint32_t fd_entry_t::size
```

Definition at line 63 of file `fat_routines.h`.

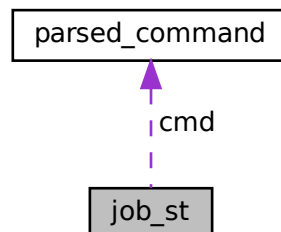
The documentation for this struct was generated from the following file:

- [SRC/fs/fat\\_routines.h](#)

### 3.3 job\_st Struct Reference

```
#include <Job.h>
```

Collaboration diagram for `job_st`:



## Public Attributes

- `jid_t` `id`
- struct `parsed_command` \* `cmd`
- `pid_t` \* `pids`
- `job_state_t` `state`
- `size_t` `num_pids`
- `pid_t` `pgid`
- `size_t` `finished_count`

### 3.3.1 Detailed Description

Definition at line 16 of file Job.h.

### 3.3.2 Member Data Documentation

#### 3.3.2.1 cmd

```
struct parsed_command* job_st::cmd
```

Definition at line 18 of file Job.h.

#### 3.3.2.2 finished\_count

```
size_t job_st::finished_count
```

Definition at line 23 of file Job.h.

#### 3.3.2.3 id

```
jid_t job_st::id
```

Definition at line 17 of file Job.h.

#### 3.3.2.4 num\_pids

```
size_t job_st::num_pids
```

Definition at line 21 of file Job.h.

### 3.3.2.5 pgid

```
pid_t job_st::pgid
```

Definition at line 22 of file Job.h.

### 3.3.2.6 pids

```
pid_t* job_st::pids
```

Definition at line 19 of file Job.h.

### 3.3.2.7 state

```
job_state_t job_st::state
```

Definition at line 20 of file Job.h.

The documentation for this struct was generated from the following file:

- SRC/shell/[Job.h](#)

## 3.4 parsed\_command Struct Reference

```
#include <parser.h>
```

### Public Attributes

- bool [is\\_background](#)
- bool [is\\_file\\_append](#)
- const char \* [stdin\\_file](#)
- const char \* [stdout\\_file](#)
- size\_t [num\\_commands](#)
- char \*\* [commands](#) []

### 3.4.1 Detailed Description

struct [parsed\\_command](#) stored all necessary information needed for penn-shell.

Definition at line 36 of file parser.h.

## 3.4.2 Member Data Documentation

### 3.4.2.1 commands

```
char** parsed_command::commands[ ]
```

Definition at line 56 of file parser.h.

### 3.4.2.2 is\_background

```
bool parsed_command::is_background
```

Definition at line 39 of file parser.h.

### 3.4.2.3 is\_file\_append

```
bool parsed_command::is_file_append
```

Definition at line 43 of file parser.h.

### 3.4.2.4 num\_commands

```
size_t parsed_command::num_commands
```

Definition at line 52 of file parser.h.

### 3.4.2.5 stdin\_file

```
const char* parsed_command::stdin_file
```

Definition at line 46 of file parser.h.

### 3.4.2.6 stdout\_file

```
const char* parsed_command::stdout_file
```

Definition at line 49 of file parser.h.

The documentation for this struct was generated from the following file:

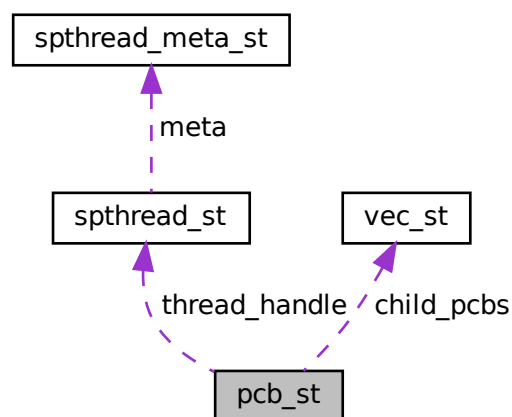
- SRC/shell/[parser.h](#)

## 3.5 pcb\_st Struct Reference

The PCB structure, which contains all the information about a process. Notably, it contains the thread handle, pid, parent pid, child pcbs, priority level, process state, command string, signals to be sent, input and output file descriptors, process status, sleeping status, and time to wake.

```
#include <kern_pcb.h>
```

Collaboration diagram for pcb\_st:



### Public Attributes

- [sptthread\\_t](#) `thread_handle`
- `pid_t` `pid`
- `pid_t` `par_pid`
- [Vec](#) `child_pcbs`
- `int` `priority`
- `char` `process_state`
- `char *` `cmd_str`
- `bool` `signals` [3]
- `int` `input_fd`
- `int` `output_fd`
- `int` `process_status`
- `bool` `is_sleeping`
- `int` `time_to_wake`
- `int` `fd_table` [`FILE_DESCRIPTOR_TABLE_SIZE`]

### 3.5.1 Detailed Description

The PCB structure, which contains all the information about a process. Notably, it contains the thread handle, pid, parent pid, child pcbs, priority level, process state, command string, signals to be sent, input and output file descriptors, process status, sleeping status, and time to wake.

Definition at line 30 of file kern\_pcb.h.

### 3.5.2 Member Data Documentation

#### 3.5.2.1 child\_pcbs

```
Vec pcb_st::child_pcbs
```

Definition at line 36 of file kern\_pcb.h.

#### 3.5.2.2 cmd\_str

```
char* pcb_st::cmd_str
```

Definition at line 42 of file kern\_pcb.h.

#### 3.5.2.3 fd\_table

```
int pcb_st::fd_table[FILE_DESCRIPTOR_TABLE_SIZE]
```

Definition at line 61 of file kern\_pcb.h.

#### 3.5.2.4 input\_fd

```
int pcb_st::input_fd
```

Definition at line 48 of file kern\_pcb.h.

#### 3.5.2.5 is\_sleeping

```
bool pcb_st::is_sleeping
```

Definition at line 58 of file kern\_pcb.h.

#### 3.5.2.6 output\_fd

```
int pcb_st::output_fd
```

Definition at line 49 of file kern\_pcb.h.

#### 3.5.2.7 par\_pid

```
pid_t pcb_st::par_pid
```

Definition at line 34 of file kern\_pcb.h.

#### 3.5.2.8 pid

```
pid_t pcb_st::pid
```

Definition at line 33 of file kern\_pcb.h.

#### 3.5.2.9 priority

```
int pcb_st::priority
```

Definition at line 38 of file kern\_pcb.h.

#### 3.5.2.10 process\_state

```
char pcb_st::process_state
```

Definition at line 39 of file kern\_pcb.h.



### 3.5.2.11 `process_status`

```
int pcb_st::process_status
```

Definition at line 51 of file `kern_pcb.h`.

### 3.5.2.12 `signals`

```
bool pcb_st::signals[3]
```

Definition at line 44 of file `kern_pcb.h`.

### 3.5.2.13 `thread_handle`

```
spthread_t pcb_st::thread_handle
```

Definition at line 31 of file `kern_pcb.h`.

### 3.5.2.14 `time_to_wake`

```
int pcb_st::time_to_wake
```

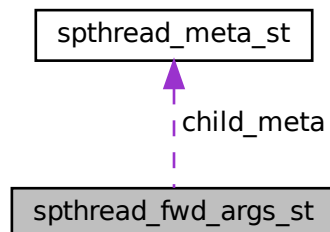
Definition at line 59 of file `kern_pcb.h`.

The documentation for this struct was generated from the following file:

- SRC/kernel/[kern\\_pcb.h](#)

## 3.6 `spthread_fwd_args_st` Struct Reference

Collaboration diagram for `spthread_fwd_args_st`:



## Public Attributes

- [pthread\\_fn](#) [actual\\_routine](#)
- void \* [actual\\_arg](#)
- bool [setup\\_done](#)
- [pthread\\_mutex\\_t](#) [setup\\_mutex](#)
- [pthread\\_cond\\_t](#) [setup\\_cond](#)
- [spthread\\_meta\\_t](#) \* [child\\_meta](#)

### 3.6.1 Detailed Description

Definition at line 22 of file `spthread.c`.

### 3.6.2 Member Data Documentation

#### 3.6.2.1 `actual_arg`

```
void* spthread_fwd_args_st::actual_arg
```

Definition at line 27 of file `spthread.c`.

#### 3.6.2.2 `actual_routine`

```
pthread_fn spthread_fwd_args_st::actual_routine
```

Definition at line 26 of file `spthread.c`.

#### 3.6.2.3 `child_meta`

```
spthread_meta_t* spthread_fwd_args_st::child_meta
```

Definition at line 40 of file `spthread.c`.

#### 3.6.2.4 `setup_cond`

```
pthread_cond_t spthread_fwd_args_st::setup_cond
```

Definition at line 37 of file `spthread.c`.

### 3.6.2.5 setup\_done

```
bool spthread_fwd_args_st::setup_done
```

Definition at line 35 of file spthread.c.

### 3.6.2.6 setup\_mutex

```
pthread_mutex_t spthread_fwd_args_st::setup_mutex
```

Definition at line 36 of file spthread.c.

The documentation for this struct was generated from the following file:

- SRC/lib/[spthread.c](#)

## 3.7 spthread\_meta\_st Struct Reference

### Public Attributes

- sigset\_t [suspend\\_set](#)
- volatile sig\_atomic\_t [state](#)
- pthread\_mutex\_t [meta\\_mutex](#)

### 3.7.1 Detailed Description

Definition at line 57 of file spthread.c.

### 3.7.2 Member Data Documentation

#### 3.7.2.1 meta\_mutex

```
pthread_mutex_t spthread_meta_st::meta_mutex
```

Definition at line 71 of file spthread.c.

### 3.7.2.2 state

```
volatile sig_atomic_t spthread_meta_st::state
```

Definition at line 68 of file spthread.c.

### 3.7.2.3 suspend\_set

```
sigset_t spthread_meta_st::suspend_set
```

Definition at line 59 of file spthread.c.

The documentation for this struct was generated from the following file:

- SRC/lib/[spthread.c](#)

## 3.8 spthread\_signal\_args\_st Struct Reference

### Public Attributes

- const int [signal](#)
- volatile sig\_atomic\_t [ack](#)
- pthread\_mutex\_t [shutup\\_mutex](#)

### 3.8.1 Detailed Description

Definition at line 46 of file spthread.c.

### 3.8.2 Member Data Documentation

#### 3.8.2.1 ack

```
volatile sig_atomic_t spthread_signal_args_st::ack
```

Definition at line 48 of file spthread.c.

### 3.8.2.2 `shutup_mutex`

```
pthread_mutex_t spthread_signal_args_st::shutup_mutex
```

Definition at line 49 of file `spthread.c`.

### 3.8.2.3 `signal`

```
const int spthread_signal_args_st::signal
```

Definition at line 47 of file `spthread.c`.

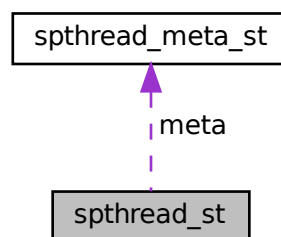
The documentation for this struct was generated from the following file:

- SRC/lib/[spthread.c](#)

## 3.9 `spthread_st` Struct Reference

```
#include <spthread.h>
```

Collaboration diagram for `spthread_st`:



### Public Attributes

- `pthread_t` [thread](#)
- `spthread_meta_t` \* [meta](#)

### 3.9.1 Detailed Description

Definition at line 28 of file `spthread.h`.

## 3.9.2 Member Data Documentation

### 3.9.2.1 meta

```
spthread_meta_t* spthread_st::meta
```

Definition at line 30 of file `spthread.h`.

### 3.9.2.2 thread

```
pthread_t spthread_st::thread
```

Definition at line 29 of file `spthread.h`.

The documentation for this struct was generated from the following file:

- SRC/lib/[spthread.h](#)

## 3.10 vec\_st Struct Reference

```
#include <Vec.h>
```

### Public Attributes

- [ptr\\_t](#) \* [data](#)
- [size\\_t](#) [length](#)
- [size\\_t](#) [capacity](#)
- [ptr\\_dtor\\_fn](#) [ele\\_dtor\\_fn](#)

### 3.10.1 Detailed Description

Definition at line 10 of file `Vec.h`.

### 3.10.2 Member Data Documentation

### 3.10.2.1 capacity

```
size_t vec_st::capacity
```

Definition at line 13 of file Vec.h.

### 3.10.2.2 data

```
ptr_t* vec_st::data
```

Definition at line 11 of file Vec.h.

### 3.10.2.3 ele\_dtor\_fn

```
ptr_dtor_fn vec_st::ele_dtor_fn
```

Definition at line 14 of file Vec.h.

### 3.10.2.4 length

```
size_t vec_st::length
```

Definition at line 12 of file Vec.h.

The documentation for this struct was generated from the following file:

- SRC/lib/[Vec.h](#)






## Chapter 4

# File Documentation

### 4.1 SRC/fs/fat\_routines.c File Reference

```
#include "fat_routines.h"
#include "../lib/pennos-errno.h"
#include "../shell/builtins.h"
#include "fs_helpers.h"
#include "fs_kfuncs.h"
#include <fcntl.h>
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/mman.h>
#include <sys/types.h>
#include <time.h>
#include <unistd.h>
```

Include dependency graph for fat\_routines.c:



### Functions

- int **mkfs** (const char \*fs\_name, int num\_blocks, int blk\_size)  
*Creates a PennFAT filesystem in the file named fs\_name at the OS-level.*
- int **mount** (const char \*fs\_name)  
*Mounts a filesystem with name fs\_name by loading its FAT into memory.*
- int **unmount** ()  
*Unmounts the current filesystem and reset variables.*
- void \* **cat** (void \*arg)  
*Concatenates and displays files.*

- void \* **ls** (void \*arg)  
*Searches root directory and lists all files in the directory.*
- void \* **touch** (void \*arg)  
*Creates files or updates timestamps.*
- void \* **mv** (void \*arg)  
*Renames files.*
- void \* **cp** (void \*arg)  
*Copies the source file to the destination.*
- void \* **rm** (void \*arg)  
*Removes files.*
- void \* **chmod** (void \*arg)  
*Changes the permissions of a file.*
- void \* **cmpctdir** (void \*arg)  
*Implements compaction of root directory.*

## 4.1.1 Function Documentation

### 4.1.1.1 cat()

```
void* cat (
    void * arg )
```

Concatenates and displays files.

This function reads the content of files and writes it to stdout or to another file. It supports reading from stdin when no input files are specified.

Usage formats:

- cat FILE ... (displays content to stdout)
- cat FILE ... -w OUTPUT\_FILE (writes content to OUTPUT\_FILE, overwriting)
- cat FILE ... -a OUTPUT\_FILE (appends content to OUTPUT\_FILE)
- cat -w OUTPUT\_FILE (reads from stdin, writes to OUTPUT\_FILE)
- cat -a OUTPUT\_FILE (reads from stdin, appends to OUTPUT\_FILE)

#### Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

#### Returns

void pointer (unused)

Definition at line 209 of file fat\_routines.c.

```

209     {
210     char** args = (char**)arg;
211
212     // verify that the file system is mounted
213     if (!is_mounted) {
214         P_ERRNO = P_EFS_NOT_MOUNTED;
215         u_perror("cat");
216         return NULL;
217     }
218
219     // early return if there is nothing after cat
220     if (args[1] == NULL) {
221         P_ERRNO = P_EINVAL;
222         u_perror("cat");
223         return NULL;
224     }
225
226     // check for output file with -w or -a flag
227     int out_fd = -1;
228     int out_mode = 0;
229
230     // scan arguments and determine output fd and output mode
231     int i;
232     for (i = 1; args[i] != NULL; i++) {
233         if (strcmp(args[i], "-w") == 0 && args[i + 1] != NULL) {
234             out_mode = F_WRITE;
235             out_fd = k_open(args[i + 1], F_WRITE);
236             if (out_fd < 0) {
237                 u_perror("cat");
238                 return NULL;
239             }
240             break;
241         } else if (strcmp(args[i], "-a") == 0 && args[i + 1] != NULL) {
242             out_mode = F_APPEND;
243             out_fd = k_open(args[i + 1], F_APPEND);
244             if (out_fd < 0) {
245                 u_perror("cat");
246                 return NULL;
247             }
248             break;
249         }
250     }
251
252     // if no output redirection found, use STDOUT
253     if (out_fd < 0) {
254         out_fd = STDOUT_FILENO;
255     }
256
257     // handle small case: cat -w OUTPUT_FILE or cat -a OUTPUT_FILE (read from stdin)
258     if ((strcmp(args[1], "-w") == 0 || strcmp(args[1], "-a") == 0) &&
259         args[2] != NULL && args[3] == NULL) {
260         char buffer[1024];
261
262         while (1) {
263             ssize_t bytes_read = k_read(STDIN_FILENO, buffer, sizeof(buffer));
264
265             if (bytes_read < 0) {
266                 u_perror("cat");
267                 if (out_fd != STDOUT_FILENO) {
268                     k_close(out_fd);
269                 }
270                 return NULL;
271             }
272
273             if (bytes_read == 0) {
274                 break;
275             }
276
277             if (k_write(out_fd, buffer, bytes_read) != bytes_read) {
278                 u_perror("cat");
279                 if (out_fd != STDOUT_FILENO) {
280                     k_close(out_fd);
281                 }
282                 return NULL;
283             }
284         }
285
286         if (out_fd != STDOUT_FILENO) {
287             k_close(out_fd);
288         }
289         return NULL;
290     }
291
292     // handle concatenating one or more files: cat FILE ... [-w/-a OUTPUT_FILE]
293     int start = 1;
294     int end = i - 1;
295

```

```

296 if (out_mode != 0) {
297     end = i - 1; // skip the output redirection arguments
298 }
299
300 // process each input file
301 for (i = start; i <= end; i++) {
302     // skip the redirection flags and their arguments
303     if (strcmp(args[i], "-w") == 0 || strcmp(args[i], "-a") == 0) {
304         i++;
305         continue;
306     }
307
308     // open the current input file
309     int in_fd = k_open(args[i], F_READ);
310     if (in_fd < 0) {
311         u_perror("cat");
312         continue;
313     }
314
315     // use lseek to get the size of in_fd
316     off_t in_fd_size = k_lseek(in_fd, 0, SEEK_END);
317     if (in_fd_size == -1) {
318         k_close(in_fd);
319         u_perror("cat");
320         continue;
321     }
322
323     // use lseek to reset position to 0 for reading
324     if (k_lseek(in_fd, 0, SEEK_SET) == -1) {
325         k_close(in_fd);
326         u_perror("cat");
327         continue;
328     }
329
330     // copy file content to output
331     char* buffer = (char*)malloc(block_size);
332     if (buffer == NULL) {
333         P_ERRNO = P_EMALLOC;
334         k_close(in_fd);
335         u_perror("cat");
336         continue;
337     }
338
339     int bytes_read;
340     ssize_t bytes_remaining = in_fd_size;
341
342     while (bytes_remaining > 0) {
343         ssize_t bytes_to_read = bytes_remaining < block_size ? bytes_remaining : block_size;
344         bytes_read = k_read(in_fd, buffer, bytes_to_read);
345
346         if (bytes_read <= 0) {
347             break;
348         }
349
350         if (k_write(out_fd, buffer, bytes_read) != bytes_read) {
351             free(buffer);
352             k_close(in_fd);
353             u_perror("cat");
354             break;
355         }
356
357         bytes_remaining -= bytes_read;
358     }
359
360     // read error
361     if (bytes_read < 0) {
362         free(buffer);
363         k_close(in_fd);
364         u_perror("cat");
365         continue;
366     }
367
368     k_close(in_fd);
369     free(buffer);
370 }
371
372 // close output file if not stdout
373 if (out_fd != STDOUT_FILENO) {
374     k_close(out_fd);
375 }
376
377 return NULL;
378 }

```

## 4.1.1.2 chmod()

```
void* chmod (
    void * arg )
```

Changes the permissions of a file.

Changes file permissions.

- chmod +x FILE (adds executable permission)
- chmod +rw FILE (adds read and write permissions)
- chmod -wx FILE (removes write and executable permissions)

Definition at line 751 of file fat\_routines.c.

```
751     {
752     char** args = (char**)arg;
753     if (!args || !args[0] || !args[1] || !args[2]) {
754         P_ERRNO = P_EINVAL;
755         return NULL;
756     }
757
758     // Parse permission string
759     const char* perm_str = args[1];
760     if (perm_str[0] != '+' && perm_str[0] != '-') {
761         P_ERRNO = P_EINVAL;
762         return NULL;
763     }
764
765     // Find the file and get its current directory entry
766     dir_entry_t dir_entry;
767     int entry_offset = find_file(args[2], &dir_entry);
768     if (entry_offset < 0) {
769         P_ERRNO = P_ENOENT;
770         return NULL;
771     }
772
773     // Calculate new permissions
774     uint8_t new_perm = dir_entry.perm;
775     int i = 1; // Start after + or -
776     while (perm_str[i] != '\0') {
777         switch (perm_str[i]) {
778             case 'r':
779                 if (perm_str[0] == '+') {
780                     new_perm |= PERM_READ;
781                 } else {
782                     new_perm &= ~PERM_READ;
783                 }
784                 break;
785             case 'w':
786                 if (perm_str[0] == '+') {
787                     new_perm |= PERM_WRITE;
788                 } else {
789                     new_perm &= ~PERM_WRITE;
790                 }
791                 break;
792             case 'x':
793                 if (perm_str[0] == '+') {
794                     new_perm |= PERM_EXEC;
795                 } else {
796                     new_perm &= ~PERM_EXEC;
797                 }
798                 break;
799             default:
800                 P_ERRNO = P_EINVAL;
801                 return NULL;
802         }
803         i++;
804     }
805
806     // Update the directory entry
807     dir_entry.perm = new_perm;
808     dir_entry.mtime = time(NULL);
809
810     // Seek to the entry's position
811     if (lseek(fs_fd, entry_offset, SEEK_SET) == -1) {
812         P_ERRNO = P_ELSEEK;
```

```

813     return NULL;
814 }
815
816 // Write the updated entry back
817 if (write(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
818     P_ERRNO = P_EWRITE;
819     return NULL;
820 }
821
822 return NULL;
823 }

```

#### 4.1.1.3 cmpctdir()

```

void* cmpctdir (
    void * arg )

```

Implements compaction of root directory.

Compacts the root directory by removing all deleted entries.

Definition at line 832 of file fat\_routines.c.

```

832     {
833     if (!is_mounted) {
834         P_ERRNO = P_EFS_NOT_MOUNTED;
835         u_perror("cmpctdir");
836         return NULL;
837     }
838
839     if (compact_directory() != 0) {
840         u_perror("cmpctdir");
841     }
842
843     return NULL;
844 }

```

#### 4.1.1.4 cp()

```

void* cp (
    void * arg )

```

Copies the source file to the destination.

Copies files.

Definition at line 621 of file fat\_routines.c.

```

621     {
622     char** args = (char**)arg;
623
624     // check that we have enough arguments
625     if (args[1] == NULL || args[2] == NULL) {
626         P_ERRNO = P_EINVAL;
627         u_perror("cp");
628         return NULL;
629     }
630
631     // cp -h SOURCE DEST
632     if (strcmp(args[1], "-h") == 0) {
633         if (args[2] == NULL || args[3] == NULL) {
634             P_ERRNO = P_EINVAL;
635             u_perror("cp");
636             return NULL;
637         }
638
639         if (copy_host_to_pennfat(args[2], args[3]) != 0) {

```

```

640     u_perror("cp");
641     return NULL;
642 }
643 return NULL;
644 }
645
646 // cp SOURCE -h DEST
647 if (args[2] != NULL && strcmp(args[2], "-h") == 0) {
648     if (args[3] == NULL) {
649         P_ERRNO = P_EINVAL;
650         u_perror("cp");
651         return NULL;
652     }
653
654     if (copy_pennfat_to_host(args[1], args[3]) != 0) {
655         u_perror("cp");
656         return NULL;
657     }
658     return NULL;
659 }
660
661 // cp SOURCE DEST
662 if ((args[1] != NULL && strcmp(args[1], "-h") != 0) &&
663     (args[2] != NULL && strcmp(args[2], "-h") != 0) && args[3] == NULL) {
664     if (copy_source_to_dest(args[1], args[2]) != 0) {
665         u_perror("cp");
666         return NULL;
667     }
668     return NULL;
669 }
670
671 P_ERRNO = P_EUNKNOWN;
672 u_perror("cp");
673 return NULL;
674 }

```

#### 4.1.1.5 ls()

```

void* ls (
    void * arg )

```

Searches root directory and lists all files in the directory.

Lists files in the current directory.

Definition at line 383 of file fat\_routines.c.

```

383     {
384     if (!is_mounted) {
385         P_ERRNO = P_EFS_NOT_MOUNTED;
386         u_perror("ls");
387         return NULL;
388     }
389
390     // start at root directory block
391     uint16_t current_block = 1;
392     int offset = 0;
393     dir_entry_t dir_entry;
394
395     while (1) {
396         // adjust pointer to beginning of current block
397         if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
398             P_ERRNO = P_ELSEEK;
399             u_perror("ls");
400             return NULL;
401         }
402
403         offset = 0;
404
405         // search current block
406         while (offset < block_size) {
407             if (read(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
408                 P_ERRNO = P_EREAD;
409                 u_perror("ls");
410                 return NULL;
411             }
412

```

```

413     // check if we've reached the end of directory
414     if (dir_entry.name[0] == 0) {
415         break;
416     }
417
418     // skip deleted entries
419     if (dir_entry.name[0] == 1 || dir_entry.name[0] == 2) {
420         offset += sizeof(dir_entry);
421         continue;
422     }
423
424     // format permission string
425     char perm_str[4] = "---";
426     if (dir_entry.perm & PERM_READ)
427         perm_str[0] = 'r';
428     if (dir_entry.perm & PERM_WRITE)
429         perm_str[1] = 'w';
430     if (dir_entry.perm & PERM_EXEC)
431         perm_str[2] = 'x';
432
433     // format time
434     struct tm* tm_info = localtime(&dir_entry.mtime);
435     char time_str[50];
436     strftime(time_str, sizeof(time_str), "%b %d %H:%M:%S %Y", tm_info);
437
438     // print entry details
439     char buffer[128];
440     int len;
441     if (dir_entry.firstBlock == 0) {
442         len = snprintf(buffer, sizeof(buffer), "  -%s- %6d %s %s\n",
443             perm_str, dir_entry.size, time_str, dir_entry.name);
444     } else {
445         len = snprintf(buffer, sizeof(buffer), "%2d -%s- %6d %s %s\n",
446             dir_entry.firstBlock, perm_str, dir_entry.size, time_str, dir_entry.name);
447     }
448
449     if (len < 0 || len >= (int)sizeof(buffer)) {
450         P_ERRNO = P_EUNKNOWN;
451         u_perror("ls");
452         return NULL;
453     }
454
455     if (k_write(STDOUT_FILENO, buffer, len) != len) {
456         P_ERRNO = P_EWRITE;
457         u_perror("ls");
458         return NULL;
459     }
460
461     offset += sizeof(dir_entry);
462 }
463
464 // move to the next block if there is one
465 if (fat[current_block] != FAT_EOF) {
466     current_block = fat[current_block];
467     continue;
468 }
469
470 // no more blocks to search
471 break;
472 }
473
474 return NULL;
475 }

```

#### 4.1.1.6 mkfs()

```

int mkfs (
    const char * fs_name,
    int num_blocks,
    int blk_size )

```

Creates a PennFAT filesystem in the file named `fs_name` at the OS-level.

Creates a PennFAT filesystem in the file named `fs_name`.

Definition at line 29 of file `fat_routines.c`.



```

29                                     {
30     // validate arguments
31     if (num_blocks < 1 || num_blocks > 32) {
32         P_ERRNO = P_EINVAL;
33         return -1;
34     }
35     if (blk_size < 0 || blk_size > 4) {
36         P_ERRNO = P_EINVAL;
37         return -1;
38     }
39
40     // determine the file system size
41     int block_sizes[] = {256, 512, 1024, 2048, 4096};
42     int actual_block_size = block_sizes[blk_size];
43     int fat_size = num_blocks * actual_block_size;
44     int fat_entries = fat_size / 2;
45     int num_data_blocks =
46         (num_blocks == 32)
47         ? fat_entries - 2
48         : fat_entries - 1; // note: first entry is reserved for metadata!
49     size_t filesystem_size = fat_size + (actual_block_size * num_data_blocks);
50
51     // create the file for the filesystem
52     int fd = open(fs_name, O_RDWR | O_CREAT | O_TRUNC, 0644);
53     if (fd == -1) {
54         P_ERRNO = P_EOPEN;
55         return -1;
56     }
57
58     // extend the file to the required size
59     if (ftruncate(fd, filesystem_size) == -1) {
60         P_ERRNO = P_EFUNC;
61         close(fd);
62         return -1;
63     }
64
65     // allocate the FAT
66     uint16_t* temp_fat = (uint16_t*)calloc(fat_entries, sizeof(uint16_t));
67     if (!temp_fat) {
68         P_ERRNO = P_EMALLOC;
69         close(fd);
70         return -1;
71     }
72
73     // initialize FAT entries to their correct values
74     temp_fat[0] = (num_blocks « 8) | blk_size;
75     temp_fat[1] = FAT_EOF;
76     for (int i = 2; i < fat_entries; i++) {
77         temp_fat[i] = FAT_FREE;
78     }
79
80     // write the FAT to the file
81     if (write(fd, temp_fat, fat_size) != fat_size) {
82         P_ERRNO = P_EWRITE;
83         free(temp_fat);
84         close(fd);
85         return -1;
86     }
87
88     // initialize the root directory + write to memory
89     uint8_t* root_dir = (uint8_t*)calloc(actual_block_size, 1);
90     if (lseek(fd, fat_size, SEEK_SET) == -1) {
91         P_ERRNO = P_ELSEEK;
92         free(temp_fat);
93         free(root_dir);
94         close(fd);
95         return -1;
96     }
97     if (write(fd, root_dir, actual_block_size) != actual_block_size) {
98         P_ERRNO = P_EWRITE;
99         free(temp_fat);
100        free(root_dir);
101        close(fd);
102        return -1;
103    }
104
105    // clean up
106    free(temp_fat);
107    free(root_dir);
108    close(fd);
109    return 0;
110 }

```

#### 4.1.1.7 mount()

```
int mount (
    const char * fs_name )
```

Mounts a filesystem with name `fs_name` by loading its FAT into memory.

Mounts the filesystem named `fs_name` by loading its FAT into memory.

Definition at line 115 of file `fat_routines.c`.

```
115     {
116     // check if a filesystem is already mounted
117     if (is_mounted) {
118         P_ERRNO = P_EBUSY;
119         return -1;
120     }
121
122     // open the file with fs_name + set the global fs_fd
123     fs_fd = open(fs_name, O_RDWR);
124     if (fs_fd == -1) {
125         P_ERRNO = P_ENOENT;
126         return -1;
127     }
128
129     // read the first two bytes to get size configuration
130     uint16_t config;
131     if (read(fs_fd, &config, sizeof(config)) != sizeof(config)) {
132         P_ERRNO = P_EREAD;
133         close(fs_fd);
134         fs_fd = -1;
135         return -1;
136     }
137
138     // extract FAT region size information
139     num_fat_blocks = (config >> 8) & 0xFF; // MSB
140     int block_size_config = config & 0xFF; // LSB
141     int block_sizes[] = {256, 512, 1024, 2048, 4096};
142     block_size = block_sizes[block_size_config];
143     fat_size = num_fat_blocks * block_size;
144
145     // map the FAT region into memory
146     if (lseek(fs_fd, 0, SEEK_SET) == -1) {
147         P_ERRNO = P_ELSEEK;
148         close(fs_fd);
149         fs_fd = -1;
150         return -1;
151     }
152
153     fat = mmap(NULL, fat_size, PROT_READ | PROT_WRITE, MAP_SHARED, fs_fd, 0);
154     if (fat == MAP_FAILED) {
155         P_ERRNO = P_EMAP;
156         close(fs_fd);
157         fs_fd = -1;
158         return -1;
159     }
160
161     init_fd_table(fd_table); // initialize the file descriptor table
162     is_mounted = true;
163     return 0;
164 }
```

#### 4.1.1.8 mv()

```
void* mv (
    void * arg )
```

Renames files.

Renames the source file to the destination name. If the destination file already exists, it will be overwritten.

Usage: `mv SOURCE DEST`

## Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

## Returns

void pointer (unused)

Definition at line 543 of file fat\_routines.c.

```

543     {
544     char** args = (char**)arg;
545
546     // verify that the file system is mounted
547     if (!is_mounted) {
548         P_ERRNO = P_EFS_NOT_MOUNTED;
549         u_perror("mv");
550         return NULL;
551     }
552
553     // check if we have both source and destination arguments
554     if (args[1] == NULL || args[2] == NULL) {
555         P_ERRNO = P_EINVAL;
556         u_perror("mv");
557         return NULL;
558     }
559
560     char* source = args[1];
561     char* dest = args[2];
562
563     // check if they're trying to rename to the same name
564     if (strcmp(source, dest) == 0) {
565         return NULL;
566     }
567
568     // check if source file exists
569     dir_entry_t source_entry;
570     int source_offset = find_file(source, &source_entry);
571     if (source_offset < 0) {
572         u_perror("mv");
573         return NULL;
574     }
575
576     // check if the destination file already exists
577     dir_entry_t dest_entry;
578     int dest_offset = find_file(dest, &dest_entry);
579
580     // destination file exists
581     if (dest_offset >= 0) {
582         // check if the destination file is currently open by any process
583         for (int i = 0; i < MAX_FDS; i++) {
584             if (fd_table[i].in_use && strcmp(fd_table[i].filename, dest) == 0) {
585                 P_ERRNO = P_EBUSY;
586                 u_perror("mv");
587                 return NULL;
588             }
589         }
590
591         // if destination file exists, delete it
592         if (mark_entry_as_deleted(&dest_entry, dest_offset) != 0) {
593             u_perror("mv");
594             return NULL;
595         }
596     }
597
598     // rename file
599     strncpy(source_entry.name, dest, sizeof(source_entry.name) - 1);
600     source_entry.name[sizeof(source_entry.name) - 1] = '\0';
601
602     // write the updated entry back to disk
603     if (lseek(fs_fd, source_offset, SEEK_SET) == -1) {
604         P_ERRNO = P_ELSEEK;
605         u_perror("mv");
606         return NULL;
607     }
608
609     if (write(fs_fd, &source_entry, sizeof(source_entry)) != sizeof(source_entry)) {
610         P_ERRNO = P_EWRITE;
611         u_perror("mv");
612         return NULL;
613     }

```

```

614
615     return NULL;
616 }

```

#### 4.1.1.9 rm()

```

void* rm (
    void * arg )

```

Removes files.

Deletes one or more files from the filesystem. Each file is processed as a separate transaction.

##### Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

##### Returns

void pointer (unused)

Definition at line 679 of file fat\_routines.c.

```

679     {
680     char** args = (char**)arg;
681
682     // verify that the file system is mounted
683     if (!is_mounted) {
684         P_ERRNO = P_EFS_NOT_MOUNTED;
685         u_perror("rm");
686         return NULL;
687     }
688
689     // check if we have any arguments
690     if (args[1] == NULL) {
691         P_ERRNO = P_EINVAL;
692         u_perror("rm");
693         return NULL;
694     }
695
696     // process each file argument
697     for (int i = 1; args[i] != NULL; i++) {
698         // find the file in the directory
699         dir_entry_t entry;
700         int entry_offset = find_file(args[i], &entry);
701
702         if (entry_offset < 0) {
703             // file doesn't exist
704             P_ERRNO = P_ENOENT;
705             u_perror("rm");
706             continue;
707         }
708
709         // check if file is currently open
710         for (int j = 0; j < MAX_FDS; j++) {
711             if (fd_table[j].in_use && strcmp(fd_table[j].filename, args[i]) == 0) {
712                 P_ERRNO = P_EBUSY;
713                 u_perror("rm");
714                 continue;
715             }
716         }
717
718         // mark the directory entry as deleted
719         if (lseek(fs_fd, entry_offset, SEEK_SET) == -1) {
720             P_ERRNO = P_ELSEEK;
721             u_perror("rm");
722             continue;
723         }
724     }

```

```

725     char deleted = 1; // mark as deleted
726     if (write(fs_fd, &deleted, sizeof(deleted)) != sizeof(deleted)) {
727         P_ERRNO = P_EWRITE;
728         u_perror("rm");
729         continue;
730     }
731
732     // free the FAT chain for this file
733     uint16_t block = entry.firstBlock;
734     while (block != FAT_FREE && block != FAT_EOF) {
735         uint16_t next_block = fat[block];
736         fat[block] = FAT_FREE;
737         block = next_block;
738     }
739 }
740
741 return NULL;
742 }

```

#### 4.1.1.10 touch()

```

void* touch (
    void * arg )

```

Creates files or updates timestamps.

Creates empty files or updates timestamps.

For each file argument, creates the file if it doesn't exist, or updates its timestamp if it already exists.

Definition at line 483 of file fat\_routines.c.

```

483     {
484     char** args = (char**)arg;
485
486     // verify that the file system is mounted
487     if (!is_mounted) {
488         P_ERRNO = P_EFS_NOT_MOUNTED;
489         u_perror("touch");
490         return NULL;
491     }
492
493     // check if we have any arguments
494     if (args[1] == NULL) {
495         P_ERRNO = P_EINVAL;
496         u_perror("touch");
497         return NULL;
498     }
499
500     // process each file argument
501     for (int i = 1; args[i] != NULL; i++) {
502         dir_entry_t entry;
503         int entry_offset = find_file(args[i], &entry);
504
505         // file exists
506         if (entry_offset >= 0) {
507             entry.mtime = time(NULL);
508
509             // write the updated entry back to the directory
510             if (lseek(fs_fd, entry_offset, SEEK_SET) == -1) {
511                 P_ERRNO = P_ELSEEK;
512                 u_perror("touch");
513                 continue;
514             }
515             if (write(fs_fd, &entry, sizeof(entry)) != sizeof(entry)) {
516                 P_ERRNO = P_EWRITE;
517                 u_perror("touch");
518                 continue;
519             }
520         } else {
521             // file doesn't exist, create a new empty file
522
523             // check if the fat is full
524             if (P_ERRNO == P_EFULL) {
525                 u_perror("touch");
526                 return NULL;

```

```

527     }
528
529     // add the file entry to root directory
530     if (add_file_entry(args[i], 0, 0, TYPE_REGULAR, PERM_READ_WRITE) == -1) {
531         u_perror("touch");
532         continue;
533     }
534 }
535 }
536
537 return NULL;
538 }

```

#### 4.1.1.11 unmount()

```
int unmount ( )
```

Unmounts the current filesystem and reset variables.

Unmounts the currently mounted filesystem.

Definition at line 169 of file fat\_routines.c.

```

169     {
170     // first check that a file system is actually mounted
171     if (!is_mounted) {
172         P_ERRNO = P_EFS_NOT_MOUNTED;
173         return -1;
174     }
175
176     // unmap the FAT
177     if (fat != NULL) {
178         if (munmap(fat, fat_size) == -1) {
179             P_ERRNO = P_EMAP;
180             return -1;
181         }
182         fat = NULL;
183     }
184
185     // close fs_fd
186     if (fs_fd != -1) {
187         if (close(fs_fd) == -1) {
188             P_ERRNO = P_ECLOSE;
189             return -1;
190         }
191         fs_fd = -1;
192     }
193
194     // reset the other globals
195     num_fat_blocks = 0;
196     block_size = 0;
197     fat_size = 0;
198     is_mounted = false;
199     return 0;
200 }

```


## 4.2 SRC/fs/fat\_routines.h File Reference

```


#include <stdint.h>
#include <time.h>

```

Include dependency graph for fat\_routines.h:



This graph shows which files directly or indirectly include this file:



## Classes

- struct [dir\\_entry\\_t](#)  
*Directory entry structure for files in the filesystem.*
- struct [fd\\_entry\\_t](#)  
*File descriptor entry structure for open files.*

## Macros

- #define [FAT\\_EOF](#) 0xFFFF
- #define [FAT\\_FREE](#) 0x0000
- #define [TYPE\\_UNKNOWN](#) 0
- #define [TYPE\\_REGULAR](#) 1
- #define [TYPE\\_DIRECTORY](#) 2
- #define [TYPE\\_SYMLINK](#) 4
- #define [PERM\\_NONE](#) 0
- #define [PERM\\_WRITE](#) 1
- #define [PERM\\_READ](#) 2
- #define [PERM\\_EXEC](#) 4
- #define [PERM\\_READ\\_WRITE](#) ([PERM\\_READ](#) | [PERM\\_WRITE](#))
- #define [PERM\\_READ\\_EXEC](#) ([PERM\\_READ](#) | [PERM\\_EXEC](#))
- #define [PERM\\_READ\\_WRITE\\_EXEC](#) ([PERM\\_READ](#) | [PERM\\_WRITE](#) | [PERM\\_EXEC](#))
- #define [F\\_READ](#) 0x01
- #define [F\\_WRITE](#) 0x02
- #define [F\\_APPEND](#) 0x04

## Functions

- int **mkfs** (const char \*fs\_name, int num\_blocks, int **block\_size**)  
*Creates a PennFAT filesystem in the file named fs\_name.*
- int **mount** (const char \*fs\_name)  
*Mounts the filesystem named fs\_name by loading its FAT into memory.*
- int **unmount** ()  
*Unmounts the currently mounted filesystem.*
- void \* **cat** (void \*arg)  
*Concatenates and displays files.*
- void \* **ls** (void \*arg)  
*Lists files in the current directory.*
- void \* **touch** (void \*arg)  
*Creates empty files or updates timestamps.*
- void \* **mv** (void \*arg)  
*Renames files.*
- void \* **cp** (void \*arg)  
*Copies files.*
- void \* **rm** (void \*arg)  
*Removes files.*
- void \* **chmod** (void \*arg)  
*Changes file permissions.*
- void \* **cmpctdir** (void \*arg)  
*Compacts the root directory by removing all deleted entries.*

## 4.2.1 Macro Definition Documentation

### 4.2.1.1 F\_APPEND

```
#define F_APPEND 0x04
```

Definition at line 37 of file fat\_routines.h.

### 4.2.1.2 F\_READ

```
#define F_READ 0x01
```

Definition at line 35 of file fat\_routines.h.



#### 4.2.1.3 F\_WRITE

```
#define F_WRITE 0x02
```

Definition at line 36 of file fat\_routines.h.

#### 4.2.1.4 FAT\_EOF

```
#define FAT_EOF 0xFFFF
```

Definition at line 16 of file fat\_routines.h.

#### 4.2.1.5 FAT\_FREE

```
#define FAT_FREE 0x0000
```

Definition at line 17 of file fat\_routines.h.

#### 4.2.1.6 PERM\_EXEC

```
#define PERM_EXEC 4
```

Definition at line 29 of file fat\_routines.h.

#### 4.2.1.7 PERM\_NONE

```
#define PERM_NONE 0
```

Definition at line 26 of file fat\_routines.h.

#### 4.2.1.8 PERM\_READ

```
#define PERM_READ 2
```

Definition at line 28 of file fat\_routines.h.

#### 4.2.1.9 PERM\_READ\_EXEC

```
#define PERM_READ_EXEC (PERM_READ | PERM_EXEC)
```

Definition at line 31 of file fat\_routines.h.

#### 4.2.1.10 PERM\_READ\_WRITE

```
#define PERM_READ_WRITE (PERM_READ | PERM_WRITE)
```

Definition at line 30 of file fat\_routines.h.

#### 4.2.1.11 PERM\_READ\_WRITE\_EXEC

```
#define PERM_READ_WRITE_EXEC (PERM_READ | PERM_WRITE | PERM_EXEC)
```

Definition at line 32 of file fat\_routines.h.

#### 4.2.1.12 PERM\_WRITE

```
#define PERM_WRITE 1
```

Definition at line 27 of file fat\_routines.h.

#### 4.2.1.13 TYPE\_DIRECTORY

```
#define TYPE_DIRECTORY 2
```

Definition at line 22 of file fat\_routines.h.

#### 4.2.1.14 TYPE\_REGULAR

```
#define TYPE_REGULAR 1
```

Definition at line 21 of file fat\_routines.h.

#### 4.2.1.15 TYPE\_SYMLINK

```
#define TYPE_SYMLINK 4
```

Definition at line 23 of file fat\_routines.h.

#### 4.2.1.16 TYPE\_UNKNOWN

```
#define TYPE_UNKNOWN 0
```

Definition at line 20 of file fat\_routines.h.

### 4.2.2 Function Documentation

#### 4.2.2.1 cat()

```
void* cat (  
    void * arg )
```

Concatenates and displays files.

This function reads the content of files and writes it to stdout or to another file. It supports reading from stdin when no input files are specified.

Usage formats:

- cat FILE ... (displays content to stdout)
- cat FILE ... -w OUTPUT\_FILE (writes content to OUTPUT\_FILE, overwriting)
- cat FILE ... -a OUTPUT\_FILE (appends content to OUTPUT\_FILE)
- cat -w OUTPUT\_FILE (reads from stdin, writes to OUTPUT\_FILE)
- cat -a OUTPUT\_FILE (reads from stdin, appends to OUTPUT\_FILE)

#### Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

#### Returns

void pointer (unused)

Definition at line 209 of file fat\_routines.c.

```

209     {
210     char** args = (char**)arg;
211
212     // verify that the file system is mounted
213     if (!is_mounted) {
214         P_ERRNO = P_EFS_NOT_MOUNTED;
215         u_perror("cat");
216         return NULL;
217     }
218
219     // early return if there is nothing after cat
220     if (args[1] == NULL) {
221         P_ERRNO = P_EINVAL;
222         u_perror("cat");
223         return NULL;
224     }
225
226     // check for output file with -w or -a flag
227     int out_fd = -1;
228     int out_mode = 0;
229
230     // scan arguments and determine output fd and output mode
231     int i;
232     for (i = 1; args[i] != NULL; i++) {
233         if (strcmp(args[i], "-w") == 0 && args[i + 1] != NULL) {
234             out_mode = F_WRITE;
235             out_fd = k_open(args[i + 1], F_WRITE);
236             if (out_fd < 0) {
237                 u_perror("cat");
238                 return NULL;
239             }
240             break;
241         } else if (strcmp(args[i], "-a") == 0 && args[i + 1] != NULL) {
242             out_mode = F_APPEND;
243             out_fd = k_open(args[i + 1], F_APPEND);
244             if (out_fd < 0) {
245                 u_perror("cat");
246                 return NULL;
247             }
248             break;
249         }
250     }
251
252     // if no output redirection found, use STDOUT
253     if (out_fd < 0) {
254         out_fd = STDOUT_FILENO;
255     }
256
257     // handle small case: cat -w OUTPUT_FILE or cat -a OUTPUT_FILE (read from stdin)
258     if ((strcmp(args[1], "-w") == 0 || strcmp(args[1], "-a") == 0) &&
259         args[2] != NULL && args[3] == NULL) {
260         char buffer[1024];
261
262         while (1) {
263             ssize_t bytes_read = k_read(STDIN_FILENO, buffer, sizeof(buffer));
264
265             if (bytes_read < 0) {
266                 u_perror("cat");
267                 if (out_fd != STDOUT_FILENO) {
268                     k_close(out_fd);
269                 }
270                 return NULL;
271             }
272
273             if (bytes_read == 0) {
274                 break;
275             }
276
277             if (k_write(out_fd, buffer, bytes_read) != bytes_read) {
278                 u_perror("cat");
279                 if (out_fd != STDOUT_FILENO) {
280                     k_close(out_fd);
281                 }
282                 return NULL;
283             }
284         }
285
286         if (out_fd != STDOUT_FILENO) {
287             k_close(out_fd);
288         }
289         return NULL;
290     }
291
292     // handle concatenating one or more files: cat FILE ... [-w/-a OUTPUT_FILE]
293     int start = 1;
294     int end = i - 1;
295

```

```

296     if (out_mode != 0) {
297         end = i - 1; // skip the output redirection arguments
298     }
299
300     // process each input file
301     for (i = start; i <= end; i++) {
302         // skip the redirection flags and their arguments
303         if (strcmp(args[i], "-w") == 0 || strcmp(args[i], "-a") == 0) {
304             i++;
305             continue;
306         }
307
308         // open the current input file
309         int in_fd = k_open(args[i], F_READ);
310         if (in_fd < 0) {
311             u_perror("cat");
312             continue;
313         }
314
315         // use lseek to get the size of in_fd
316         off_t in_fd_size = k_lseek(in_fd, 0, SEEK_END);
317         if (in_fd_size == -1) {
318             k_close(in_fd);
319             u_perror("cat");
320             continue;
321         }
322
323         // use lseek to reset position to 0 for reading
324         if (k_lseek(in_fd, 0, SEEK_SET) == -1) {
325             k_close(in_fd);
326             u_perror("cat");
327             continue;
328         }
329
330         // copy file content to output
331         char* buffer = (char*)malloc(block_size);
332         if (buffer == NULL) {
333             P_ERRNO = P_EMALLOC;
334             k_close(in_fd);
335             u_perror("cat");
336             continue;
337         }
338
339         int bytes_read;
340         ssize_t bytes_remaining = in_fd_size;
341
342         while (bytes_remaining > 0) {
343             ssize_t bytes_to_read = bytes_remaining < block_size ? bytes_remaining : block_size;
344             bytes_read = k_read(in_fd, buffer, bytes_to_read);
345
346             if (bytes_read <= 0) {
347                 break;
348             }
349
350             if (k_write(out_fd, buffer, bytes_read) != bytes_read) {
351                 free(buffer);
352                 k_close(in_fd);
353                 u_perror("cat");
354                 break;
355             }
356
357             bytes_remaining -= bytes_read;
358         }
359
360         // read error
361         if (bytes_read < 0) {
362             free(buffer);
363             k_close(in_fd);
364             u_perror("cat");
365             continue;
366         }
367
368         k_close(in_fd);
369         free(buffer);
370     }
371
372     // close output file if not stdout
373     if (out_fd != STDOUT_FILENO) {
374         k_close(out_fd);
375     }
376
377     return NULL;
378 }

```

#### 4.2.2.2 chmod()

```
void* chmod (
    void * arg )
```

Changes file permissions.

Modifies the permissions of the specified file.

Usage formats:

- chmod +x FILE (adds executable permission)
- chmod +rw FILE (adds read and write permissions)
- chmod -wx FILE (removes write and executable permissions)

##### Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

##### Returns

void pointer (unused)

Changes file permissions.

- chmod +x FILE (adds executable permission)
- chmod +rw FILE (adds read and write permissions)
- chmod -wx FILE (removes write and executable permissions)

Definition at line 751 of file fat\_routines.c.

```
751     {
752     char** args = (char**)arg;
753     if (!args || !args[0] || !args[1] || !args[2]) {
754         P_ERRNO = P_EINVAL;
755         return NULL;
756     }
757
758     // Parse permission string
759     const char* perm_str = args[1];
760     if (perm_str[0] != '+' && perm_str[0] != '-') {
761         P_ERRNO = P_EINVAL;
762         return NULL;
763     }
764
765     // Find the file and get its current directory entry
766     dir_entry_t dir_entry;
767     int entry_offset = find_file(args[2], &dir_entry);
768     if (entry_offset < 0) {
769         P_ERRNO = P_ENOENT;
770         return NULL;
771     }
772
773     // Calculate new permissions
774     uint8_t new_perm = dir_entry.perm;
775     int i = 1; // Start after + or -
776     while (perm_str[i] != '\0') {
777         switch (perm_str[i]) {
778             case 'r':
779                 if (perm_str[0] == '+') {
780                     new_perm |= PERM_READ;
```

```

781         } else {
782             new_perm &= ~PERM_READ;
783         }
784         break;
785     case 'w':
786         if (perm_str[0] == '+') {
787             new_perm |= PERM_WRITE;
788         } else {
789             new_perm &= ~PERM_WRITE;
790         }
791         break;
792     case 'x':
793         if (perm_str[0] == '+') {
794             new_perm |= PERM_EXEC;
795         } else {
796             new_perm &= ~PERM_EXEC;
797         }
798         break;
799     default:
800         P_ERRNO = P_EINVAL;
801         return NULL;
802     }
803     i++;
804 }
805
806 // Update the directory entry
807 dir_entry.perm = new_perm;
808 dir_entry.mtime = time(NULL);
809
810 // Seek to the entry's position
811 if (lseek(fs_fd, entry_offset, SEEK_SET) == -1) {
812     P_ERRNO = P_ELSEEK;
813     return NULL;
814 }
815
816 // Write the updated entry back
817 if (write(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
818     P_ERRNO = P_EWRITE;
819     return NULL;
820 }
821
822 return NULL;
823 }

```

### 4.2.2.3 cmpctdir()

```

void* cmpctdir (
    void * arg )

```

Compacts the root directory by removing all deleted entries.

#### Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

#### Returns

void pointer (unused)

Compacts the root directory by removing all deleted entries.

Definition at line 832 of file fat\_routines.c.

```

832     {
833     if (!is_mounted) {
834         P_ERRNO = P_EFS_NOT_MOUNTED;
835         u_perror("cmpctdir");
836         return NULL;
837     }
838 }

```

```

839  if (compact_directory() != 0) {
840      u_perror("cmpctdir");
841  }
842
843  return NULL;
844 }

```

#### 4.2.2.4 cp()

```

void* cp (
    void * arg )

```

Copies files.

Copies the source file to the destination. If the destination file already exists, it will be overwritten.

Usage formats:

- cp SOURCE DEST (copies within PennFAT)
- cp -h SOURCE DEST (copies from host OS to PennFAT)
- cp SOURCE -h DEST (copies from PennFAT to host OS)

#### Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

#### Returns

return 0 on success, -1 on error

Copies files.

Definition at line 621 of file fat\_routines.c.

```

621  {
622  char** args = (char**)arg;
623
624  // check that we have enough arguments
625  if (args[1] == NULL || args[2] == NULL) {
626      P_ERRNO = P_EINVAL;
627      u_perror("cp");
628      return NULL;
629  }
630
631  // cp -h SOURCE DEST
632  if (strcmp(args[1], "-h") == 0) {
633      if (args[2] == NULL || args[3] == NULL) {
634          P_ERRNO = P_EINVAL;
635          u_perror("cp");
636          return NULL;
637      }
638
639      if (copy_host_to_pennfat(args[2], args[3]) != 0) {
640          u_perror("cp");
641          return NULL;
642      }
643      return NULL;
644  }
645
646  // cp SOURCE -h DEST
647  if (args[2] != NULL && strcmp(args[2], "-h") == 0) {

```



```

648     if (args[3] == NULL) {
649         P_ERRNO = P_EINVAL;
650         u_perror("cp");
651         return NULL;
652     }
653
654     if (copy_pennfat_to_host(args[1], args[3]) != 0) {
655         u_perror("cp");
656         return NULL;
657     }
658     return NULL;
659 }
660
661 // cp SOURCE DEST
662 if ((args[1] != NULL && strcmp(args[1], "-h") != 0) &&
663     (args[2] != NULL && strcmp(args[2], "-h") != 0) && args[3] == NULL) {
664     if (copy_source_to_dest(args[1], args[2]) != 0) {
665         u_perror("cp");
666         return NULL;
667     }
668     return NULL;
669 }
670
671 P_ERRNO = P_EUNKNOWN;
672 u_perror("cp");
673 return NULL;
674 }

```

#### 4.2.2.5 ls()

```

void* ls (
    void * arg )

```

Lists files in the current directory.

This function displays information about files in the current directory, including block number, permissions, size, and name.

##### Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

##### Returns

0 on success, -1 on error

Lists files in the current directory.

Definition at line 383 of file fat\_routines.c.

```

383     {
384     if (!is_mounted) {
385         P_ERRNO = P_EFS_NOT_MOUNTED;
386         u_perror("ls");
387         return NULL;
388     }
389
390     // start at root directory block
391     uint16_t current_block = 1;
392     int offset = 0;
393     dir_entry_t dir_entry;
394
395     while (1) {
396         // adjust pointer to beginning of current block
397         if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
398             P_ERRNO = P_ELSEEK;
399             u_perror("ls");
400             return NULL;

```

```

401     }
402
403     offset = 0;
404
405     // search current block
406     while (offset < block_size) {
407         if (read(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
408             P_ERRNO = P_EREAD;
409             u_perror("ls");
410             return NULL;
411         }
412
413         // check if we've reached the end of directory
414         if (dir_entry.name[0] == 0) {
415             break;
416         }
417
418         // skip deleted entries
419         if (dir_entry.name[0] == 1 || dir_entry.name[0] == 2) {
420             offset += sizeof(dir_entry);
421             continue;
422         }
423
424         // format permission string
425         char perm_str[4] = "---";
426         if (dir_entry.perm & PERM_READ)
427             perm_str[0] = 'r';
428         if (dir_entry.perm & PERM_WRITE)
429             perm_str[1] = 'w';
430         if (dir_entry.perm & PERM_EXEC)
431             perm_str[2] = 'x';
432
433         // format time
434         struct tm* tm_info = localtime(&dir_entry.mtime);
435         char time_str[50];
436         strftime(time_str, sizeof(time_str), "%b %d %H:%M:%S %Y", tm_info);
437
438         // print entry details
439         char buffer[128];
440         int len;
441         if (dir_entry.firstBlock == 0) {
442             len = snprintf(buffer, sizeof(buffer), "  -%-6d %s %s\n",
443                           perm_str, dir_entry.size, time_str, dir_entry.name);
444         } else {
445             len = snprintf(buffer, sizeof(buffer), "%2d -%-6d %s %s\n",
446                           dir_entry.firstBlock, perm_str, dir_entry.size, time_str, dir_entry.name);
447         }
448
449         if (len < 0 || len >= (int)sizeof(buffer)) {
450             P_ERRNO = P_EUNKNOWN;
451             u_perror("ls");
452             return NULL;
453         }
454
455         if (k_write(STDOUT_FILENO, buffer, len) != len) {
456             P_ERRNO = P_EWRITE;
457             u_perror("ls");
458             return NULL;
459         }
460
461         offset += sizeof(dir_entry);
462     }
463
464     // move to the next block if there is one
465     if (fat[current_block] != FAT_EOF) {
466         current_block = fat[current_block];
467         continue;
468     }
469
470     // no more blocks to search
471     break;
472 }
473
474 return NULL;
475 }

```

#### 4.2.2.6 mkfs()

```

int mkfs (
    const char * fs_name,

```

```
int num_blocks,
int blk_size )
```

Creates a PennFAT filesystem in the file named `fs_name`.

This function initializes a new PennFAT filesystem with the specified parameters. The number of blocks in the FAT ranges from 1 through 32, and the block size is determined by `block_size` (0=256B, 1=512B, 2=1024B, 3=2048B, 4=4096B).

#### Parameters

<i>fs_name</i>	The name of the file to create the filesystem in.
<i>num_blocks</i>	The number of blocks in the FAT region (1-32).
<i>block_size</i>	The block size configuration (0-4).

Creates a PennFAT filesystem in the file named `fs_name`.

Definition at line 29 of file `fat_routines.c`.

```
29
30 // validate arguments
31 if (num_blocks < 1 || num_blocks > 32) {
32     P_ERRNO = P_EINVAL;
33     return -1;
34 }
35 if (blk_size < 0 || blk_size > 4) {
36     P_ERRNO = P_EINVAL;
37     return -1;
38 }
39
40 // determine the file system size
41 int block_sizes[] = {256, 512, 1024, 2048, 4096};
42 int actual_block_size = block_sizes[blk_size];
43 int fat_size = num_blocks * actual_block_size;
44 int fat_entries = fat_size / 2;
45 int num_data_blocks =
46     (num_blocks == 32)
47     ? fat_entries - 2
48     : fat_entries - 1; // note: first entry is reserved for metadata!
49 size_t filesystem_size = fat_size + (actual_block_size * num_data_blocks);
50
51 // create the file for the filesystem
52 int fd = open(fs_name, O_RDWR | O_CREAT | O_TRUNC, 0644);
53 if (fd == -1) {
54     P_ERRNO = P_EOPEN;
55     return -1;
56 }
57
58 // extend the file to the required size
59 if (ftruncate(fd, filesystem_size) == -1) {
60     P_ERRNO = P_EFUNC;
61     close(fd);
62     return -1;
63 }
64
65 // allocate the FAT
66 uint16_t* temp_fat = (uint16_t*)calloc(fat_entries, sizeof(uint16_t));
67 if (!temp_fat) {
68     P_ERRNO = P_EMALLOC;
69     close(fd);
70     return -1;
71 }
72
73 // initialize FAT entries to their correct values
74 temp_fat[0] = (num_blocks << 8) | blk_size;
75 temp_fat[1] = FAT_EOF;
76 for (int i = 2; i < fat_entries; i++) {
77     temp_fat[i] = FAT_FREE;
78 }
79
80 // write the FAT to the file
81 if (write(fd, temp_fat, fat_size) != fat_size) {
82     P_ERRNO = P_EWRITE;
83     free(temp_fat);
84     close(fd);
85     return -1;
86 }
```

```

87
88 // initialize the root directory + write to memory
89 uint8_t* root_dir = (uint8_t*)calloc(actual_block_size, 1);
90 if (lseek(fd, fat_size, SEEK_SET) == -1) {
91     P_ERRNO = P_ELSEEK;
92     free(temp_fat);
93     free(root_dir);
94     close(fd);
95     return -1;
96 }
97 if (write(fd, root_dir, actual_block_size) != actual_block_size) {
98     P_ERRNO = P_EWRITE;
99     free(temp_fat);
100     free(root_dir);
101     close(fd);
102     return -1;
103 }
104
105 // clean up
106 free(temp_fat);
107 free(root_dir);
108 close(fd);
109 return 0;
110 }

```

#### 4.2.2.7 mount()

```

int mount (
    const char * fs_name )

```

Mounts the filesystem named `fs_name` by loading its FAT into memory.

This function loads the filesystem's FAT into memory for subsequent operations. Only one filesystem can be mounted at a time.

##### Parameters

<code>fs_name</code>	The name of the filesystem file to mount.
----------------------	---

##### Returns

0 on success, -1 on failure with `P_ERRNO` set.

Mounts the filesystem named `fs_name` by loading its FAT into memory.

Definition at line 115 of file `fat_routines.c`.

```

115 {
116 // check if a filesystem is already mounted
117 if (is_mounted) {
118     P_ERRNO = P_EBUSY;
119     return -1;
120 }
121
122 // open the file with fs_name + set the global fs_fd
123 fs_fd = open(fs_name, O_RDWR);
124 if (fs_fd == -1) {
125     P_ERRNO = P_ENOENT;
126     return -1;
127 }
128
129 // read the first two bytes to get size configuration
130 uint16_t config;
131 if (read(fs_fd, &config, sizeof(config)) != sizeof(config)) {
132     P_ERRNO = P_EREAD;
133     close(fs_fd);
134     fs_fd = -1;
135     return -1;

```

```

136 }
137
138 // extract FAT region size information
139 num_fat_blocks = (config >> 8) & 0xFF; // MSB
140 int block_size_config = config & 0xFF; // LSB
141 int block_sizes[] = {256, 512, 1024, 2048, 4096};
142 block_size = block_sizes[block_size_config];
143 fat_size = num_fat_blocks * block_size;
144
145 // map the FAT region into memory
146 if (lseek(fs_fd, 0, SEEK_SET) == -1) {
147     P_ERRNO = P_ELSEEK;
148     close(fs_fd);
149     fs_fd = -1;
150     return -1;
151 }
152
153 fat = mmap(NULL, fat_size, PROT_READ | PROT_WRITE, MAP_SHARED, fs_fd, 0);
154 if (fat == MAP_FAILED) {
155     P_ERRNO = P_EMAP;
156     close(fs_fd);
157     fs_fd = -1;
158     return -1;
159 }
160
161 init_fd_table(fd_table); // initialize the file descriptor table
162 is_mounted = true;
163 return 0;
164 }

```

#### 4.2.2.8 mv()

```

void* mv (
    void * arg )

```

Renames files.

Renames the source file to the destination name. If the destination file already exists, it will be overwritten.

Usage: mv SOURCE DEST

##### Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

##### Returns

void pointer (unused)

Definition at line 543 of file fat\_routines.c.

```

543 {
544     char** args = (char**)arg;
545
546     // verify that the file system is mounted
547     if (!is_mounted) {
548         P_ERRNO = P_EFS_NOT_MOUNTED;
549         u_perror("mv");
550         return NULL;
551     }
552
553     // check if we have both source and destination arguments
554     if (args[1] == NULL || args[2] == NULL) {
555         P_ERRNO = P_EINVAL;
556         u_perror("mv");
557         return NULL;
558     }
559 }

```

```

560 char* source = args[1];
561 char* dest = args[2];
562
563 // check if they're trying to rename to the same name
564 if (strcmp(source, dest) == 0) {
565     return NULL;
566 }
567
568 // check if source file exists
569 dir_entry_t source_entry;
570 int source_offset = find_file(source, &source_entry);
571 if (source_offset < 0) {
572     u_perror("mv");
573     return NULL;
574 }
575
576 // check if the destination file already exists
577 dir_entry_t dest_entry;
578 int dest_offset = find_file(dest, &dest_entry);
579
580 // destination file exists
581 if (dest_offset >= 0) {
582     // check if the destination file is currently open by any process
583     for (int i = 0; i < MAX_FDS; i++) {
584         if (fd_table[i].in_use && strcmp(fd_table[i].filename, dest) == 0) {
585             P_ERRNO = P_EBUSY;
586             u_perror("mv");
587             return NULL;
588         }
589     }
590
591     // if destination file exists, delete it
592     if (mark_entry_as_deleted(&dest_entry, dest_offset) != 0) {
593         u_perror("mv");
594         return NULL;
595     }
596 }
597
598 // rename file
599 strncpy(source_entry.name, dest, sizeof(source_entry.name) - 1);
600 source_entry.name[sizeof(source_entry.name) - 1] = '\0';
601
602 // write the updated entry back to disk
603 if (lseek(fs_fd, source_offset, SEEK_SET) == -1) {
604     P_ERRNO = P_ELSEEK;
605     u_perror("mv");
606     return NULL;
607 }
608
609 if (write(fs_fd, &source_entry, sizeof(source_entry)) != sizeof(source_entry)) {
610     P_ERRNO = P_EWRITE;
611     u_perror("mv");
612     return NULL;
613 }
614
615 return NULL;
616 }

```

#### 4.2.2.9 rm()

```

void* rm (
    void * arg )

```

Removes files.

Deletes one or more files from the filesystem. Each file is processed as a separate transaction.

##### Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

## Returns

void pointer (unused)

Definition at line 679 of file fat\_routines.c.

```

679     {
680     char** args = (char**)arg;
681
682     // verify that the file system is mounted
683     if (!is_mounted) {
684         P_ERRNO = P_EFS_NOT_MOUNTED;
685         u_perror("rm");
686         return NULL;
687     }
688
689     // check if we have any arguments
690     if (args[1] == NULL) {
691         P_ERRNO = P_EINVAL;
692         u_perror("rm");
693         return NULL;
694     }
695
696     // process each file argument
697     for (int i = 1; args[i] != NULL; i++) {
698         // find the file in the directory
699         dir_entry_t entry;
700         int entry_offset = find_file(args[i], &entry);
701
702         if (entry_offset < 0) {
703             // file doesn't exist
704             P_ERRNO = P_ENOENT;
705             u_perror("rm");
706             continue;
707         }
708
709         // check if file is currently open
710         for (int j = 0; j < MAX_FDS; j++) {
711             if (fd_table[j].in_use && strcmp(fd_table[j].filename, args[i]) == 0) {
712                 P_ERRNO = P_EBUSY;
713                 u_perror("rm");
714                 continue;
715             }
716         }
717
718         // mark the directory entry as deleted
719         if (lseek(fs_fd, entry_offset, SEEK_SET) == -1) {
720             P_ERRNO = P_ELSEEK;
721             u_perror("rm");
722             continue;
723         }
724
725         char deleted = 1; // mark as deleted
726         if (write(fs_fd, &deleted, sizeof(deleted)) != sizeof(deleted)) {
727             P_ERRNO = P_EWRITE;
728             u_perror("rm");
729             continue;
730         }
731
732         // free the FAT chain for this file
733         uint16_t block = entry.firstBlock;
734         while (block != FAT_FREE && block != FAT_EOF) {
735             uint16_t next_block = fat[block];
736             fat[block] = FAT_FREE;
737             block = next_block;
738         }
739     }
740
741     return NULL;
742 }

```

#### 4.2.2.10 touch()

```

void* touch (
    void * arg )

```

Creates empty files or updates timestamps.

For each file argument, creates the file if it doesn't exist, or updates its timestamp if it already exists.

## Parameters

<i>arg</i>	Arguments array (command line arguments)
------------	--

## Returns

void pointer (unused)

Creates empty files or updates timestamps.

For each file argument, creates the file if it doesn't exist, or updates its timestamp if it already exists.

Definition at line 483 of file fat\_routines.c.

```

483     {
484     char** args = (char**)arg;
485
486     // verify that the file system is mounted
487     if (!is_mounted) {
488         P_ERRNO = P_EFS_NOT_MOUNTED;
489         u_perror("touch");
490         return NULL;
491     }
492
493     // check if we have any arguments
494     if (args[1] == NULL) {
495         P_ERRNO = P_EINVAL;
496         u_perror("touch");
497         return NULL;
498     }
499
500     // process each file argument
501     for (int i = 1; args[i] != NULL; i++) {
502         dir_entry_t entry;
503         int entry_offset = find_file(args[i], &entry);
504
505         // file exists
506         if (entry_offset >= 0) {
507             entry.mtime = time(NULL);
508
509             // write the updated entry back to the directory
510             if (lseek(fs_fd, entry_offset, SEEK_SET) == -1) {
511                 P_ERRNO = P_ELSEEK;
512                 u_perror("touch");
513                 continue;
514             }
515             if (write(fs_fd, &entry, sizeof(entry)) != sizeof(entry)) {
516                 P_ERRNO = P_EWRITE;
517                 u_perror("touch");
518                 continue;
519             }
520         } else {
521             // file doesn't exist, create a new empty file
522
523             // check if the fat is full
524             if (P_ERRNO == P_EFULL) {
525                 u_perror("touch");
526                 return NULL;
527             }
528
529             // add the file entry to root directory
530             if (add_file_entry(args[i], 0, 0, TYPE_REGULAR, PERM_READ_WRITE) == -1) {
531                 u_perror("touch");
532                 continue;
533             }
534         }
535     }
536
537     return NULL;
538 }

```



#### 4.2.2.11 unmount()

```
int unmount ( )
```

Unmounts the currently mounted filesystem.

This function flushes any pending changes and unmounts the filesystem.

##### Returns

0 on success, -1 on failure with P\_ERRNO set.

Unmounts the currently mounted filesystem.

Definition at line 169 of file fat\_routines.c.


```
169 {
170     // first check that a file system is actually mounted
171     if (!is_mounted) {
172         P_ERRNO = P_EFS_NOT_MOUNTED;
173         return -1;
174     }
175
176     // unmap the FAT
177     if (fat != NULL) {
178         if (munmap(fat, fat_size) == -1) {
179             P_ERRNO = P_EMAP;
180             return -1;
181         }
182         fat = NULL;
183     }
184
185     // close fs_fd
186     if (fs_fd != -1) {
187         if (close(fs_fd) == -1) {
188             P_ERRNO = P_ECLOSE;
189             return -1;
190         }
191         fs_fd = -1;
192     }
193
194     // reset the other globals
195     num_fat_blocks = 0;
196     block_size = 0;
197     fat_size = 0;
198     is_mounted = false;
199     return 0;
200 }
```

## 4.3 SRC/fs/fs\_helpers.c File Reference

```
#include "fs_helpers.h"
#include "fat_routines.h"
#include "fs_kfuncs.h"
#include "lib/pennos-errno.h"
#include "shell/builtins.h"
#include <fcntl.h>
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/mman.h>
#include <sys/types.h>
#include <time.h>
```

```
#include <unistd.h>
```

Include dependency graph for fs\_helpers.c:



## Functions

- void [init\\_fd\\_table](#) ([fd\\_entry\\_t](#) \*[fd\\_table](#))  
*Initializes the global kernel-level file descriptor table.*
- int [get\\_free\\_fd](#) ([fd\\_entry\\_t](#) \*[fd\\_table](#))  
*Gets a free file descriptor.*
- int [increment\\_fd\\_ref\\_count](#) (int [fd](#))  
*Increments the reference count of a file descriptor.*
- int [decrement\\_fd\\_ref\\_count](#) (int [fd](#))  
*Decrements the reference count of a file descriptor.*
- int [has\\_executable\\_permission](#) (int [fd](#))  
*Checks if a file has executable permissions.*
- [uint16\\_t](#) [allocate\\_block](#) ()  
*Allocates a block.*
- int [find\\_file](#) (const char \*[filename](#), [dir\\_entry\\_t](#) \*[entry](#))  
*Searches for a file in the root directory.*
- int [add\\_file\\_entry](#) (const char \*[filename](#), [uint32\\_t](#) [size](#), [uint16\\_t](#) [first\\_block](#), [uint8\\_t](#) [type](#), [uint8\\_t](#) [perm](#))  
*Adds a file to the root directory.*
- int [mark\\_entry\\_as\\_deleted](#) ([dir\\_entry\\_t](#) \*[entry](#), int [absolute\\_offset](#))  
*Marks a file entry as deleted and frees its blocks.*
- int [copy\\_host\\_to\\_pennfat](#) (const char \*[host\\_filename](#), const char \*[pennfat\\_filename](#))  
*Copies data from host OS file to the PennFAT file.*
- int [copy\\_pennfat\\_to\\_host](#) (const char \*[pennfat\\_filename](#), const char \*[host\\_filename](#))  
*Copies data from PennFAT file to host OS file.*
- int [copy\\_source\\_to\\_dest](#) (const char \*[source\\_filename](#), const char \*[dest\\_filename](#))  
*Copies data from source file to destination file.*
- int [compact\\_directory](#) ()  
*Compacts a directory.*

## Variables

- int [fs\\_fd](#) = -1
- int [block\\_size](#) = 0
- int [num\\_fat\\_blocks](#) = 0
- int [fat\\_size](#) = 0
- [uint16\\_t](#) \* [fat](#) = NULL
- bool [is\\_mounted](#) = false
- int [MAX\\_FDS](#) = 100
- [fd\\_entry\\_t](#) [fd\\_table](#) [100]

## 4.3.1 Function Documentation

### 4.3.1.1 add\_file\_entry()

```
int add_file_entry (
    const char * filename,
    uint32_t size,
    uint16_t first_block,
    uint8_t type,
    uint8_t perm )
```

Adds a file to the root directory.

Adds a new file entry to the root directory.

Definition at line 266 of file fs\_helpers.c.

```
270     {
271     if (!is_mounted) {
272         P_ERRNO = P_EFS_NOT_MOUNTED;
273         return -1;
274     }
275
276     // check if file already exists
277     dir_entry_t existing;
278     if (find_file(filename, &existing) >= 0) {
279         P_ERRNO = P_EEXIST;
280         return -1;
281     }
282
283     // start with root directory block (block 1)
284     uint16_t current_block = 1;
285     int offset = 0;
286     dir_entry_t dir_entry;
287
288     while (1) {
289         // position at the start of current block of the root directory
290         if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
291             P_ERRNO = P_ELSEEK;
292             return -1;
293         }
294
295         // reset offset for new block
296         offset = 0;
297
298         // search current block for free slot
299         while (offset < block_size) {
300             if (read(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
301                 P_ERRNO = P_EREAD;
302                 return -1;
303             }
304
305             // found a free slot
306             if (dir_entry.name[0] == 0 || dir_entry.name[0] == 1) {
307                 // initialize the new entry
308                 memset(&dir_entry, 0, sizeof(dir_entry));
309                 strncpy(dir_entry.name, filename, 31);
310                 dir_entry.size = size;
311                 dir_entry.firstBlock = first_block;
312                 dir_entry.type = type;
313                 dir_entry.perm = perm;
314                 dir_entry.mtime = time(NULL);
315
316                 // write the entry
317                 if (lseek(fs_fd, fat_size + (current_block - 1) * block_size + offset, SEEK_SET) == -1) {
318                     P_ERRNO = P_ELSEEK;
319                     return -1;
320                 }
321                 if (write(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
322                     P_ERRNO = P_EWRITE;
323                     return -1;
324                 }
325
326                 return offset;
```

```

327     }
328
329     offset += sizeof(dir_entry);
330 }
331
332 // current block is full, check if there's a next block
333 if (fat[current_block] != FAT_EOF) {
334     current_block = fat[current_block];
335     continue;
336 }
337
338 // allocate a new block for the root directory
339 uint16_t new_block = allocate_block();
340 if (new_block == 0) {
341     P_ERRNO = P_EFULL;
342     return -1;
343 }
344
345 // chain the new block
346 fat[current_block] = new_block;
347 fat[new_block] = FAT_EOF;
348
349 // initialize new block
350 uint8_t* zero_block = calloc(block_size, 1);
351 if (!zero_block) {
352     P_ERRNO = P_EINVAL;
353     return -1;
354 }
355
356 // write this new block to the file system
357 if (lseek(fs_fd, fat_size + (new_block - 1) * block_size, SEEK_SET) == -1) {
358     P_ERRNO = P_ELSEEK;
359     free(zero_block);
360     return -1;
361 }
362 if (write(fs_fd, zero_block, block_size) != block_size) {
363     P_ERRNO = P_EWRITE;
364     free(zero_block);
365     return -1;
366 }
367
368 free(zero_block);
369
370 // initialize the new entry
371 memset(&dir_entry, 0, sizeof(dir_entry));
372 strncpy(dir_entry.name, filename, 31);
373 dir_entry.size = size;
374 dir_entry.firstBlock = first_block;
375 dir_entry.type = type;
376 dir_entry.perm = perm;
377 dir_entry.mtime = time(NULL);
378
379 // write the new entry at the start of the new block in the file system
380 if (lseek(fs_fd, fat_size + (new_block - 1) * block_size, SEEK_SET) == -1) {
381     P_ERRNO = P_ELSEEK;
382     return -1;
383 }
384 if (write(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
385     P_ERRNO = P_EWRITE;
386     return -1;
387 }
388
389 return 0;
390 }
391 }

```

#### 4.3.1.2 allocate\_block()

```
uint16_t allocate_block ( )
```

Allocates a block.

Allocates a free block in the FAT.

If no block found, we try compacting the directory.

Definition at line 167 of file fs\_helpers.c.

```

167     {
168     for (int i = 2; i < fat_size / 2; i++) {
169         if (fat[i] == FAT_FREE) {
170             fat[i] = FAT_EOF;
171             return i;
172         }
173     }
174
175     if (compact_directory() == 0) {
176         for (int i = 2; i < fat_size / 2; i++) {
177             if (fat[i] == FAT_FREE) {
178                 fat[i] = FAT_EOF;
179                 return i;
180             }
181         }
182     }
183
184     return 0;
185 }

```

#### 4.3.1.3 compact\_directory()

```
int compact_directory ( )
```

Compacts a directory.

Compacts the root directory by removing all deleted entries.

Definition at line 692 of file fs\_helpers.c.

```

692     {
693     if (!is_mounted) {
694         P_ERRNO = P_EFS_NOT_MOUNTED;
695         return -1;
696     }
697
698     // buffer for temp storage of a block
699     uint8_t* dir_buffer = malloc(block_size);
700     if (!dir_buffer) {
701         P_ERRNO = P_EMALLOC;
702         return -1;
703     }
704
705     // start at root directory
706     uint16_t current_block = 1;
707     int dir_entries_count = 0;
708     int deleted_entries_count = 0;
709
710     // calculate number of entries and deleted entries in the root directory
711     while (current_block != FAT_EOF) {
712         if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
713             P_ERRNO = P_ELSEEK;
714             free(dir_buffer);
715             return -1;
716         }
717
718         if (read(fs_fd, dir_buffer, block_size) != block_size) {
719             P_ERRNO = P_EREAD;
720             free(dir_buffer);
721             return -1;
722         }
723
724         // count entries and deleted entries in this block
725         for (int offset = 0; offset < block_size; offset += sizeof(dir_entry_t)) {
726             dir_entry_t* entry = (dir_entry_t*)(dir_buffer + offset);
727
728             // check if we've reached the end of directory
729             if (entry->name[0] == 0) {
730                 break;
731             }
732
733             dir_entries_count++;
734
735             // check if it's a deleted entry
736             if (entry->name[0] == 1) {
737                 deleted_entries_count++;
738             }

```

```

739     }
740
741     // move onto next block, if there is one
742     if (fat[current_block] != FAT_EOF) {
743         current_block = fat[current_block];
744     } else {
745         break;
746     }
747 }
748
749 // if no deleted entries, no compaction needed
750 if (deleted_entries_count == 0) {
751     free(dir_buffer);
752     return 0;
753 }
754
755 // allocate space for all valid entries
756 dir_entry_t* all_entries = malloc(dir_entries_count * sizeof(dir_entry_t));
757 if (!all_entries) {
758     P_ERRNO = P_EMALLOC;
759     free(dir_buffer);
760     return -1;
761 }
762
763 // read all entries into the buffer, skipping deleted ones
764 current_block = 1;
765 int valid_entry_idx = 0;
766
767 while (current_block != FAT_EOF) {
768     if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
769         P_ERRNO = P_ELSEEK;
770         free(dir_buffer);
771         free(all_entries);
772         return -1;
773     }
774
775     if (read(fs_fd, dir_buffer, block_size) != block_size) {
776         P_ERRNO = P_EREAD;
777         free(dir_buffer);
778         free(all_entries);
779         return -1;
780     }
781
782     // process entries in this block
783     for (int offset = 0; offset < block_size; offset += sizeof(dir_entry_t)) {
784         dir_entry_t* entry = (dir_entry_t*)(dir_buffer + offset);
785
786         // check if we've reached the end of directory
787         if (entry->name[0] == 0) {
788             break;
789         }
790
791         // skip deleted entries
792         if (entry->name[0] == 1) {
793             continue;
794         }
795
796         // copy valid entry to our array
797         memcpy(&all_entries[valid_entry_idx++], entry, sizeof(dir_entry_t));
798     }
799
800     // move to the next block
801     if (fat[current_block] != FAT_EOF) {
802         current_block = fat[current_block];
803     } else {
804         break;
805     }
806 }
807
808 // rewrite the directory with only valid entries
809 current_block = 1;
810 int entries_per_block = block_size / sizeof(dir_entry_t);
811 int blocks_needed = (valid_entry_idx + entries_per_block - 1) / entries_per_block;
812
813 // clean up any excess directory blocks in the FAT chain
814 uint16_t next_block = fat[current_block];
815 if (blocks_needed == 1) {
816     // only need one block, free all others
817     while (next_block != FAT_EOF) {
818         uint16_t temp = fat[next_block];
819         fat[next_block] = FAT_FREE;
820         next_block = temp;
821     }
822     fat[current_block] = FAT_EOF;
823 } else {
824     // navigate through needed blocks
825     int block_count = 1;

```

```

826     uint16_t prev_block = current_block;
827
828     while (block_count < blocks_needed) {
829         if (next_block == FAT_EOF) {
830             // need to allocate a new block
831             uint16_t new_block = allocate_block();
832             if (new_block == 0) {
833                 P_ERRNO = P_EFULL;
834                 free(dir_buffer);
835                 free(all_entries);
836                 return -1;
837             }
838             fat[prev_block] = new_block;
839             next_block = new_block;
840         }
841
842         prev_block = next_block;
843         next_block = fat[next_block];
844         block_count++;
845     }
846
847     // free any excess blocks
848     fat[prev_block] = FAT_EOF;
849     while (next_block != FAT_EOF) {
850         uint16_t temp = fat[next_block];
851         fat[next_block] = FAT_FREE;
852         next_block = temp;
853     }
854 }
855
856 // write the valid entries back to the directory blocks
857 current_block = 1;
858 int entries_written = 0;
859
860 while (entries_written < valid_entry_idx) {
861     if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
862         P_ERRNO = P_ELSEEK;
863         free(dir_buffer);
864         free(all_entries);
865         return -1;
866     }
867
868     memset(dir_buffer, 0, block_size);
869
870     // copy entries to the buffer
871     int entries_in_this_block = 0;
872     while (entries_written < valid_entry_idx && entries_in_this_block < entries_per_block) {
873         memcpy(dir_buffer + (entries_in_this_block * sizeof(dir_entry_t)),
874             &all_entries[entries_written],
875             sizeof(dir_entry_t));
876         entries_written++;
877         entries_in_this_block++;
878     }
879
880     // write the buffer to the file system
881     if (write(fs_fd, dir_buffer, block_size) != block_size) {
882         P_ERRNO = P_EINVAL;
883         free(dir_buffer);
884         free(all_entries);
885         return -1;
886     }
887
888     // move to the next block if needed
889     if (entries_written < valid_entry_idx) {
890         current_block = fat[current_block];
891     }
892 }
893
894 free(dir_buffer);
895 free(all_entries);
896 return 0;
897 }

```

#### 4.3.1.4 copy\_host\_to\_pennfat()

```

int copy_host_to_pennfat (
    const char * host_filename,
    const char * pennfat_filename )

```

Copies data from host OS file to the PennFAT file.

Copies a file from the host OS to the PennFAT filesystem.

Definition at line 434 of file fs\_helpers.c.

```

435                                     {
436     if (!is_mounted) {
437         P_ERRNO = P_EFS_NOT_MOUNTED;
438         return -1;
439     }
440
441     // open the host file
442     int host_fd = open(host_filename, O_RDONLY);
443     if (host_fd == -1) {
444         P_ERRNO = P_EOPEN;
445         return -1;
446     }
447
448     // determine file size by seeking to the end and getting position
449     off_t host_file_size_in_bytes = lseek(host_fd, 0, SEEK_END);
450     if (host_file_size_in_bytes == -1) {
451         P_ERRNO = P_ELSEEK;
452         close(host_fd);
453         return -1;
454     }
455
456     // go back to beginning of file for reading
457     if (lseek(host_fd, 0, SEEK_SET) == -1) {
458         P_ERRNO = P_ELSEEK;
459         close(host_fd);
460         return -1;
461     }
462
463     // open the destination file in PennFAT
464     int pennfat_fd = k_open(pennfat_filename, F_WRITE);
465     if (pennfat_fd < 0) {
466         close(host_fd);
467         return -1;
468     }
469
470     // copy the data into this buffer
471     uint8_t* buffer = (uint8_t*)malloc(block_size);
472     if (!buffer) {
473         P_ERRNO = P_EMALLOC;
474         k_close(pennfat_fd);
475         close(host_fd);
476         return -1;
477     }
478
479     uint32_t bytes_remaining = host_file_size_in_bytes;
480     ssize_t bytes_read;
481
482     // read from host file
483     while (bytes_remaining > 0) {
484         // ensure bytes to read never exceeds the block size
485         ssize_t bytes_to_read = bytes_remaining < block_size ? bytes_remaining : block_size;
486         bytes_read = read(host_fd, buffer, bytes_to_read);
487
488         if (bytes_read <= 0) {
489             break;
490         }
491
492         // write to pennfat_fd using k_write
493         if (k_write(pennfat_fd, (const char*)buffer, bytes_read) != bytes_read) {
494             free(buffer);
495             k_close(pennfat_fd);
496             close(host_fd);
497             return -1;
498         }
499
500         bytes_remaining -= bytes_read;
501     }
502
503     // check for read error
504     if (bytes_read < 0) {
505         P_ERRNO = P_EREAD;
506         free(buffer);
507         k_close(pennfat_fd);
508         close(host_fd);
509         return -1;
510     }
511
512     // otherwise, cleanup and return success
513     free(buffer);
514     k_close(pennfat_fd);

```



```

515     close(host_fd);
516     return 0;
517 }

```

#### 4.3.1.5 copy\_pennfat\_to\_host()

```

int copy_pennfat_to_host (
    const char * pennfat_filename,
    const char * host_filename )

```

Copies data from PennFAT file to host OS file.

Copies a file from the PennFAT filesystem to the host OS.

Definition at line 522 of file fs\_helpers.c.

```

523     {
524     if (!is_mounted) {
525         P_ERRNO = P_EFS_NOT_MOUNTED;
526         return -1;
527     }
528
529     // open the PennFAT file
530     int pennfat_fd = k_open(pennfat_filename, F_READ);
531     if (pennfat_fd < 0) {
532         return -1;
533     }
534
535     // get the pennfat file size
536     off_t pennfat_file_size_in_bytes = k_lseek(pennfat_fd, 0, SEEK_END);
537     if (pennfat_file_size_in_bytes == -1) {
538         k_close(pennfat_fd);
539         return -1;
540     }
541
542     // go back to beginning of file for reading
543     if (k_lseek(pennfat_fd, 0, SEEK_SET) == -1) {
544         k_close(pennfat_fd);
545         return -1;
546     }
547
548     // open the host file
549     int host_fd = open(host_filename, O_WRONLY | O_CREAT | O_TRUNC, 0644);
550     if (host_fd == -1) {
551         P_ERRNO = P_EOPEN;
552         k_close(pennfat_fd);
553         return -1;
554     }
555
556     // allocate buffer for data transfer
557     char* buffer = (char*)malloc(block_size);
558     if (!buffer) {
559         P_ERRNO = P_EMALLOC;
560         k_close(pennfat_fd);
561         close(host_fd);
562         return -1;
563     }
564
565     uint32_t bytes_remaining = pennfat_file_size_in_bytes;
566     ssize_t bytes_read;
567
568     // read from PennFAT file and write to host file
569     while (bytes_remaining > 0) {
570         // ensure bytes to read never exceeds the block size
571         ssize_t bytes_to_read = bytes_remaining < block_size ? bytes_remaining : block_size;
572         bytes_read = k_read(pennfat_fd, buffer, bytes_to_read);
573
574         if (bytes_read <= 0) {
575             break;
576         }
577
578         if (write(host_fd, buffer, bytes_read) != bytes_read) {
579             P_ERRNO = P_EINVAL;
580             free(buffer);
581             close(host_fd);
582             k_close(pennfat_fd);
583             return -1;

```

```

584     }
585
586     bytes_remaining -= bytes_read;
587 }
588
589 // check for read error
590 if (bytes_read < 0) {
591     P_ERRNO = P_EREAD;
592     free(buffer);
593     close(host_fd);
594     k_close(pennfat_fd);
595     return -1;
596 }
597
598 // otherwise, cleanup and return success
599 free(buffer);
600 close(host_fd);
601 k_close(pennfat_fd);
602 return 0;
603 }

```

#### 4.3.1.6 copy\_source\_to\_dest()

```

int copy_source_to_dest (
    const char * source_filename,
    const char * dest_filename )

```

Copies data from source file to destination file.

Copies a file from a source file to a destination file.

Definition at line 608 of file fs\_helpers.c.

```

609                                     {
610     if (!is_mounted) {
611         P_ERRNO = P_EFS_NOT_MOUNTED;
612         return -1;
613     }
614
615     // open the source file
616     int source_fd = k_open(source_filename, F_READ);
617     if (source_fd < 0) {
618         return -1;
619     }
620
621     // get the source file size
622     off_t source_file_size_in_bytes = k_lseek(source_fd, 0, SEEK_END);
623     if (source_file_size_in_bytes == -1) {
624         k_close(source_fd);
625         return -1;
626     }
627
628     // move to the beginning of the source file for reading
629     if (k_lseek(source_fd, 0, SEEK_SET) < 0) {
630         k_close(source_fd);
631         return -1;
632     }
633
634     // open the destination file
635     int dest_fd = k_open(dest_filename, F_WRITE);
636     if (dest_fd < 0) {
637         k_close(source_fd);
638         return -1;
639     }
640
641     // read from source to destination
642     char* buffer = (char*)malloc(block_size);
643     if (!buffer) {
644         P_ERRNO = P_EMALLOC;
645         k_close(source_fd);
646         k_close(dest_fd);
647         return -1;
648     }
649
650     uint32_t bytes_remaining = source_file_size_in_bytes;
651     ssize_t bytes_read;
652

```

```

653 while (bytes_remaining > 0) {
654     // make sure the bytes to read doesn't exceed block size
655     ssize_t bytes_to_read = bytes_remaining < block_size ? bytes_remaining : block_size;
656     bytes_read = k_read(source_fd, buffer, bytes_to_read);
657
658     if (bytes_read <= 0) {
659         break;
660     }
661
662     if (k_write(dest_fd, buffer, bytes_read) != bytes_read) {
663         free(buffer);
664         k_close(source_fd);
665         k_close(dest_fd);
666         return -1;
667     }
668 }
669
670 // check for read error
671 if (bytes_read < 0) {
672     free(buffer);
673     k_close(source_fd);
674     k_close(dest_fd);
675     return -1;
676 }
677
678 // otherwise, cleanup and return success
679 free(buffer);
680 k_close(source_fd);
681 k_close(dest_fd);
682 return 0;
683 }

```

#### 4.3.1.7 decrement\_fd\_ref\_count()

```

int decrement_fd_ref_count (
    int fd )

```

Decrements the reference count of a file descriptor.

If reference count reaches 0, flush field values.

Definition at line 107 of file fs\_helpers.c.

```

107 {
108     if (fd < 0 || fd >= MAX_FDS) {
109         P_ERRNO = P_EBADF;
110         return -1;
111     }
112
113     if (!fd_table[fd].in_use) {
114         P_ERRNO = P_EBADF;
115         return -1;
116     }
117
118     fd_table[fd].ref_count--;
119     if (fd_table[fd].ref_count == 0) {
120         fd_table[fd].in_use = 0;
121         memset(fd_table[fd].filename, 0, sizeof(fd_table[fd].filename));
122         fd_table[fd].size = 0;
123         fd_table[fd].first_block = 0;
124         fd_table[fd].position = 0;
125         fd_table[fd].mode = 0;
126     }
127     return fd_table[fd].ref_count;
128 }

```

#### 4.3.1.8 find\_file()

```
int find_file (
    const char * filename,
    dir_entry_t * entry )
```

Searches for a file in the root directory.

Retrieves the file's absolute offset in the filesystem.

Definition at line 192 of file fs\_helpers.c.

```
192                                     {
193     if (!is_mounted) {
194         P_ERRNO = P_EFS_NOT_MOUNTED;
195         return -1;
196     }
197
198     // Start with root directory block (block 1)
199     uint16_t current_block = 1;
200     int offset_in_block = 0;
201     int absolute_offset = 0;
202     dir_entry_t dir_entry;
203
204     while (1) {
205         // Position at the start of current block
206         if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
207             P_ERRNO = P_ELSEEK;
208             return -1;
209         }
210
211         // reset offset for new block
212         offset_in_block = 0;
213
214         // calculate the absolute offset
215         absolute_offset = fat_size + (current_block - 1) * block_size;
216
217         // search current block
218         while (offset_in_block < block_size) {
219             if (read(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
220                 P_ERRNO = P_EREAD;
221                 return -1;
222             }
223
224             // check if we've reached the end of directory
225             if (dir_entry.name[0] == 0) {
226                 break;
227             }
228
229             // check if this is a deleted entry
230             if (dir_entry.name[0] == 1 || dir_entry.name[0] == 2) {
231                 offset_in_block += sizeof(dir_entry);
232                 absolute_offset += sizeof(dir_entry);
233                 continue;
234             }
235
236             // check if we found the file
237             if (strcmp(dir_entry.name, filename) == 0) {
238                 if (entry) {
239                     memcpy(entry, &dir_entry, sizeof(dir_entry));
240                 }
241                 return absolute_offset; // return the absolute file offset
242             }
243
244             offset_in_block += sizeof(dir_entry);
245             absolute_offset += sizeof(dir_entry);
246         }
247
248         // if we've reached the end of the current block, check if there's a next block
249         if (fat[current_block] != FAT_EOF) {
250             current_block = fat[current_block];
251             continue;
252         }
253
254         // no more blocks to search
255         break;
256     }
257
258     // file not found
259     P_ERRNO = P_ENOENT;
260     return -1;
261 }
```

#### 4.3.1.9 get\_free\_fd()

```
int get_free_fd (
    fd_entry_t * fd_table )
```

Gets a free file descriptor.

Finds the first available file descriptor in the table.

Definition at line 77 of file fs\_helpers.c.

```
77 {
78     for (int i = 3; i < MAX_FDS; i++) {
79         if (!fd_table[i].in_use) {
80             return i;
81         }
82     }
83     return -1;
84 }
```

#### 4.3.1.10 has\_executable\_permission()

```
int has_executable_permission (
    int fd )
```

Checks if a file has executable permissions.

Checks if a file has executable permissions in the PennFAT filesystem.

Definition at line 133 of file fs\_helpers.c.

```
133 {
134     // check if fs is mounted
135     if (!is_mounted) {
136         P_ERRNO = P_EFS_NOT_MOUNTED;
137         return -1;
138     }
139
140     // validate fd argument
141     if (fd < 0 || fd >= MAX_FDS) {
142         P_ERRNO = P_EINVAL;
143         return -1;
144     }
145
146
147     // determine whether the file exists
148     dir_entry_t entry;
149     int entry_offset = find_file(fd_table[fd].filename, &entry);
150     if (entry_offset < 0) {
151         return -1;
152     }
153
154     // if it exists, get its permission
155     if (entry.perm & PERM_EXEC) {
156         return 1;
157     }
158
159     return 0;
160 }
```

#### 4.3.1.11 increment\_fd\_ref\_count()

```
int increment_fd_ref_count (
    int fd )
```

Increments the reference count of a file descriptor.

**Parameters**

<i>fd</i>	file descriptor to increment
-----------	------------------------------

**Returns**

new reference count, or -1 on error

Definition at line 89 of file fs\_helpers.c.

```

89                                     {
90     if (fd < 0 || fd >= MAX_FDS) {
91         P_ERRNO = P_EBADF;
92         return -1;
93     }
94     if (!fd_table[fd].in_use) {
95         P_ERRNO = P_EBADF;
96         return -1;
97     }
98     fd_table[fd].ref_count++;
99     return fd_table[fd].ref_count;
100 }
```

**4.3.1.12 init\_fd\_table()**

```

void init_fd_table (
    fd_entry_t * fd_table )
```

Initializes the global kernel-level file descriptor table.

Initializes all entries in the file descriptor table to not in use.

Definition at line 43 of file fs\_helpers.c.

```

43                                     {
44     // STDIN (fd 0)
45     fd_table[0].in_use = 1;
46     fd_table[0].ref_count = 1;
47     strncpy(fd_table[0].filename, "<stdin>", 31);
48     fd_table[0].mode = F_READ;
49
50     // STDOUT (fd 1)
51     fd_table[1].in_use = 1;
52     strncpy(fd_table[1].filename, "<stdout>", 31);
53     fd_table[1].mode = F_WRITE; // write-only
54     fd_table[1].ref_count = 1;
55
56     // STDERR (fd 2)
57     fd_table[2].in_use = 1;
58     strncpy(fd_table[2].filename, "<stderr>", 31);
59     fd_table[2].mode = F_WRITE; // write-only
60     fd_table[2].ref_count = 1;
61
62     // other file descriptors (fd 3 and above)
63     for (int i = 3; i < MAX_FDS; i++) {
64         fd_table[i].in_use = 0;
65         fd_table[i].ref_count = 0;
66         memset(fd_table[i].filename, 0, sizeof(fd_table[i].filename));
67         fd_table[i].size = 0;
68         fd_table[i].first_block = 0;
69         fd_table[i].position = 0;
70         fd_table[i].mode = 0;
71     }
72 }
```

#### 4.3.1.13 mark\_entry\_as\_deleted()

```
int mark_entry_as_deleted (
    dir_entry_t * entry,
    int absolute_offset )
```

Marks a file entry as deleted and frees its blocks.

Marks a file entry as deleted and frees its blocks in the FAT.

Definition at line 396 of file fs\_helpers.c.

```
396 {
397     if (!is_mounted || entry == NULL || absolute_offset < 0) {
398         P_ERRNO = P_EINVAL;
399         return -1;
400     }
401
402     // free the blocks
403     uint16_t current_block = entry->firstBlock;
404     while (current_block != FAT_FREE && current_block != FAT_EOF) {
405         uint16_t next_block = fat[current_block];
406         fat[current_block] = FAT_FREE;
407         current_block = next_block;
408     }
409
410     // mark the entry as deleted in the root directory
411     dir_entry_t deleted_entry = *entry;
412     deleted_entry.name[0] = 1;
413     if (lseek(fs_fd, absolute_offset, SEEK_SET) == -1) {
414         P_ERRNO = P_ELSEEK;
415         return -1;
416     }
417     if (write(fs_fd, &deleted_entry, sizeof(deleted_entry)) != sizeof(deleted_entry)) {
418         P_ERRNO = P_EINVAL;
419         return -1;
420     }
421
422     // mark the passed entry as deleted
423     entry->name[0] = 1;
424     return 0;
425 }
```

### 4.3.2 Variable Documentation

#### 4.3.2.1 block\_size

```
int block_size = 0
```

Definition at line 28 of file fs\_helpers.c.

#### 4.3.2.2 fat

```
uint16_t* fat = NULL
```

Definition at line 31 of file fs\_helpers.c.

#### 4.3.2.3 fat\_size

```
int fat_size = 0
```

Definition at line 30 of file fs\_helpers.c.

#### 4.3.2.4 fd\_table

```
fd_entry_t fd_table[100]
```

Definition at line 34 of file fs\_helpers.c.

#### 4.3.2.5 fs\_fd

```
int fs_fd = -1
```

Definition at line 27 of file fs\_helpers.c.

#### 4.3.2.6 is\_mounted

```
bool is_mounted = false
```

Definition at line 32 of file fs\_helpers.c.

#### 4.3.2.7 MAX\_FDS

```
int MAX_FDS = 100
```

Definition at line 33 of file fs\_helpers.c.

#### 4.3.2.8 num\_fat\_blocks


```
int num_fat_blocks = 0
```

Definition at line 29 of file fs\_helpers.c.




## 4.4 SRC/fs/fs\_helpers.h File Reference

```
#include <stdint.h>
#include "fat_routines.h"
Include dependency graph for fs_helpers.h:
```



This graph shows which files directly or indirectly include this file:



## Functions

- void `init_fd_table` (`fd_entry_t *fd_table`)  
*Initializes all entries in the file descriptor table to not in use.*
- int `get_free_fd` (`fd_entry_t *fd_table`)  
*Finds the first available file descriptor in the table.*
- int `increment_fd_ref_count` (int fd)  
*Increments the reference count of a file descriptor.*
- int `decrement_fd_ref_count` (int fd)  
*Decrements the reference count of a file descriptor.*
- int `has_executable_permission` (int fd)  
*Checks if a file has executable permissions in the PennFAT filesystem.*
- uint16\_t `allocate_block` ()  
*Allocates a free block in the FAT.*
- int `find_file` (const char \*filename, `dir_entry_t *entry`)  
*Searches for a file in the root directory.*
- int `add_file_entry` (const char \*filename, uint32\_t size, uint16\_t first\_block, uint8\_t type, uint8\_t perm)

- Adds a new file entry to the root directory.*

  - int `mark_entry_as_deleted` (`dir_entry_t` \*entry, int offset)

*Marks a file entry as deleted and frees its blocks in the FAT.*
- int `copy_host_to_pennfat` (const char \*host\_filename, const char \*pennfat\_filename)

*Copies a file from the host OS to the PennFAT filesystem.*
- int `copy_pennfat_to_host` (const char \*pennfat\_filename, const char \*host\_filename)

*Copies a file from the PennFAT filesystem to the host OS.*
- int `copy_source_to_dest` (const char \*source\_filename, const char \*dest\_filename)

*Copies a file from a source file to a destination file.*
- int `compact_directory` ()

*Compacts the root directory by removing all deleted entries.*

## Variables

- int `fs_fd`
- int `block_size`
- int `num_fat_blocks`
- int `fat_size`
- uint16\_t \* `fat`
- bool `is_mounted`
- int `MAX_FDS`
- `fd_entry_t` `fd_table` [100]

## 4.4.1 Function Documentation

### 4.4.1.1 add\_file\_entry()

```
int add_file_entry (
    const char * filename,
    uint32_t size,
    uint16_t first_block,
    uint8_t type,
    uint8_t perm )
```

Adds a new file entry to the root directory.

#### Parameters

<i>filename</i>	name of the file to add
<i>size</i>	size of the file in bytes
<i>first_block</i>	block number of the first block of the file
<i>type</i>	file type (regular, directory, etc.)
<i>perm</i>	file permissions

## Returns

offset of the new entry in the directory if successful, -1 on error

Adds a new file entry to the root directory.

Definition at line 266 of file fs\_helpers.c.

```

270     {
271     if (!is_mounted) {
272         P_ERRNO = P_EFS_NOT_MOUNTED;
273         return -1;
274     }
275
276     // check if file already exists
277     dir_entry_t existing;
278     if (find_file(filename, &existing) >= 0) {
279         P_ERRNO = P_EEXIST;
280         return -1;
281     }
282
283     // start with root directory block (block 1)
284     uint16_t current_block = 1;
285     int offset = 0;
286     dir_entry_t dir_entry;
287
288     while (1) {
289         // position at the start of current block of the root directory
290         if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
291             P_ERRNO = P_ELSEEK;
292             return -1;
293         }
294
295         // reset offset for new block
296         offset = 0;
297
298         // search current block for free slot
299         while (offset < block_size) {
300             if (read(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
301                 P_ERRNO = P_ERead;
302                 return -1;
303             }
304
305             // found a free slot
306             if (dir_entry.name[0] == 0 || dir_entry.name[0] == 1) {
307                 // initialize the new entry
308                 memset(&dir_entry, 0, sizeof(dir_entry));
309                 strncpy(dir_entry.name, filename, 31);
310                 dir_entry.size = size;
311                 dir_entry.firstBlock = first_block;
312                 dir_entry.type = type;
313                 dir_entry.perm = perm;
314                 dir_entry.mtime = time(NULL);
315
316                 // write the entry
317                 if (lseek(fs_fd, fat_size + (current_block - 1) * block_size + offset, SEEK_SET) == -1) {
318                     P_ERRNO = P_ELSEEK;
319                     return -1;
320                 }
321                 if (write(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
322                     P_ERRNO = P_EWRITE;
323                     return -1;
324                 }
325
326                 return offset;
327             }
328
329             offset += sizeof(dir_entry);
330         }
331
332         // current block is full, check if there's a next block
333         if (fat[current_block] != FAT_EOF) {
334             current_block = fat[current_block];
335             continue;
336         }
337
338         // allocate a new block for the root directory
339         uint16_t new_block = allocate_block();
340         if (new_block == 0) {
341             P_ERRNO = P_EFULL;
342             return -1;
343         }
344
345         // chain the new block
346         fat[current_block] = new_block;

```

```

347     fat[new_block] = FAT_EOF;
348
349     // initialize new block
350     uint8_t* zero_block = calloc(block_size, 1);
351     if (!zero_block) {
352         P_ERRNO = P_EINVAL;
353         return -1;
354     }
355
356     // write this new block to the file system
357     if (lseek(fs_fd, fat_size + (new_block - 1) * block_size, SEEK_SET) == -1) {
358         P_ERRNO = P_ELSEEK;
359         free(zero_block);
360         return -1;
361     }
362     if (write(fs_fd, zero_block, block_size) != block_size) {
363         P_ERRNO = P_EWRITE;
364         free(zero_block);
365         return -1;
366     }
367
368     free(zero_block);
369
370     // initialize the new entry
371     memset(&dir_entry, 0, sizeof(dir_entry));
372     strncpy(dir_entry.name, filename, 31);
373     dir_entry.size = size;
374     dir_entry.firstBlock = first_block;
375     dir_entry.type = type;
376     dir_entry.perm = perm;
377     dir_entry.mtime = time(NULL);
378
379     // write the new entry at the start of the new block in the file system
380     if (lseek(fs_fd, fat_size + (new_block - 1) * block_size, SEEK_SET) == -1) {
381         P_ERRNO = P_ELSEEK;
382         return -1;
383     }
384     if (write(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
385         P_ERRNO = P_EWRITE;
386         return -1;
387     }
388
389     return 0;
390 }
391 }

```

#### 4.4.1.2 allocate\_block()

```
uint16_t allocate_block ( )
```

Allocates a free block in the FAT.

##### Returns

block number of the allocated block, or 0 if no free blocks available

Allocates a free block in the FAT.

If no block found, we try compacting the directory.

Definition at line 167 of file fs\_helpers.c.

```

167     {
168     for (int i = 2; i < fat_size / 2; i++) {
169         if (fat[i] == FAT_FREE) {
170             fat[i] = FAT_EOF;
171             return i;
172         }
173     }
174
175     if (compact_directory() == 0) {
176         for (int i = 2; i < fat_size / 2; i++) {
177             if (fat[i] == FAT_FREE) {
178                 fat[i] = FAT_EOF;

```

```

179         return i;
180     }
181 }
182 }
183
184 return 0;
185 }

```

#### 4.4.1.3 compact\_directory()

```
int compact_directory ( )
```

Compacts the root directory by removing all deleted entries.

##### Returns

0 on success, -1 on error

Compacts the root directory by removing all deleted entries.

Definition at line 692 of file fs\_helpers.c.

```

692     {
693     if (!is_mounted) {
694         P_ERRNO = P_EFS_NOT_MOUNTED;
695         return -1;
696     }
697
698     // buffer for temp storage of a block
699     uint8_t* dir_buffer = malloc(block_size);
700     if (!dir_buffer) {
701         P_ERRNO = P_EMALLOC;
702         return -1;
703     }
704
705     // start at root directory
706     uint16_t current_block = 1;
707     int dir_entries_count = 0;
708     int deleted_entries_count = 0;
709
710     // calculate number of entries and deleted entries in the root directory
711     while (current_block != FAT_EOF) {
712         if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
713             P_ERRNO = P_ELSEEK;
714             free(dir_buffer);
715             return -1;
716         }
717
718         if (read(fs_fd, dir_buffer, block_size) != block_size) {
719             P_ERRNO = P_EREAD;
720             free(dir_buffer);
721             return -1;
722         }
723
724         // count entries and deleted entries in this block
725         for (int offset = 0; offset < block_size; offset += sizeof(dir_entry_t)) {
726             dir_entry_t* entry = (dir_entry_t*)(dir_buffer + offset);
727
728             // check if we've reached the end of directory
729             if (entry->name[0] == 0) {
730                 break;
731             }
732
733             dir_entries_count++;
734
735             // check if it's a deleted entry
736             if (entry->name[0] == 1) {
737                 deleted_entries_count++;
738             }
739         }
740
741         // move onto next block, if there is one
742         if (fat[current_block] != FAT_EOF) {
743             current_block = fat[current_block];
744         } else {

```

```

745     break;
746 }
747 }
748
749 // if no deleted entries, no compaction needed
750 if (deleted_entries_count == 0) {
751     free(dir_buffer);
752     return 0;
753 }
754
755 // allocate space for all valid entries
756 dir_entry_t* all_entries = malloc(dir_entries_count * sizeof(dir_entry_t));
757 if (!all_entries) {
758     P_ERRNO = P_EMALLOC;
759     free(dir_buffer);
760     return -1;
761 }
762
763 // read all entries into the buffer, skipping deleted ones
764 current_block = 1;
765 int valid_entry_idx = 0;
766
767 while (current_block != FAT_EOF) {
768     if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
769         P_ERRNO = P_ELSEEK;
770         free(dir_buffer);
771         free(all_entries);
772         return -1;
773     }
774
775     if (read(fs_fd, dir_buffer, block_size) != block_size) {
776         P_ERRNO = P_EREAD;
777         free(dir_buffer);
778         free(all_entries);
779         return -1;
780     }
781
782     // process entries in this block
783     for (int offset = 0; offset < block_size; offset += sizeof(dir_entry_t)) {
784         dir_entry_t* entry = (dir_entry_t*)(dir_buffer + offset);
785
786         // check if we've reached the end of directory
787         if (entry->name[0] == 0) {
788             break;
789         }
790
791         // skip deleted entries
792         if (entry->name[0] == 1) {
793             continue;
794         }
795
796         // copy valid entry to our array
797         memcpy(&all_entries[valid_entry_idx++], entry, sizeof(dir_entry_t));
798     }
799
800     // move to the next block
801     if (fat[current_block] != FAT_EOF) {
802         current_block = fat[current_block];
803     } else {
804         break;
805     }
806 }
807
808 // rewrite the directory with only valid entries
809 current_block = 1;
810 int entries_per_block = block_size / sizeof(dir_entry_t);
811 int blocks_needed = (valid_entry_idx + entries_per_block - 1) / entries_per_block;
812
813 // clean up any excess directory blocks in the FAT chain
814 uint16_t next_block = fat[current_block];
815 if (blocks_needed == 1) {
816     // only need one block, free all others
817     while (next_block != FAT_EOF) {
818         uint16_t temp = fat[next_block];
819         fat[next_block] = FAT_FREE;
820         next_block = temp;
821     }
822     fat[current_block] = FAT_EOF;
823 } else {
824     // navigate through needed blocks
825     int block_count = 1;
826     uint16_t prev_block = current_block;
827
828     while (block_count < blocks_needed) {
829         if (next_block == FAT_EOF) {
830             // need to allocate a new block
831             uint16_t new_block = allocate_block();

```

```

832         if (new_block == 0) {
833             P_ERRNO = P_EFULL;
834             free(dir_buffer);
835             free(all_entries);
836             return -1;
837         }
838         fat[prev_block] = new_block;
839         next_block = new_block;
840     }
841
842     prev_block = next_block;
843     next_block = fat[next_block];
844     block_count++;
845 }
846
847 // free any excess blocks
848 fat[prev_block] = FAT_EOF;
849 while (next_block != FAT_EOF) {
850     uint16_t temp = fat[next_block];
851     fat[next_block] = FAT_FREE;
852     next_block = temp;
853 }
854 }
855
856 // write the valid entries back to the directory blocks
857 current_block = 1;
858 int entries_written = 0;
859
860 while (entries_written < valid_entry_idx) {
861     if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
862         P_ERRNO = P_ELSEEK;
863         free(dir_buffer);
864         free(all_entries);
865         return -1;
866     }
867
868     memset(dir_buffer, 0, block_size);
869
870     // copy entries to the buffer
871     int entries_in_this_block = 0;
872     while (entries_written < valid_entry_idx && entries_in_this_block < entries_per_block) {
873         memcpy(dir_buffer + (entries_in_this_block * sizeof(dir_entry_t)),
874             &all_entries[entries_written],
875             sizeof(dir_entry_t));
876         entries_written++;
877         entries_in_this_block++;
878     }
879
880     // write the buffer to the file system
881     if (write(fs_fd, dir_buffer, block_size) != block_size) {
882         P_ERRNO = P_EINVAL;
883         free(dir_buffer);
884         free(all_entries);
885         return -1;
886     }
887
888     // move to the next block if needed
889     if (entries_written < valid_entry_idx) {
890         current_block = fat[current_block];
891     }
892 }
893
894 free(dir_buffer);
895 free(all_entries);
896 return 0;
897 }

```

#### 4.4.1.4 copy\_host\_to\_pennfat()

```

int copy_host_to_pennfat (
    const char * host_filename,
    const char * pennfat_filename )

```

Copies a file from the host OS to the PennFAT filesystem.

## Parameters

<i>host_filename</i>	path to the file on the host OS
<i>pennfat_filename</i>	name to give the file in PennFAT

## Returns

0 on success, -1 on error

Copies a file from the host OS to the PennFAT filesystem.

Definition at line 434 of file fs\_helpers.c.

```

435                                     {
436     if (!is_mounted) {
437         P_ERRNO = P_EFS_NOT_MOUNTED;
438         return -1;
439     }
440
441     // open the host file
442     int host_fd = open(host_filename, O_RDONLY);
443     if (host_fd == -1) {
444         P_ERRNO = P_EOPEN;
445         return -1;
446     }
447
448     // determine file size by seeking to the end and getting position
449     off_t host_file_size_in_bytes = lseek(host_fd, 0, SEEK_END);
450     if (host_file_size_in_bytes == -1) {
451         P_ERRNO = P_ELSEEK;
452         close(host_fd);
453         return -1;
454     }
455
456     // go back to beginning of file for reading
457     if (lseek(host_fd, 0, SEEK_SET) == -1) {
458         P_ERRNO = P_ELSEEK;
459         close(host_fd);
460         return -1;
461     }
462
463     // open the destination file in PennFAT
464     int pennfat_fd = k_open(pennfat_filename, F_WRITE);
465     if (pennfat_fd < 0) {
466         close(host_fd);
467         return -1;
468     }
469
470     // copy the data into this buffer
471     uint8_t* buffer = (uint8_t*)malloc(block_size);
472     if (!buffer) {
473         P_ERRNO = P_EMALLOC;
474         k_close(pennfat_fd);
475         close(host_fd);
476         return -1;
477     }
478
479     uint32_t bytes_remaining = host_file_size_in_bytes;
480     ssize_t bytes_read;
481
482     // read from host file
483     while (bytes_remaining > 0) {
484         // ensure bytes to read never exceeds the block size
485         ssize_t bytes_to_read = bytes_remaining < block_size ? bytes_remaining : block_size;
486         bytes_read = read(host_fd, buffer, bytes_to_read);
487
488         if (bytes_read <= 0) {
489             break;
490         }
491
492         // write to pennfat_fd using k_write
493         if (k_write(pennfat_fd, (const char*)buffer, bytes_read) != bytes_read) {
494             free(buffer);
495             k_close(pennfat_fd);
496             close(host_fd);
497             return -1;
498         }
499
500         bytes_remaining -= bytes_read;
501     }

```



```

502
503 // check for read error
504 if (bytes_read < 0) {
505     P_ERRNO = P_EREAD;
506     free(buffer);
507     k_close(pennfat_fd);
508     close(host_fd);
509     return -1;
510 }
511
512 // otherwise, cleanup and return success
513 free(buffer);
514 k_close(pennfat_fd);
515 close(host_fd);
516 return 0;
517 }

```

#### 4.4.1.5 copy\_pennfat\_to\_host()

```

int copy_pennfat_to_host (
    const char * pennfat_filename,
    const char * host_filename )

```

Copies a file from the PennFAT filesystem to the host OS.

##### Parameters

<i>pennfat_filename</i>	name of the file in PennFAT
<i>host_filename</i>	path to save the file on the host OS

##### Returns

0 on success, -1 on error

Copies a file from the PennFAT filesystem to the host OS.

Definition at line 522 of file fs\_helpers.c.

```

523
524 if (!is_mounted) {
525     P_ERRNO = P_EFS_NOT_MOUNTED;
526     return -1;
527 }
528
529 // open the PennFAT file
530 int pennfat_fd = k_open(pennfat_filename, F_READ);
531 if (pennfat_fd < 0) {
532     return -1;
533 }
534
535 // get the pennfat file size
536 off_t pennfat_file_size_in_bytes = k_lseek(pennfat_fd, 0, SEEK_END);
537 if (pennfat_file_size_in_bytes == -1) {
538     k_close(pennfat_fd);
539     return -1;
540 }
541
542 // go back to beginning of file for reading
543 if (k_lseek(pennfat_fd, 0, SEEK_SET) == -1) {
544     k_close(pennfat_fd);
545     return -1;
546 }
547
548 // open the host file
549 int host_fd = open(host_filename, O_WRONLY | O_CREAT | O_TRUNC, 0644);
550 if (host_fd == -1) {
551     P_ERRNO = P_EOPEN;
552     k_close(pennfat_fd);

```

```

553     return -1;
554 }
555
556 // allocate buffer for data transfer
557 char* buffer = (char*)malloc(block_size);
558 if (!buffer) {
559     P_ERRNO = P_EMALLOC;
560     k_close(pennfat_fd);
561     close(host_fd);
562     return -1;
563 }
564
565 uint32_t bytes_remaining = pennfat_file_size_in_bytes;
566 ssize_t bytes_read;
567
568 // read from PennFAT file and write to host file
569 while (bytes_remaining > 0) {
570     // ensure bytes to read never exceeds the block size
571     ssize_t bytes_to_read = bytes_remaining < block_size ? bytes_remaining : block_size;
572     bytes_read = k_read(pennfat_fd, buffer, bytes_to_read);
573
574     if (bytes_read <= 0) {
575         break;
576     }
577
578     if (write(host_fd, buffer, bytes_read) != bytes_read) {
579         P_ERRNO = P_EINVAL;
580         free(buffer);
581         close(host_fd);
582         k_close(pennfat_fd);
583         return -1;
584     }
585
586     bytes_remaining -= bytes_read;
587 }
588
589 // check for read error
590 if (bytes_read < 0) {
591     P_ERRNO = P_EREAD;
592     free(buffer);
593     close(host_fd);
594     k_close(pennfat_fd);
595     return -1;
596 }
597
598 // otherwise, cleanup and return success
599 free(buffer);
600 close(host_fd);
601 k_close(pennfat_fd);
602 return 0;
603 }

```

#### 4.4.1.6 copy\_source\_to\_dest()

```

int copy_source_to_dest (
    const char * source_filename,
    const char * dest_filename )

```

Copies a file from a source file to a destination file.

##### Parameters

<i>source_filename</i>	name of the source filename
<i>dest_filename</i>	name of the destination filename

##### Returns

0 on success, -1 on error

Copies a file from a source file to a destination file.

Definition at line 608 of file fs\_helpers.c.

```

609                                     {
610     if (!is_mounted) {
611         P_ERRNO = P_EFS_NOT_MOUNTED;
612         return -1;
613     }
614
615     // open the source file
616     int source_fd = k_open(source_filename, F_READ);
617     if (source_fd < 0) {
618         return -1;
619     }
620
621     // get the source file size
622     off_t source_file_size_in_bytes = k_lseek(source_fd, 0, SEEK_END);
623     if (source_file_size_in_bytes == -1) {
624         k_close(source_fd);
625         return -1;
626     }
627
628     // move to the beginning of the source file for reading
629     if (k_lseek(source_fd, 0, SEEK_SET) < 0) {
630         k_close(source_fd);
631         return -1;
632     }
633
634     // open the destination file
635     int dest_fd = k_open(dest_filename, F_WRITE);
636     if (dest_fd < 0) {
637         k_close(source_fd);
638         return -1;
639     }
640
641     // read from source to destination
642     char* buffer = (char*)malloc(block_size);
643     if (!buffer) {
644         P_ERRNO = P_EMALLOC;
645         k_close(source_fd);
646         k_close(dest_fd);
647         return -1;
648     }
649
650     uint32_t bytes_remaining = source_file_size_in_bytes;
651     ssize_t bytes_read;
652
653     while (bytes_remaining > 0) {
654         // make sure the bytes to read doesn't exceed block size
655         ssize_t bytes_to_read = bytes_remaining < block_size ? bytes_remaining : block_size;
656         bytes_read = k_read(source_fd, buffer, bytes_to_read);
657
658         if (bytes_read <= 0) {
659             break;
660         }
661
662         if (k_write(dest_fd, buffer, bytes_read) != bytes_read) {
663             free(buffer);
664             k_close(source_fd);
665             k_close(dest_fd);
666             return -1;
667         }
668     }
669
670     // check for read error
671     if (bytes_read < 0) {
672         free(buffer);
673         k_close(source_fd);
674         k_close(dest_fd);
675         return -1;
676     }
677
678     // otherwise, cleanup and return success
679     free(buffer);
680     k_close(source_fd);
681     k_close(dest_fd);
682     return 0;
683 }

```

#### 4.4.1.7 decrement\_fd\_ref\_count()

```

int decrement_fd_ref_count (
    int fd )

```

Decrements the reference count of a file descriptor.

#### Parameters

<i>fd</i>	file descriptor to decrement
-----------	------------------------------

#### Returns

new reference count, or -1 on error

If reference count reaches 0, flush field values.

Definition at line 107 of file `fs_helpers.c`.

```

107     {
108     if (fd < 0 || fd >= MAX_FDS) {
109         P_ERRNO = P_EBADF;
110         return -1;
111     }
112
113     if (!fd_table[fd].in_use) {
114         P_ERRNO = P_EBADF;
115         return -1;
116     }
117
118     fd_table[fd].ref_count--;
119     if (fd_table[fd].ref_count == 0) {
120         fd_table[fd].in_use = 0;
121         memset(fd_table[fd].filename, 0, sizeof(fd_table[fd].filename));
122         fd_table[fd].size = 0;
123         fd_table[fd].first_block = 0;
124         fd_table[fd].position = 0;
125         fd_table[fd].mode = 0;
126     }
127     return fd_table[fd].ref_count;
128 }

```

#### 4.4.1.8 find\_file()

```

int find_file (
    const char * filename,
    dir_entry_t * entry )

```

Searches for a file in the root directory.

#### Parameters

<i>filename</i>	name of the file to find
<i>entry</i>	pointer to store the directory entry if found

#### Returns

offset of the entry in the directory if found, -1 if not found

Retrieves the file's absolute offset in the filesystem.

Definition at line 192 of file `fs_helpers.c`.

```

192     {
193     if (!is_mounted) {

```

```

194     P_ERRNO = P_EFS_NOT_MOUNTED;
195     return -1;
196 }
197
198 // Start with root directory block (block 1)
199 uint16_t current_block = 1;
200 int offset_in_block = 0;
201 int absolute_offset = 0;
202 dir_entry_t dir_entry;
203
204 while (1) {
205     // Position at the start of current block
206     if (lseek(fs_fd, fat_size + (current_block - 1) * block_size, SEEK_SET) == -1) {
207         P_ERRNO = P_ELSEEK;
208         return -1;
209     }
210
211     // reset offset for new block
212     offset_in_block = 0;
213
214     // calculate the absolute offset
215     absolute_offset = fat_size + (current_block - 1) * block_size;
216
217     // search current block
218     while (offset_in_block < block_size) {
219         if (read(fs_fd, &dir_entry, sizeof(dir_entry)) != sizeof(dir_entry)) {
220             P_ERRNO = P_EREAD;
221             return -1;
222         }
223
224         // check if we've reached the end of directory
225         if (dir_entry.name[0] == 0) {
226             break;
227         }
228
229         // check if this is a deleted entry
230         if (dir_entry.name[0] == 1 || dir_entry.name[0] == 2) {
231             offset_in_block += sizeof(dir_entry);
232             absolute_offset += sizeof(dir_entry);
233             continue;
234         }
235
236         // check if we found the file
237         if (strcmp(dir_entry.name, filename) == 0) {
238             if (entry) {
239                 memcpy(entry, &dir_entry, sizeof(dir_entry));
240             }
241             return absolute_offset; // return the absolute file offset
242         }
243
244         offset_in_block += sizeof(dir_entry);
245         absolute_offset += sizeof(dir_entry);
246     }
247
248     // if we've reached the end of the current block, check if there's a next block
249     if (fat[current_block] != FAT_EOF) {
250         current_block = fat[current_block];
251         continue;
252     }
253
254     // no more blocks to search
255     break;
256 }
257
258 // file not found
259 P_ERRNO = P_ENOENT;
260 return -1;
261 }

```

#### 4.4.1.9 get\_free\_fd()

```

int get_free_fd (
    fd_entry_t * fd_table )

```

Finds the first available file descriptor in the table.

**Parameters**

<i>fd_table</i>	pointer to the file descriptor table to search
-----------------	--

**Returns**

index of the first free file descriptor, or -1 if none available

Finds the first available file descriptor in the table.

Definition at line 77 of file fs\_helpers.c.

```

77
78     for (int i = 3; i < MAX_FDS; i++) {
79         if (!fd_table[i].in_use) {
80             return i;
81         }
82     }
83     return -1;
84 }
```

**4.4.1.10 has\_executable\_permission()**

```

int has_executable_permission (
    int fd )
```

Checks if a file has executable permissions in the PennFAT filesystem.

**Parameters**

<i>fd</i>	The fd of the file to check.
-----------	------------------------------

**Returns**

1 if the file has executable permissions, 0 if it doesn't, -1 if an error occurred.

Checks if a file has executable permissions in the PennFAT filesystem.

Definition at line 133 of file fs\_helpers.c.

```

133
134     // check if fs is mounted
135     if (!is_mounted) {
136         P_ERRNO = P_EFS_NOT_MOUNTED;
137         return -1;
138     }
139
140     // validate fd argument
141     if (fd < 0 || fd >= MAX_FDS) {
142         P_ERRNO = P_EINVAL;
143         return -1;
144     }
145
146
147     // determine whether the file exists
148     dir_entry_t entry;
149     int entry_offset = find_file(fd_table[fd].filename, &entry);
150     if (entry_offset < 0) {
151         return -1;
152     }
153
154     // if it exists, get its permission
```

```

155     if (entry.perm & PERM_EXEC) {
156         return 1;
157     }
158
159     return 0;
160 }

```

#### 4.4.1.11 increment\_fd\_ref\_count()

```

int increment_fd_ref_count (
    int fd )

```

Increments the reference count of a file descriptor.

##### Parameters

<i>fd</i>	file descriptor to increment
-----------	------------------------------

##### Returns

new reference count, or -1 on error

Definition at line 89 of file fs\_helpers.c.

```

89     {
90     if (fd < 0 || fd >= MAX_FDS) {
91         P_ERRNO = P_EBADF;
92         return -1;
93     }
94     if (!fd_table[fd].in_use) {
95         P_ERRNO = P_EBADF;
96         return -1;
97     }
98     fd_table[fd].ref_count++;
99     return fd_table[fd].ref_count;
100 }

```

#### 4.4.1.12 init\_fd\_table()

```

void init_fd_table (
    fd_entry_t * fd_table )

```

Initializes all entries in the file descriptor table to not in use.

##### Parameters

<i>fd_table</i>	pointer to the file descriptor table to initialize
-----------------	--

Initializes all entries in the file descriptor table to not in use.

Definition at line 43 of file fs\_helpers.c.

```

43     {
44     // STDIN (fd 0)
45     fd_table[0].in_use = 1;

```

```

46 fd_table[0].ref_count = 1;
47 strncpy(fd_table[0].filename, "<stdin>", 31);
48 fd_table[0].mode = F_READ;
49
50 // STDOUT (fd 1)
51 fd_table[1].in_use = 1;
52 strncpy(fd_table[1].filename, "<stdout>", 31);
53 fd_table[1].mode = F_WRITE; // write-only
54 fd_table[1].ref_count = 1;
55
56 // STDERR (fd 2)
57 fd_table[2].in_use = 1;
58 strncpy(fd_table[2].filename, "<stderr>", 31);
59 fd_table[2].mode = F_WRITE; // write-only
60 fd_table[2].ref_count = 1;
61
62 // other file descriptors (fd 3 and above)
63 for (int i = 3; i < MAX_FDS; i++) {
64     fd_table[i].in_use = 0;
65     fd_table[i].ref_count = 0;
66     memset(fd_table[i].filename, 0, sizeof(fd_table[i].filename));
67     fd_table[i].size = 0;
68     fd_table[i].first_block = 0;
69     fd_table[i].position = 0;
70     fd_table[i].mode = 0;
71 }
72 }

```

#### 4.4.1.13 mark\_entry\_as\_deleted()

```

int mark_entry_as_deleted (
    dir_entry_t * entry,
    int absolute_offset )

```

Marks a file entry as deleted and frees its blocks in the FAT.

This function takes a directory entry and its offset in the directory, marks it as deleted in the directory, and frees all blocks in its FAT chain.

##### Parameters

<i>entry</i>	the entry struct of the file to mark as deleted.
<i>offset</i>	the offset of the entry in the directory

##### Returns

0 on success, -1 on error

Marks a file entry as deleted and frees its blocks in the FAT.

Definition at line 396 of file fs\_helpers.c.

```

396
397 if (!is_mounted || entry == NULL || absolute_offset < 0) {
398     P_ERRNO = P_EINVAL;
399     return -1;
400 }
401
402 // free the blocks
403 uint16_t current_block = entry->firstBlock;
404 while (current_block != FAT_FREE && current_block != FAT_EOF) {
405     uint16_t next_block = fat[current_block];
406     fat[current_block] = FAT_FREE;
407     current_block = next_block;
408 }
409
410 // mark the entry as deleted in the root directory

```



```
411     dir_entry_t deleted_entry = *entry;
412     deleted_entry.name[0] = 1;
413     if (lseek(fs_fd, absolute_offset, SEEK_SET) == -1) {
414         P_ERRNO = P_ELSEEK;
415         return -1;
416     }
417     if (write(fs_fd, &deleted_entry, sizeof(deleted_entry)) != sizeof(deleted_entry)) {
418         P_ERRNO = P_EINVAL;
419         return -1;
420     }
421
422     // mark the passed entry as deleted
423     entry->name[0] = 1;
424     return 0;
425 }
```

## 4.4.2 Variable Documentation

### 4.4.2.1 block\_size

```
int block_size [extern]
```

Definition at line 28 of file fs\_helpers.c.

### 4.4.2.2 fat

```
uint16_t* fat [extern]
```

Definition at line 31 of file fs\_helpers.c.

### 4.4.2.3 fat\_size

```
int fat_size [extern]
```

Definition at line 30 of file fs\_helpers.c.

### 4.4.2.4 fd\_table

```
fd_entry_t fd_table[100] [extern]
```

Definition at line 34 of file fs\_helpers.c.

#### 4.4.2.5 fs\_fd

```
int fs_fd [extern]
```

Definition at line 27 of file fs\_helpers.c.

#### 4.4.2.6 is\_mounted

```
bool is_mounted [extern]
```

Definition at line 32 of file fs\_helpers.c.

#### 4.4.2.7 MAX\_FDS

```
int MAX_FDS [extern]
```

Definition at line 33 of file fs\_helpers.c.

#### 4.4.2.8 num\_fat\_blocks


```
int num_fat_blocks [extern]
```

Definition at line 29 of file fs\_helpers.c.

### 4.5 SRC/fs/fs\_kfuncs.c File Reference

```
#include "fs_kfuncs.h"
#include "../kernel/kern_pcb.h"
#include "../kernel/kern_sys_calls.h"
#include "../kernel/signal.h"
#include "../lib/pennos-errno.h"
#include "fat_routines.h"
#include "fs_helpers.h"
#include "fs_syscalls.h"
#include <errno.h>
#include <fcntl.h>
#include <stdbool.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <time.h>
#include <unistd.h>
```

Include dependency graph for fs\_kfuncs.c:



## Functions

- `int k_open` (const char \*fname, int mode)  
*Kernel-level call to open a file.*
- `int k_read` (int fd, char \*buf, int n)  
*Kernel-level call to read a file.*
- `int k_write` (int fd, const char \*str, int n)  
*Kernel-level call to write to a file.*
- `int k_close` (int fd)  
*Kernel-level call to close a file.*
- `int k_unlink` (const char \*fname)  
*Kernel-level call to remove a file.*
- `int k_lseek` (int fd, int offset, int whence)  
*Kernel-level call to re-position a file offset.*
- `void format_file_info` (dir\_entry\_t \*entry, char \*buffer)
- `int k_ls` (const char \*filename)  
*Kernel-level call to list files.*

## Variables

- `pcb_t * current_running_pcb`
- `pid_t current_fg_pid`

### 4.5.1 Function Documentation

#### 4.5.1.1 format\_file\_info()

```
void format_file_info (
    dir_entry_t * entry,
    char * buffer )
```

Definition at line 666 of file fs\_kfuncs.c.

```
666                                     {
667     // convert permissions to string
668     char perms[4] = "---";
669     if (entry->perm & PERM_READ)
670         perms[0] = 'r';
671     if (entry->perm & PERM_WRITE)
672         perms[1] = 'w';
673     if (entry->perm & PERM_READ_EXEC & ~PERM_READ)
674         perms[2] = 'x';
675
676     // convert time to string
677     char time_str[20];
678     struct tm* tm = localtime(&entry->mtime);
679     strftime(time_str, sizeof(time_str), "%b %d %H:%M", tm);
680
681     // format the output string
682     snprintf(buffer, 256, "%4d %s %6d %s %s\n", entry->firstBlock, perms,
683             entry->size, time_str, entry->name);
684 }
```

#### 4.5.1.2 k\_close()

```
int k_close (
    int fd )
```

Kernel-level call to close a file.

Closes an open file.

Definition at line 528 of file fs\_kfuncs.c.

```
528     {
529     /*if (fd == STDIN_FILENO || fd == STDOUT_FILENO || fd == STDERR_FILENO) {
530         P_ERRNO = P_EINVAL;
531         return -1;
532     }*/
533
534     // validate the file descriptor
535     if (fd < 0 || fd >= MAX_FDS || !fd_table[fd].in_use) {
536         P_ERRNO = P_EBADF;
537         return -1;
538     }
539
540     // ensure any pending changes are written to disk
541     // update the directory entry with the current file size
542     dir_entry_t entry;
543     int file_offset = find_file(fd_table[fd].filename, &entry);
544
545     if (file_offset >= 0) {
546         // update file size if it changed
547         if (entry.size != fd_table[fd].size) {
548             entry.size = fd_table[fd].size;
549             entry.mtime = time(NULL);
550
551             if (lseek(fs_fd, file_offset, SEEK_SET) != -1) {
552                 write(fs_fd, &entry, sizeof(entry));
553             }
554         }
555     }
556
557     // decrement the reference count
558     decrement_fd_ref_count(fd);
559
560     return 0;
561 }
```

#### 4.5.1.3 k\_ls()

```
int k_ls (
    const char * filename )
```

Kernel-level call to list files.

Lists files or file information.

Definition at line 689 of file fs\_kfuncs.c.

```
689     {
690     if (!is_mounted) {
691         P_ERRNO = P_EFS_NOT_MOUNTED;
692         return -1;
693     }
694
695     // start with root directory block
696     uint16_t current_block = 1;
697     dir_entry_t entry;
698     uint32_t offset_in_block = 0;
699
700     // if filename is null, list all files in the current directory
701     if (filename == NULL) {
702         while (current_block != FAT_EOF) {
703             // calculate absolute offset in filesystem
704             off_t abs_offset =
```

```

705         fat_size + (current_block - 1) * block_size + offset_in_block;
706
707     // read directory entry
708     if (lseek(fs_fd, abs_offset, SEEK_SET) == -1) {
709         P_ERRNO = P_ELSEEK;
710         return -1;
711     }
712     if (read(fs_fd, &entry, sizeof(entry)) != sizeof(entry)) {
713         P_ERRNO = P_EREAD;
714         return -1;
715     }
716
717     // check for end of directory
718     if (entry.name[0] == 0)
719         break;
720
721     // skip deleted entries
722     if (entry.name[0] == 1 || entry.name[0] == 2) {
723         offset_in_block += sizeof(entry);
724         // check if we need to move to next block
725         if (offset_in_block + sizeof(entry) > block_size) {
726             current_block = fat[current_block];
727             offset_in_block = 0;
728         }
729         continue;
730     }
731
732     // format and write entry information
733     char output_buffer[256];
734     format_file_info(&entry, output_buffer);
735     if (k_write(STDOUT_FILENO, output_buffer, strlen(output_buffer)) < 0) {
736         return -1;
737     }
738
739     // move to next entry
740     offset_in_block += sizeof(entry);
741     // check if we need to move to next block
742     if (offset_in_block + sizeof(entry) > block_size) {
743         current_block = fat[current_block];
744         offset_in_block = 0;
745     }
746 }
747 } else {
748     // find and display specific file
749     int file_offset = find_file(filename, &entry);
750     if (file_offset < 0) {
751         P_ERRNO = P_ENOENT;
752         return -1;
753     }
754
755     char output_buffer[256];
756     format_file_info(&entry, output_buffer);
757     if (k_write(STDOUT_FILENO, output_buffer, strlen(output_buffer)) < 0) {
758         return -1;
759     }
760 }
761
762 return 0;
763 }

```

#### 4.5.1.4 k\_lseek()

```

int k_lseek (
    int fd,
    int offset,
    int whence )

```

Kernel-level call to re-position a file offset.

Repositions the file offset of an open file.

Definition at line 622 of file fs\_kfuncs.c.

```

622     {
623     // standard file descriptors don't support lseek
624     if (fd == STDIN_FILENO || fd == STDOUT_FILENO || fd == STDERR_FILENO) {

```

```

625     P_ERRNO = P_EINVAL;
626     return -1;
627 }
628
629 // validate the file descriptor
630 if (fd < 0 || fd >= MAX_FDS || !fd_table[fd].in_use) {
631     P_ERRNO = P_EBADF;
632     return -1;
633 }
634
635 // calculate new position based on whence
636 int32_t new_position;
637
638 switch (whence) {
639     case SEEK_SET:
640         new_position = offset;
641         break;
642     case SEEK_CUR:
643         new_position = fd_table[fd].position + offset;
644         break;
645     case SEEK_END:
646         new_position = fd_table[fd].size + offset;
647         break;
648     default:
649         P_ERRNO = P_EINVAL;
650         return -1;
651 }
652
653 // check if new position is valid
654 if (new_position < 0) {
655     P_ERRNO = P_EINVAL;
656     return -1;
657 }
658
659 // update file position
660 fd_table[fd].position = new_position;
661
662 return new_position;
663 }

```

#### 4.5.1.5 k\_open()

```

int k_open (
    const char * fname,
    int mode )

```

Kernel-level call to open a file.

Opens a file with the specified mode.

Definition at line 33 of file fs\_kfuncs.c.

```

33 {
34     // validate arguments
35     if (fname == NULL || *fname == '\0') {
36         P_ERRNO = P_EINVAL;
37         return -1;
38     }
39     if ((mode & (F_READ | F_WRITE | F_APPEND)) == 0) {
40         P_ERRNO = P_EINVAL;
41         return -1;
42     }
43
44     // check if the file system is mounted
45     if (!is_mounted) {
46         P_ERRNO = P_EFS_NOT_MOUNTED;
47         return -1;
48     }
49
50     // get a free file descriptor
51     int fd = get_free_fd(fd_table);
52     if (fd < 0) {
53         P_ERRNO = P_EFULL; // no free file descriptors
54         return -1;
55     }
56
57     // check if the file exists

```

```

58  dir_entry_t entry;
59  int file_offset = find_file(fname, &entry);
60
61  // file exists
62  if (file_offset >= 0) {
63      // check if the file is already open in write mode by another descriptor
64      if ((mode & (F_WRITE | F_APPEND)) != 0) {
65          for (int i = 0; i < MAX_FDS; i++) {
66              if (i != fd && fd_table[i].in_use &&
67                  strcmp(fd_table[i].filename, fname) == 0 &&
68                      (fd_table[i].mode & (F_WRITE | F_APPEND)) != 0) {
69                  P_ERRNO = P_EBUSY; // file is already open for writing
70                  return -1;
71              }
72          }
73      }
74
75      // fill in the file descriptor entry
76      fd_table[fd].in_use = 1;
77      fd_table[fd].ref_count++;
78      strncpy(fd_table[fd].filename, fname, 31);
79      fd_table[fd].filename[31] = '\0';
80      fd_table[fd].size = entry.size;
81      fd_table[fd].first_block = entry.firstBlock;
82      fd_table[fd].mode = mode;
83
84      // set the initial position
85      if (mode & F_APPEND) {
86          fd_table[fd].position = entry.size;
87      } else {
88          fd_table[fd].position = 0;
89      }
90
91      // if mode includes F_WRITE and not F_APPEND, truncate the file
92      if ((mode & F_WRITE) && !(mode & F_APPEND)) {
93          // free all blocks except the first one
94          uint16_t block = entry.firstBlock;
95          uint16_t next_block;
96
97          if (block != 0 && block != FAT_EOF) {
98              next_block = fat[block];
99              fat[block] = FAT_EOF; // terminate the chain at the first block
100              block = next_block;
101
102              // free the rest of the chain
103              while (block != 0 && block != FAT_EOF) {
104                  next_block = fat[block];
105                  fat[block] = FAT_FREE;
106                  block = next_block;
107              }
108          }
109
110          // update file size to 0
111          fd_table[fd].size = 0;
112          entry.size = 0;
113          entry.mtime = time(NULL);
114
115          // update the file system with the truncated file
116          if (lseek(fs_fd, file_offset, SEEK_SET) == -1) {
117              P_ERRNO = P_ELSEEK;
118              return -1;
119          }
120          if (write(fs_fd, &entry, sizeof(entry)) != sizeof(entry)) {
121              P_ERRNO = P_EWRITE;
122              return -1;
123          }
124      }
125      // file doesn't exist
126
127      // we can only create it if we are reading the file
128      if (!(mode & F_WRITE)) {
129          P_ERRNO = P_ENOENT;
130          return -1;
131      }
132
133      // allocate the first block
134      uint16_t first_block = allocate_block();
135      if (first_block == 0) {
136          P_ERRNO = P_EFULL;
137          return -1;
138      }
139
140      // create a new file entry
141      if (add_file_entry(fname, 0, first_block, TYPE_REGULAR, PERM_READ_WRITE) ==
142          -1) {
143          // error code already set by add_file_entry
144

```

```

145     fat[first_block] = FAT_FREE;
146     return -1;
147 }
148
149 // fill in the file descriptor entry
150 fd_table[fd].in_use = 1;
151 fd_table[fd].ref_count++;
152 strncpy(fd_table[fd].filename, fname, 31);
153 fd_table[fd].filename[31] = '\0';
154 fd_table[fd].size = 0;
155 fd_table[fd].first_block = first_block;
156 fd_table[fd].position = 0;
157 fd_table[fd].mode = mode;
158 }
159
160 return fd;
161 }

```

#### 4.5.1.6 k\_read()

```

int k_read (
    int fd,
    char * buf,
    int n )

```

Kernel-level call to read a file.

Reads data from an open file.

Definition at line 166 of file fs\_kfuncs.c.

```

166     {
167     // handle terminal control (if doesn't control, send a STOP signal)
168     if (fd == STDIN_FILENO && current_running_pcb != NULL) {
169         if (current_running_pcb->pid != current_fg_pid) {
170             s_kill(current_running_pcb->pid, P_SIGSTOP);
171         }
172     }
173
174     // handle standard input
175     if (fd == STDIN_FILENO) {
176         return read(STDIN_FILENO, buf, n);
177     }
178
179     // validate inputs
180     if (fd < 0 || fd >= MAX_FDS || !fd_table[fd].in_use) {
181         P_ERRNO = P_EBADF;
182         return -1;
183     }
184     if (buf == NULL || n < 0) {
185         P_ERRNO = P_EINVAL;
186         return -1;
187     }
188     if (n == 0) {
189         return 0;
190     }
191
192     // check if we're at EOF already
193     if (fd_table[fd].position >= fd_table[fd].size) {
194         return 0;
195     }
196
197     // determine how many bytes we can actually read
198     uint32_t bytes_to_read = n;
199     if (fd_table[fd].position + bytes_to_read > fd_table[fd].size) {
200         bytes_to_read = fd_table[fd].size - fd_table[fd].position;
201     }
202
203     // find the block containing the current position
204     uint16_t current_block = fd_table[fd].first_block;
205     uint32_t block_index = fd_table[fd].position / block_size;
206     uint32_t block_offset = fd_table[fd].position % block_size;
207
208     // navigate to the correct block in the chain
209     for (uint32_t i = 0; i < block_index; i++) {
210         if (current_block == 0 || current_block == FAT_EOF) {

```



```

211         // unexpected end of chain
212         P_ERRNO = P_EINVAL;
213         return -1;
214     }
215     current_block = fat[current_block];
216 }
217
218 // now we're at the right block, start reading
219 uint32_t bytes_read = 0;
220
221 while (bytes_read < bytes_to_read) {
222     // how much data can we read from the current block
223     uint32_t bytes_left_in_block = block_size - block_offset;
224     uint32_t bytes_to_read_now =
225         (bytes_to_read - bytes_read) < bytes_left_in_block
226         ? (bytes_to_read - bytes_read)
227         : bytes_left_in_block;
228
229     // seek to the right position in the file
230     if (lseek(fs_fd, fat_size + (current_block - 1) * block_size + block_offset,
231             SEEK_SET) == -1) {
232         P_ERRNO = P_ELSEEK;
233         if (bytes_read > 0) {
234             fd_table[fd].position += bytes_read;
235             return bytes_read;
236         }
237         return -1;
238     }
239
240     // read the data from the file
241     ssize_t read_result = read(fs_fd, buf + bytes_read, bytes_to_read_now);
242     if (read_result <= 0) {
243         P_ERRNO = P_EREAD;
244         // if we already read some data, return that count
245         if (bytes_read > 0) {
246             fd_table[fd].position += bytes_read;
247             return bytes_read;
248         }
249         return -1;
250     }
251
252     bytes_read += read_result;
253     block_offset += read_result;
254
255     // if we've read all data from this block and still have more to read, go to
256     // the next block
257     if (block_offset == block_size && bytes_read < bytes_to_read) {
258         if (current_block == FAT_EOF) {
259             // unexpected end of chain
260             break;
261         }
262         current_block = fat[current_block];
263         block_offset = 0;
264     }
265
266     // if we read less than expected, we might have hit EOF
267     if (read_result < bytes_to_read_now) {
268         break;
269     }
270 }
271
272 // update file position
273 fd_table[fd].position += bytes_read;
274
275 return bytes_read;
276 }

```

#### 4.5.1.7 k\_unlink()

```

int k_unlink (
    const char * fname )

```

Kernel-level call to remove a file.

Removes a file from the file system.

Definition at line 566 of file fs\_kfuncs.c.

```

566     {
567     if (fname == NULL || *fname == '\\0') {
568         P_ERRNO = P_EINVAL;
569         return -1;
570     }
571
572     if (!is_mounted) {
573         P_ERRNO = P_EFS_NOT_MOUNTED;
574         return -1;
575     }
576
577     // check if file is currently open by any process
578     for (int i = 0; i < MAX_FDS; i++) {
579         if (fd_table[i].in_use && strcmp(fd_table[i].filename, fname) == 0) {
580             P_ERRNO = P_EBUSY;
581             return -1;
582         }
583     }
584
585     // find the file in directory
586     dir_entry_t entry;
587     int file_offset = find_file(fname, &entry);
588     if (file_offset < 0) {
589         P_ERRNO = P_ENOENT;
590         return -1;
591     }
592
593     // mark the directory entry as deleted (set first byte to 1)
594     entry.name[0] = 1;
595
596     // write the modified directory entry back
597     if (lseek(fs_fd, file_offset, SEEK_SET) == -1) {
598         P_ERRNO = P_ELSEEK;
599         return -1;
600     }
601     if (write(fs_fd, &entry, sizeof(entry)) != sizeof(entry)) {
602         P_ERRNO = P_EWRITE;
603         return -1;
604     }
605
606     // free all blocks in the file chain
607     uint16_t current_block = entry.firstBlock;
608     uint16_t next_block;
609
610     while (current_block != FAT_FREE && current_block != FAT_EOF) {
611         next_block = fat[current_block];
612         fat[current_block] = FAT_FREE;
613         current_block = next_block;
614     }
615
616     return 0;
617 }

```

#### 4.5.1.8 k\_write()

```

int k_write (
    int fd,
    const char * str,
    int n )

```

Kernel-level call to write to a file.

Writes data to an open file.

Definition at line 281 of file fs\_kfuncs.c.

```

281     {
282     // handle standard output and error
283     if (fd == STDOUT_FILENO) {
284         return write(STDOUT_FILENO, str, n);
285     }
286     if (fd == STDERR_FILENO) {
287         return write(STDERR_FILENO, str, n);
288     }
289
290     // validate inputs

```

```

291 if (fd < 0 || fd >= MAX_FDS || !fd_table[fd].in_use) {
292     P_ERRNO = P_EBADF;
293     return -1;
294 }
295 if (str == NULL || n < 0) {
296     P_ERRNO = P_EINVAL;
297     return -1;
298 }
299 if (n == 0) {
300     return 0;
301 }
302
303 // check if filesystem is mounted and FAT is valid
304 if (!is_mounted || fat == NULL) {
305     P_ERRNO = P_EFS_NOT_MOUNTED;
306     return -1;
307 }
308
309 // get file information
310 uint16_t current_block = fd_table[fd].first_block;
311 uint32_t current_position = fd_table[fd].position;
312
313 // create a local buffer for block data
314 char* block_buffer = (char*)malloc(block_size);
315 if (block_buffer == NULL) {
316     P_ERRNO = P_EMALLOC;
317     return -1;
318 }
319
320 // calculate initial block position
321 uint32_t block_index = current_position / block_size;
322 uint32_t block_offset = current_position % block_size;
323
324 // if the file doesn't have a first block yet, allocate one
325 if (current_block == 0) {
326     current_block = allocate_block();
327     if (current_block == 0) {
328         P_ERRNO = P_EFULL;
329         free(block_buffer);
330         return -1;
331     }
332     fd_table[fd].first_block = current_block;
333 }
334
335 // navigate to the appropriate block
336 uint16_t prev_block = 0;
337 for (uint32_t i = 0; i < block_index; i++) {
338     if (current_block == 0 || current_block == FAT_EOF ||
339         current_block >= fat_size / 2) {
340         // reached the end of chain prematurely, need to allocate a new block
341         uint16_t new_block = allocate_block();
342         if (new_block == 0) {
343             P_ERRNO = P_EFULL;
344             free(block_buffer);
345             return -1;
346         }
347
348         // update the chain
349         if (prev_block != 0 && prev_block < fat_size / 2) {
350             fat[prev_block] = new_block;
351         } else {
352             // if there's no previous block, this must be the first one
353             fd_table[fd].first_block = new_block;
354         }
355
356         current_block = new_block;
357     }
358
359     prev_block = current_block;
360
361     // validate the block number before accessing FAT
362     if (current_block >= fat_size / 2) {
363         P_ERRNO = P_EINVAL;
364         free(block_buffer);
365         return -1;
366     }
367
368     current_block = fat[current_block];
369 }
370
371 // if we ended up without a valid block, go back to the last valid one
372 if (current_block == 0 || current_block == FAT_EOF ||
373     current_block >= fat_size / 2) {
374     if (prev_block != 0 && prev_block < fat_size / 2) {
375         uint16_t new_block = allocate_block();
376         if (new_block == 0) {
377             P_ERRNO = P_EFULL;

```

```

378         free(block_buffer);
379         return -1;
380     }
381
382     fat[prev_block] = new_block;
383     current_block = new_block;
384 } else {
385     P_ERRNO = P_EINVAL;
386     free(block_buffer);
387     return -1;
388 }
389 }
390
391 // start writing data
392 uint32_t bytes_written = 0;
393
394 while (bytes_written < n) {
395     // validate current block
396     if (current_block == 0 || current_block == FAT_EOF ||
397         current_block >= fat_size / 2) {
398         P_ERRNO = P_EINVAL;
399         break;
400     }
401
402     // how much can we write to this block
403     uint32_t space_in_block = block_size - block_offset;
404     uint32_t bytes_to_write = (n - bytes_written) < space_in_block
405         ? (n - bytes_written)
406         : space_in_block;
407
408     // position in filesystem
409     off_t block_position = fat_size + (current_block - 1) * block_size;
410
411     // if we're not writing a full block or not starting at the beginning, we
412     // need to read-modify-write
413     if (bytes_to_write < block_size || block_offset > 0) {
414         // read the current block
415         if (lseek(fs_fd, block_position, SEEK_SET) == -1) {
416             P_ERRNO = P_ELSEEK;
417             break;
418         }
419
420         // read the current block data
421         ssize_t read_result = read(fs_fd, block_buffer, block_size);
422         if (read_result < 0) {
423             P_ERRNO = P_EREAD;
424             break;
425         }
426
427         // copy the new data into the block buffer
428         memcpy(block_buffer + block_offset, str + bytes_written, bytes_to_write);
429
430         // seek back to write the modified block
431         if (lseek(fs_fd, block_position, SEEK_SET) == -1) {
432             P_ERRNO = P_ELSEEK;
433             break;
434         }
435
436         // write the full block back
437         ssize_t write_result = write(fs_fd, block_buffer, block_size);
438         if (write_result != block_size) {
439             P_ERRNO = P_EWRITE;
440             // we might have a partial write, but that's hard to handle correctly
441             break;
442         }
443     } else {
444         // we're writing a full block from the beginning
445         if (lseek(fs_fd, block_position, SEEK_SET) == -1) {
446             P_ERRNO = P_ELSEEK;
447             break;
448         }
449
450         ssize_t write_result = write(fs_fd, str + bytes_written, bytes_to_write);
451         if (write_result != bytes_to_write) {
452             P_ERRNO = P_EWRITE;
453             break;
454         }
455     }
456
457     // update counters
458     bytes_written += bytes_to_write;
459     block_offset = (block_offset + bytes_to_write) % block_size;
460
461     // if we've filled this block and still have more to write, go to the next
462     // block
463     if (block_offset == 0 && bytes_written < n) {
464         // validate current block before accessing FAT

```

```

465     if (current_block >= fat_size / 2) {
466         P_ERRNO = P_EINVAL;
467         break;
468     }
469
470     // check if there's a next block
471     if (fat[current_block] == FAT_EOF) {
472         // allocate a new block
473         uint16_t new_block = allocate_block();
474         if (new_block == 0) {
475             P_ERRNO = P_EFULL;
476             break;
477         }
478
479         // Update the FAT safely
480         if (current_block < fat_size / 2) {
481             fat[current_block] = new_block;
482         } else {
483             P_ERRNO = P_EINVAL;
484             break;
485         }
486
487         current_block = new_block;
488     } else {
489         current_block = fat[current_block];
490     }
491 }
492 }
493
494 // free the block buffer
495 free(block_buffer);
496
497 // update file position
498 fd_table[fd].position += bytes_written;
499
500 // update file size if needed
501 if (fd_table[fd].position > fd_table[fd].size) {
502     fd_table[fd].size = fd_table[fd].position;
503 }
504
505 // update the directory entry
506 dir_entry_t entry;
507 int dir_offset = find_file(fd_table[fd].filename, &entry);
508 if (dir_offset >= 0) {
509     entry.size = fd_table[fd].size;
510     entry.mtime = time(NULL);
511
512     if (lseek(fs_fd, dir_offset, SEEK_SET) == -1) {
513         P_ERRNO = P_ELSEEK;
514         return -1;
515     }
516     if (write(fs_fd, &entry, sizeof(entry)) != sizeof(entry)) {
517         P_ERRNO = P_EWRITE;
518         return -1;
519     }
520 }
521
522 return bytes_written;
523 }

```

## 4.5.2 Variable Documentation

### 4.5.2.1 current\_fg\_pid

pid\_t current\_fg\_pid [extern]

Definition at line 31 of file kern\_sys\_calls.c.

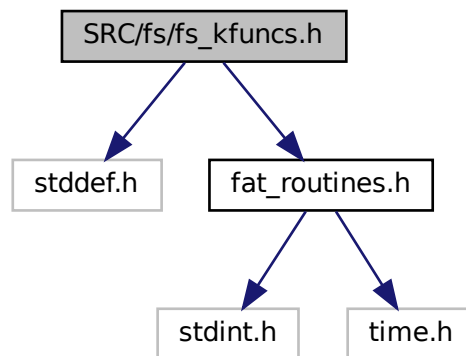
#### 4.5.2.2 current\_running\_pcb

```
pcb_t* current_running_pcb [extern]
```


Definition at line 38 of file scheduler.c.

## 4.6 SRC/fs/fs\_kfuncs.h File Reference

```
#include <stddef.h>
#include "fat_routines.h"
Include dependency graph for fs_kfuncs.h:
```



This graph shows which files directly or indirectly include this file:



## Macros

- #define `SEEK_SET` 0
- #define `SEEK_CUR` 1
- #define `SEEK_END` 2

## Functions

- int [k\\_open](#) (const char \*fname, int mode)  
*Opens a file with the specified mode.*
- int [k\\_read](#) (int fd, char \*buf, int n)  
*Reads data from an open file.*
- int [k\\_write](#) (int fd, const char \*str, int n)  
*Writes data to an open file.*
- int [k\\_close](#) (int fd)  
*Closes an open file.*
- int [k\\_unlink](#) (const char \*fname)  
*Removes a file from the file system.*
- int [k\\_lseek](#) (int fd, int offset, int whence)  
*Repositions the file offset of an open file.*
- int [k\\_ls](#) (const char \*filename)  
*Lists files or file information.*

## 4.6.1 Macro Definition Documentation

### 4.6.1.1 SEEK\_CUR

```
#define SEEK_CUR 1
```

Definition at line 17 of file fs\_kfuncs.h.

### 4.6.1.2 SEEK\_END

```
#define SEEK_END 2
```

Definition at line 18 of file fs\_kfuncs.h.

### 4.6.1.3 SEEK\_SET

```
#define SEEK_SET 0
```

Definition at line 16 of file fs\_kfuncs.h.

## 4.6.2 Function Documentation

### 4.6.2.1 k\_close()

```
int k_close (  
    int fd )
```

Closes an open file.

This is a kernel-level function that closes an open file and releases the associated file descriptor. Any unsaved changes are flushed to disk.

**Parameters**

<i>fd</i>	File descriptor of the open file.
-----------	-----------------------------------

**Returns**

0 on success, -1 on error with P\_ERRNO set. Possible error codes:

- P\_EBADF: Invalid file descriptor.

Closes an open file.

Definition at line 528 of file fs\_kfuncs.c.

```

528     {
529     /*if (fd == STDIN_FILENO || fd == STDOUT_FILENO || fd == STDERR_FILENO) {
530         P_ERRNO = P_EINVAL;
531         return -1;
532     }*/
533
534     // validate the file descriptor
535     if (fd < 0 || fd >= MAX_FDS || !fd_table[fd].in_use) {
536         P_ERRNO = P_EBADF;
537         return -1;
538     }
539
540     // ensure any pending changes are written to disk
541     // update the directory entry with the current file size
542     dir_entry_t entry;
543     int file_offset = find_file(fd_table[fd].filename, &entry);
544
545     if (file_offset >= 0) {
546         // update file size if it changed
547         if (entry.size != fd_table[fd].size) {
548             entry.size = fd_table[fd].size;
549             entry.mtime = time(NULL);
550
551             if (lseek(fs_fd, file_offset, SEEK_SET) != -1) {
552                 write(fs_fd, &entry, sizeof(entry));
553             }
554         }
555     }
556
557     // decrement the reference count
558     decrement_fd_ref_count(fd);
559
560     return 0;
561 }
```

**4.6.2.2 k\_ls()**

```

int k_ls (
    const char * filename )
```

Lists files or file information.

This is a kernel-level function that provides directory listing functionality. If filename is NULL or refers to a directory, it lists all files in that directory. If filename refers to a specific file, it displays detailed information about that file.

**Parameters**

<i>filename</i>	The name of the file or directory to list, or NULL for the current directory.
-----------------	---



## Returns

0 on success, -1 on error with P\_ERRNO set. Possible error codes:

- P\_ENOENT: Specified file or directory doesn't exist.

Lists files or file information.

Definition at line 689 of file fs\_kfuncs.c.

```

689     {
690     if (!is_mounted) {
691         P_ERRNO = P_EFS_NOT_MOUNTED;
692         return -1;
693     }
694
695     // start with root directory block
696     uint16_t current_block = 1;
697     dir_entry_t entry;
698     uint32_t offset_in_block = 0;
699
700     // if filename is null, list all files in the current directory
701     if (filename == NULL) {
702         while (current_block != FAT_EOF) {
703             // calculate absolute offset in filesystem
704             off_t abs_offset =
705                 fat_size + (current_block - 1) * block_size + offset_in_block;
706
707             // read directory entry
708             if (lseek(fs_fd, abs_offset, SEEK_SET) == -1) {
709                 P_ERRNO = P_ELSEEK;
710                 return -1;
711             }
712             if (read(fs_fd, &entry, sizeof(entry)) != sizeof(entry)) {
713                 P_ERRNO = P_EREAD;
714                 return -1;
715             }
716
717             // check for end of directory
718             if (entry.name[0] == 0)
719                 break;
720
721             // skip deleted entries
722             if (entry.name[0] == 1 || entry.name[0] == 2) {
723                 offset_in_block += sizeof(entry);
724                 // check if we need to move to next block
725                 if (offset_in_block + sizeof(entry) > block_size) {
726                     current_block = fat[current_block];
727                     offset_in_block = 0;
728                 }
729                 continue;
730             }
731
732             // format and write entry information
733             char output_buffer[256];
734             format_file_info(&entry, output_buffer);
735             if (k_write(STDOUT_FILENO, output_buffer, strlen(output_buffer)) < 0) {
736                 return -1;
737             }
738
739             // move to next entry
740             offset_in_block += sizeof(entry);
741             // check if we need to move to next block
742             if (offset_in_block + sizeof(entry) > block_size) {
743                 current_block = fat[current_block];
744                 offset_in_block = 0;
745             }
746         }
747     } else {
748         // find and display specific file
749         int file_offset = find_file(filename, &entry);
750         if (file_offset < 0) {
751             P_ERRNO = P_ENOENT;
752             return -1;
753         }
754
755         char output_buffer[256];
756         format_file_info(&entry, output_buffer);
757         if (k_write(STDOUT_FILENO, output_buffer, strlen(output_buffer)) < 0) {
758             return -1;
759         }
760     }
761     return 0;
762 }
763 }
```

### 4.6.2.3 k\_lseek()

```
int k_lseek (
    int fd,
    int offset,
    int whence )
```

Repositions the file offset of an open file.

This is a kernel-level function that changes the current position within an open file. The interpretation of the offset depends on the whence parameter.

#### Parameters

<i>fd</i>	File descriptor of the open file.
<i>offset</i>	The offset in bytes to set the position to.
<i>whence</i>	How to interpret the offset: <ul style="list-style-type: none"> <li>• SEEK_SET (0): Offset is from the beginning of the file.</li> <li>• SEEK_CUR (1): Offset is from the current position.</li> <li>• SEEK_END (2): Offset is from the end of the file.</li> </ul>

#### Returns

The new offset location on success, -1 on error with P\_ERRNO set. Possible error codes:

- P\_EBADF: Invalid file descriptor.
- P\_EINVAL: Invalid whence or the resulting position would be negative.

Repositions the file offset of an open file.

Definition at line 622 of file fs\_kfuncs.c.

```
622                                     {
623     // standard file descriptors don't support lseek
624     if (fd == STDIN_FILENO || fd == STDOUT_FILENO || fd == STDERR_FILENO) {
625         P_ERRNO = P_EINVAL;
626         return -1;
627     }
628
629     // validate the file descriptor
630     if (fd < 0 || fd >= MAX_FDS || !fd_table[fd].in_use) {
631         P_ERRNO = P_EBADF;
632         return -1;
633     }
634
635     // calculate new position based on whence
636     int32_t new_position;
637
638     switch (whence) {
639         case SEEK_SET:
640             new_position = offset;
641             break;
642         case SEEK_CUR:
643             new_position = fd_table[fd].position + offset;
644             break;
645         case SEEK_END:
646             new_position = fd_table[fd].size + offset;
647             break;
648         default:
649             P_ERRNO = P_EINVAL;
650             return -1;
651     }
652
653     // check if new position is valid
654     if (new_position < 0) {
```

```

655     P_ERRNO = P_EINVAL;
656     return -1;
657 }
658
659 // update file position
660 fd_table[fd].position = new_position;
661
662 return new_position;
663 }

```

#### 4.6.2.4 k\_open()

```

int k_open (
    const char * fname,
    int mode )

```

Opens a file with the specified mode.

This is a kernel-level function that opens a file and returns a file descriptor. The file is created if it doesn't exist and the mode includes F\_WRITE. If the file exists and F\_APPEND is specified, the file position is set to the end.

##### Parameters

<i>fname</i>	The name of the file to open.
<i>mode</i>	A combination of F_READ, F_WRITE, and F_APPEND.

##### Returns

A non-negative file descriptor on success, -1 on error with P\_ERRNO set. Possible error codes:

- P\_ENOENT: File doesn't exist and F\_READ only.
- P\_EFULL: Cannot create file (file system full).
- P\_EINVAL: Invalid mode or filename.

Opens a file with the specified mode.

Definition at line 33 of file fs\_kfuncs.c.

```

33                                     {
34     // validate arguments
35     if (fname == NULL || *fname == '\0') {
36         P_ERRNO = P_EINVAL;
37         return -1;
38     }
39     if ((mode & (F_READ | F_WRITE | F_APPEND)) == 0) {
40         P_ERRNO = P_EINVAL;
41         return -1;
42     }
43
44     // check if the file system is mounted
45     if (!is_mounted) {
46         P_ERRNO = P_EFS_NOT_MOUNTED;
47         return -1;
48     }
49
50     // get a free file descriptor
51     int fd = get_free_fd(fd_table);
52     if (fd < 0) {
53         P_ERRNO = P_EFULL; // no free file descriptors
54         return -1;
55     }
56
57     // check if the file exists
58     dir_entry_t entry;

```

```

59 int file_offset = find_file(fname, &entry);
60
61 // file exists
62 if (file_offset >= 0) {
63     // check if the file is already open in write mode by another descriptor
64     if ((mode & (F_WRITE | F_APPEND)) != 0) {
65         for (int i = 0; i < MAX_FDS; i++) {
66             if (i != fd && fd_table[i].in_use &&
67                 strcmp(fd_table[i].filename, fname) == 0 &&
68                 (fd_table[i].mode & (F_WRITE | F_APPEND)) != 0) {
69                 P_ERRNO = P_EBUSY; // file is already open for writing
70                 return -1;
71             }
72         }
73     }
74
75     // fill in the file descriptor entry
76     fd_table[fd].in_use = 1;
77     fd_table[fd].ref_count++;
78     strncpy(fd_table[fd].filename, fname, 31);
79     fd_table[fd].filename[31] = '\0';
80     fd_table[fd].size = entry.size;
81     fd_table[fd].first_block = entry.firstBlock;
82     fd_table[fd].mode = mode;
83
84     // set the initial position
85     if (mode & F_APPEND) {
86         fd_table[fd].position = entry.size;
87     } else {
88         fd_table[fd].position = 0;
89     }
90
91     // if mode includes F_WRITE and not F_APPEND, truncate the file
92     if ((mode & F_WRITE) && !(mode & F_APPEND)) {
93         // free all blocks except the first one
94         uint16_t block = entry.firstBlock;
95         uint16_t next_block;
96
97         if (block != 0 && block != FAT_EOF) {
98             next_block = fat[block];
99             fat[block] = FAT_EOF; // terminate the chain at the first block
100             block = next_block;
101
102             // free the rest of the chain
103             while (block != 0 && block != FAT_EOF) {
104                 next_block = fat[block];
105                 fat[block] = FAT_FREE;
106                 block = next_block;
107             }
108         }
109
110         // update file size to 0
111         fd_table[fd].size = 0;
112         entry.size = 0;
113         entry.mtime = time(NULL);
114
115         // update the file system with the truncated file
116         if (lseek(fs_fd, file_offset, SEEK_SET) == -1) {
117             P_ERRNO = P_ELSEEK;
118             return -1;
119         }
120         if (write(fs_fd, &entry, sizeof(entry)) != sizeof(entry)) {
121             P_ERRNO = P_EWRITE;
122             return -1;
123         }
124     }
125 } else {
126     // file doesn't exist
127
128     // we can only create it if we are reading the file
129     if (!(mode & F_WRITE)) {
130         P_ERRNO = P_ENOENT;
131         return -1;
132     }
133
134     // allocate the first block
135     uint16_t first_block = allocate_block();
136     if (first_block == 0) {
137         P_ERRNO = P_EFULL;
138         return -1;
139     }
140
141     // create a new file entry
142     if (add_file_entry(fname, 0, first_block, TYPE_REGULAR, PERM_READ_WRITE) ==
143         -1) {
144         // error code already set by add_file_entry
145         fat[first_block] = FAT_FREE;

```

```

146     return -1;
147 }
148
149 // fill in the file descriptor entry
150 fd_table[fd].in_use = 1;
151 fd_table[fd].ref_count++;
152 strncpy(fd_table[fd].filename, fname, 31);
153 fd_table[fd].filename[31] = '\0';
154 fd_table[fd].size = 0;
155 fd_table[fd].first_block = first_block;
156 fd_table[fd].position = 0;
157 fd_table[fd].mode = mode;
158 }
159
160 return fd;
161 }

```

#### 4.6.2.5 k\_read()

```

int k_read (
    int fd,
    char * buf,
    int n )

```

Reads data from an open file.

This is a kernel-level function that reads up to *n* bytes from an open file into the provided buffer. The file position is advanced by the number of bytes read.

##### Parameters

<i>fd</i>	File descriptor of the open file.
<i>n</i>	Maximum number of bytes to read.
<i>buf</i>	Buffer to store the read data.

##### Returns

The number of bytes read on success, -1 on error with *P\_ERRNO* set. Possible error codes:

- *P\_EBADF*: Invalid file descriptor.
- *P\_EINVAL*: Invalid buffer or count.

Reads data from an open file.

Definition at line 166 of file *fs\_kfuncs.c*.

```

166     {
167     // handle terminal control (if doesn't control, send a STOP signal)
168     if (fd == STDIN_FILENO && current_running_pcb != NULL) {
169         if (current_running_pcb->pid != current_fg_pid) {
170             s_kill(current_running_pcb->pid, P_SIGSTOP);
171         }
172     }
173
174     // handle standard input
175     if (fd == STDIN_FILENO) {
176         return read(STDIN_FILENO, buf, n);
177     }
178
179     // validate inputs
180     if (fd < 0 || fd >= MAX_FDS || !fd_table[fd].in_use) {
181         P_ERRNO = P_EBADF;
182         return -1;
183     }

```

```

184     if (buf == NULL || n < 0) {
185         P_ERRNO = P_EINVAL;
186         return -1;
187     }
188     if (n == 0) {
189         return 0;
190     }
191
192     // check if we're at EOF already
193     if (fd_table[fd].position >= fd_table[fd].size) {
194         return 0;
195     }
196
197     // determine how many bytes we can actually read
198     uint32_t bytes_to_read = n;
199     if (fd_table[fd].position + bytes_to_read > fd_table[fd].size) {
200         bytes_to_read = fd_table[fd].size - fd_table[fd].position;
201     }
202
203     // find the block containing the current position
204     uint16_t current_block = fd_table[fd].first_block;
205     uint32_t block_index = fd_table[fd].position / block_size;
206     uint32_t block_offset = fd_table[fd].position % block_size;
207
208     // navigate to the correct block in the chain
209     for (uint32_t i = 0; i < block_index; i++) {
210         if (current_block == 0 || current_block == FAT_EOF) {
211             // unexpected end of chain
212             P_ERRNO = P_EINVAL;
213             return -1;
214         }
215         current_block = fat[current_block];
216     }
217
218     // now we're at the right block, start reading
219     uint32_t bytes_read = 0;
220
221     while (bytes_read < bytes_to_read) {
222         // how much data can we read from the current block
223         uint32_t bytes_left_in_block = block_size - block_offset;
224         uint32_t bytes_to_read_now =
225             (bytes_to_read - bytes_read) < bytes_left_in_block
226             ? (bytes_to_read - bytes_read)
227             : bytes_left_in_block;
228
229         // seek to the right position in the file
230         if (lseek(fs_fd, fat_size + (current_block - 1) * block_size + block_offset,
231             SEEK_SET) == -1) {
232             P_ERRNO = P_ELSEEK;
233             if (bytes_read > 0) {
234                 fd_table[fd].position += bytes_read;
235                 return bytes_read;
236             }
237             return -1;
238         }
239
240         // read the data from the file
241         ssize_t read_result = read(fs_fd, buf + bytes_read, bytes_to_read_now);
242         if (read_result <= 0) {
243             P_ERRNO = P_EREAD;
244             // if we already read some data, return that count
245             if (bytes_read > 0) {
246                 fd_table[fd].position += bytes_read;
247                 return bytes_read;
248             }
249             return -1;
250         }
251
252         bytes_read += read_result;
253         block_offset += read_result;
254
255         // if we've read all data from this block and still have more to read, go to
256         // the next block
257         if (block_offset == block_size && bytes_read < bytes_to_read) {
258             if (current_block == FAT_EOF) {
259                 // unexpected end of chain
260                 break;
261             }
262             current_block = fat[current_block];
263             block_offset = 0;
264         }
265
266         // if we read less than expected, we might have hit EOF
267         if (read_result < bytes_to_read_now) {
268             break;
269         }
270     }

```

```

271
272 // update file position
273 fd_table[fd].position += bytes_read;
274
275 return bytes_read;
276 }

```

#### 4.6.2.6 k\_unlink()

```

int k_unlink (
    const char * fname )

```

Removes a file from the file system.

This is a kernel-level function that deletes the specified file from the file system. The file must not be open by any process.

##### Parameters

<i>fname</i>	The name of the file to remove.
--------------	---------------------------------

##### Returns

0 on success, -1 on error with P\_ERRNO set. Possible error codes:

- P\_ENOENT: File doesn't exist.
- P\_EBUSY: File is still open by some process.

Removes a file from the file system.

Definition at line 566 of file fs\_kfuncs.c.

```

566
567 if (fname == NULL || *fname == '\0') {
568     P_ERRNO = P_EINVAL;
569     return -1;
570 }
571
572 if (!is_mounted) {
573     P_ERRNO = P_EFS_NOT_MOUNTED;
574     return -1;
575 }
576
577 // check if file is currently open by any process
578 for (int i = 0; i < MAX_FDS; i++) {
579     if (fd_table[i].in_use && strcmp(fd_table[i].filename, fname) == 0) {
580         P_ERRNO = P_EBUSY;
581         return -1;
582     }
583 }
584
585 // find the file in directory
586 dir_entry_t entry;
587 int file_offset = find_file(fname, &entry);
588 if (file_offset < 0) {
589     P_ERRNO = P_ENOENT;
590     return -1;
591 }
592
593 // mark the directory entry as deleted (set first byte to 1)
594 entry.name[0] = 1;
595
596 // write the modified directory entry back
597 if (lseek(fs_fd, file_offset, SEEK_SET) == -1) {
598     P_ERRNO = P_ELSEEK;
599     return -1;
600 }

```

```

601  if (write(fs_fd, &entry, sizeof(entry)) != sizeof(entry)) {
602      P_ERRNO = P_EWRITE;
603      return -1;
604  }
605
606  // free all blocks in the file chain
607  uint16_t current_block = entry.firstBlock;
608  uint16_t next_block;
609
610  while (current_block != FAT_FREE && current_block != FAT_EOF) {
611      next_block = fat[current_block];
612      fat[current_block] = FAT_FREE;
613      current_block = next_block;
614  }
615
616  return 0;
617 }

```

#### 4.6.2.7 k\_write()

```

int k_write (
    int fd,
    const char * str,
    int n )

```

Writes data to an open file.

This is a kernel-level function that writes *n* bytes from the provided buffer to an open file. The file position is advanced by the number of bytes written. If necessary, the file is extended.

##### Parameters

<i>fd</i>	File descriptor of the open file.
<i>str</i>	Buffer containing the data to write.
<i>n</i>	Number of bytes to write.

##### Returns

The number of bytes written on success, -1 on error with P\_ERRNO set. Possible error codes:

- P\_EBADF: Invalid file descriptor.
- P\_EINVAL: Invalid buffer or count.
- P\_EFULL: File system is full.

Writes data to an open file.

Definition at line 281 of file fs\_kfuncs.c.

```

281  {
282      // handle standard output and error
283      if (fd == STDOUT_FILENO) {
284          return write(STDOUT_FILENO, str, n);
285      }
286      if (fd == STDERR_FILENO) {
287          return write(STDERR_FILENO, str, n);
288      }
289
290      // validate inputs
291      if (fd < 0 || fd >= MAX_FDS || !fd_table[fd].in_use) {
292          P_ERRNO = P_EBADF;
293          return -1;
294      }
295      if (str == NULL || n < 0) {

```



```

296     P_ERRNO = P_EINVAL;
297     return -1;
298 }
299 if (n == 0) {
300     return 0;
301 }
302
303 // check if filesystem is mounted and FAT is valid
304 if (!is_mounted || fat == NULL) {
305     P_ERRNO = P_EFS_NOT_MOUNTED;
306     return -1;
307 }
308
309 // get file information
310 uint16_t current_block = fd_table[fd].first_block;
311 uint32_t current_position = fd_table[fd].position;
312
313 // create a local buffer for block data
314 char* block_buffer = (char*)malloc(block_size);
315 if (block_buffer == NULL) {
316     P_ERRNO = P_EALLOC;
317     return -1;
318 }
319
320 // calculate initial block position
321 uint32_t block_index = current_position / block_size;
322 uint32_t block_offset = current_position % block_size;
323
324 // if the file doesn't have a first block yet, allocate one
325 if (current_block == 0) {
326     current_block = allocate_block();
327     if (current_block == 0) {
328         P_ERRNO = P_EFULL;
329         free(block_buffer);
330         return -1;
331     }
332     fd_table[fd].first_block = current_block;
333 }
334
335 // navigate to the appropriate block
336 uint16_t prev_block = 0;
337 for (uint32_t i = 0; i < block_index; i++) {
338     if (current_block == 0 || current_block == FAT_EOF ||
339         current_block >= fat_size / 2) {
340         // reached the end of chain prematurely, need to allocate a new block
341         uint16_t new_block = allocate_block();
342         if (new_block == 0) {
343             P_ERRNO = P_EFULL;
344             free(block_buffer);
345             return -1;
346         }
347
348         // update the chain
349         if (prev_block != 0 && prev_block < fat_size / 2) {
350             fat[prev_block] = new_block;
351         } else {
352             // if there's no previous block, this must be the first one
353             fd_table[fd].first_block = new_block;
354         }
355
356         current_block = new_block;
357     }
358
359     prev_block = current_block;
360
361     // validate the block number before accessing FAT
362     if (current_block >= fat_size / 2) {
363         P_ERRNO = P_EINVAL;
364         free(block_buffer);
365         return -1;
366     }
367
368     current_block = fat[current_block];
369 }
370
371 // if we ended up without a valid block, go back to the last valid one
372 if (current_block == 0 || current_block == FAT_EOF ||
373     current_block >= fat_size / 2) {
374     if (prev_block != 0 && prev_block < fat_size / 2) {
375         uint16_t new_block = allocate_block();
376         if (new_block == 0) {
377             P_ERRNO = P_EFULL;
378             free(block_buffer);
379             return -1;
380         }
381
382         fat[prev_block] = new_block;

```

```

383     current_block = new_block;
384 } else {
385     P_ERRNO = P_EINVAL;
386     free(block_buffer);
387     return -1;
388 }
389 }
390
391 // start writing data
392 uint32_t bytes_written = 0;
393
394 while (bytes_written < n) {
395     // validate current block
396     if (current_block == 0 || current_block == FAT_EOF ||
397         current_block >= fat_size / 2) {
398         P_ERRNO = P_EINVAL;
399         break;
400     }
401
402     // how much can we write to this block
403     uint32_t space_in_block = block_size - block_offset;
404     uint32_t bytes_to_write = (n - bytes_written) < space_in_block
405         ? (n - bytes_written)
406         : space_in_block;
407
408     // position in filesystem
409     off_t block_position = fat_size + (current_block - 1) * block_size;
410
411     // if we're not writing a full block or not starting at the beginning, we
412     // need to read-modify-write
413     if (bytes_to_write < block_size || block_offset > 0) {
414         // read the current block
415         if (lseek(fs_fd, block_position, SEEK_SET) == -1) {
416             P_ERRNO = P_ELSEEK;
417             break;
418         }
419
420         // read the current block data
421         ssize_t read_result = read(fs_fd, block_buffer, block_size);
422         if (read_result < 0) {
423             P_ERRNO = P_EREAD;
424             break;
425         }
426
427         // copy the new data into the block buffer
428         memcpy(block_buffer + block_offset, str + bytes_written, bytes_to_write);
429
430         // seek back to write the modified block
431         if (lseek(fs_fd, block_position, SEEK_SET) == -1) {
432             P_ERRNO = P_ELSEEK;
433             break;
434         }
435
436         // write the full block back
437         ssize_t write_result = write(fs_fd, block_buffer, block_size);
438         if (write_result != block_size) {
439             P_ERRNO = P_EWRITE;
440             // we might have a partial write, but that's hard to handle correctly
441             break;
442         }
443     } else {
444         // we're writing a full block from the beginning
445         if (lseek(fs_fd, block_position, SEEK_SET) == -1) {
446             P_ERRNO = P_ELSEEK;
447             break;
448         }
449
450         ssize_t write_result = write(fs_fd, str + bytes_written, bytes_to_write);
451         if (write_result != bytes_to_write) {
452             P_ERRNO = P_EWRITE;
453             break;
454         }
455     }
456
457     // update counters
458     bytes_written += bytes_to_write;
459     block_offset = (block_offset + bytes_to_write) % block_size;
460
461     // if we've filled this block and still have more to write, go to the next
462     // block
463     if (block_offset == 0 && bytes_written < n) {
464         // validate current block before accessing FAT
465         if (current_block >= fat_size / 2) {
466             P_ERRNO = P_EINVAL;
467             break;
468         }
469     }

```

```

470     // check if there's a next block
471     if (fat[current_block] == FAT_EOF) {
472         // allocate a new block
473         uint16_t new_block = allocate_block();
474         if (new_block == 0) {
475             P_ERRNO = P_EFULL;
476             break;
477         }
478
479         // Update the FAT safely
480         if (current_block < fat_size / 2) {
481             fat[current_block] = new_block;
482         } else {
483             P_ERRNO = P_EINVAL;
484             break;
485         }
486
487         current_block = new_block;
488     } else {
489         current_block = fat[current_block];
490     }
491 }
492 }
493
494 // free the block buffer
495 free(block_buffer);
496
497 // update file position
498 fd_table[fd].position += bytes_written;
499
500 // update file size if needed
501 if (fd_table[fd].position > fd_table[fd].size) {
502     fd_table[fd].size = fd_table[fd].position;
503 }
504
505 // update the directory entry
506 dir_entry_t entry;
507 int dir_offset = find_file(fd_table[fd].filename, &entry);
508 if (dir_offset >= 0) {
509     entry.size = fd_table[fd].size;
510     entry.mtime = time(NULL);
511
512     if (lseek(fs_fd, dir_offset, SEEK_SET) == -1) {
513         P_ERRNO = P_ELSEEK;
514         return -1;
515     }
516     if (write(fs_fd, &entry, sizeof(entry)) != sizeof(entry)) {
517         P_ERRNO = P_EWRITE;
518         return -1;
519     }
520 }
521
522 return bytes_written;
523 }

```


## 4.7 SRC/fs/fs\_syscalls.c File Reference

```

#include "fs_syscalls.h"
#include "fs_kfuncs.h"
#include "../lib/pennos-errno.h"

```

Include dependency graph for fs\_syscalls.c:



## Functions

- `int s_open (const char *fname, int mode)`  
*System call to open a file.*
- `int s_read (int fd, char *buf, int n)`  
*System call to read from a file.*
- `int s_write (int fd, const char *str, int n)`  
*System call to write to a file.*
- `int s_close (int fd)`  
*System call to close a file.*
- `int s_unlink (const char *fname)`  
*System call to remove a file.*
- `int s_lseek (int fd, int offset, int whence)`  
*System call to reposition the file offset.*
- `int s_ls (const char *filename)`  
*System call to list files.*

### 4.7.1 Function Documentation

#### 4.7.1.1 s\_close()

```
int s_close (
    int fd )
```

System call to close a file.

Closes an open file descriptor.

This is a wrapper around the kernel function k\_close.

Definition at line 42 of file fs\_syscalls.c.

```
42     {
43     return k_close(fd);
44 }
```

#### 4.7.1.2 s\_ls()

```
int s_ls (
    const char * filename )
```

System call to list files.

Lists files in the current directory or displays file information.

This is a wrapper around the kernel function k\_ls.

Definition at line 69 of file fs\_syscalls.c.

```
69     {
70     return k_ls(filename);
71 }
```

#### 4.7.1.3 s\_lseek()

```
int s_lseek (
    int fd,
    int offset,
    int whence )
```

System call to reposition the file offset.

Repositions the file offset of an open file.

This is a wrapper around the kernel function k\_lseek.

Definition at line 60 of file fs\_syscalls.c.

```
60     {
61     return k_lseek(fd, offset, whence);
62 }
```

#### 4.7.1.4 s\_open()

```
int s_open (
    const char * fname,
    int mode )
```

System call to open a file.

Opens a file with the specified access mode.

This is a wrapper around the kernel function k\_open.

Definition at line 15 of file fs\_syscalls.c.

```
15                                     {
16     return k_open(fname, mode);
17 }
```

#### 4.7.1.5 s\_read()

```
int s_read (
    int fd,
    char * buf,
    int n )
```

System call to read from a file.

Reads data from an open file.

This is a wrapper around the kernel function k\_read.

Definition at line 24 of file fs\_syscalls.c.

```
24                                     {
25     return k_read(fd, buf, n);
26 }
```

#### 4.7.1.6 s\_unlink()

```
int s_unlink (
    const char * fname )
```

System call to remove a file.

Removes a file from the file system.

This is a wrapper around the kernel function k\_unlink.

Definition at line 51 of file fs\_syscalls.c.

```
51                                     {
52     return k_unlink(fname);
53 }
```

## 4.7.1.7 s\_write()

```
int s_write (
    int fd,
    const char * str,
    int n )
```

System call to write to a file.

Writes data to an open file.

This is a wrapper around the kernel function k\_write.


Definition at line 33 of file fs\_syscalls.c.

```
33 {
34     return k_write(fd, str, n);
35 }
```


## 4.8 SRC/fs/fs\_syscalls.h File Reference

```
#include <stddef.h>
```

Include dependency graph for fs\_syscalls.h:



This graph shows which files directly or indirectly include this file:



## Macros

- #define `STDIN_FILENO` 0
- #define `STDOUT_FILENO` 1
- #define `STDERR_FILENO` 2

## Functions

- int `s_open` (const char \*fname, int mode)  
*Opens a file with the specified access mode.*
- int `s_read` (int fd, char \*buf, int n)  
*Reads data from an open file.*
- int `s_write` (int fd, const char \*str, int n)  
*Writes data to an open file.*
- int `s_close` (int fd)  
*Closes an open file descriptor.*
- int `s_unlink` (const char \*fname)  
*Removes a file from the file system.*
- int `s_lseek` (int fd, int offset, int whence)  
*Repositions the file offset of an open file.*
- int `s_ls` (const char \*filename)  
*Lists files in the current directory or displays file information.*

## 4.8.1 Macro Definition Documentation

### 4.8.1.1 STDERR\_FILENO

```
#define STDERR_FILENO 2
```

Definition at line 18 of file fs\_syscalls.h.

### 4.8.1.2 STDIN\_FILENO

```
#define STDIN_FILENO 0
```

Definition at line 16 of file fs\_syscalls.h.

### 4.8.1.3 STDOUT\_FILENO

```
#define STDOUT_FILENO 1
```

Definition at line 17 of file fs\_syscalls.h.

## 4.8.2 Function Documentation

### 4.8.2.1 s\_close()

```
int s_close (  
    int fd )
```

Closes an open file descriptor.

This function closes the file descriptor fd, making it available for reuse. If this is the last reference to the underlying file, any necessary cleanup is performed.



#### Parameters

<i>fd</i>	The file descriptor to close.
-----------	-------------------------------

#### Returns

On success, returns 0. On error, returns -1 and sets P\_ERRNO appropriately:

- P\_EBADF: *fd* is not a valid file descriptor.

Closes an open file descriptor.

This is a wrapper around the kernel function `k_close`.

Definition at line 42 of file `fs_syscalls.c`.

```
42     {  
43         return k_close(fd);  
44     }
```

#### 4.8.2.2 s\_ls()

```
int s_ls (  
        const char * filename )
```

Lists files in the current directory or displays file information.

If *filename* is NULL, this function lists all files in the current directory. If *filename* refers to a specific file, it displays detailed information about that file.

#### Parameters

<i>filename</i>	The name of the file to get information about, or NULL to list all files.
-----------------	---

#### Returns

On success, returns 0. On error, returns -1 and sets P\_ERRNO appropriately:

- P\_ENOENT: The specified file does not exist.

Lists files in the current directory or displays file information.

This is a wrapper around the kernel function `k_ls`.

Definition at line 69 of file `fs_syscalls.c`.

```
69     {  
70         return k_ls(filename);  
71     }
```

#### 4.8.2.3 s\_lseek()

```
int s_lseek (
    int fd,
    int offset,
    int whence )
```

Repositions the file offset of an open file.

This function repositions the offset of the file descriptor `fd` to the argument `offset` according to the directive `whence`.

##### Parameters

<i>fd</i>	The file descriptor of an open file.
<i>offset</i>	The offset in bytes.
<i>whence</i>	Specifies the reference position: <ul style="list-style-type: none"><li>• <code>SEEK_SET</code> (0): The offset is set relative to the start of the file.</li><li>• <code>SEEK_CUR</code> (1): The offset is set relative to the current position.</li><li>• <code>SEEK_END</code> (2): The offset is set relative to the end of the file.</li></ul>

##### Returns

On success, returns the resulting offset from the beginning of the file. On error, returns -1 and sets `P_ERRNO` appropriately:

- `P_EBADF`: `fd` is not a valid file descriptor.
- `P_EINVAL`: `whence` is not valid or the resulting offset would be negative.

Repositions the file offset of an open file.

This is a wrapper around the kernel function `k_lseek`.

Definition at line 60 of file `fs_syscalls.c`.

```
60 {
61     return k_lseek(fd, offset, whence);
62 }
```

#### 4.8.2.4 s\_open()

```
int s_open (
    const char * fname,
    int mode )
```

Opens a file with the specified access mode.

This function provides a user-level interface to the kernel's file open operation. It opens the specified file with the given access mode and returns a file descriptor that can be used in subsequent operations on the file.

#### Parameters

<i>fname</i>	The name of the file to open.
<i>mode</i>	A combination of F_READ, F_WRITE, and F_APPEND.

#### Returns

On success, returns a non-negative integer representing the file descriptor. On error, returns -1 and sets P\_ERRNO appropriately:

- P\_ENOENT: The file does not exist and F\_READ was specified.
- P\_EINVAL: Invalid parameters (NULL filename or invalid mode).
- P\_EFULL: No space left on device or file descriptor table is full.

Opens a file with the specified access mode.

This is a wrapper around the kernel function k\_open.

Definition at line 15 of file fs\_syscalls.c.

```
15 {  
16     return k_open(fname, mode);  
17 }
```

#### 4.8.2.5 s\_read()

```
int s_read (  
    int fd,  
    char * buf,  
    int n )
```

Reads data from an open file.

This function reads up to n bytes from the file associated with the file descriptor fd into the buffer starting at buf. The file offset is advanced by the number of bytes read.

#### Parameters

<i>fd</i>	The file descriptor of an open file.
<i>n</i>	The maximum number of bytes to read.
<i>buf</i>	The buffer to store the read data.

#### Returns

On success, returns the number of bytes read (0 indicates end of file). On error, returns -1 and sets P\_ERRNO appropriately:

- P\_EBADF: fd is not a valid file descriptor or is not open for reading.
- P\_EINVAL: Invalid parameters (NULL buffer or negative count).

Reads data from an open file.

This is a wrapper around the kernel function `k_read`.

Definition at line 24 of file `fs_syscalls.c`.

```
24                                     {  
25     return k_read(fd, buf, n);  
26 }
```

#### 4.8.2.6 `s_unlink()`

```
int s_unlink (  
    const char * fname )
```

Removes a file from the file system.

This function removes the specified file from the file system. If the file is currently open, the behavior depends on the implementation.

##### Parameters

<i>fname</i>	The name of the file to remove.
--------------	---------------------------------

##### Returns

On success, returns 0. On error, returns -1 and sets `P_ERRNO` appropriately:

- `P_ENOENT`: The file does not exist.
- `P_EBUSY`: The file is currently in use.
- `P_EINVAL`: Invalid parameter (NULL filename).

Removes a file from the file system.

This is a wrapper around the kernel function `k_unlink`.

Definition at line 51 of file `fs_syscalls.c`.

```
51                                     {  
52     return k_unlink(fname);  
53 }
```

#### 4.8.2.7 `s_write()`

```
int s_write (  
    int fd,  
    const char * str,  
    int n )
```

Writes data to an open file.

This function writes up to `n` bytes from the buffer starting at `str` to the file associated with the file descriptor `fd`. The file offset is advanced by the number of bytes written.

## Parameters

<i>fd</i>	The file descriptor of an open file.
<i>str</i>	The buffer containing the data to be written.
<i>n</i>	The number of bytes to write.

## Returns

On success, returns the number of bytes written. On error, returns -1 and sets P\_ERRNO appropriately:

- P\_EBADF: *fd* is not a valid file descriptor or is not open for writing.
- P\_EINVAL: Invalid parameters (NULL buffer or negative count).
- P\_EFULL: No space left on device.

Writes data to an open file.

This is a wrapper around the kernel function `k_write`.

Definition at line 33 of file `fs_syscalls.c`.

```

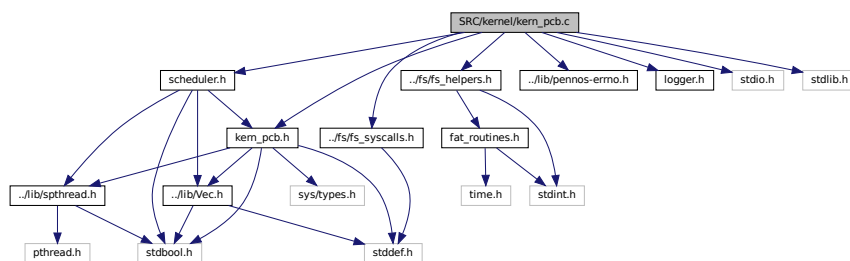
33                                     {
34     return k_write(fd, str, n);
35 }
```

## 4.9 SRC/kernel/kern\_pcb.c File Reference

```

#include "kern_pcb.h"
#include "../fs/fs_helpers.h"
#include "../fs/fs_syscalls.h"
#include "../lib/pennos-errno.h"
#include "logger.h"
#include "scheduler.h"
#include "stdio.h"
#include "stdlib.h"
```

Include dependency graph for `kern_pcb.c`:



## Functions

- void `free_pcb` (void \*pcb)  
*Free resources associated with a PCB.*
- `pcb_t * create_pcb` (pid\_t pid, pid\_t par\_pid, int priority, int input\_fd, int output\_fd)  
*Initializes a PCB with the given parameters.*
- void `remove_child_in_parent` (pcb\_t \*parent, pcb\_t \*child)  
*Removes a child PCB from its parent's child list.*
- `pcb_t * k_proc_create` (pcb\_t \*parent, int priority)  
*Creates a new process. If the parent is NULL, it creates the init process.*
- void `k_proc_cleanup` (pcb\_t \*proc)  
*Cleans up a process by removing it from its parent's child list, removing its children, decrementing file descriptor reference counts, closing files, and freeing the PCB.*

## Variables

- int `next_pid` = 2
- `Vec` `current_pcb`s
- `pcb_t *` `current_running_pcb`

### 4.9.1 Function Documentation

#### 4.9.1.1 create\_pcb()

```
pcb_t* create_pcb (
    pid_t pid,
    pid_t par_pid,
    int priority,
    int input_fd,
    int output_fd )
```

Initializes a PCB with the given parameters.

Creates a new PCB and initializes its fields. Notably, the thread handle and cmd are left out. It's up to the user to assign them post-call.

Definition at line 41 of file kern\_pcb.c.

```
45     {
46     pcb_t* ret_pcb = malloc(sizeof(pcb_t));
47     if (ret_pcb == NULL) {
48         perror("malloc failed for PCB creation");
49         return NULL;
50     }
51
52     ret_pcb->pid = pid;
53     ret_pcb->par_pid = par_pid;
54     ret_pcb->priority = priority;
55     ret_pcb->process_state = 'R'; // running by default
56     ret_pcb->input_fd = input_fd;
57     ret_pcb->output_fd = output_fd;
58     ret_pcb->process_status = 0; // default status
59
60     ret_pcb->child_pcb = vec_new(0, NULL); // NULL deconstructor prevents
61                                           // double free
62
63     for (int i = 0; i < 3; i++) {
64         ret_pcb->signals[i] = false;
65     }
66
67     ret_pcb->is_sleeping = false;
68     ret_pcb->time_to_wake = -1; // default to not sleeping
69
70     return ret_pcb;
71 }
```

### 4.9.1.2 free\_pcb()

```
void free_pcb (
    void * pcb )
```

Free resources associated with a PCB.

Frees all malloced memory associated with the PCB. Note that this will destroy children too, so be careful when using it. In particular, make sure to remove any children pcbs you want to preserve.

Definition at line 29 of file kern\_pcb.c.

```
29 {
30     pcb_t* casted_pcb = (pcb_t*)pcb;
31
32     free(casted_pcb->cmd_str);
33     vec_destroy(&casted_pcb->child_pcbs); // will free any remaining
34                                           // children too!
35     free(casted_pcb);
36 }
```

### 4.9.1.3 k\_proc\_cleanup()

```
void k_proc_cleanup (
    pcb_t * proc )
```

Cleans up a process by removing it from its parent's child list, removing its children, decrementing file descriptor reference counts, closing files, and freeing the PCB.

Clean up a terminated/finished thread's resources. This may include freeing the PCB, handling children, etc. If a child is orphaned, the INIT process becomes its parent.

Definition at line 150 of file kern\_pcb.c.

```
150 {
151     // if proc has parent (i.e. isn't init) then remove it from parent's child
152     // list
153     pcb_t* par_pcb = get_pcb_in_queue(&current_pcbs, proc->par_pid);
154     if (par_pcb != NULL) {
155         remove_child_in_parent(par_pcb, proc);
156     } else {
157         P_ERRNO = P_ENULL;
158         return;
159     }
160
161     // if proc has children, remove them and assign them to init parent
162     if (vec_len(&proc->child_pcbs) > 0) {
163         // retrieve the init process
164         pcb_t* init_pcb =
165             get_pcb_in_queue(&current_pcbs, 1); // init process has pid 1
166
167         while (vec_len(&proc->child_pcbs) > 0) {
168             pcb_t* curr_child = vec_get(&proc->child_pcbs, 0);
169             vec_push_back(&init_pcb->child_pcbs, curr_child);
170             vec_erase_no_deletor(&proc->child_pcbs, 0); // don't free in erase
171             curr_child->par_pid = 1; // update parent to init (pid 1)
172             log_generic_event('O', curr_child->pid, curr_child->priority,
173                             curr_child->cmd_str);
174         }
175     }
176
177     // decr reference counts + close files if necessary
178     for (int i = 0; i < FILE_DESCRIPTOR_TABLE_SIZE; i++) {
179         if (proc->fd_table[i] != -1) {
180             if (decrement_fd_ref_count(proc->fd_table[i]) == 0) {
181                 s_close(
182                     proc->fd_table[i]); // close the fd since no other process using
183             }
184         }
185     }
186
187     // cancel + join this thread
```

```

188  pthread_cancel(proc->thread_handle);
189  pthread_continue(proc->thread_handle);
190  pthread_suspend(proc->thread_handle);
191  pthread_join(proc->thread_handle, NULL);
192
193  // delete this process from any queue it's in + free it
194  delete_process_from_all_queues(proc);
195  free_pcb(proc);
196 }

```

#### 4.9.1.4 k\_proc\_create()

```

pcb_t* k_proc_create (
    pcb_t * parent,
    int priority )

```

Creates a new process. If the parent is NULL, it creates the init process.

Create a new child process, inheriting applicable properties from the parent. Also inserts the created child into the correct scheduler queue based on its priority.

Definition at line 93 of file kern\_pcb.c.

```

93  {
94  if (parent == NULL) { // init creation case
95      pcb_t* init = create_pcb(1, 0, 0, 0, 1);
96      if (init == NULL) {
97          P_ERRNO = P_ENULL;
98          return NULL;
99      }
100     init->fd_table[0] = STDIN_FILENO;
101     init->fd_table[1] = STDOUT_FILENO;
102     init->fd_table[2] = STDERR_FILENO;
103     for (int i = 3; i < FILE_DESCRIPTOR_TABLE_SIZE; i++) {
104         init->fd_table[i] = -1;
105     }
106
107     increment_fd_ref_count(STDIN_FILENO);
108     increment_fd_ref_count(STDOUT_FILENO);
109     increment_fd_ref_count(STDERR_FILENO);
110
111     current_running_pcb = init;
112     put_pcb_into_correct_queue(init);
113     vec_push_back(&current_pcb, init);
114     return init;
115 }
116
117 pcb_t* child = create_pcb(next_pid++, parent->pid, priority, parent->input_fd,
118                          parent->output_fd);
119 if (child == NULL) {
120     P_ERRNO = P_ENULL;
121     return NULL;
122 }
123
124 // copy parent's fd table
125 for (int i = 0; i < FILE_DESCRIPTOR_TABLE_SIZE; i++) {
126     child->fd_table[i] = parent->fd_table[i];
127 }
128
129 for (int i = 0; i < FILE_DESCRIPTOR_TABLE_SIZE; i++) {
130     if (child->fd_table[i] != -1) {
131         increment_fd_ref_count(child->fd_table[i]);
132     }
133 }
134
135 // update parent as needed
136 vec_push_back(&parent->child_pcb, child);
137
138 // add to appropriate queue
139 put_pcb_into_correct_queue(child);
140 vec_push_back(&current_pcb, child);
141
142 return child;
143 }

```



#### 4.9.1.5 remove\_child\_in\_parent()

```
void remove_child_in_parent (
    pcb_t * parent,
    pcb_t * child )
```

Removes a child PCB from its parent's child list.

Given a parent, removes the child from the parent's child vector if its exists. Notably, it does not free the child but simply removes it via the `vec_erase_no_deletor` function.

Definition at line 76 of file `kern_pcb.c`.

```
76
77     for (int i = 0; i < vec_len(&parent->child_pcb); i++) {
78         pcb_t* curr_child = (pcb_t*)vec_get(&parent->child_pcb, i);
79         if (curr_child->pid == child->pid) {
80             vec_erase_no_deletor(&parent->child_pcb, i);
81             return;
82         }
83     }
84 }
```

### 4.9.2 Variable Documentation

#### 4.9.2.1 current\_pcb

```
Vec current_pcb [extern]
```

Definition at line 30 of file `scheduler.c`.

#### 4.9.2.2 current\_running\_pcb

```
pcb_t* current_running_pcb [extern]
```

Definition at line 38 of file `scheduler.c`.

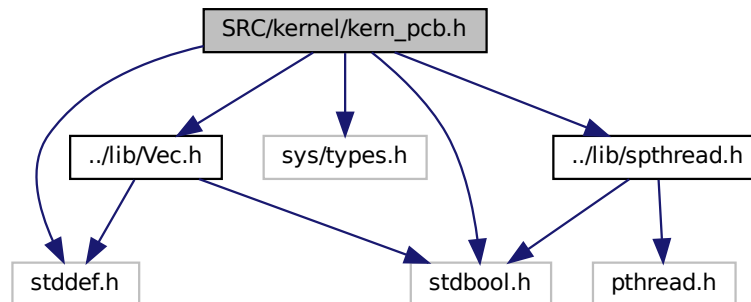
#### 4.9.2.3 next\_pid

```
int next_pid = 2
```


Definition at line 16 of file `kern_pcb.c`.

## 4.10 SRC/kernel/kern\_pcb.h File Reference

```
#include <stddef.h>
#include <stdbool.h>
#include <sys/types.h>
#include "../lib/spthread.h"
#include "../lib/Vec.h"
Include dependency graph for kern_pcb.h:
```



This graph shows which files directly or indirectly include this file:



## Classes

- struct [pcb\\_st](#)

The PCB structure, which contains all the information about a process. Notably, it contains the thread handle, pid, parent pid, child pcbs, priority level, process state, command string, signals to be sent, input and output file descriptors, process status, sleeping status, and time to wake.

## Macros

- #define [FILE\\_DESCRIPTOR\\_TABLE\\_SIZE](#) 100

## Typedefs

- typedef struct [pcb\\_st](#) [pcb\\_t](#)

The PCB structure, which contains all the information about a process. Notably, it contains the thread handle, pid, parent pid, child pcbs, priority level, process state, command string, signals to be sent, input and output file descriptors, process status, sleeping status, and time to wake.

## Functions

- `pcb_t * create_pcb` (`pid_t pid`, `pid_t par_pid`, `int priority`, `int input_fd`, `int output_fd`)  
*Creates a new PCB and initializes its fields. Notably, the thread handle and cmd are left out. It's up to the user to assign them post-call.*
- `void free_pcb` (`void *pcb`)  
*Frees all malloced memory associated with the PCB. Note that this will destroy children too, so be careful when using it. In particular, make sure to remove any children pcbs you want to preserve.*
- `void remove_child_in_parent` (`pcb_t *parent`, `pcb_t *child`)  
*Given a parent, removes the child from the parent's child vector if its exists. Notably, it does not free the child but simply removes it via the `vec_erase_no_deletor` function.*
- `pcb_t * k_proc_create` (`pcb_t *parent`, `int priority`)  
*Create a new child process, inheriting applicable properties from the parent. Also inserts the created child into the correct scheduler queue based on its priority.*
- `void k_proc_cleanup` (`pcb_t *proc`)  
*Clean up a terminated/finished thread's resources. This may include freeing the PCB, handling children, etc. If a child is orphaned, the INIT process becomes its parent.*

## 4.10.1 Macro Definition Documentation

### 4.10.1.1 FILE\_DESCRIPTOR\_TABLE\_SIZE

```
#define FILE_DESCRIPTOR_TABLE_SIZE 100
```

Definition at line 17 of file kern\_pcb.h.

## 4.10.2 Typedef Documentation

### 4.10.2.1 pcb\_t

```
typedef struct pcb_st pcb_t
```

The PCB structure, which contains all the information about a process. Notably, it contains the thread handle, pid, parent pid, child pcbs, priority level, process state, command string, signals to be sent, input and output file descriptors, process status, sleeping status, and time to wake.

## 4.10.3 Function Documentation

### 4.10.3.1 create\_pcb()

```
pcb_t* create_pcb (
    pid_t pid,
    pid_t par_pid,
    int priority,
    int input_fd,
    int output_fd )
```

Creates a new PCB and initializes its fields. Notably, the thread handle and cmd are left out. It's up to the user to assign them post-call.

**Parameters**

<i>pid</i>	the new process id
<i>par_pid</i>	the parent process id
<i>priority</i>	the priority level (0,1,2)
<i>input_fd</i>	input fd
<i>output_fd</i>	output fd

**Returns**

pointer to the newly created and malloced PCB or NULL if failure

Creates a new PCB and initializes its fields. Notably, the thread handle and cmd are left out. It's up to the user to assign them post-call.

Definition at line 41 of file kern\_pcb.c.

```

45 {
46     pcb_t* ret_pcb = malloc(sizeof(pcb_t));
47     if (ret_pcb == NULL) {
48         perror("malloc failed for PCB creation");
49         return NULL;
50     }
51
52     ret_pcb->pid = pid;
53     ret_pcb->par_pid = par_pid;
54     ret_pcb->priority = priority;
55     ret_pcb->process_state = 'R'; // running by default
56     ret_pcb->input_fd = input_fd;
57     ret_pcb->output_fd = output_fd;
58     ret_pcb->process_status = 0; // default status
59
60     ret_pcb->child_pcb = vec_new(0, NULL); // NULL deconstructor prevents
61                                           // double free
62
63     for (int i = 0; i < 3; i++) {
64         ret_pcb->signals[i] = false;
65     }
66
67     ret_pcb->is_sleeping = false;
68     ret_pcb->time_to_wake = -1; // default to not sleeping
69
70     return ret_pcb;
71 }
```

**4.10.3.2 free\_pcb()**

```

void free_pcb (
    void * pcb )
```

Frees all malloced memory associated with the PCB. Note that this will destroy children too, so be careful when using it. In particular, make sure to remove any children pcbs you want to preserve.

**Parameters**

<i>pcb</i>	Pointer to the PCB to be freed, NULL if error
------------	---

Frees all malloced memory associated with the PCB. Note that this will destroy children too, so be careful when using it. In particular, make sure to remove any children pcbs you want to preserve.

Definition at line 29 of file kern\_pcb.c.

```

29      {
30  pcb_t* casted_pcb = (pcb_t*)pcb;
31
32  free(casted_pcb->cmd_str);
33  vec_destroy(&casted_pcb->child_pcbs); // will free any remaining
34                                     // children too!
35  free(casted_pcb);
36  }
```

#### 4.10.3.3 k\_proc\_cleanup()

```

void k_proc_cleanup (
    pcb_t * proc )
```

Clean up a terminated/finished thread's resources. This may include freeing the PCB, handling children, etc. If a child is orphaned, the INIT process becomes its parent.

##### Parameters

<i>proc</i>	a pcb ptr to the terminated/finished thread
-------------	---

Clean up a terminated/finished thread's resources. This may include freeing the PCB, handling children, etc. If a child is orphaned, the INIT process becomes its parent.

Definition at line 150 of file kern\_pcb.c.

```

150      {
151  // if proc has parent (i.e. isn't init) then remove it from parent's child
152  // list
153  pcb_t* par_pcb = get_pcb_in_queue(&current_pcbs, proc->par_pid);
154  if (par_pcb != NULL) {
155      remove_child_in_parent(par_pcb, proc);
156  } else {
157      P_ERRNO = P_ENULL;
158      return;
159  }
160
161  // if proc has children, remove them and assign them to init parent
162  if (vec_len(&proc->child_pcbs) > 0) {
163      // retrieve the init process
164      pcb_t* init_pcb =
165          get_pcb_in_queue(&current_pcbs, 1); // init process has pid 1
166
167      while (vec_len(&proc->child_pcbs) > 0) {
168          pcb_t* curr_child = vec_get(&proc->child_pcbs, 0);
169          vec_push_back(&init_pcb->child_pcbs, curr_child);
170          vec_erase_no_deletor(&proc->child_pcbs, 0); // don't free in erase
171          curr_child->par_pid = 1; // update parent to init (pid 1)
172          log_generic_event('O', curr_child->pid, curr_child->priority,
173                          curr_child->cmd_str);
174      }
175  }
176
177  // decr reference counts + close files if necessary
178  for (int i = 0; i < FILE_DESCRIPTOR_TABLE_SIZE; i++) {
179      if (proc->fd_table[i] != -1) {
180          if (decrement_fd_ref_count(proc->fd_table[i]) == 0) {
181              s_close(
182                  proc->fd_table[i]); // close the fd since no other process using
183          }
184      }
185  }
186
187  // cancel + join this thread
188  spthread_cancel(proc->thread_handle);
189  spthread_continue(proc->thread_handle);
190  spthread_suspend(proc->thread_handle);
191  spthread_join(proc->thread_handle, NULL);
192
193  // delete this process from any queue it's in + free it
194  delete_process_from_all_queues(proc);
```

```

195     free_pcb(proc);
196 }

```

#### 4.10.3.4 k\_proc\_create()

```

pcb_t* k_proc_create (
    pcb_t * parent,
    int priority )

```

Create a new child process, inheriting applicable properties from the parent. Also inserts the created child into the correct scheduler queue based on its priority.

##### Parameters

<i>parent</i>	a pointer to the parent pcb
<i>priority</i>	the priority of the child, usually 1 but exceptions like shell exist

##### Returns

Reference to the child PCB or NULL if error

Create a new child process, inheriting applicable properties from the parent. Also inserts the created child into the correct scheduler queue based on its priority.

Definition at line 93 of file kern\_pcb.c.

```

93     {
94         if (parent == NULL) { // init creation case
95             pcb_t* init = create_pcb(1, 0, 0, 0, 1);
96             if (init == NULL) {
97                 P_ERRNO = P_ENULL;
98                 return NULL;
99             }
100             init->fd_table[0] = STDIN_FILENO;
101             init->fd_table[1] = STDOUT_FILENO;
102             init->fd_table[2] = STDERR_FILENO;
103             for (int i = 3; i < FILE_DESCRIPTOR_TABLE_SIZE; i++) {
104                 init->fd_table[i] = -1;
105             }
106             increment_fd_ref_count(STDIN_FILENO);
107             increment_fd_ref_count(STDOUT_FILENO);
108             increment_fd_ref_count(STDERR_FILENO);
109             current_running_pcb = init;
110             put_pcb_into_correct_queue(init);
111             vec_push_back(&current_pcb, init);
112             return init;
113         }
114         pcb_t* child = create_pcb(next_pid++, parent->pid, priority, parent->input_fd,
115                                 parent->output_fd);
116         if (child == NULL) {
117             P_ERRNO = P_ENULL;
118             return NULL;
119         }
120         // copy parent's fd table
121         for (int i = 0; i < FILE_DESCRIPTOR_TABLE_SIZE; i++) {
122             child->fd_table[i] = parent->fd_table[i];
123         }
124         for (int i = 0; i < FILE_DESCRIPTOR_TABLE_SIZE; i++) {
125             if (child->fd_table[i] != -1) {
126                 increment_fd_ref_count(child->fd_table[i]);
127             }
128         }
129     }

```

```

134
135 // update parent as needed
136 vec_push_back(&parent->child_pcb, child);
137
138 // add to appropriate queue
139 put_pcb_into_correct_queue(child);
140 vec_push_back(&current_pcb, child);
141
142 return child;
143 }

```

#### 4.10.3.5 remove\_child\_in\_parent()

```

void remove_child_in_parent (
    pcb_t * parent,
    pcb_t * child )

```

Given a parent, removes the child from the parent's child vector if its exists. Notably, it does not free the child but simply removes it via the `vec_erase_no_deletor` function.

##### Parameters

<i>parent</i>	a ptr to the parent pcb with the child list
<i>child</i>	a ptr to the child pcb that we'd like to remove

Given a parent, removes the child from the parent's child vector if its exists. Notably, it does not free the child but simply removes it via the `vec_erase_no_deletor` function.

Definition at line 76 of file `kern_pcb.c`.

```

76
77 for (int i = 0; i < vec_len(&parent->child_pcb); i++) {
78     pcb_t* curr_child = (pcb_t*)vec_get(&parent->child_pcb, i);
79     if (curr_child->pid == child->pid) {
80         vec_erase_no_deletor (&parent->child_pcb, i);
81         return;
82     }
83 }
84 }

```

## 4.11 SRC/kernel/kern\_sys\_calls.c File Reference

```

#include "kern_sys_calls.h"
#include <stdlib.h>
#include <string.h>
#include "../fs/fs_kfuncs.h"
#include "../lib/Vec.h"
#include "../lib/pennos-errno.h"
#include "../shell/builtins.h"
#include "../shell/shell.h"
#include "kern_pcb.h"
#include "logger.h"
#include "scheduler.h"
#include "signal.h"
#include "../fs/fs_syscalls.h"

```





## 4.11.1 Function Documentation

### 4.11.1.1 delete\_from\_explicit\_queue()

```
void delete_from_explicit_queue (
    Vec * queue_to_delete_from,
    int pid )
```

Deletes a PCB from the specified explicit queue based on its PID.

Helper function that deletes the given PCB from the explicit queue passed in. Notably, it does not free the PCB but instead uses `vec_erase_no_deletor` to remove it from the queue.

Definition at line 107 of file `kern_sys_calls.c`.

```
107 {
108     int index = determine_index_in_queue(queue_to_delete_from, pid);
109     if (index != -1) {
110         vec_erase_no_deletor(queue_to_delete_from, index);
111     }
112 }
```

### 4.11.1.2 delete\_from\_queue()

```
void delete_from_queue (
    int queue_id,
    int pid )
```

Deletes a PCB from the specified queue based on its PID.

Deletes the PCB with the specified PID from one of the priority queues, selected by the provided `queue_id` (0, 1, or 2).

Definition at line 88 of file `kern_sys_calls.c`.

```
88 {
89     Vec* queue = NULL;
90     if (queue_id == 0) {
91         queue = &zero_priority_queue;
92     } else if (queue_id == 1) {
93         queue = &one_priority_queue;
94     } else {
95         queue = &two_priority_queue;
96     }
97
98     int index = determine_index_in_queue(queue, pid);
99     if (index != -1) {
100         vec_erase_no_deletor(queue, index);
101     }
102 }
```

#### 4.11.1.3 determine\_index\_in\_queue()

```
int determine_index_in_queue (
    Vec * queue,
    int pid )
```

Determines the index of a PCB in a given queue.

Given a thread pid and Vec\* queue, this helper function determines the vector index of the thread/pid in the queue. If the thread/pid is not found, it returns -1.

Definition at line 40 of file kern\_sys\_calls.c.

```
40
41 for (int i = 0; i < vec_len(queue); i++) {
42     pcb_t* curr_pcb = vec_get(queue, i);
43     if (curr_pcb->pid == pid) {
44         return i;
45     }
46 }
47
48 return -1; // not found
49 }
```

#### 4.11.1.4 init\_func()

```
void* init_func (
    void * input )
```

The function that runs the shell process.

The init process function. It spawns the shell process and reaps zombie children.

Definition at line 117 of file kern\_sys\_calls.c.

```
117 {
118     char* shell_argv[] = {"shell", NULL};
119     s_spawn(shell, shell_argv, STDIN_FILENO, STDOUT_FILENO);
120
121     // continuously wait for and reap zombie children
122     while (true) {
123         int status;
124         s_waitpid(-1, &status, false);
125     }
126
127     return NULL; // should never reach
128 }
```

#### 4.11.1.5 move\_pcb\_correct\_queue()

```
void move_pcb_correct_queue (
    int prev_priority,
    int new_priority,
    pcb_t * curr_pcb )
```

Moves a PCB from its previous priority queue to its new priority queue.

Given a thread's previous priority, this helper checks if the thread is present in that priority's queue, removes it from that queue if so, and then puts it into the new priority level's queue.

Definition at line 54 of file kern\_sys\_calls.c.

```

56                                     {
57     Vec* prev_queue;
58     Vec* new_queue;
59
60     if (prev_priority == 0) {
61         prev_queue = &zero_priority_queue;
62     } else if (prev_priority == 1) {
63         prev_queue = &one_priority_queue;
64     } else {
65         prev_queue = &two_priority_queue;
66     }
67
68     if (new_priority == 0) {
69         new_queue = &zero_priority_queue;
70     } else if (new_priority == 1) {
71         new_queue = &one_priority_queue;
72     } else {
73         new_queue = &two_priority_queue;
74     }
75
76     // delete from prev_queue, if it's present at all
77     int ind = determine_index_in_queue(prev_queue, curr_pcb->pid);
78     if (ind != -1) {
79         vec_erase_no_deletor(prev_queue, ind);
80     }
81
82     vec_push_back(new_queue, curr_pcb);
83 }
```

#### 4.11.1.6 s\_cleanup\_init\_process()

```
void s_cleanup_init_process ( )
```

Cleans up Init's resources.

Wrapper system-level function to be called in pennos's main method to clean up the init process.

Definition at line 157 of file kern\_sys\_calls.c.

```

157                                     {
158     k_proc_cleanup(get_pcb_in_queue(&current_pcb, 1));
159 }
```

#### 4.11.1.7 s\_echo()

```
void* s_echo (
    void * arg )
```

System-level wrapper for the shell built-in command "echo".

##### Parameters

<i>arg</i>	the pass along arguments to the u_echo function
------------	---

##### Returns

NULL, dummy return value

Definition at line 358 of file kern\_sys\_calls.c.

```

358     {
359     char** argv = (char**)arg;
360     if (argv[1] == NULL) { // no args case
361         s_exit();
362         return NULL;
363     }
364
365     int i = 1; // words after "echo"
366     while (argv[i] != NULL) { // while the arg isn't NULL
367         s_write(current_running_pcb->output_fd, argv[i], strlen(argv[i]));
368         s_write(current_running_pcb->output_fd, " ", 1);
369         i++;
370     }
371
372     s_write(current_running_pcb->output_fd, "\n", 1);
373     return NULL;
374 }

```

#### 4.11.1.8 s\_exit()

```

void s_exit (
    void )

```

Exits the current process and cleans up its resources.

Unconditionally exit the calling process.

Definition at line 293 of file kern\_sys\_calls.c.

```

293     {
294     // Set process state to zombie
295     current_running_pcb->process_state = 'Z';
296     current_running_pcb->process_status = 20; // EXITED_NORMALLY
297
298     // Log the exit
299     log_generic_event('E', current_running_pcb->pid,
300                     current_running_pcb->priority,
301                     current_running_pcb->cmd_str);
302
303     delete_from_queue(current_running_pcb->priority, current_running_pcb->pid);
304
305     log_generic_event('Z', current_running_pcb->pid,
306                     current_running_pcb->priority,
307                     current_running_pcb->cmd_str);
308 }

```

#### 4.11.1.9 s\_kill()

```

int s_kill (
    pid_t pid,
    int signal )

```

Sends a signal to a process with specified pid.

Send a signal to a particular process.

Definition at line 279 of file kern\_sys\_calls.c.

```

279     {
280     pcb_t* pcb_with_pid = get_pcb_in_queue(&current_pcb, pid);
281     if (pcb_with_pid == NULL) {
282         return -1; // pid not found case
283     }
284
285     pcb_with_pid->signals[signal] = true; // signal flagged
286     log_generic_event('S', pid, pcb_with_pid->priority, pcb_with_pid->cmd_str);
287     return 0;
288 }

```

## 4.11.1.10 s\_nice()

```
int s_nice (
    pid_t pid,
    int priority )
```

Sets the priority of a process with specified pid.

Set the priority of the specified thread.

Definition at line 313 of file kern\_sys\_calls.c.

```
313 {
314     if (priority < 0 || priority > 2) { // error check
315         return -1;
316     }
317
318     pcb_t* curr_pcb = get_pcb_in_queue(&current_pcb, pid);
319     if (curr_pcb != NULL) { // found + exists
320         move_pcb_correct_queue(curr_pcb->priority, priority, curr_pcb);
321         log_nice_event(pid, curr_pcb->priority, priority, curr_pcb->cmd_str);
322         curr_pcb->priority = priority;
323         return 0;
324     }
325     return -1; // pid not found
326 }
327 }
```

## 4.11.1.11 s\_ps()

```
void* s_ps (
    void * arg )
```

System-level wrapper for the shell built-in command "ps".

## Parameters

<i>arg</i>	the pass along arguments to the u_ps function
------------	---

## Returns

NULL, dummy return value

Definition at line 379 of file kern\_sys\_calls.c.

```
379 {
380     char pid_top[] = "PID\tPPID\tPRI\tSTAT\tCMD\n";
381     s_write(current_running_pcb->output_fd, pid_top, strlen(pid_top));
382     for (int i = 0; i < vec_len(&current_pcb); i++) {
383         pcb_t* curr_pcb = (pcb_t*)vec_get(&current_pcb, i);
384         char buffer[100];
385         snprintf(buffer, sizeof(buffer), "%d\t%d\t%d\t%c\t%s\n", curr_pcb->pid,
386             curr_pcb->par_pid, curr_pcb->priority, curr_pcb->process_state,
387             curr_pcb->cmd_str);
388         s_write(current_running_pcb->output_fd, buffer, strlen(buffer));
389     }
390     return NULL;
391 }
```

#### 4.11.1.12 s\_sleep()

```
void s_sleep (
    unsigned int ticks )
```

Suspends the current process for a specified number of ticks.

Suspends execution of the calling proces for a specified number of clock ticks.

Definition at line 332 of file kern\_sys\_calls.c.

```
332     {
333     if (ticks <= 0) {
334         P_ERRNO = P_EINVAL;
335         return;
336     }
337
338     // block current process, set state to sleep
339     current_running_pcb->process_state = 'B';
340     current_running_pcb->is_sleeping = true;
341     current_running_pcb->time_to_wake = tick_counter + ticks;
342     log_generic_event('B', current_running_pcb->pid,
343                     current_running_pcb->priority,
344                     current_running_pcb->cmd_str);
345     if (spthread_suspend(current_running_pcb->thread_handle) != 0) { // give scheduler control
346         perror("Error in spthread_suspend in s_sleep call");
347     }
348 }
```

#### 4.11.1.13 s\_spawn()

```
pid_t s_spawn (
    void (*)(void *) func,
    char * argv[],
    int fd0,
    int fd1 )
```

Spawns a child process with the given function and arguments.

Create a child process that executes the function `func`. The child will retain some attributes of the parent.

Definition at line 164 of file kern\_sys\_calls.c.

```
164     {
165     pcb_t* child;
166     if (strcmp(argv[0], "shell") == 0) {
167         child = k_proc_create(current_running_pcb, 0);
168     } else {
169         child = k_proc_create(current_running_pcb, 1);
170     }
171
172     if (child == NULL) {
173         P_ERRNO = P_ENULL;
174         return -1;
175     }
176
177     spthread_t thread_handle;
178
179     if (spthread_create(&thread_handle, NULL, func, argv) != 0) {
180         perror("Error in spthread_create in s_spawn call");
181     }
182
183     child->cmd_str = strdup(argv[0]);
184     child->thread_handle = thread_handle;
185     child->input_fd = fd0;
186     child->output_fd = fd1;
187     child->fd_table[0] = fd0;
188     child->fd_table[1] = fd1;
189
190     log_generic_event('C', child->pid, child->priority, child->cmd_str);
191
192     return child->pid;
193 }
```

## 4.11.1.14 s\_spawn\_init()

```
pid_t s_spawn_init ( )
```

Creates the init process and spawns the shell process.

Similar to s\_spawn except only called when you want to spawn the init process. It will create the init process and also spawn in the shell.

Definition at line 137 of file kern\_sys\_calls.c.

```

137     {
138     pcb_t* init = k_proc_create(NULL, 0);
139     if (init == NULL) {
140         P_ERRNO = P_ENULL;
141         return -1;
142     }
143
144     spthread_t thread_handle;
145     if (spthread_create(&thread_handle, NULL, init_func, NULL) != 0) {
146         perror("Error in spthread_create in s_spawn_init call");
147     }
148
149     init->cmd_str = strdup("init");
150     init->thread_handle = thread_handle;
151     return init->pid;
152 }
```

## 4.11.1.15 s\_waitpid()

```

pid_t s_waitpid (
    pid_t pid,
    int * wstatus,
    bool nohang )
```

Waits for a child of the calling process.

Wait on a child of the calling process, until it changes state. If nohang is true, this will not block the calling process and return immediately.

Definition at line 198 of file kern\_sys\_calls.c.

```

198     {
199     pcb_t* parent = current_running_pcb;
200     if (parent == NULL) {
201         return -1;
202     }
203
204     // if no children, return -1
205     bool has_child = false;
206     for (int i = 0; i < vec_len(&current_pcb); i++) {
207         pcb_t* child = vec_get(&current_pcb, i);
208         if (child->par_pid == parent->pid) {
209             has_child = true;
210             break;
211         }
212     }
213     if (!has_child) {
214         return -1;
215     }
216
217     // Scan the zombie queue first for terminated children.
218     for (int i = 0; i < vec_len(&zombie_queue); i++) {
219         pcb_t* child = vec_get(&zombie_queue, i);
220         if ((pid == -1 || child->pid == pid) && child->par_pid == parent->pid) {
221             if (wstatus != NULL) {
222                 *wstatus = child->process_status;
223             }
224             log_generic_event('W', child->pid, child->priority, child->cmd_str);
225             vec_erase_no_deleter(&zombie_queue, i);
226             delete_from_explicit_queue(&parent->child_pcb, child->pid);
227             k_proc_cleanup(child);

```

```

228     return child->pid;
229 }
230 }
231
232 // If nohang is true, return immediately if no child has exited
233 if (nohang) {
234     return 0;
235 }
236
237 // Block the parent until a child exits
238 delete_from_queue(parent->priority, parent->pid);
239 parent->process_state = 'B';
240 log_generic_event('B', parent->pid, parent->priority, parent->cmd_str);
241
242 while (true) {
243     // Scan the zombie queue first for terminated children.
244     for (int i = 0; i < vec_len(&zombie_queue); i++) {
245         pcb_t* child = vec_get(&zombie_queue, i);
246         if ((pid == -1 || child->pid == pid) && child->par_pid == parent->pid) {
247             if (wstatus != NULL) {
248                 *wstatus = child->process_status;
249             }
250             log_generic_event('W', child->pid, child->priority, child->cmd_str);
251             vec_erase_no_deletor(&zombie_queue, i);
252             delete_from_explicit_queue(&parent->child_pcb, child->pid);
253             k_proc_cleanup(child);
254             return child->pid;
255         }
256     }
257
258     // scan children of current running process for non-terminated state changes
259     for (int i = 0; i < vec_len(&parent->child_pcb); i++) {
260         pcb_t* child = vec_get(&parent->child_pcb, i);
261         if ((pid == -1 || child->pid == pid) && (child->process_status == 21 || child->process_status ==
262 23)) { // signaled
263             if (wstatus != NULL) {
264                 *wstatus = child->process_status;
265             }
266             log_generic_event('W', child->pid, child->priority, child->cmd_str);
267             child->process_status = 0; // reset status
268             return child->pid;
269         }
270     }
271
272     // If we get here, something went wrong
273     return -1;
274 }

```

## 4.11.2 Variable Documentation

### 4.11.2.1 current\_fg\_pid

pid\_t current\_fg\_pid = 2

Definition at line 31 of file kern\_sys\_calls.c.

### 4.11.2.2 current\_pcb

Vec current\_pcb [extern]

Definition at line 30 of file scheduler.c.



#### 4.11.2.3 current\_running\_pcb

`pcb_t* current_running_pcb [extern]`

Definition at line 38 of file scheduler.c.

#### 4.11.2.4 one\_priority\_queue

`Vec one_priority_queue [extern]`

Definition at line 25 of file scheduler.c.

#### 4.11.2.5 sleep\_blocked\_queue

`Vec sleep_blocked_queue [extern]`

Definition at line 28 of file scheduler.c.

#### 4.11.2.6 tick\_counter

`int tick_counter [extern]`

Definition at line 35 of file scheduler.c.

#### 4.11.2.7 two\_priority\_queue

`Vec two_priority_queue [extern]`

Definition at line 26 of file scheduler.c.

#### 4.11.2.8 zero\_priority\_queue

`Vec zero_priority_queue [extern]`

Definition at line 24 of file scheduler.c.

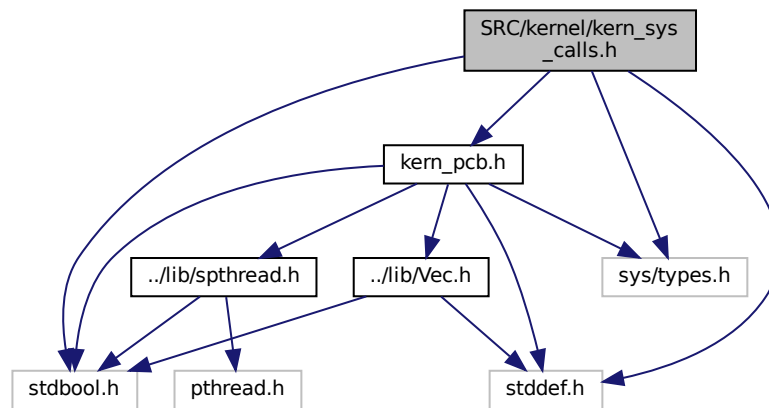
#### 4.11.2.9 zombie\_queue

`Vec zombie_queue [extern]`


Definition at line 27 of file scheduler.c.

## 4.12 SRC/kernel/kern\_sys\_calls.h File Reference

```
#include <stdbool.h>
#include <stddef.h>
#include <sys/types.h>
#include "kern_pcb.h"
Include dependency graph for kern_sys_calls.h:
```



This graph shows which files directly or indirectly include this file:



## Functions

- `int determine_index_in_queue (Vec *queue, int pid)`  
 Given a thread `pid` and `Vec* queue`, this helper function determines the vector index of the thread/`pid` in the queue. If the thread/`pid` is not found, it returns -1.
- `void move_pcb_correct_queue (int prev_priority, int new_priority, pcb_t *curr_pcb)`  
 Given a thread's previous priority, this helper checks if the thread is present in that priority's queue, removes it from that queue if so, and then puts it into the new priority level's queue.

- void `delete_from_queue` (int queue\_id, int pid)  
*Deletes the PCB with the specified PID from one of the priority queues, selected by the provided queue\_id (0, 1, or 2).*
- void `delete_from_explicit_queue` (Vec \*queue\_to\_delete\_from, int pid)  
*Helper function that deletes the given PCB from the explicit queue passed in. Notably, it does not free the PCB but instead uses vec\_erase\_no\_deletor to remove it from the queue.*
- void \* `init_func` (void \*input)  
*The init process function. It spawns the shell process and reaps zombie children.*
- pid\_t `s_spawn_init` ()  
*Similar to s\_spawn except only called when you want to spawn the init process. It will create the init process and also spawn in the shell.*
- void `s_cleanup_init_process` ()  
*Wrapper system-level function to be called in pennos's main method to clean up the init process.*
- pid\_t `s_spawn` (void \*(\*func)(void \*), char \*argv[], int fd0, int fd1)  
*Create a child process that executes the function func. The child will retain some attributes of the parent.*
- pid\_t `s_waitpid` (pid\_t pid, int \*wstatus, bool nohang)  
*Wait on a child of the calling process, until it changes state. If nohang is true, this will not block the calling process and return immediately.*
- int `s_kill` (pid\_t pid, int signal)  
*Send a signal to a particular process.*
- void `s_exit` (void)  
*Unconditionally exit the calling process.*
- int `s_nice` (pid\_t pid, int priority)  
*Set the priority of the specified thread.*
- void `s_sleep` (unsigned int ticks)  
*Suspends execution of the calling proces for a specified number of clock ticks.*
- void \* `s_echo` (void \*arg)  
*System-level wrapper for the shell built-in command "echo".*
- void \* `s_ps` (void \*arg)  
*System-level wrapper for the shell built-in command "ps".*

## 4.12.1 Function Documentation

### 4.12.1.1 delete\_from\_explicit\_queue()

```
void delete_from_explicit_queue (
    Vec * queue_to_delete_from,
    int pid )
```

Helper function that deletes the given PCB from the explicit queue passed in. Notably, it does not free the PCB but instead uses vec\_erase\_no\_deletor to remove it from the queue.

#### Parameters

<i>queue_to_delete_from</i>	ptr to Vec* queue to delete from
<i>pid</i>	the pid of the PCB to delete

Helper function that deletes the given PCB from the explicit queue passed in. Notably, it does not free the PCB but instead uses `vec_erase_no_deletor` to remove it from the queue.

Definition at line 107 of file `kern_sys_calls.c`.

```

107
108     int index = determine_index_in_queue(queue_to_delete_from, pid);
109     if (index != -1) {
110         vec_erase_no_deletor(queue_to_delete_from, index);
111     }
112 }
```

#### 4.12.1.2 delete\_from\_queue()

```

void delete_from_queue (
    int queue_id,
    int pid )
```

Deletes the PCB with the specified PID from one of the priority queues, selected by the provided `queue_id` (0, 1, or 2).

##### Parameters

<i>queue_id</i>	An integer representing the queue: 0 for <code>zero_priority_queue</code> , 1 for <code>one_priority_queue</code> , or 2 for <code>two_priority_queue</code> .
<i>pid</i>	The PID of the PCB to be removed.

Deletes the PCB with the specified PID from one of the priority queues, selected by the provided `queue_id` (0, 1, or 2).

Definition at line 88 of file `kern_sys_calls.c`.

```

88
89     Vec* queue = NULL;
90     if (queue_id == 0) {
91         queue = &zero_priority_queue;
92     } else if (queue_id == 1) {
93         queue = &one_priority_queue;
94     } else {
95         queue = &two_priority_queue;
96     }
97
98     int index = determine_index_in_queue(queue, pid);
99     if (index != -1) {
100         vec_erase_no_deletor(queue, index);
101     }
102 }
```

#### 4.12.1.3 determine\_index\_in\_queue()

```

int determine_index_in_queue (
    Vec * queue,
    int pid )
```

Given a thread `pid` and `Vec* queue`, this helper function determines the vector index of the thread/`pid` in the queue. If the thread/`pid` is not found, it returns -1.

## Parameters

<i>queue</i>	pointer to the vector queue that may contain the thread/pid
<i>pid</i>	the thread's pid

## Returns

the index of the thread/pid in the queue, or -1 if not found

Given a thread pid and Vec\* queue, this helper function determines the vector index of the thread/pid in the queue. If the thread/pid is not found, it returns -1.

Definition at line 40 of file kern\_sys\_calls.c.

```

40                                     {
41     for (int i = 0; i < vec_len(queue); i++) {
42         pcb_t* curr_pcb = vec_get(queue, i);
43         if (curr_pcb->pid == pid) {
44             return i;
45         }
46     }
47
48     return -1; // not found
49 }
```

## 4.12.1.4 init\_func()

```

void* init_func (
    void * input )
```

The init process function. It spawns the shell process and reaps zombie children.

## Parameters

<i>input</i>	unused but needed for typing reasons
--------------	--------------------------------------

## Returns

irrelevant return value because never supposed to return

The init process function. It spawns the shell process and reaps zombie children.

Definition at line 117 of file kern\_sys\_calls.c.

```

117                                     {
118     char* shell_argv[] = {"shell", NULL};
119     s_spawn(shell, shell_argv, STDIN_FILENO, STDOUT_FILENO);
120
121     // continuously wait for and reap zombie children
122     while (true) {
123         int status;
124         s_waitpid(-1, &status, false);
125     }
126
127     return NULL; // should never reach
128 }
```

#### 4.12.1.5 move\_pcb\_correct\_queue()

```
void move_pcb_correct_queue (
    int prev_priority,
    int new_priority,
    pcb_t * curr_pcb )
```

Given a thread's previous priority, this helper checks if the thread is present in that priority's queue, removes it from that queue if so, and then puts it into the new priority level's queue.

##### Parameters

<i>prev_priority</i>	thread's previous priority
<i>new_priority</i>	thread's new priority
<i>curr_pcb</i>	pointer to the thread's PCB

##### Precondition

assumes the *prev\_priority* and *new\_priority* falls in integers [0, 2]

Given a thread's previous priority, this helper checks if the thread is present in that priority's queue, removes it from that queue if so, and then puts it into the new priority level's queue.

Definition at line 54 of file kern\_sys\_calls.c.

```
56                                     {
57     Vec* prev_queue;
58     Vec* new_queue;
59
60     if (prev_priority == 0) {
61         prev_queue = &zero_priority_queue;
62     } else if (prev_priority == 1) {
63         prev_queue = &one_priority_queue;
64     } else {
65         prev_queue = &two_priority_queue;
66     }
67
68     if (new_priority == 0) {
69         new_queue = &zero_priority_queue;
70     } else if (new_priority == 1) {
71         new_queue = &one_priority_queue;
72     } else {
73         new_queue = &two_priority_queue;
74     }
75
76     // delete from prev_queue, if it's present at all
77     int ind = determine_index_in_queue(prev_queue, curr_pcb->pid);
78     if (ind != -1) {
79         vec_erase_no_deletor(prev_queue, ind);
80     }
81
82     vec_push_back(new_queue, curr_pcb);
83 }
```

#### 4.12.1.6 s\_cleanup\_init\_process()

```
void s_cleanup_init_process ( )
```

Wrapper system-level function to be called in pennos's main method to clean up the init process.

Wrapper system-level function to be called in pennos's main method to clean up the init process.

Definition at line 157 of file kern\_sys\_calls.c.

```
157                                     {
158     k_proc_cleanup(get_pcb_in_queue(&current_pcb, 1));
159 }
```

### 4.12.1.7 s\_echo()

```
void* s_echo (
    void * arg )
```

System-level wrapper for the shell built-in command "echo".

#### Parameters

<i>arg</i>	the pass along arguments to the u_echo function
------------	---

#### Returns

NULL, dummy return value

Definition at line 358 of file kern\_sys\_calls.c.

```
358     {
359     char** argv = (char**)arg;
360     if (argv[1] == NULL) { // no args case
361         s_exit();
362         return NULL;
363     }
364
365     int i = 1; // words after "echo"
366     while (argv[i] != NULL) { // while the arg isn't NULL
367         s_write(current_running_pcb->output_fd, argv[i], strlen(argv[i]));
368         s_write(current_running_pcb->output_fd, " ", 1);
369         i++;
370     }
371
372     s_write(current_running_pcb->output_fd, "\n", 1);
373     return NULL;
374 }
```

### 4.12.1.8 s\_exit()

```
void s_exit (
    void )
```

Unconditionally exit the calling process.

Unconditionally exit the calling process.

Definition at line 293 of file kern\_sys\_calls.c.

```
293     {
294     // Set process state to zombie
295     current_running_pcb->process_state = 'Z';
296     current_running_pcb->process_status = 20; // EXITED_NORMALLY
297
298     // Log the exit
299     log_generic_event('E', current_running_pcb->pid,
300                     current_running_pcb->priority,
301                     current_running_pcb->cmd_str);
302
303     delete_from_queue(current_running_pcb->priority, current_running_pcb->pid);
304
305     log_generic_event('Z', current_running_pcb->pid,
306                     current_running_pcb->priority,
307                     current_running_pcb->cmd_str);
308 }
```

#### 4.12.1.9 `s_kill()`

```
int s_kill (
    pid_t pid,
    int signal )
```

Send a signal to a particular process.



## Parameters

<i>pid</i>	Process ID of the target proces.
<i>signal</i>	Signal number to be sent 0 = P_SIGSTOP, 1 = P_SIGCONT, 2 = P_SIGTERM

## Returns

0 on success, -1 on error.

Send a signal to a particular process.

Definition at line 279 of file kern\_sys\_calls.c.

```

279     {
280     pcb_t* pcb_with_pid = get_pcb_in_queue(&current_pcb, pid);
281     if (pcb_with_pid == NULL) {
282         return -1; // pid not found case
283     }
284
285     pcb_with_pid->signals[signal] = true; // signal flagged
286     log_generic_event('S', pid, pcb_with_pid->priority, pcb_with_pid->cmd_str);
287     return 0;
288 }
```

## 4.12.1.10 s\_nice()

```

int s_nice (
    pid_t pid,
    int priority )
```

Set the priority of the specified thread.

## Parameters

<i>pid</i>	Process ID of the target thread.
<i>priority</i>	The new priority value of the thread (0, 1, or 2)

## Returns

0 on success, -1 on failure.

Set the priority of the specified thread.

Definition at line 313 of file kern\_sys\_calls.c.

```

313     {
314     if (priority < 0 || priority > 2) { // error check
315         return -1;
316     }
317
318     pcb_t* curr_pcb = get_pcb_in_queue(&current_pcb, pid);
319     if (curr_pcb != NULL) { // found + exists
320         move_pcb_correct_queue(curr_pcb->priority, priority, curr_pcb);
321         log_nice_event(pid, curr_pcb->priority, priority, curr_pcb->cmd_str);
322         curr_pcb->priority = priority;
323         return 0;
324     }
325
326     return -1; // pid not found
327 }
```

#### 4.12.1.11 s\_ps()

```
void* s_ps (
    void * arg )
```

System-level wrapper for the shell built-in command "ps".

##### Parameters

<i>arg</i>	the pass along arguments to the u_ps function
------------	---

##### Returns

NULL, dummy return value

Definition at line 379 of file kern\_sys\_calls.c.

```
379     {
380     char pid_top[] = "PID\tPPID\tPRI\tSTAT\tCMD\n";
381     s_write(current_running_pcb->output_fd, pid_top, strlen(pid_top));
382     for (int i = 0; i < vec_len(&current_pcbs); i++) {
383         pcb_t* curr_pcb = (pcb_t*)vec_get(&current_pcbs, i);
384         char buffer[100];
385         snprintf(buffer, sizeof(buffer), "%d\t%d\t%d\t%c\t%s\n", curr_pcb->pid,
386             curr_pcb->par_pid, curr_pcb->priority, curr_pcb->process_state,
387             curr_pcb->cmd_str);
388         s_write(current_running_pcb->output_fd, buffer, strlen(buffer));
389     }
390     return NULL;
391 }
```

#### 4.12.1.12 s\_sleep()

```
void s_sleep (
    unsigned int ticks )
```

Suspends execution of the calling proces for a specified number of clock ticks.

This function is analogous to `sleep(3)` in Linux, with the behavior that the system clock continues to tick even if the call is interrupted. The sleep can be interrupted by a `P_SIGTERM` signal, after which the function will return prematurely.

##### Parameters

<i>ticks</i>	Duration of the sleep in system clock ticks. Must be greater than 0.
--------------	--

Suspends execution of the calling proces for a specified number of clock ticks.

Definition at line 332 of file kern\_sys\_calls.c.

```
332     {
333     if (ticks <= 0) {
334         P_ERRNO = P_EINVAL;
335         return;
336     }
337
338     // block current process, set state to sleep
339     current_running_pcb->process_state = 'B';
340     current_running_pcb->is_sleeping = true;
```

```

341     current_running_pcb->time_to_wake = tick_counter + ticks;
342     log_generic_event('B', current_running_pcb->pid,
343                     current_running_pcb->priority,
344                     current_running_pcb->cmd_str);
345     if (spthread_suspend(current_running_pcb->thread_handle) != 0) { // give scheduler control
346         perror("Error in pthread_suspend in s_sleep call");
347     }
348 }

```

#### 4.12.1.13 s\_spawn()

```

pid_t s_spawn (
    void (*)(void *) func,
    char * argv[],
    int fd0,
    int fd1 )

```

Create a child process that executes the function `func`. The child will retain some attributes of the parent.

##### Parameters

<i>func</i>	Function to be executed by the child process.
<i>argv</i>	Null-terminated array of args, including the command name as <code>argv[0]</code> .
<i>fd0</i>	Input file descriptor.
<i>fd1</i>	Output file descriptor.

##### Returns

`pid_t` The process ID of the created child process or -1 on error

Create a child process that executes the function `func`. The child will retain some attributes of the parent.

Definition at line 164 of file `kern_sys_calls.c`.

```

164
165     pcb_t* child;
166     if (strcmp(argv[0], "shell") == 0) {
167         child = k_proc_create(current_running_pcb, 0);
168     } else {
169         child = k_proc_create(current_running_pcb, 1);
170     }
171
172     if (child == NULL) {
173         P_ERRNO = P_ENULL;
174         return -1;
175     }
176
177     pthread_t thread_handle;
178
179     if (pthread_create(&thread_handle, NULL, func, argv) != 0) {
180         perror("Error in pthread_create in s_spawn call");
181     }
182
183     child->cmd_str = strdup(argv[0]);
184     child->thread_handle = thread_handle;
185     child->input_fd = fd0;
186     child->output_fd = fd1;
187     child->fd_table[0] = fd0;
188     child->fd_table[1] = fd1;
189
190     log_generic_event('C', child->pid, child->priority, child->cmd_str);
191
192     return child->pid;
193 }

```

#### 4.12.1.14 s\_spawn\_init()

```
pid_t s_spawn_init ( )
```

Similar to `s_spawn` except only called when you want to spawn the init process. It will create the init process and also spawn in the shell.

##### Returns

the `pid_t` of the created process on success or -1 on error

Similar to `s_spawn` except only called when you want to spawn the init process. It will create the init process and also spawn in the shell.

Definition at line 137 of file `kern_sys_calls.c`.

```
137     {
138     pcb_t* init = k_proc_create(NULL, 0);
139     if (init == NULL) {
140         P_ERRNO = P_ENULL;
141         return -1;
142     }
143
144     spthread_t thread_handle;
145     if (spthread_create(&thread_handle, NULL, init_func, NULL) != 0){
146         perror("Error in spthread_create in s_spawn_init call");
147     }
148
149     init->cmd_str = strdup("init");
150     init->thread_handle = thread_handle;
151     return init->pid;
152 }
```

#### 4.12.1.15 s\_waitpid()

```
pid_t s_waitpid (
    pid_t pid,
    int * wstatus,
    bool nohang )
```

Wait on a child of the calling process, until it changes state. If `nohang` is true, this will not block the calling process and return immediately.

##### Parameters

<i>pid</i>	Process ID of the child to wait for.
<i>wstatus</i>	Pointer to an integer variable where the status will be stored.
<i>nohang</i>	If true, return immediately if no child has exited.

##### Returns

`pid_t` The process ID of the child which has changed state on success, -1 on error.

Wait on a child of the calling process, until it changes state. If `nohang` is true, this will not block the calling process and return immediately.

Definition at line 198 of file `kern_sys_calls.c`.

```


198                                     {
199     pcb_t* parent = current_running_pcb;
200     if (parent == NULL) {
201         return -1;
202     }
203
204     // if no children, return -1
205     bool has_child = false;
206     for (int i = 0; i < vec_len(&current_pcbs); i++) {
207         pcb_t* child = vec_get(&current_pcbs, i);
208         if (child->par_pid == parent->pid) {
209             has_child = true;
210             break;
211         }
212     }
213     if (!has_child) {
214         return -1;
215     }
216
217     // Scan the zombie queue first for terminated children.
218     for (int i = 0; i < vec_len(&zombie_queue); i++) {
219         pcb_t* child = vec_get(&zombie_queue, i);
220         if ((pid == -1 || child->pid == pid) && child->par_pid == parent->pid) {
221             if (wstatus != NULL) {
222                 *wstatus = child->process_status;
223             }
224             log_generic_event('W', child->pid, child->priority, child->cmd_str);
225             vec_erase_no_deletor(&zombie_queue, i);
226             delete_from_explicit_queue(&parent->child_pcbs, child->pid);
227             k_proc_cleanup(child);
228             return child->pid;
229         }
230     }
231
232     // If nohang is true, return immediately if no child has exited
233     if (nohang) {
234         return 0;
235     }
236
237     // Block the parent until a child exits
238     delete_from_queue(parent->priority, parent->pid);
239     parent->process_state = 'B';
240     log_generic_event('B', parent->pid, parent->priority, parent->cmd_str);
241
242     while (true) {
243         // Scan the zombie queue first for terminated children.
244         for (int i = 0; i < vec_len(&zombie_queue); i++) {
245             pcb_t* child = vec_get(&zombie_queue, i);
246             if ((pid == -1 || child->pid == pid) && child->par_pid == parent->pid) {
247                 if (wstatus != NULL) {
248                     *wstatus = child->process_status;
249                 }
250                 log_generic_event('W', child->pid, child->priority, child->cmd_str);
251                 vec_erase_no_deletor(&zombie_queue, i);
252                 delete_from_explicit_queue(&parent->child_pcbs, child->pid);
253                 k_proc_cleanup(child);
254                 return child->pid;
255             }
256         }
257
258         // scan children of current running process for non-terminated state changes
259         for (int i = 0; i < vec_len(&parent->child_pcbs); i++) {
260             pcb_t* child = vec_get(&parent->child_pcbs, i);
261             if ((pid == -1 || child->pid == pid) && (child->process_status == 21 || child->process_status ==
262 23)) { // signaled
263                 if (wstatus != NULL) {
264                     *wstatus = child->process_status;
265                 }
266                 log_generic_event('W', child->pid, child->priority, child->cmd_str);
267                 child->process_status = 0; // reset status
268                 return child->pid;
269             }
270         }
271
272         // If we get here, something went wrong
273         return -1;
274     }

```

## 4.13 SRC/kernel/logger.c File Reference

```
#include "logger.h"
```

```
#include <stdio.h>
#include <unistd.h>
Include dependency graph for logger.c:
```



## Functions

- void [log\\_scheduling\\_event](#) (int pid, int queue\_num, char \*process\_name)  
*Logs a scheduling event i.e. the scheduling of a process for this clock tick.*
- void [log\\_generic\\_event](#) (char event\_type, int pid, int nice\_value, char \*process\_name)  
*Logs a non-nice, non-scheduling event (i.e any event that follows the EVENT PID NICE\_VALUE PROCESS\_NAME format)*
- void [log\\_nice\\_event](#) (int pid, int old\_nice\_value, int new\_nice\_value, char \*process\_name)  
*Logs a nice event, which is the adjusting of a process's nice value.*

### 4.13.1 Function Documentation

#### 4.13.1.1 log\_generic\_event()

```
void log_generic_event (
    char event_type,
    int pid,
    int nice_value,
    char * process_name )
```

Logs a non-nice, non-scheduling event (i.e any event that follows the EVENT PID NICE\_VALUE PROCESS\_NAME format)

#### Parameters

<i>event_type</i>	the type of event, defined by: 'C' = CREATE, 'S' = SIGNALLED, 'E' = EXITED, 'Z' = ZOMBIE, 'O' = ORPHAN, 'W' = WAITED 'B' = BLOCKED, 'U' = UNBLOCKED 's' = STOPPED, 'c' = CONTINUED (notably lower-cased)
<i>pid</i>	process pid
<i>nice_value</i>	process nice value
<i>process_name</i>	string containing process name

**Precondition**

assumes event\_type matches one of the above characters

**Postcondition**

will perror if the write fails

Definition at line 14 of file logger.c.

```

14                                     {
15     char* operation;
16
17     switch(event_type) {
18     case 'C':
19         operation = "CREATE";
20         break;
21     case 'S':
22         operation = "SIGNALLED";
23         break;
24     case 'E':
25         operation = "EXITED";
26         break;
27     case 'Z':
28         operation = "ZOMBIE";
29         break;
30     case 'O':
31         operation = "ORPHAN";
32         break;
33     case 'W':
34         operation = "WAITED";
35         break;
36     case 'B':
37         operation = "BLOCKED";
38         break;
39     case 'U':
40         operation = "UNBLOCKED";
41         break;
42     case 's':
43         operation = "STOPPED";
44         break;
45     default:
46         operation = "CONTINUED";
47         break;
48     }
49
50     char buffer[200];
51     int str_len = snprintf(buffer, sizeof(buffer), "[%d]\t%s\t%d\t%d\t%s\n", tick_counter, operation,
pid, nice_value, process_name);
52     if (write(log_fd, buffer, str_len) == -1) {
53         perror("error in writing to the log file for generic event");
54     }
55 }
```

**4.13.1.2 log\_nice\_event()**

```

void log_nice_event (
    int pid,
    int old_nice_value,
    int new_nice_value,
    char * process_name )
```

Logs a nice event, which is the adjusting of a process's nice value.

**Parameters**

<i>pid</i>	process pid
<i>old_nice_value</i>	old nice value
<i>new_nice_value</i>	new nice value
<i>process_name</i>	string containing process name

**Postcondition**

will perror if the write fails

Definition at line 57 of file logger.c.

```

57                                     {
58     char buffer[200];
59     int str_len = snprintf(buffer, sizeof(buffer), "[%d]\tNICE\t%d\t%d\t%d\t%s\n", tick_counter, pid,
        old_nice_value, new_nice_value, process_name);
60     if (write(log_fd, buffer, str_len) == -1) {
61         perror("error in writing to the log file for nice event");
62     }
63 }
```

**4.13.1.3 log\_scheduling\_event()**

```

void log_scheduling_event (
    int pid,
    int queue_num,
    char * process_name )
```

Logs a scheduling event i.e. the scheduling of a process for this clock tick.

**Parameters**

<i>pid</i>	pid of the process being scheduled
<i>queue_num</i>	the priority queue num of the process
<i>process_name</i>	string containing scheduled process's name

**Postcondition**

will perror if the write fails


Definition at line 6 of file logger.c.

```

6                                     {
7     char buffer[200];
8     int str_len = snprintf(buffer, sizeof(buffer), "[%d]\tSCHEDULE\t%d\t%d\t%s\n", tick_counter, pid,
        queue_num, process_name);
9     if (write(log_fd, buffer, str_len) == -1) {
10         perror("error in writing to the log file for scheduling event");
11     }
12 }
```

**4.14 SRC/kernel/logger.h File Reference**

This graph shows which files directly or indirectly include this file:





## Functions

- void `log_scheduling_event` (int pid, int queue\_num, char \*process\_name)  
*Logs a scheduling event i.e. the scheduling of a process for this clock tick.*
- void `log_generic_event` (char event\_type, int pid, int nice\_value, char \*process\_name)  
*Logs a non-nice, non-scheduling event (i.e any event that follows the EVENT PID NICE\_VALUE PROCESS\_NAME format)*
- void `log_nice_event` (int pid, int old\_nice\_value, int new\_nice\_value, char \*process\_name)  
*Logs a nice event, which is the adjusting of a process's nice value.*

## Variables

- int `tick_counter`
- int `log_fd`

### 4.14.1 Function Documentation

#### 4.14.1.1 `log_generic_event()`

```
void log_generic_event (
    char event_type,
    int pid,
    int nice_value,
    char * process_name )
```

Logs a non-nice, non-scheduling event (i.e any event that follows the EVENT PID NICE\_VALUE PROCESS\_NAME format)

##### Parameters

<i>event_type</i>	the type of event, defined by: 'C' = CREATE, 'S' = SIGNALLED, 'E' = EXITED, 'Z' = ZOMBIE, 'O' = ORPHAN, 'W' = WAITED 'B' = BLOCKED, 'U' = UNBLOCKED 's' = STOPPED, 'c' = CONTINUED (notably lower-cased)
<i>pid</i>	process pid
<i>nice_value</i>	process nice value
<i>process_name</i>	string containing process name

##### Precondition

assumes event\_type matches one of the above characters

##### Postcondition

will perror if the write fails

Definition at line 14 of file logger.c.

```

14                                     {
15     char* operation;
16
17     switch(event_type) {
18         case 'C':
19             operation = "CREATE";
20             break;
21         case 'S':
22             operation = "SIGNALLED";
23             break;
24         case 'E':
25             operation = "EXITED";
26             break;
27         case 'Z':
28             operation = "ZOMBIE";
29             break;
30         case 'O':
31             operation = "ORPHAN";
32             break;
33         case 'W':
34             operation = "WAITED";
35             break;
36         case 'B':
37             operation = "BLOCKED";
38             break;
39         case 'U':
40             operation = "UNBLOCKED";
41             break;
42         case 's':
43             operation = "STOPPED";
44             break;
45         default:
46             operation = "CONTINUED";
47             break;
48     }
49
50     char buffer[200];
51     int str_len = snprintf(buffer, sizeof(buffer), "[%d]\t%s\t%d\t%d\t%s\n", tick_counter, operation,
52                             pid, nice_value, process_name);
53     if (write(log_fd, buffer, str_len) == -1) {
54         perror("error in writing to the log file for generic event");
55     }

```

#### 4.14.1.2 log\_nice\_event()

```

void log_nice_event (
    int pid,
    int old_nice_value,
    int new_nice_value,
    char * process_name )

```

Logs a nice event, which is the adjusting of a process's nice value.

##### Parameters

<i>pid</i>	process pid
<i>old_nice_value</i>	old nice value
<i>new_nice_value</i>	new nice value
<i>process_name</i>	string containing process name

##### Postcondition

will perror if the write fails

Definition at line 57 of file logger.c.

```

57                                     {
58     char buffer[200];
59     int str_len = snprintf(buffer, sizeof(buffer), "[%d]\tNICE\t%d\t%d\t%d\t%s\n", tick_counter, pid,
        old_nice_value, new_nice_value, process_name);
60     if (write(log_fd, buffer, str_len) == -1) {
61         perror("error in writing to the log file for nice event");
62     }
63 }

```

#### 4.14.1.3 log\_scheduling\_event()

```

void log_scheduling_event (
    int pid,
    int queue_num,
    char * process_name )

```

Logs a scheduling event i.e. the scheduling of a process for this clock tick.

##### Parameters

<i>pid</i>	pid of the process being scheduled
<i>queue_num</i>	the priority queue num of the process
<i>process_name</i>	string containing scheduled process's name

##### Postcondition

will perror if the write fails

Definition at line 6 of file logger.c.

```

6                                     {
7     char buffer[200];
8     int str_len = snprintf(buffer, sizeof(buffer), "[%d]\tSCHEDULE\t%d\t%d\t%s\n", tick_counter, pid,
        queue_num, process_name);
9     if (write(log_fd, buffer, str_len) == -1) {
10         perror("error in writing to the log file for scheduling event");
11     }
12 }

```

### 4.14.2 Variable Documentation

#### 4.14.2.1 log\_fd

```
int log_fd [extern]
```

Definition at line 36 of file scheduler.c.

#### 4.14.2.2 tick\_counter

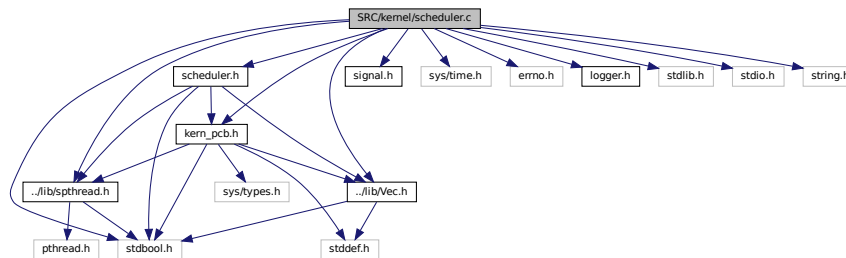
```
int tick_counter [extern]
```

Definition at line 35 of file scheduler.c.

### 4.15 SRC/kernel/scheduler.c File Reference

```
#include "scheduler.h"
#include <signal.h>
#include <stdbool.h>
#include <sys/time.h>
#include "../lib/Vec.h"
#include "../lib/spthread.h"
#include "errno.h"
#include "kern_pcb.h"
#include "logger.h"
#include "stdlib.h"
#include <stdio.h>
#include <string.h>
```

Include dependency graph for scheduler.c:



### Functions

- void [initialize\\_scheduler\\_queues](#) ()  
*Initializes the scheduler queues.*
- void [free\\_scheduler\\_queues](#) ()  
*Frees the scheduler queues.*
- int [generate\\_next\\_priority](#) ()  
*Generates the next priority for scheduling based on the defined probabilities.*
- [pcb\\_t \\*](#) [get\\_next\\_pcb](#) (int priority)  
*Gets the next PCB from the specified priority queue.*
- void [put\\_pcb\\_into\\_correct\\_queue](#) ([pcb\\_t \\*](#)pcb)  
*Puts the given PCB into the correct queue based on its priority and state.*
- void [delete\\_process\\_from\\_particular\\_queue](#) ([pcb\\_t \\*](#)pcb, [Vec \\*](#)queue)  
*Deletes the given PCB from the specified queue.*
- void [delete\\_process\\_from\\_all\\_queues\\_except\\_current](#) ([pcb\\_t \\*](#)pcb)  
*Deletes the given PCB from all queues except the current one.*
- void [delete\\_process\\_from\\_all\\_queues](#) ([pcb\\_t \\*](#)pcb)

- Deletes the given PCB from all queues.*

  - `pcb_t * get_pcb_in_queue (Vec *queue, pid_t pid)`  
*Gets the PCB with the specified PID from the given queue.*
- `bool child_in_zombie_queue (pcb_t *parent)`  
*Checks if the given parent PCB has any children in the zombie queue.*
- `bool child_with_changed_process_status (pcb_t *parent)`  
*Checks if the given parent PCB has any children with a changed process status.*
- `void alarm_handler (int signum)`  
*Signal handler for SIGALRM.*
- `void handle_signal (pcb_t *pcb, int signal)`  
*Handles the specified signal for the given PCB.*
- `void s_shutdown_pennos (void)`  
*Shuts down the scheduler and cleans up resources.*
- `void scheduler ()`  
*The main scheduler function for PennOS.*

## Variables

- `Vec zero_priority_queue`
- `Vec one_priority_queue`
- `Vec two_priority_queue`
- `Vec zombie_queue`
- `Vec sleep_blocked_queue`
- `Vec current_pcb`
- `int tick_counter = 0`
- `int log_fd`
- `pcb_t * current_running_pcb`
- `int curr_priority_arr_index = 0`
- `int det_priorities_arr [19]`

## 4.15.1 Function Documentation

### 4.15.1.1 alarm\_handler()

```
void alarm_handler (
    int signum )
```

Signal handler for SIGALRM.

Handles the alarm signal.

Definition at line 225 of file scheduler.c.

```
225                                     {
226     tick_counter++;
227 }
```

#### 4.15.1.2 child\_in\_zombie\_queue()

```
bool child_in_zombie_queue (
    pcb_t * parent )
```

Checks if the given parent PCB has any children in the zombie queue.

Checks if a child of the given parent process is in the zombie queue.

Definition at line 199 of file scheduler.c.

```
199                                     {
200     for (int i = 0; i < vec_len(&zombie_queue); i++) {
201         pcb_t* child = vec_get(&zombie_queue, i);
202         if (child->par_pid == parent->pid) {
203             return true;
204         }
205     }
206     return false;
207 }
```

#### 4.15.1.3 child\_with\_changed\_process\_status()

```
bool child_with_changed_process_status (
    pcb_t * parent )
```

Checks if the given parent PCB has any children with a changed process status.

Checks if a child of the given parent process has a changed process status.

Definition at line 212 of file scheduler.c.

```
212                                     {
213     for (int i = 0; i < vec_len(&current_pcbs); i++) {
214         pcb_t* child = vec_get(&current_pcbs, i);
215         if (child->par_pid == parent->pid && child->process_status != 0) {
216             return true;
217         }
218     }
219     return false;
220 }
```

#### 4.15.1.4 delete\_process\_from\_all\_queues()

```
void delete_process_from_all_queues (
    pcb_t * pcb )
```

Deletes the given PCB from all queues.

Searches through all of the scheduler's queues and deletes the the given pcb from all of them. Notably, it does not free the pcb via calling vec\_erase\_no\_deletor instead of vec\_erase. If a particular queue does not contain the pcb, nothing occurs.

Definition at line 177 of file scheduler.c.

```
177                                     {
178     delete_process_from_all_queues_except_current(pcb);
179     delete_process_from_particular_queue(pcb, &current_pcbs);
180 }
```

#### 4.15.1.5 delete\_process\_from\_all\_queues\_except\_current()

```
void delete_process_from_all_queues_except_current (
    pcb_t * pcb )
```

Deletes the given PCB from all queues except the current one.

Searches through all of the scheduler's queues except the one containing all of the current processes and deletes the given pcb from all of them. Notably, it does not free the pcb via calling `vec_erase_no_deletor` instead of `vec_erase`. If a particular queue does not contain the pcb, nothing occurs.

Definition at line 166 of file scheduler.c.

```
166 {
167     delete_process_from_particular_queue(pcb, &zero_priority_queue);
168     delete_process_from_particular_queue(pcb, &one_priority_queue);
169     delete_process_from_particular_queue(pcb, &two_priority_queue);
170     delete_process_from_particular_queue(pcb, &zombie_queue);
171     delete_process_from_particular_queue(pcb, &sleep_blocked_queue);
172 }
```

#### 4.15.1.6 delete\_process\_from\_particular\_queue()

```
void delete_process_from_particular_queue (
    pcb_t * pcb,
    Vec * queue )
```

Deletes the given PCB from the specified queue.

Given a queue in the form of a vector, searches through it for the given pcb and deletes it from the queue if found. Notably, it does not free the pcb. Instead, the implementation calls `vec_erase_no_deletor` instead of `vec_erase`. If the pcb isn't in the queue, this function does nothing.

Definition at line 153 of file scheduler.c.

```
153 {
154     for (int i = 0; i < vec_len(queue); i++) {
155         pcb_t* curr_pcb = vec_get(queue, i);
156         if (curr_pcb->pid == pcb->pid) {
157             vec_erase_no_deletor(queue, i);
158             return;
159         }
160     }
161 }
```

#### 4.15.1.7 free\_scheduler\_queues()

```
void free_scheduler_queues ( )
```

Frees the scheduler queues.

Frees the scheduler queues. This function should be called when the scheduler is no longer needed.

Definition at line 66 of file scheduler.c.

```
66 {
67     vec_destroy(&zero_priority_queue);
68     vec_destroy(&one_priority_queue);
69     vec_destroy(&two_priority_queue);
70     vec_destroy(&zombie_queue);
71     vec_destroy(&sleep_blocked_queue);
72     vec_destroy(&current_pcb);
73 }
```

#### 4.15.1.8 generate\_next\_priority()

```
int generate_next_priority ( )
```

Generates the next priority for scheduling based on the defined probabilities.

Deterministically chooses an integer from 0, 1, 2 at the prescribed probabilities. In particular, 0 is output 1.5x more than 1, which is output 1.5x more than 2. Notably, it accounts for cases where some of the queues are empty. If all queues are empty, it'll return -1.

Definition at line 82 of file scheduler.c.

```
82     {
83     // check if all queues are empty
84     if (vec_is_empty(&zero_priority_queue) && vec_is_empty(&one_priority_queue) &&
85         vec_is_empty(&two_priority_queue)) {
86         return -1;
87     }
88
89     int priorities_attempted = 0;
90     while (priorities_attempted < 19) {
91         int curr_pri = det_priorities_arr[curr_priority_arr_index];
92         curr_priority_arr_index = (curr_priority_arr_index + 1) % 19;
93         if (curr_pri == 0 && !vec_is_empty(&zero_priority_queue)) {
94             priorities_attempted++;
95             return 0;
96         } else if (curr_pri == 1 && !vec_is_empty(&one_priority_queue)) {
97             priorities_attempted++;
98             return 1;
99         } else if (curr_pri == 2 && !vec_is_empty(&two_priority_queue)) {
100             priorities_attempted++;
101             return 2;
102         }
103     }
104
105     return -1; // should never reach
106 }
```

#### 4.15.1.9 get\_next\_pcb()

```
pcb_t* get_next_pcb (
    int priority )
```

Gets the next PCB from the specified priority queue.

Returns the next PCB in the queue of the specified priority. or NULL if that queue is empty. Notably, it removes the PCB from the queue.

Definition at line 111 of file scheduler.c.

```
111     {
112     if (priority == -1) { // all queues empty
113         return NULL;
114     }
115
116     pcb_t* next_pcb = NULL;
117     if (priority == 0) {
118         next_pcb = vec_get(&zero_priority_queue, 0);
119         vec_erase_no_deletor(&zero_priority_queue, 0);
120     } else if (priority == 1) {
121         next_pcb = vec_get(&one_priority_queue, 0);
122         vec_erase_no_deletor(&one_priority_queue, 0);
123     } else if (priority == 2) {
124         next_pcb = vec_get(&two_priority_queue, 0);
125         vec_erase_no_deletor(&two_priority_queue, 0);
126     }
127
128     return next_pcb;
129 }
```



## 4.15.1.10 get\_pcb\_in\_queue()

```
pcb_t* get_pcb_in_queue (
    Vec * queue,
    pid_t pid )
```

Gets the PCB with the specified PID from the given queue.

Given a queue, searches for a particular pid inside that queue and, if found, returns the pcb\_t\* associated with that pid.

Definition at line 185 of file scheduler.c.

```
185 {
186     for (int i = 0; i < vec_len(queue); i++) {
187         pcb_t* curr_pcb = vec_get(queue, i);
188         if (curr_pcb->pid == pid) {
189             return curr_pcb;
190         }
191     }
192     return NULL;
193 }
194 }
```

## 4.15.1.11 handle\_signal()

```
void handle_signal (
    pcb_t * pcb,
    int signal )
```

Handles the specified signal for the given PCB.

Handles a signal for a given process.

Definition at line 232 of file scheduler.c.

```
232 {
233     switch (signal) {
234         case 0: // P_SIGSTOP
235             if (pcb->process_state == 'R' || pcb->process_state == 'B') {
236                 pcb->process_state = 'S';
237                 log_generic_event('s', pcb->pid, pcb->priority, pcb->cmd_str);
238                 delete_process_from_all_queues_except_current(pcb);
239                 pcb->process_status = 21; // STOPPED_BY_SIG
240             }
241             pcb->signals[0] = false;
242             break;
243         case 1: // P_SIGCONT
244             if (pcb->process_state == 'S') { // Only continue if stopped
245                 pcb->process_state = 'R';
246                 log_generic_event('c', pcb->pid, pcb->priority, pcb->cmd_str);
247                 delete_process_from_all_queues_except_current(pcb);
248                 put_pcb_into_correct_queue(pcb);
249                 pcb->process_status = 23; // Reset status
250             }
251             pcb->signals[1] = false;
252             break;
253         case 2: // P_SIGTERM
254             if (pcb->process_state != 'Z') { // Don't terminate if already zombie
255                 pcb->process_state = 'Z';
256                 pcb->process_status = 22; // TERM_BY_SIG
257                 log_generic_event('Z', pcb->pid, pcb->priority, pcb->cmd_str);
258                 delete_process_from_all_queues_except_current(pcb);
259                 put_pcb_into_correct_queue(pcb);
260                 pcb->process_status = 22; // TERM_BY_SIG
261             }
262             pcb->signals[2] = false;
263             break;
264     }
265 }
```

#### 4.15.1.12 initialize\_scheduler\_queues()

```
void initialize_scheduler_queues ( )
```

Initializes the scheduler queues.

Initializes the scheduler queues. This function should be called before any other scheduler functions are called.

##### Note

The destructors for the queues are set to NULL to prevent double freeing when exiting PennOS.

Definition at line 54 of file scheduler.c.

```
54 {
55     zero_priority_queue = vec_new(0, NULL);
56     one_priority_queue = vec_new(0, NULL);
57     two_priority_queue = vec_new(0, NULL);
58     zombie_queue = vec_new(0, NULL);
59     sleep_blocked_queue = vec_new(0, NULL);
60     current_pcb = vec_new(0, free_pcb);
61 }
```

#### 4.15.1.13 put\_pcb\_into\_correct\_queue()

```
void put_pcb_into_correct_queue (
    pcb_t * pcb )
```

Puts the given PCB into the correct queue based on its priority and state.

Puts the given pcb struct pointer into its appropriate queue. Notably, it solely uses the pcb's internal fields to determine the correct queue (priority and state).

Definition at line 134 of file scheduler.c.

```
134 {
135     if (pcb->process_state == 'R') {
136         if (pcb->priority == 0) {
137             vec_push_back(&zero_priority_queue, pcb);
138         } else if (pcb->priority == 1) {
139             vec_push_back(&one_priority_queue, pcb);
140         } else if (pcb->priority == 2) {
141             vec_push_back(&two_priority_queue, pcb);
142         }
143     } else if (pcb->process_state == 'Z') {
144         vec_push_back(&zombie_queue, pcb);
145     } else if (pcb->process_state == 'B' || pcb->is_sleeping) {
146         vec_push_back(&sleep_blocked_queue, pcb);
147     }
148 }
```

#### 4.15.1.14 s\_shutdown\_pennos()

```
void s_shutdown_pennos (
    void )
```

Shuts down the scheduler and cleans up resources.

Shuts down the PennOS scheduler.

Definition at line 270 of file scheduler.c.

```
270 {
271     scheduling_done = true;
272 }
```

## 4.15.1.15 scheduler()

```
void scheduler ( )
```

The main scheduler function for PennOS.

This function manages process scheduling, signal handling, and timer-based preemption. It ensures that processes are executed based on their priority and handles signals for both the currently running process and other processes.

Definition at line 277 of file scheduler.c.

```

277     {
278     int curr_priority_queue_num;
279
280     // mask for while scheduler is waiting for alarm
281     sigset_t suspend_set;
282     sigfillset(&suspend_set);
283     sigdelset(&suspend_set, SIGALRM);
284
285     // ensure sigarlm doesn't terminate the process
286     struct sigaction act = (struct sigaction){
287         .sa_handler = alarm_handler,
288         .sa_mask = suspend_set,
289         .sa_flags = SA_RESTART,
290     };
291     sigaction(SIGALRM, &act, NULL);
292
293     // make sure SIGALRM is unblocked
294     sigset_t alarm_set;
295     sigemptyset(&alarm_set);
296     sigaddset(&alarm_set, SIGALRM);
297     pthread_sigmask(SIG_UNBLOCK, &alarm_set, NULL);
298
299     struct itimerval it;
300     it.it_interval = (struct timeval){.tv_usec = hundred_millisec};
301     it.it_value = it.it_interval;
302     setitimer(ITIMER_REAL, &it, NULL);
303
304     while (!scheduling_done) {
305         // handle signals for the currently running process
306         if (current_running_pcb != NULL) {
307             for (int i = 0; i < 3; i++) {
308                 if (current_running_pcb->signals[i]) {
309                     handle_signal(current_running_pcb, i);
310                     // If process was terminated, don't continue scheduling it
311                     if (current_running_pcb->process_state != 'R') {
312                         current_running_pcb = NULL;
313                         break;
314                     }
315                 }
316             }
317         }
318
319         // handle signals for all other processes (currently running or not)
320         for (int i = 0; i < vec_len(&current_pcb); i++) {
321             pcb_t* curr_pcb = vec_get(&current_pcb, i);
322             for (int j = 0; j < 3; j++) {
323                 if (curr_pcb->signals[j]) {
324                     handle_signal(curr_pcb, j);
325                 }
326             }
327         }
328
329         // Check sleep/blocked queue to move processes back to scheduable queues
330         for (int i = 0; i < vec_len(&sleep_blocked_queue); i++) {
331             pcb_t* blocked_proc = vec_get(&sleep_blocked_queue, i);
332             bool make_runnable = false;
333             if (blocked_proc->is_sleeping &&
334                 blocked_proc->time_to_wake == tick_counter) {
335                 blocked_proc->is_sleeping = false;
336                 blocked_proc->time_to_wake = -1;
337                 blocked_proc->signals[2] = false; // Unlikely, but reset signal
338                 make_runnable = true;
339             } else if (blocked_proc->is_sleeping &&
340                 blocked_proc->signals[2]) { // P_SIGTERM received
341                 blocked_proc->is_sleeping = false;
342                 blocked_proc->process_state = 'Z';
343                 blocked_proc->process_status = 22; // TERM_BY_SIG
344                 blocked_proc->signals[2] = false;
345                 delete_process_from_all_queues_except_current(blocked_proc);
346                 put_pcb_into_correct_queue(blocked_proc);
347                 log_generic_event('Z', blocked_proc->pid, blocked_proc->priority,
348                                 blocked_proc->cmd_str);

```

```

349     i--;
350 } else if (child_in_zombie_queue(blocked_proc)) {
351     make_runnable = true;
352 } else if (child_with_changed_process_status(blocked_proc)) {
353     make_runnable = true;
354 }
355
356 if (make_runnable) {
357     blocked_proc->process_state = 'R';
358     vec_erase_no_deletor(&sleep_blocked_queue, i);
359     delete_process_from_all_queues_except_current(blocked_proc);
360     put_pcb_into_correct_queue(blocked_proc);
361     log_generic_event('U', blocked_proc->pid, blocked_proc->priority,
362                     blocked_proc->cmd_str);
363     i--;
364 }
365 }
366
367 curr_priority_queue_num = generate_next_priority();
368
369 current_running_pcb = get_next_pcb(curr_priority_queue_num);
370 if (current_running_pcb == NULL) {
371     sigsuspend(&suspend_set); // idle until signal received
372     continue;
373 }
374
375 log_scheduling_event(current_running_pcb->pid, curr_priority_queue_num,
376                     current_running_pcb->cmd_str);
377
378 if (spthread_continue(current_running_pcb->thread_handle) != 0 &&
379     errno != EINTR) {
380     perror("spthread_continue failed in scheduler");
381 }
382 sigsuspend(&suspend_set);
383 if (spthread_suspend(current_running_pcb->thread_handle) != 0 &&
384     errno != EINTR) {
385     perror("spthread_suspend failed in scheduler");
386 }
387 put_pcb_into_correct_queue(current_running_pcb);
388 }
389 }

```

## 4.15.2 Variable Documentation

### 4.15.2.1 curr\_priority\_arr\_index

```
int curr_priority_arr_index = 0
```

Definition at line 40 of file scheduler.c.

### 4.15.2.2 current\_pcb

```
Vec current_pcb
```

Definition at line 30 of file scheduler.c.

#### 4.15.2.3 current\_running\_pcb

```
pcb_t* current_running_pcb
```

Definition at line 38 of file scheduler.c.

#### 4.15.2.4 det\_priorities\_arr

```
int det_priorities_arr[19]
```

**Initial value:**

```
= {0, 1, 2, 0, 0, 1, 0, 1, 2, 0,  
    0, 1, 2, 0, 1, 0, 0, 1, 2}
```

Definition at line 41 of file scheduler.c.

#### 4.15.2.5 log\_fd

```
int log_fd
```

Definition at line 36 of file scheduler.c.

#### 4.15.2.6 one\_priority\_queue

```
Vec one_priority_queue
```

Definition at line 25 of file scheduler.c.

#### 4.15.2.7 sleep\_blocked\_queue

```
Vec sleep_blocked_queue
```

Definition at line 28 of file scheduler.c.

#### 4.15.2.8 tick\_counter

```
int tick_counter = 0
```

Definition at line 35 of file scheduler.c.

#### 4.15.2.9 two\_priority\_queue

[Vec](#) two\_priority\_queue

Definition at line 26 of file scheduler.c.

#### 4.15.2.10 zero\_priority\_queue

[Vec](#) zero\_priority\_queue

Definition at line 24 of file scheduler.c.

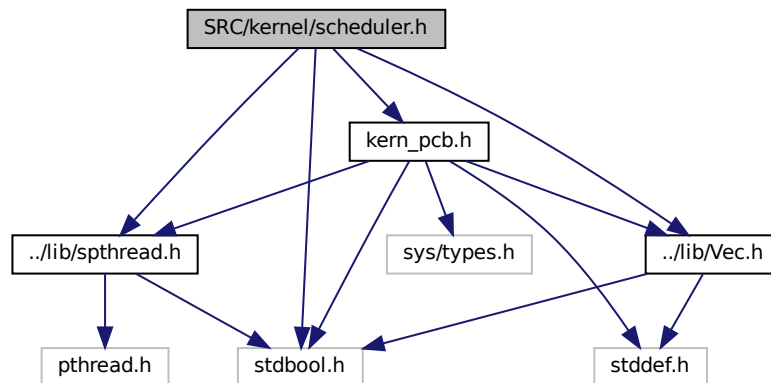
#### 4.15.2.11 zombie\_queue

[Vec](#) zombie\_queue

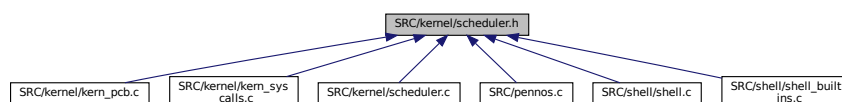
Definition at line 27 of file scheduler.c.

### 4.16 SRC/kernel/scheduler.h File Reference

```
#include <stdbool.h>
#include "../lib/Vec.h"
#include "../lib/spthread.h"
#include "kern_pcb.h"
Include dependency graph for scheduler.h:
```



This graph shows which files directly or indirectly include this file:



## Functions

- void `initialize_scheduler_queues` ()  
*Initializes the scheduler queues. This function should be called before any other scheduler functions are called.*
- void `free_scheduler_queues` ()  
*Frees the scheduler queues. This function should be called when the scheduler is no longer needed.*
- int `generate_next_priority` ()  
*Deterministically chooses an integer from 0, 1, 2 at the prescribed probabilities. In particular, 0 is output 1.5x more than 1, which is output 1.5x more than 2. Notably, it accounts for cases where some of the queues are empty. If all queues are empty, it'll return -1.*
- `pcb_t` \* `get_next_pcb` (int priority)  
*Returns the next PCB in the queue of the specified priority. or NULL if that queue is empty. Notably, it removes the PCB from the queue.*
- void `put_pcb_into_correct_queue` (`pcb_t` \*pcb)  
*Puts the given pcb struct pointer into its appropriate queue. Notably, it solely uses the pcb's internal fields to determine the correct queue (priority and state).*
- void `delete_process_from_particular_queue` (`pcb_t` \*pcb, `Vec` \*queue)  
*Given a queue in the form of a vector, searches through it for the given pcb and deletes it from the queue if found. Notably, it does not free the pcb. Instead, the implementation calls `vec_erase_no_deletor` instead of `vec_erase`. If the pcb isn't in the queue, this function does nothing.*
- void `delete_process_from_all_queues_except_current` (`pcb_t` \*pcb)  
*Searches through all of the scheduler's queues except the one containing all of the current processes and deletes the given pcb from all of them. Notably, it does not free the pcb via calling `vec_erase_no_deletor` instead of `vec_erase`. If a particular queue does not contain the pcb, nothing occurs.*
- void `delete_process_from_all_queues` (`pcb_t` \*pcb)  
*Searches through all of the scheduler's queues and deletes the the given pcb from all of them. Notably, it does not free the pcb via calling `vec_erase_no_deletor` instead of `vec_erase`. If a particular queue does not contain the pcb, nothing occurs.*
- `pcb_t` \* `get_pcb_in_queue` (`Vec` \*queue, `pid_t` pid)  
*Given a queue, searches for a particular pid inside that queue and, if found, returns the `pcb_t`\* associated with that pid.*
- bool `child_in_zombie_queue` (`pcb_t` \*parent)  
*Checks if a child of the given parent process is in the zombie queue.*
- bool `child_with_changed_process_status` (`pcb_t` \*parent)  
*Checks if a child of the given parent process has a changed process status.*
- void `alarm_handler` (int signum)  
*Handles the alarm signal.*
- void `handle_signal` (`pcb_t` \*pcb, int signal)  
*Handles a signal for a given process.*
- void `scheduler` ()  
*The main scheduler function for PennOS.*
- void `s_shutdown_pennos` ()  
*Shuts down the PennOS scheduler.*

### 4.16.1 Function Documentation

#### 4.16.1.1 alarm\_handler()

```
void alarm_handler (
    int signum )
```

Handles the alarm signal.

This function is triggered when the alarm signal is received. It increments the global tick counter, which is used for scheduling and timing purposes.

**Parameters**

<i>signum</i>	The signal number (unused in this implementation).
---------------	--

Handles the alarm signal.

Definition at line 225 of file scheduler.c.

```

225     {
226     tick_counter++;
227 }
```

**4.16.1.2 child\_in\_zombie\_queue()**

```

bool child_in_zombie_queue (
    pcb_t * parent )
```

Checks if a child of the given parent process is in the zombie queue.

This function iterates through the zombie queue to determine if any process in the queue has the given parent process as its parent.

**Parameters**

<i>parent</i>	A pointer to the parent PCB.
---------------	------------------------------

**Returns**

true if a child of the parent is in the zombie queue, false otherwise.

Checks if a child of the given parent process is in the zombie queue.

Definition at line 199 of file scheduler.c.

```

199     {
200     for (int i = 0; i < vec_len(&zombie_queue); i++) {
201         pcb_t* child = vec_get(&zombie_queue, i);
202         if (child->par_pid == parent->pid) {
203             return true;
204         }
205     }
206     return false;
207 }
```

**4.16.1.3 child\_with\_changed\_process\_status()**

```

bool child_with_changed_process_status (
    pcb_t * parent )
```

Checks if a child of the given parent process has a changed process status.

This function iterates through the current PCBs to determine if any child of the given parent process has a non-zero process status, indicating a change.



## Parameters

<i>parent</i>	A pointer to the parent PCB.
---------------	------------------------------

## Returns

true if a child of the parent has a changed process status, false otherwise.

Checks if a child of the given parent process has a changed process status.

Definition at line 212 of file scheduler.c.

```

212                                     {
213     for (int i = 0; i < vec_len(&current_pcb); i++) {
214         pcb_t* child = vec_get(&current_pcb, i);
215         if (child->par_pid == parent->pid && child->process_status != 0) {
216             return true;
217         }
218     }
219     return false;
220 }
```

## 4.16.1.4 delete\_process\_from\_all\_queues()

```

void delete_process_from_all_queues (
    pcb_t * pcb )
```

Searches through all of the scheduler's queues and deletes the the given pcb from all of them. Notably, it does not free the pcb via calling vec\_erase\_no\_deletor instead of vec\_erase. If a particular queue does not contain the pcb, nothing occurs.

## Parameters

<i>pcb</i>	a pointer to the pcb with the pid to delete
------------	---

Searches through all of the scheduler's queues and deletes the the given pcb from all of them. Notably, it does not free the pcb via calling vec\_erase\_no\_deletor instead of vec\_erase. If a particular queue does not contain the pcb, nothing occurs.

Definition at line 177 of file scheduler.c.

```

177                                     {
178     delete_process_from_all_queues_except_current(pcb);
179     delete_process_from_particular_queue(pcb, &current_pcb);
180 }
```

## 4.16.1.5 delete\_process\_from\_all\_queues\_except\_current()

```

void delete_process_from_all_queues_except_current (
    pcb_t * pcb )
```

Searches through all of the scheduler's queues except the one containing all of the current processes and deletes the given pcb from all of them. Notably, it does not free the pcb via calling vec\_erase\_no\_deletor instead of vec\_erase. If a particular queue does not contain the pcb, nothing occurs.

## Parameters

<i>pcb</i>	a pointer to the pcb with the pid to delete
------------	---

Searches through all of the scheduler's queues except the one containing all of the current processes and deletes the given pcb from all of them. Notably, it does not free the pcb via calling `vec_erase_no_deletor` instead of `vec_erase`. If a particular queue does not contain the pcb, nothing occurs.

Definition at line 166 of file `scheduler.c`.

```

166
167     delete_process_from_particular_queue(pcb, &zero_priority_queue);
168     delete_process_from_particular_queue(pcb, &one_priority_queue);
169     delete_process_from_particular_queue(pcb, &two_priority_queue);
170     delete_process_from_particular_queue(pcb, &zombie_queue);
171     delete_process_from_particular_queue(pcb, &sleep_blocked_queue);
172 }
```

#### 4.16.1.6 delete\_process\_from\_particular\_queue()

```

void delete_process_from_particular_queue (
    pcb_t * pcb,
    Vec * queue )
```

Given a queue in the form of a vector, searches through it for the given pcb and deletes it from the queue if found. Notably, it does not free the pcb. Instead, the implementation calls `vec_erase_no_deletor` instead of `vec_erase`. If the pcb isn't in the queue, this function does nothing.

Given a queue in the form of a vector, searches through it for the given pcb and deletes it from the queue if found. Notably, it does not free the pcb. Instead, the implementation calls `vec_erase_no_deletor` instead of `vec_erase`. If the pcb isn't in the queue, this function does nothing.

Definition at line 153 of file `scheduler.c`.

```

153
154     for (int i = 0; i < vec_len(queue); i++) {
155         pcb_t* curr_pcb = vec_get(queue, i);
156         if (curr_pcb->pid == pcb->pid) {
157             vec_erase_no_deletor(queue, i);
158             return;
159         }
160     }
161 }
```

#### 4.16.1.7 free\_scheduler\_queues()

```

void free_scheduler_queues ( )
```

Frees the scheduler queues. This function should be called when the scheduler is no longer needed.

Frees the scheduler queues. This function should be called when the scheduler is no longer needed.

Definition at line 66 of file `scheduler.c`.

```

66     {
67         vec_destroy(&zero_priority_queue);
68         vec_destroy(&one_priority_queue);
69         vec_destroy(&two_priority_queue);
70         vec_destroy(&zombie_queue);
71         vec_destroy(&sleep_blocked_queue);
72         vec_destroy(&current_pcb);
73 }
```

#### 4.16.1.8 generate\_next\_priority()

```
int generate_next_priority ( )
```

Deterministically chooses an integer from 0, 1, 2 at the prescribed probabilities. In particular, 0 is output 1.5x more than 1, which is output 1.5x more than 2. Notably, it accounts for cases where some of the queues are empty. If all queues are empty, it'll return -1.

##### Precondition

assumes that at least one of the scheduler queues is non-empty

##### Returns

int 0, 1, or 2 for priority or -1 to signify that all queues are empty

Deterministically chooses an integer from 0, 1, 2 at the prescribed probabilities. In particular, 0 is output 1.5x more than 1, which is output 1.5x more than 2. Notably, it accounts for cases where some of the queues are empty. If all queues are empty, it'll return -1.

Definition at line 82 of file scheduler.c.

```
82  {
83  // check if all queues are empty
84  if (vec_is_empty(&zero_priority_queue) && vec_is_empty(&one_priority_queue) &&
85      vec_is_empty(&two_priority_queue)) {
86      return -1;
87  }
88
89  int priorities_attempted = 0;
90  while (priorities_attempted < 19) {
91      int curr_pri = det_priorities_arr[curr_priority_arr_index];
92      curr_priority_arr_index = (curr_priority_arr_index + 1) % 19;
93      if (curr_pri == 0 && !vec_is_empty(&zero_priority_queue)) {
94          priorities_attempted++;
95          return 0;
96      } else if (curr_pri == 1 && !vec_is_empty(&one_priority_queue)) {
97          priorities_attempted++;
98          return 1;
99      } else if (curr_pri == 2 && !vec_is_empty(&two_priority_queue)) {
100          priorities_attempted++;
101          return 2;
102      }
103  }
104
105  return -1; // should never reach
106 }
```

#### 4.16.1.9 get\_next\_pcb()

```
pcb_t* get_next_pcb (
    int priority )
```

Returns the next PCB in the queue of the specified priority. or NULL if that queue is empty. Notably, it removes the PCB from the queue.

##### Parameters

<i>priority</i>	queue priority to get next PCB from, or -1 if none
-----------------	--

**Returns**

a ptr to the next pcb struct in queue or NULL if the queue is empty

Returns the next PCB in the queue of the specified priority. or NULL if that queue is empty. Notably, it removes the PCB from the queue.

Definition at line 111 of file scheduler.c.

```

111     {
112     if (priority == -1) { // all queues empty
113         return NULL;
114     }
115
116     pcb_t* next_pcb = NULL;
117     if (priority == 0) {
118         next_pcb = vec_get(&zero_priority_queue, 0);
119         vec_erase_no_deletor(&zero_priority_queue, 0);
120     } else if (priority == 1) {
121         next_pcb = vec_get(&one_priority_queue, 0);
122         vec_erase_no_deletor(&one_priority_queue, 0);
123     } else if (priority == 2) {
124         next_pcb = vec_get(&two_priority_queue, 0);
125         vec_erase_no_deletor(&two_priority_queue, 0);
126     }
127
128     return next_pcb;
129 }
```

**4.16.1.10 get\_pcb\_in\_queue()**

```

pcb_t* get_pcb_in_queue (
    Vec * queue,
    pid_t pid )
```

Given a queue, searches for a particular pid inside that queue and, if found, returns the pcb\_t\* associated with that pid.

**Parameters**

<i>queue</i>	the queue of pcb_t* ptrs to search
<i>pid</i>	the pid to search for

**Returns**

a ptr to the pcb w/ the desired pid if found, NULL otherwise

Given a queue, searches for a particular pid inside that queue and, if found, returns the pcb\_t\* associated with that pid.

Definition at line 185 of file scheduler.c.

```

185     {
186     for (int i = 0; i < vec_len(queue); i++) {
187         pcb_t* curr_pcb = vec_get(queue, i);
188         if (curr_pcb->pid == pid) {
189             return curr_pcb;
190         }
191     }
192
193     return NULL;
194 }
```

## 4.16.1.11 handle\_signal()

```
void handle_signal (
    pcb_t * pcb,
    int signal )
```

Handles a signal for a given process.

This function processes a signal sent to a process and updates its state accordingly. Supported signals include:

- P\_SIGSTOP: Stops the process.
- P\_SIGCONT: Continues a stopped process.
- P\_SIGTERM: Terminates the process.

## Parameters

<i>pcb</i>	A pointer to the PCB of the process receiving the signal.
<i>signal</i>	The signal to handle (0 for P_SIGSTOP, 1 for P_SIGCONT, 2 for P_SIGTERM).

Handles a signal for a given process.

Definition at line 232 of file scheduler.c.

```
232 {
233     switch (signal) {
234         case 0: // P_SIGSTOP
235             if (pcb->process_state == 'R' || pcb->process_state == 'B') {
236                 pcb->process_state = 'S';
237                 log_generic_event('s', pcb->pid, pcb->priority, pcb->cmd_str);
238                 delete_process_from_all_queues_except_current(pcb);
239                 pcb->process_status = 21; // STOPPED_BY_SIG
240             }
241             pcb->signals[0] = false;
242             break;
243         case 1: // P_SIGCONT
244             if (pcb->process_state == 'S') { // Only continue if stopped
245                 pcb->process_state = 'R';
246                 log_generic_event('c', pcb->pid, pcb->priority, pcb->cmd_str);
247                 delete_process_from_all_queues_except_current(pcb);
248                 put_pcb_into_correct_queue(pcb);
249                 pcb->process_status = 23; // Reset status
250             }
251             pcb->signals[1] = false;
252             break;
253         case 2: // P_SIGTERM
254             if (pcb->process_state != 'Z') { // Don't terminate if already zombie
255                 pcb->process_state = 'Z';
256                 pcb->process_status = 22; // TERM_BY_SIG
257                 log_generic_event('Z', pcb->pid, pcb->priority, pcb->cmd_str);
258                 delete_process_from_all_queues_except_current(pcb);
259                 put_pcb_into_correct_queue(pcb);
260                 pcb->process_status = 22; // TERM_BY_SIG
261             }
262             pcb->signals[2] = false;
263             break;
264     }
265 }
```

## 4.16.1.12 initialize\_scheduler\_queues()

```
void initialize_scheduler_queues ( )
```

Initializes the scheduler queues. This function should be called before any other scheduler functions are called.

Initializes the scheduler queues. This function should be called before any other scheduler functions are called.

**Note**

The destructors for the queues are set to NULL to prevent double freeing when exiting PennOS.

Definition at line 54 of file scheduler.c.

```

54      {
55  zero_priority_queue = vec_new(0, NULL);
56  one_priority_queue = vec_new(0, NULL);
57  two_priority_queue = vec_new(0, NULL);
58  zombie_queue = vec_new(0, NULL);
59  sleep_blocked_queue = vec_new(0, NULL);
60  current_pcb = vec_new(0, free_pcb);
61  }
```

**4.16.1.13 put\_pcb\_into\_correct\_queue()**

```

void put_pcb_into_correct_queue (
    pcb_t * pcb )
```

Puts the given pcb struct pointer into its appropriate queue. Notably, it solely uses the pcb's internal fields to determine the correct queue (priority and state).

Puts the given pcb struct pointer into its appropriate queue. Notably, it solely uses the pcb's internal fields to determine the correct queue (priority and state).

Definition at line 134 of file scheduler.c.

```

134      {
135  if (pcb->process_state == 'R') {
136  if (pcb->priority == 0) {
137  vec_push_back(&zero_priority_queue, pcb);
138  } else if (pcb->priority == 1) {
139  vec_push_back(&one_priority_queue, pcb);
140  } else if (pcb->priority == 2) {
141  vec_push_back(&two_priority_queue, pcb);
142  }
143  } else if (pcb->process_state == 'Z') {
144  vec_push_back(&zombie_queue, pcb);
145  } else if (pcb->process_state == 'B' || pcb->is_sleeping) {
146  vec_push_back(&sleep_blocked_queue, pcb);
147  }
148  }
```

**4.16.1.14 s\_shutdown\_pennos()**

```

void s_shutdown_pennos (
    void )
```

Shuts down the PennOS scheduler.

This function sets the scheduling\_done flag to true, signaling the scheduler to terminate its loop and shut down.

Shuts down the PennOS scheduler.

Definition at line 270 of file scheduler.c.

```

270      {
271  scheduling_done = true;
272  }
```

## 4.16.1.15 scheduler()

```
void scheduler ( )
```

The main scheduler function for PennOS.

This function manages process scheduling, signal handling, and timer-based preemption. It ensures that processes are executed based on their priority and handles signals for both the currently running process and other processes.

Definition at line 277 of file scheduler.c.

```

277     {
278     int curr_priority_queue_num;
279
280     // mask for while scheduler is waiting for alarm
281     sigset_t suspend_set;
282     sigfillset(&suspend_set);
283     sigdelset(&suspend_set, SIGALRM);
284
285     // ensure sigarlm doesn't terminate the process
286     struct sigaction act = (struct sigaction){
287         .sa_handler = alarm_handler,
288         .sa_mask = suspend_set,
289         .sa_flags = SA_RESTART,
290     };
291     sigaction(SIGALRM, &act, NULL);
292
293     // make sure SIGALRM is unblocked
294     sigset_t alarm_set;
295     sigemptyset(&alarm_set);
296     sigaddset(&alarm_set, SIGALRM);
297     pthread_sigmask(SIG_UNBLOCK, &alarm_set, NULL);
298
299     struct itimerval it;
300     it.it_interval = (struct timeval){.tv_usec = hundred_millisec};
301     it.it_value = it.it_interval;
302     setitimer(ITIMER_REAL, &it, NULL);
303
304     while (!scheduling_done) {
305         // handle signals for the currently running process
306         if (current_running_pcb != NULL) {
307             for (int i = 0; i < 3; i++) {
308                 if (current_running_pcb->signals[i]) {
309                     handle_signal(current_running_pcb, i);
310                     // If process was terminated, don't continue scheduling it
311                     if (current_running_pcb->process_state != 'R') {
312                         current_running_pcb = NULL;
313                         break;
314                     }
315                 }
316             }
317         }
318
319         // handle signals for all other processes (currently running or not)
320         for (int i = 0; i < vec_len(&current_pcb); i++) {
321             pcb_t* curr_pcb = vec_get(&current_pcb, i);
322             for (int j = 0; j < 3; j++) {
323                 if (curr_pcb->signals[j]) {
324                     handle_signal(curr_pcb, j);
325                 }
326             }
327         }
328
329         // Check sleep/blocked queue to move processes back to scheduable queues
330         for (int i = 0; i < vec_len(&sleep_blocked_queue); i++) {
331             pcb_t* blocked_proc = vec_get(&sleep_blocked_queue, i);
332             bool make_runnable = false;
333             if (blocked_proc->is_sleeping &&
334                 blocked_proc->time_to_wake == tick_counter) {
335                 blocked_proc->is_sleeping = false;
336                 blocked_proc->time_to_wake = -1;
337                 blocked_proc->signals[2] = false; // Unlikely, but reset signal
338                 make_runnable = true;
339             } else if (blocked_proc->is_sleeping &&
340                 blocked_proc->signals[2]) { // P_SIGTERM received
341                 blocked_proc->is_sleeping = false;
342                 blocked_proc->process_state = 'Z';
343                 blocked_proc->process_status = 22; // TERM_BY_SIG
344                 blocked_proc->signals[2] = false;
345                 delete_process_from_all_queues_except_current(blocked_proc);
346                 put_pcb_into_correct_queue(blocked_proc);
347                 log_generic_event('Z', blocked_proc->pid, blocked_proc->priority,
348                                 blocked_proc->cmd_str);

```

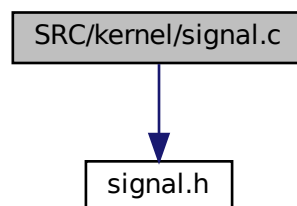
```

349     i--;
350 } else if (child_in_zombie_queue(blocked_proc)) {
351     make_runnable = true;
352 } else if (child_with_changed_process_status(blocked_proc)) {
353     make_runnable = true;
354 }
355
356 if (make_runnable) {
357     blocked_proc->process_state = 'R';
358     vec_erase_no_deletor(&sleep_blocked_queue, i);
359     delete_process_from_all_queues_except_current(blocked_proc);
360     put_pcb_into_correct_queue(blocked_proc);
361     log_generic_event('U', blocked_proc->pid, blocked_proc->priority,
362                     blocked_proc->cmd_str);
363     i--;
364 }
365 }
366
367 curr_priority_queue_num = generate_next_priority();
368
369 current_running_pcb = get_next_pcb(curr_priority_queue_num);
370 if (current_running_pcb == NULL) {
371     sigsuspend(&suspend_set); // idle until signal received
372     continue;
373 }
374
375 log_scheduling_event(current_running_pcb->pid, curr_priority_queue_num,
376                     current_running_pcb->cmd_str);
377
378 if (spthread_continue(current_running_pcb->thread_handle) != 0 &&
379     errno != EINTR) {
380     perror("spthread_continue failed in scheduler");
381 }
382 sigsuspend(&suspend_set);
383 if (spthread_suspend(current_running_pcb->thread_handle) != 0 &&
384     errno != EINTR) {
385     perror("spthread_suspend failed in scheduler");
386 }
387 put_pcb_into_correct_queue(current_running_pcb);
388 }
389 }

```

## 4.17 SRC/kernel/signal.c File Reference


#include <signal.h>  
 Include dependency graph for signal.c:





## 4.18 SRC/kernel/signal.h File Reference

This graph shows which files directly or indirectly include this file:



### Macros

- `#define P_SIGSTOP 0`  
*Signals for PennOS.*
- `#define P_SIGCONT 1`
- `#define P_SIGTERM 2`
- `#define EXITED_NORMALLY 20`  
*Status definitions.*
- `#define STOPPED_BY_SIG 21`
- `#define TERM_BY_SIG 22`
- `#define CONT_BY_SIG 23`
- `#define P_WIFEXITED(status) ((status) == EXITED_NORMALLY)`  
*User-level macros for waitpid status.*
- `#define P_WIFSTOPPED(status) ((status) == STOPPED_BY_SIG)`
- `#define P_WIFSIGNALED(status) ((status) == TERM_BY_SIG)`

### 4.18.1 Macro Definition Documentation

#### 4.18.1.1 CONT\_BY\_SIG

```
#define CONT_BY_SIG 23
```

Definition at line 17 of file signal.h.

#### 4.18.1.2 EXITED\_NORMALLY

```
#define EXITED_NORMALLY 20
```

Status definitions.

Definition at line 14 of file signal.h.

#### 4.18.1.3 P\_SIGCONT

```
#define P_SIGCONT 1
```

Definition at line 8 of file signal.h.

#### 4.18.1.4 P\_SIGSTOP

```
#define P_SIGSTOP 0
```

Signals for PennOS.

Definition at line 7 of file signal.h.

#### 4.18.1.5 P\_SIGTERM

```
#define P_SIGTERM 2
```

Definition at line 9 of file signal.h.

#### 4.18.1.6 P\_WIFEXITED

```
#define P_WIFEXITED(  
    status ) ((status) == EXITED_NORMALLY)
```

User-level macros for waitpid status.

Definition at line 22 of file signal.h.

#### 4.18.1.7 P\_WIFSIGNALED

```
#define P_WIFSIGNALED(  
    status ) ((status) == TERM_BY_SIG)
```

Definition at line 24 of file signal.h.

## 4.18.1.8 P\_WIFSTOPPED

```
#define P_WIFSTOPPED (
    status ) ((status) == STOPPED_BY_SIG)
```

Definition at line 23 of file signal.h.

## 4.18.1.9 STOPPED\_BY\_SIG

```
#define STOPPED_BY_SIG 21
```

Definition at line 15 of file signal.h.

## 4.18.1.10 TERM\_BY\_SIG

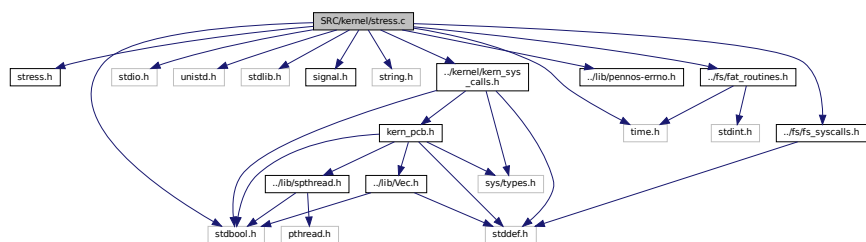
```
#define TERM_BY_SIG 22
```

Definition at line 16 of file signal.h.

## 4.19 SRC/kernel/stress.c File Reference

```
#include "stress.h"
#include <stdbool.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <signal.h>
#include <string.h>
#include <time.h>
#include "../kernel/kern_sys_calls.h"
#include "../lib/pennos-errno.h"
#include "../fs/fs_syscalls.h"
#include "../fs/fat_routines.h"
```

Include dependency graph for stress.c:



## Functions

- void \* [hang](#) (void \*arg)
- void \* [nohang](#) (void \*arg)
- void \* [recur](#) (void \*arg)
- void \* [crash](#) (void \*arg)

### 4.19.1 Function Documentation

#### 4.19.1.1 crash()

```
void* crash (
    void * arg )
```

Definition at line 233 of file stress.c.

```
233     {
234     // This one only works on a file system big enough to hold 5480 bytes
235     crash_main();
236     s_exit();
237     return NULL;
238 }
```

#### 4.19.1.2 hang()

```
void* hang (
    void * arg )
```

Definition at line 215 of file stress.c.

```
215     {
216     spawn(false);
217     s_exit();
218     return NULL;
219 }
```

#### 4.19.1.3 nohang()

```
void* nohang (
    void * arg )
```

Definition at line 221 of file stress.c.

```
221     {
222     spawn(true);
223     s_exit();
224     return NULL;
225 }
```

#### 4.19.1.4 recur()

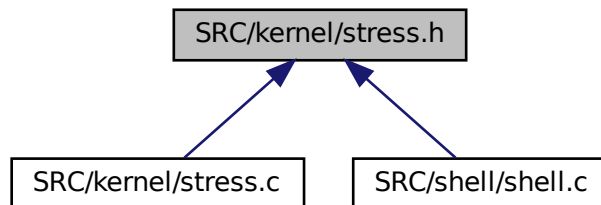
```
void* recur (
    void * arg )
```

Definition at line 227 of file stress.c.

```
227     {
228     spawn_r (NULL);
229     s_exit();
230     return NULL;
231 }
```

## 4.20 SRC/kernel/stress.h File Reference

This graph shows which files directly or indirectly include this file:



### Functions

- void \* [hang](#) (void \*)
- void \* [nohang](#) (void \*)
- void \* [recur](#) (void \*)
- void \* [crash](#) (void \*)

### 4.20.1 Function Documentation

#### 4.20.1.1 crash()

```
void* crash (
    void * arg )
```

Definition at line 233 of file stress.c.

```
233     {
234     // This one only works on a file system big enough to hold 5480 bytes
235     crash_main();
236     s_exit();
237     return NULL;
238 }
```

#### 4.20.1.2 hang()

```
void* hang (
    void * arg )
```

Definition at line 215 of file stress.c.

```
215     {
216     spawn(false);
217     s_exit();
218     return NULL;
219 }
```

#### 4.20.1.3 nohang()

```
void* nohang (
    void * arg )
```

Definition at line 221 of file stress.c.

```
221     {
222     spawn(true);
223     s_exit();
224     return NULL;
225 }
```

#### 4.20.1.4 recur()

```
void* recur (
    void * arg )
```

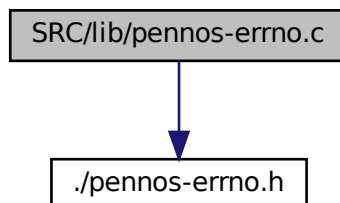
Definition at line 227 of file stress.c.

```
227     {
228     spawn_r(NULL);
229     s_exit();
230     return NULL;
231 }
```

## 4.21 SRC/lib/pennos-errno.c File Reference

```
#include " ./pennos-errno.h"
```

Include dependency graph for pennos-errno.c:





## Variables

- int [P\\_ERRNO](#)

### 4.22.1 Macro Definition Documentation

#### 4.22.1.1 P\_EBADF

```
#define P_EBADF 2
```

Definition at line 8 of file pennos-errno.h.

#### 4.22.1.2 P\_EBUSY

```
#define P_EBUSY 6
```

Definition at line 12 of file pennos-errno.h.

#### 4.22.1.3 P\_ECLOSE

```
#define P_ECLOSE 19
```

Definition at line 25 of file pennos-errno.h.

#### 4.22.1.4 P\_ECOMMAND

```
#define P_ECOMMAND 21
```

Definition at line 27 of file pennos-errno.h.

#### 4.22.1.5 P\_EEXIST

```
#define P_EEXIST 5
```

Definition at line 11 of file pennos-errno.h.



#### 4.22.1.6 P\_EFS\_NOT\_MOUNTED

```
#define P_EFS_NOT_MOUNTED 8
```

Definition at line 14 of file pennos-errno.h.

#### 4.22.1.7 P\_EFULL

```
#define P_EFULL 7
```

Definition at line 13 of file pennos-errno.h.

#### 4.22.1.8 P\_EFUNC

```
#define P_EFUNC 14
```

Definition at line 20 of file pennos-errno.h.

#### 4.22.1.9 P\_EINTR

```
#define P_EINTR 9
```

Definition at line 15 of file pennos-errno.h.

#### 4.22.1.10 P\_EINVAL

```
#define P_EINVAL 4
```

Definition at line 10 of file pennos-errno.h.

#### 4.22.1.11 P\_ELSEEK

```
#define P_ELSEEK 12
```

Definition at line 18 of file pennos-errno.h.

#### 4.22.1.12 P\_EMALLOC

```
#define P_EMALLOC 16
```

Definition at line 22 of file pennos-errno.h.

#### 4.22.1.13 P\_EMAP

```
#define P_EMAP 13
```

Definition at line 19 of file pennos-errno.h.

#### 4.22.1.14 P\_ENOENT

```
#define P_ENOENT 1
```

Definition at line 7 of file pennos-errno.h.

#### 4.22.1.15 P\_ENULL

```
#define P_ENULL 10
```

Definition at line 16 of file pennos-errno.h.

#### 4.22.1.16 P\_EOPEN

```
#define P_EOPEN 15
```

Definition at line 21 of file pennos-errno.h.

#### 4.22.1.17 P\_EPARSE

```
#define P_EPARSE 20
```

Definition at line 26 of file pennos-errno.h.

#### 4.22.1.18 P\_EPERM

```
#define P_EPERM 3
```

Definition at line 9 of file pennos-errno.h.

#### 4.22.1.19 P\_EREAD

```
#define P_EREAD 11
```

Definition at line 17 of file pennos-errno.h.

#### 4.22.1.20 P\_ESIGNAL

```
#define P_ESIGNAL 17
```

Definition at line 23 of file pennos-errno.h.

#### 4.22.1.21 P\_EUNKNOWN

```
#define P_EUNKNOWN 99
```

Definition at line 30 of file pennos-errno.h.

#### 4.22.1.22 P\_EWRITE

```
#define P_EWRITE 18
```

Definition at line 24 of file pennos-errno.h.

#### 4.22.1.23 P\_INITFAIL

```
#define P_INITFAIL 23
```

Definition at line 29 of file pennos-errno.h.

#### 4.22.1.24 P\_NEEDF

```
#define P_NEEDF 22
```

Definition at line 28 of file pennos-errno.h.

### 4.22.2 Variable Documentation

#### 4.22.2.1 P\_ERRNO


```
int P_ERRNO [extern]
```

Definition at line 8 of file pennos-errno.c.

## 4.23 SRC/lib/spthread.c File Reference

```
#include <errno.h>
#include <pthread.h>
#include <signal.h>
#include <stdbool.h>
#include <stdlib.h>
#include <unistd.h>
#include "../spthread.h"
#include <string.h>
#include <stdio.h>
```

Include dependency graph for spthread.c:



## Classes

- struct [spthread\\_fwd\\_args\\_st](#)
- struct [spthread\\_signal\\_args\\_st](#)
- struct [spthread\\_meta\\_st](#)

## Macros

- `#define _GNU_SOURCE`
- `#define MILLISEC_IN_NANO 100000`
- `#define SPTHREAD_RUNNING_STATE 0`
- `#define SPTHREAD_SUSPENDED_STATE 1`
- `#define SPTHREAD_TERMINATED_STATE 2`
- `#define SPTHREAD_SIG_SUSPEND -1`
- `#define SPTHREAD_SIG_CONTINUE -2`

## Typedefs

- `typedef void (* pthread_fn) (void *)`
- `typedef struct pthread_fwd_args_st pthread_fwd_args`
- `typedef struct pthread_signal_args_st pthread_signal_args`
- `typedef struct pthread_meta_st pthread_meta_t`

## Functions

- `int pthread_create (pthread_t *thread, const pthread_attr_t *attr, pthread_fn start_routine, void *arg)`
- `int pthread_suspend (pthread_t thread)`
- `int pthread_suspend_self ()`
- `int pthread_continue (pthread_t thread)`
- `int pthread_cancel (pthread_t thread)`
- `bool pthread_self (pthread_t *thread)`
- `int pthread_join (pthread_t thread, void **retval)`
- `void pthread_exit (void *status)`
- `bool pthread_equal (pthread_t first, pthread_t second)`
- `int pthread_disable_interrupts_self ()`
- `int pthread_enable_interrupts_self ()`

### 4.23.1 Macro Definition Documentation

#### 4.23.1.1 \_GNU\_SOURCE

```
#define _GNU_SOURCE
```

Definition at line 1 of file `spthread.c`.

#### 4.23.1.2 MILLISEC\_IN\_NANO

```
#define MILLISEC_IN_NANO 100000
```

Definition at line 12 of file `spthread.c`.

#### 4.23.1.3 SPTHREAD\_RUNNING\_STATE

```
#define SPTHREAD_RUNNING_STATE 0
```

Definition at line 76 of file spthread.c.

#### 4.23.1.4 SPTHREAD\_SIG\_CONTINUE

```
#define SPTHREAD_SIG_CONTINUE -2
```

Definition at line 85 of file spthread.c.

#### 4.23.1.5 SPTHREAD\_SIG\_SUSPEND

```
#define SPTHREAD_SIG_SUSPEND -1
```

Definition at line 84 of file spthread.c.

#### 4.23.1.6 SPTHREAD\_SUSPENDED\_STATE

```
#define SPTHREAD_SUSPENDED_STATE 1
```

Definition at line 77 of file spthread.c.

#### 4.23.1.7 SPTHREAD\_TERMINATED\_STATE

```
#define SPTHREAD_TERMINATED_STATE 2
```

Definition at line 78 of file spthread.c.

### 4.23.2 Typedef Documentation

#### 4.23.2.1 pthread\_fn

```
typedef void*(* pthread_fn) (void *)
```

Definition at line 20 of file spthread.c.

#### 4.23.2.2 `spthread_fwd_args`

```
typedef struct spthread_fwd_args_st spthread_fwd_args
```

#### 4.23.2.3 `spthread_meta_t`

```
typedef struct spthread_meta_st spthread_meta_t
```

#### 4.23.2.4 `spthread_signal_args`

```
typedef struct spthread_signal_args_st spthread_signal_args
```

### 4.23.3 Function Documentation

#### 4.23.3.1 `spthread_cancel()`

```
int spthread_cancel (  
    spthread_t thread )
```

Definition at line 293 of file `spthread.c`.

```
293  
294     return pthread_cancel(thread.thread);  
295 }
```

#### 4.23.3.2 `spthread_continue()`

```
int spthread_continue (  
    spthread_t thread )
```

Definition at line 241 of file `spthread.c`.

```
241  
242     pthread_t pself = pthread_self();  
243  
244     if (pthread_equal(pself, thread.thread) != 0) {  
245         // I am already running... so just return 0  
246         my_meta->state = SPTHREAD_RUNNING_STATE;  
247         return 0;  
248     }  
249  
250     spthread_signal_args args = (spthread_signal_args){  
251         .signal = SPTHREAD_SIG_CONTINUE,  
252         .ack = 0,  
253     };  
254     pthread_mutex_init(&args.shutup_mutex, NULL);  
255  
256     int ret = pthread_sigqueue(thread.thread, SIGPTH,  
257                             (union sigval){  
258                                 .sival_ptr = &args,
```

```

259             });
260     if (ret != 0) {
261         pthread_mutex_destroy(&args.shutup_mutex);
262         // handles the case where the thread is already dead.
263         return ret;
264     }
265
266     // wait for our signal to be ack'd
267
268     // setting up args to nanosleep
269     const struct timespec t = (struct timespec){
270         .tv_nsec = MILLISEC_IN_NANO,
271     };
272
273     pthread_mutex_lock(&args.shutup_mutex);
274     while (args.ack != 1) {
275         // wait for a mili second
276         pthread_mutex_unlock(&args.shutup_mutex);
277
278         nanosleep(&t, NULL);
279
280         // fprintf(stderr, "susp checking...\n");
281         pthread_mutex_lock(&args.shutup_mutex);
282
283         if (thread.meta->state == SPTHREAD_TERMINATED_STATE) {
284             // child called exit, can break
285             break;
286         }
287     }
288     pthread_mutex_unlock(&args.shutup_mutex);
289     pthread_mutex_destroy(&args.shutup_mutex);
290     return ret;
291 }

```

#### 4.23.3.3 pthread\_create()

```

int pthread_create (
    pthread_t * thread,
    const pthread_attr_t * attr,
    pthread_fn start_routine,
    void * arg )

```

Definition at line 114 of file pthread.c.

```

117     {
118         pthread_meta_t* child_meta = malloc(sizeof(pthread_meta_t));
119         if (child_meta == NULL) {
120             return EAGAIN;
121         }
122
123         pthread_fwd_args* fwd_args = malloc(sizeof(pthread_fwd_args));
124         if (fwd_args == NULL) {
125             free(child_meta);
126             return EAGAIN;
127         }
128         *fwd_args = (pthread_fwd_args){
129             .actual_routine = start_routine,
130             .actual_arg = arg,
131             .setup_done = false,
132             .child_meta = child_meta,
133         };
134
135         int ret = pthread_mutex_init(&(fwd_args->setup_mutex), NULL);
136         if (ret != 0) {
137             free(child_meta);
138             free(fwd_args);
139             return EAGAIN;
140         }
141
142         ret = pthread_cond_init(&(fwd_args->setup_cond), NULL);
143         if (ret != 0) {
144             free(child_meta);
145             pthread_mutex_destroy(&(fwd_args->setup_mutex));
146             free(fwd_args);
147             return EAGAIN;
148         }
149
150         pthread_t pthread;

```



```

151  int result = pthread_create(&pthread, attr, spthread_start, fwd_args);
152
153  pthread_mutex_lock(&(fwd_args->setup_mutex));
154  while (fwd_args->setup_done == false) {
155      pthread_cond_wait(&(fwd_args->setup_cond), &(fwd_args->setup_mutex));
156  }
157  pthread_mutex_unlock(&(fwd_args->setup_mutex));
158
159  pthread_cond_destroy(&(fwd_args->setup_cond));
160  pthread_mutex_destroy(&(fwd_args->setup_mutex));
161  free(fwd_args);
162
163  *thread = (spthread_t){
164      .thread = pthread,
165      .meta = child_meta,
166  };
167
168  return result;
169 }

```

#### 4.23.3.4 spthread\_disable\_interrupts\_self()

```
int spthread_disable_interrupts_self ( )
```

Definition at line 326 of file spthread.c.

```

326  {
327      sigset_t block_set;
328      int res = sigemptyset(&block_set);
329      if (res != 0) {
330          return res;
331      }
332      res = sigaddset(&block_set, SIGPTHD);
333      if (res != 0) {
334          return res;
335      }
336      res = pthread_sigmask(SIG_BLOCK, &block_set, NULL);
337      if (res != 0) {
338          return res;
339      }
340      return 0;
341 }

```

#### 4.23.3.5 spthread\_enable\_interrupts\_self()

```
int spthread_enable_interrupts_self ( )
```

Definition at line 345 of file spthread.c.

```

345  {
346      sigset_t block_set;
347      int res = sigemptyset(&block_set);
348      if (res != 0) {
349          return res;
350      }
351      res = sigaddset(&block_set, SIGPTHD);
352      if (res != 0) {
353          return res;
354      }
355      res = pthread_sigmask(SIG_UNBLOCK, &block_set, NULL);
356      if (res != 0) {
357          return res;
358      }
359      return 0;
360 }

```

#### 4.23.3.6 pthread\_equal()

```
bool pthread_equal (
    pthread_t first,
    pthread_t second )
```

Definition at line 322 of file pthread.c.

```
322                                     {
323     return pthread_equal(first.thread, second.thread) && (first.meta == second.meta);
324 }
```

#### 4.23.3.7 pthread\_exit()

```
void pthread_exit (
    void * status )
```

Definition at line 315 of file pthread.c.

```
315                                     {
316     // necessary cleanup is registered
317     // in a cleanup routine
318     // that is pushed at start of an pthread
319     pthread_exit(status);
320 }
```

#### 4.23.3.8 pthread\_join()

```
int pthread_join (
    pthread_t thread,
    void ** retval )
```

Definition at line 308 of file pthread.c.

```
308                                     {
309     int res = pthread_join(thread.thread, retval);
310     pthread_mutex_destroy(&thread.meta->meta_mutex);
311     free(thread.meta);
312     return res;
313 }
```

#### 4.23.3.9 pthread\_self()

```
bool pthread_self (
    pthread_t * thread )
```

Definition at line 297 of file pthread.c.

```
297                                     {
298     if (my_meta == NULL) {
299         return false;
300     }
301     *thread = (pthread_t){
302         .thread = pthread_self(),
303         .meta = my_meta,
304     };
305     return true;
306 }
```

4.23.3.10 `spthread_suspend()`

```
int spthread_suspend (
    spthread_t thread )
```

Definition at line 171 of file `spthread.c`.

```
171 {
172     pthread_t pself = pthread_self();
173
174     if (pthread_equal(pself, thread.thread) != 0) {
175         return spthread_suspend_self();
176     }
177
178     spthread_signal_args args = (spthread_signal_args){
179         .signal = SPTHREAD_SIG_SUSPEND,
180         .ack = 0,
181     };
182     pthread_mutex_init(&args.shutup_mutex, NULL);
183
184     int ret = pthread_sigqueue(thread.thread, SIGPTH,
185                               (union sigval){
186                                   .sival_ptr = &args,
187                               });
188     if (ret != 0) {
189         pthread_mutex_destroy(&args.shutup_mutex);
190         // handles the case where the thread is already dead.
191         return ret;
192     }
193
194     // wait for our signal to be ack'd
195
196     // setting up args to nanosleep
197     const struct timespec t = (struct timespec){
198         .tv_nsec = MILLISEC_IN_NANO,
199     };
200
201     nanosleep(&t, NULL);
202
203     pthread_mutex_lock(&args.shutup_mutex);
204     while (args.ack != 1) {
205         // wait for a mili second
206         pthread_mutex_unlock(&args.shutup_mutex);
207
208         nanosleep(&t, NULL);
209
210         // fprintf(stderr, "susp checking...\n");
211         pthread_mutex_lock(&args.shutup_mutex);
212
213         if (thread.meta->state == SPTHREAD_TERMINATED_STATE) {
214             // child called exit, can break
215             break;
216         }
217     }
218
219     pthread_mutex_unlock(&args.shutup_mutex);
220
221     pthread_mutex_destroy(&args.shutup_mutex);
222     return ret;
223 }
```

4.23.3.11 `spthread_suspend_self()`

```
int spthread_suspend_self ( )
```

Definition at line 225 of file `spthread.c`.


```
225 {
226     spthread_t self;
227     bool am_sp = spthread_self(&self);
228     if (!am_sp) {
229         return ESRCH;
230     }
231
232     my_meta->state = SPTHREAD_SUSPENDED_STATE;
233
234     do {
235         sigsuspend(&my_meta->suspend_set);
236     } while (my_meta->state == SPTHREAD_SUSPENDED_STATE);
237
238     return 0;
239 }
```

## 4.24 SRC/lib/spthread.h File Reference


```
#include <pthread.h>
```

```
#include <stdbool.h>
```

Include dependency graph for spthread.h:



This graph shows which files directly or indirectly include this file:



### Classes

- struct [spthread\\_st](#)

### Macros

- #define [SIGPTHD](#) SIGUSR1

### Typedefs

- typedef struct [spthread\\_meta\\_st](#) [spthread\\_meta\\_t](#)
- typedef struct [spthread\\_st](#) [spthread\\_t](#)

## Functions

- int `spthread_create` (`spthread_t` \*thread, const pthread\_attr\_t \*attr, void \*(\*start\_routine)(void \*), void \*arg)
- int `spthread_suspend` (`spthread_t` thread)
- int `spthread_suspend_self` ()
- int `spthread_continue` (`spthread_t` thread)
- int `spthread_cancel` (`spthread_t` thread)
- bool `spthread_self` (`spthread_t` \*thread)
- int `spthread_join` (`spthread_t` thread, void \*\*retval)
- void `spthread_exit` (void \*status)
- bool `spthread_equal` (`spthread_t` first, `spthread_t` second)
- int `spthread_disable_interrupts_self` ()
- int `spthread_enable_interrupts_self` ()

### 4.24.1 Macro Definition Documentation

#### 4.24.1.1 SIGPTHD

```
#define SIGPTHD SIGUSR1
```

Definition at line 19 of file `spthread.h`.

### 4.24.2 Typedef Documentation

#### 4.24.2.1 `spthread_meta_t`

```
typedef struct spthread_meta_st spthread_meta_t
```

Definition at line 1 of file `spthread.h`.

#### 4.24.2.2 `spthread_t`

```
typedef struct spthread_st spthread_t
```

### 4.24.3 Function Documentation

#### 4.24.3.1 `spthread_cancel()`

```
int pthread_cancel (
    pthread_t thread )
```

Definition at line 293 of file `spthread.c`.

```
293 {
294     return pthread_cancel(thread.thread);
295 }
```

#### 4.24.3.2 `spthread_continue()`

```
int pthread_continue (
    pthread_t thread )
```

Definition at line 241 of file `spthread.c`.

```
241 {
242     pthread_t pself = pthread_self();
243
244     if (pthread_equal(pself, thread.thread) != 0) {
245         // I am already running... so just return 0
246         my_meta->state = SPTHREAD_RUNNING_STATE;
247         return 0;
248     }
249
250     pthread_signal_args args = (pthread_signal_args){
251         .signal = SPTHREAD_SIG_CONTINUE,
252         .ack = 0,
253     };
254     pthread_mutex_init(&args.shutup_mutex, NULL);
255
256     int ret = pthread_sigqueue(thread.thread, SIGPTH,
257                               (union sigval){
258                                   .sival_ptr = &args,
259                               });
260     if (ret != 0) {
261         pthread_mutex_destroy(&args.shutup_mutex);
262         // handles the case where the thread is already dead.
263         return ret;
264     }
265
266     // wait for our signal to be ack'd
267
268     // setting up args to nanosleep
269     const struct timespec t = (struct timespec){
270         .tv_nsec = MILLISEC_IN_NANO,
271     };
272
273     pthread_mutex_lock(&args.shutup_mutex);
274     while (args.ack != 1) {
275         // wait for a mili second
276         pthread_mutex_unlock(&args.shutup_mutex);
277
278         nanosleep(&t, NULL);
279
280         // fprintf(stderr, "susp checking...\n");
281         pthread_mutex_lock(&args.shutup_mutex);
282
283         if (thread.meta->state == SPTHREAD_TERMINATED_STATE) {
284             // child called exit, can break
285             break;
286         }
287     }
288     pthread_mutex_unlock(&args.shutup_mutex);
289     pthread_mutex_destroy(&args.shutup_mutex);
290     return ret;
291 }
```

## 4.24.3.3 pthread\_create()

```
int pthread_create (
    pthread_t * thread,
    const pthread_attr_t * attr,
    void (*)(void *) start_routine,
    void * arg )
```

Definition at line 114 of file `spthread.c`.

```
117     {
118     pthread_meta_t* child_meta = malloc(sizeof(pthread_meta_t));
119     if (child_meta == NULL) {
120         return EAGAIN;
121     }
122
123     pthread_fwd_args* fwd_args = malloc(sizeof(pthread_fwd_args));
124     if (fwd_args == NULL) {
125         free(child_meta);
126         return EAGAIN;
127     }
128     *fwd_args = (pthread_fwd_args){
129         .actual_routine = start_routine,
130         .actual_arg = arg,
131         .setup_done = false,
132         .child_meta = child_meta,
133     };
134
135     int ret = pthread_mutex_init(&(fwd_args->setup_mutex), NULL);
136     if (ret != 0) {
137         free(child_meta);
138         free(fwd_args);
139         return EAGAIN;
140     }
141
142     ret = pthread_cond_init(&(fwd_args->setup_cond), NULL);
143     if (ret != 0) {
144         free(child_meta);
145         pthread_mutex_destroy(&(fwd_args->setup_mutex));
146         free(fwd_args);
147         return EAGAIN;
148     }
149
150     pthread_t pthread;
151     int result = pthread_create(&pthread, attr, spthread_start, fwd_args);
152
153     pthread_mutex_lock(&(fwd_args->setup_mutex));
154     while (fwd_args->setup_done == false) {
155         pthread_cond_wait(&(fwd_args->setup_cond), &(fwd_args->setup_mutex));
156     }
157     pthread_mutex_unlock(&(fwd_args->setup_mutex));
158
159     pthread_cond_destroy(&(fwd_args->setup_cond));
160     pthread_mutex_destroy(&(fwd_args->setup_mutex));
161     free(fwd_args);
162
163     *thread = (pthread_t){
164         .thread = pthread,
165         .meta = child_meta,
166     };
167
168     return result;
169 }
```

## 4.24.3.4 pthread\_disable\_interrupts\_self()

```
int pthread_disable_interrupts_self ( )
```

Definition at line 326 of file `spthread.c`.

```
326     {
327     sigset_t block_set;
328     int res = sigemptyset(&block_set);
329     if (res != 0) {
330         return res;
331     }
```

```

332  res = sigaddset(&block_set, SIGPTHD);
333  if (res != 0) {
334      return res;
335  }
336  res = pthread_sigmask(SIG_BLOCK, &block_set, NULL);
337  if (res != 0) {
338      return res;
339  }
340  return 0;
341 }

```

#### 4.24.3.5 `spthread_enable_interrupts_self()`

```
int spthread_enable_interrupts_self ( )
```

Definition at line 345 of file `spthread.c`.

```

345  {
346      sigset_t block_set;
347      int res = sigemptyset(&block_set);
348      if (res != 0) {
349          return res;
350      }
351      res = sigaddset(&block_set, SIGPTHD);
352      if (res != 0) {
353          return res;
354      }
355      res = pthread_sigmask(SIG_UNBLOCK, &block_set, NULL);
356      if (res != 0) {
357          return res;
358      }
359      return 0;
360 }

```

#### 4.24.3.6 `spthread_equal()`

```

bool spthread_equal (
    pthread_t first,
    pthread_t second )

```

Definition at line 322 of file `spthread.c`.

```

322  {
323      return pthread_equal(first.thread, second.thread) && (first.meta == second.meta);
324 }

```

#### 4.24.3.7 `spthread_exit()`

```

void spthread_exit (
    void * status )

```

Definition at line 315 of file `spthread.c`.

```

315  {
316      // necessary cleanup is registered
317      // in a cleanup routine
318      // that is pushed at start of an spthread
319      pthread_exit(status);
320 }

```



## 4.24.3.8 pthread\_join()

```
int pthread_join (
    pthread_t thread,
    void ** retval )
```

Definition at line 308 of file `spthread.c`.

```
308
309 int res = pthread_join(thread.thread, retval);
310 pthread_mutex_destroy(&thread.meta->meta_mutex);
311 free(thread.meta);
312 return res;
313 }
```

## 4.24.3.9 pthread\_self()

```
bool pthread_self (
    pthread_t * thread )
```

Definition at line 297 of file `spthread.c`.

```
297
298 if (my_meta == NULL) {
299     return false;
300 }
301 *thread = (pthread_t){
302     .thread = pthread_self(),
303     .meta = my_meta,
304 };
305 return true;
306 }
```

## 4.24.3.10 pthread\_suspend()

```
int pthread_suspend (
    pthread_t thread )
```

Definition at line 171 of file `spthread.c`.

```
171 pthread_t pself = pthread_self();
172
173 if (pthread_equal(pself, thread.thread) != 0) {
174     return pthread_suspend_self();
175 }
176
177 pthread_signal_args args = (pthread_signal_args){
178     .signal = SPTHREAD_SIG_SUSPEND,
179     .ack = 0,
180 };
181 pthread_mutex_init(&args.shutup_mutex, NULL);
182
183 int ret = pthread_sigqueue(thread.thread, SIGPTH,
184                             (union sigval){
185                                 .sival_ptr = &args,
186                             });
187 if (ret != 0) {
188     pthread_mutex_destroy(&args.shutup_mutex);
189     // handles the case where the thread is already dead.
190     return ret;
191 }
192
193 // wait for our signal to be ack'd
194
195 // setting up args to nanosleep
196 const struct timespec t = (struct timespec){
197     .tv_nsec = MILLISEC_IN_NANO,
```

```

199     };
200
201     nanosleep(&t, NULL);
202
203     pthread_mutex_lock(&args.shutup_mutex);
204     while (args.ack != 1) {
205         // wait for a mili second
206         pthread_mutex_unlock(&args.shutup_mutex);
207
208
209         nanosleep(&t, NULL);
210
211         // fprintf(stderr, "susp checking...\n");
212         pthread_mutex_lock(&args.shutup_mutex);
213
214         if (thread.meta->state == SPTHREAD_TERMINATED_STATE) {
215             // child called exit, can break
216             break;
217         }
218     }
219     pthread_mutex_unlock(&args.shutup_mutex);
220
221     pthread_mutex_destroy(&args.shutup_mutex);
222     return ret;
223 }

```

#### 4.24.3.11 spthread\_suspend\_self()

```
int spthread_suspend_self ( )
```

Definition at line 225 of file spthread.c.

```

225     {
226         spthread_t self;
227         bool am_sp = spthread_self(&self);
228         if (!am_sp) {
229             return ESRCH;
230         }
231
232         my_meta->state = SPTHREAD_SUSPENDED_STATE;
233
234         do {
235             sigsuspend(&my_meta->suspend_set);
236         } while (my_meta->state == SPTHREAD_SUSPENDED_STATE);
237
238         return 0;
239     }

```

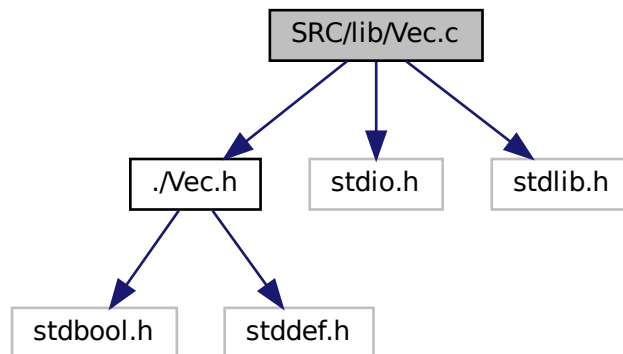
## 4.25 SRC/lib/Vec.c File Reference

```

#include "../Vec.h"
#include <stdio.h>
#include <stdlib.h>

```

Include dependency graph for Vec.c:



## Functions

- [Vec](#) `vec_new` (`size_t` initial\_capacity, `ptr_dtor_fn` ele\_dtor\_fn)
- void `vec_destroy` (`Vec` \*self)
- void `vec_clear` (`Vec` \*self)
- void `vec_resize` (`Vec` \*self, `size_t` new\_capacity)
- void `vec_erase` (`Vec` \*self, `size_t` index)
- void `vec_erase_no_deletor` (`Vec` \*self, `size_t` index)
- void `vec_insert` (`Vec` \*self, `size_t` index, `ptr_t` new\_ele)
- bool `vec_pop_back` (`Vec` \*self)
- void `vec_push_back` (`Vec` \*self, `ptr_t` new\_ele)
- void `vec_set` (`Vec` \*self, `size_t` index, `ptr_t` new\_ele)
- `ptr_t` `vec_get` (`Vec` \*self, `size_t` index)

### 4.25.1 Function Documentation

#### 4.25.1.1 `vec_clear()`

```
void vec_clear (
    Vec * self )
```

Erases all elements from the container. After this, the length of the vector is zero. Capacity of the vector is unchanged.

##### Parameters

<i>self</i>	a pointer to the vector we want to clear.
-------------	---

**Precondition**

Assumes self points to a valid vector.

**Postcondition**

The removed elements are destructed (cleaned up).

Definition at line 34 of file Vec.c.

```

34     {
35     if (self->ele_dtor_fn) {
36         for (int i = 0; i < self->length; i++) {
37             self->ele_dtor_fn(self->data[i]);
38         }
39     }
40
41     self->length = 0;
42 }
```

**4.25.1.2 vec\_destroy()**

```

void vec_destroy (
    Vec * self )
```

Destruct the vector. All elements are destructed and storage is deallocated. Must set capacity and length to zero. Data is set to NULL.

**Parameters**

<i>self</i>	a pointer to the vector we want to destruct.
-------------	--

**Precondition**

Assumes self points to a valid vector.

**Postcondition**

The removed elements are destructed (cleaned up) and data storage deallocated.

Definition at line 17 of file Vec.c.

```

17     {
18     if (self->ele_dtor_fn) {
19         for (int i = 0; i < self->length; i++) {
20             self->ele_dtor_fn(self->data[i]);
21         }
22     }
23     free(self->data);
24 }
```

**4.25.1.3 vec\_erase()**

```

void vec_erase (
    Vec * self,
    size_t index )
```

Erases an element at the specified valid location in the container

## Parameters

<i>self</i>	a pointer to the vector we want to erase from.
<i>index</i>	the index of the element we want to erase at. Elements after this index are "shifted" down one position.

## Precondition

Assumes *self* points to a valid vector. If the index is  $\geq$  *self*->length then this function will call perror

Definition at line 63 of file Vec.c.

```

63                                     {
64     if (index >= self->length) {
65         perror("vec_erase: index >= vec length");
66     }
67
68     if (self->ele_dtor_fn) {
69         self->ele_dtor_fn(self->data[index]);
70     }
71
72     for (unsigned int i = index; i < self->length - 1; i++) {
73         self->data[i] = self->data[i + 1];
74     }
75
76     self->length--;
77 }
```

## 4.25.1.4 vec\_erase\_no\_deletor()

```

void vec_erase_no_deletor (
    Vec * self,
    size_t index )
```

Erases an element at the specified location in the container, except it does not call the element deleter function on the element.

## Parameters

<i>self</i>	a pointer to the vector we want to erase from
<i>index</i>	the index of the element we want to erase at

Definition at line 79 of file Vec.c.

```

79                                     {
80     if (index >= self->length) {
81         perror("vec_erase: index >= vec length");
82     }
83
84     for (unsigned int i = index; i < self->length - 1; i++) {
85         self->data[i] = self->data[i + 1];
86     }
87
88     self->length--;
89 }
```

## 4.25.1.5 vec\_get()

```

ptr_t vec_get (
    Vec * self,
    size_t index )
```

Gets the specified element of the Vec

#### Parameters

<i>self</i>	a pointer to the vector who's element we want to get.
<i>index</i>	the index of the element to get.

#### Returns

the element at the specified index.

#### Precondition

Assumes self points to a valid vector. If the index is  $\geq$  self->length then this function will call perror

Definition at line 152 of file Vec.c.

```

152                                     {
153     if (index >= self->length) {
154         perror("vec_get: index greater than length");
155     }
156     return self->data[index];
157 }
```

#### 4.25.1.6 vec\_insert()

```

void vec_insert (
    Vec * self,
    size_t index,
    ptr_t new_ele )
```

Inserts an element at the specified location in the container

#### Parameters

<i>self</i>	a pointer to the vector we want to insert into.
<i>index</i>	the index of the element we want to insert at. Elements at this index and after it are "shifted" up one position. If index is equal to the length, then we insert at the end of the vector.
<i>new_ele</i>	the value we want to insert

#### Precondition

Assumes self points to a valid vector. If the index is  $>$  self->length then this function will call perror

#### Postcondition

If after the operation the new length is greater than the old capacity then a reallocation takes place and all elements are copied over. Capacity is doubled. Any pointers to elements prior to this reallocation are invalidated.

Definition at line 91 of file Vec.c.

```

91                                     {
92     if (index > self->length) {
93         perror("vec_insert: index > vec length");
94     }
95
96     if (index == self->length) { // Insertion at end = Adding at end
97         vec_push_back(self, new_ele);
98     } else { // Inserting not at the end
99         // Vector is full
100         if (self->length == self->capacity) {
101             vec_resize(self, self->capacity * 2);
102         }
103         // Insertion + Displacement
104         for (unsigned int i = self->length; i > index; i--) {
105             self->data[i] = self->data[i - 1];
106         }
107         self->data[index] = new_ele;
108
109         self->length++;
110     }
111 }

```

#### 4.25.1.7 vec\_new()

```

Vec vec_new (
    size_t initial_capacity,
    ptr_dtor_fn ele_dtor_fn )

```

Creates a new empty Vec(tor) with the specified initial\_capacity and specified function to clean up elements in the vector.

##### Parameters

<i>initial_capacity</i>	the initial capacity of the newly created vector, non negative
<i>ele_dtor_fn</i>	a function pointer to a function that takes in a ptr_t (a vector element) and cleans it up. This is commonly just free but custom functions can be passed in. NULL can also be passed in to specify that there is no cleanup function that needs to be called on each element.

##### Returns

a newly created vector with specified capacity, 0 length and the specified element destructor (cleanup) function.

##### Postcondition

if memory allocation fails, the function will perror

##### Definition at line 5 of file Vec.c.

```

5                                     {
6     Vec vector;
7     vector.capacity = initial_capacity;
8     vector.data = malloc(sizeof(void*) * initial_capacity);
9     if (vector.data == NULL) {
10         perror("vec_new: malloc failed");
11     }
12     vector.ele_dtor_fn = ele_dtor_fn;
13     vector.length = 0;
14     return vector;
15 }

```

#### 4.25.1.8 vec\_pop\_back()

```
bool vec_pop_back (
    Vec * self )
```

Removes and destroys the last element of the Vec

##### Parameters

<i>self</i>	a pointer to the vector we are popping.
-------------	---

##### Returns

true iff an element was removed.

##### Precondition

Assumes self points to a valid vector.

##### Postcondition

The capacity of self stays the same. The removed element is destructed (cleaned up) as specified by the `dtor_fn` provided in `vec_new`.

Definition at line 113 of file `Vec.c`.

```
113 {
114     if (self->length == 0) {
115         return false;
116     }
117     if (self->ele_dtor_fn) {
118         self->ele_dtor_fn(self->data[--self->length]);
119     } else {
120         self->length--;
121     }
122     return true;
123 }
```

#### 4.25.1.9 vec\_push\_back()

```
void vec_push_back (
    Vec * self,
    ptr_t new_ele )
```

Appends the given element to the end of the Vec

##### Parameters

<i>self</i>	a pointer to the vector we are pushing onto
<i>new_ele</i>	the value we want to add to the end of the container



**Precondition**

Assumes self points to a valid vector.

**Postcondition**

If a resize is needed and it fails, then this function will call perror

If after the operation the new length is greater than the old capacity then a reallocation takes place and all elements are copied over. Capacity is doubled. If initial capacity is zero, it is resized to capacity 1. Any pointers to elements prior to this reallocation are invalidated.

Definition at line 125 of file Vec.c.

```

125                                     {
126     if (self->capacity == self->length) {
127         if (self->capacity == 0) {
128             vec_resize(self, 1);
129         } else {
130             vec_resize(self, self->capacity * 2);
131         }
132     }
133
134     if (self->capacity == self->length) {
135         perror("vec_push_back: resize failed");
136     }
137
138     // The array is 0 indexed
139     self->data[self->length++] = new_ele;
140 }
```

**4.25.1.10 vec\_resize()**

```

void vec_resize (
    Vec * self,
    size_t new_capacity )
```

Resizes the container to a new specified capacity. Does nothing if new\_capacity <= self->length

**Parameters**

<i>self</i>	a pointer to the vector we want to resize.
<i>new_capacity</i>	the new capacity of the vector.

**Precondition**

Assumes self points to a valid vector.

**Postcondition**

If a resize takes place, then a reallocation takes place and all elements are copied over. Any pointers to elements prior to this reallocation are invalidated.

The removed elements are destructed (cleaned up).

Definition at line 44 of file Vec.c.

```

44                                     {
45     if (new_capacity * sizeof(void*) < new_capacity) {
46         perror("vec_resize: new capacity too large");
```

```

47  }
48  if (new_capacity > self->length) {
49      self->capacity = new_capacity;
50      ptr_t* new_data = malloc(sizeof(void*) * self->capacity);
51
52      // Copy over old elements
53      for (int i = 0; i < self->length; i++) {
54          new_data[i] = self->data[i];
55      }
56
57      free(self->data);
58
59      self->data = new_data;
60  }
61 }

```

#### 4.25.1.11 vec\_set()

```

void vec_set (
    Vec * self,
    size_t index,
    ptr_t new_ele )

```

Sets the specified element of the Vec to the specified value

##### Parameters

<i>self</i>	a pointer to the vector who's element we want to set.
<i>index</i>	the index of the element to set.
<i>new_ele</i>	the value we want to set the element at that index to

##### Returns

the element at the specified index.

##### Precondition

Assumes self points to a valid vector. If the index is  $\geq$  self->length then this function will call perror

Definition at line 142 of file Vec.c.

```

142                                     {
143     if (index >= self->length) {
144         perror("vec_set: idx >= len");
145     }
146     if (self->ele_dtor_fn) {
147         self->ele_dtor_fn(self->data[index]);
148     }
149     self->data[index] = new_ele;
150 }

```


## 4.26 SRC/lib/Vec.h File Reference

```


#include <stdbool.h>
#include <stddef.h>

```

Include dependency graph for Vec.h:



This graph shows which files directly or indirectly include this file:



## Classes

- struct [vec\\_st](#)

## Macros

- #define [vec\\_capacity](#)(vec) ((vec)->capacity)
- #define [vec\\_len](#)(vec) ((vec)->length)
- #define [vec\\_is\\_empty](#)(vec) ((vec)->length == 0)

## Typedefs

- typedef void \* [ptr\\_t](#)
- typedef void(\* [ptr\\_dtor\\_fn](#)) ([ptr\\_t](#))
- typedef struct [vec\\_st](#) [Vec](#)

## Functions

- [Vec](#) [vec\\_new](#) (size\_t initial\_capacity, [ptr\\_dtor\\_fn](#) ele\_dtor\_fn)
- [ptr\\_t](#) [vec\\_get](#) ([Vec](#) \*self, size\_t index)
- void [vec\\_set](#) ([Vec](#) \*self, size\_t index, [ptr\\_t](#) new\_ele)
- void [vec\\_push\\_back](#) ([Vec](#) \*self, [ptr\\_t](#) new\_ele)
- bool [vec\\_pop\\_back](#) ([Vec](#) \*self)
- void [vec\\_insert](#) ([Vec](#) \*self, size\_t index, [ptr\\_t](#) new\_ele)
- void [vec\\_erase](#) ([Vec](#) \*self, size\_t index)
- void [vec\\_erase\\_no\\_deletor](#) ([Vec](#) \*self, size\_t index)
- void [vec\\_resize](#) ([Vec](#) \*self, size\_t new\_capacity)
- void [vec\\_clear](#) ([Vec](#) \*self)
- void [vec\\_destroy](#) ([Vec](#) \*self)

## 4.26.1 Macro Definition Documentation

### 4.26.1.1 `vec_capacity`

```
#define vec_capacity(  
    vec ) ((vec)->capacity)
```

Returns the current capacity of the Vec Written as a function-like macro

#### Parameters

<code>vec,a</code>	pointer to the vector we want to grab the capacity of.
--------------------	--

Definition at line 40 of file Vec.h.

### 4.26.1.2 `vec_is_empty`

```
#define vec_is_empty(  
    vec ) ((vec)->length == 0)
```

Checks if the Vec is empty written as a function-like macro

#### Parameters

<code>vec,a</code>	pointer to the vector we want to check emptiness of.
--------------------	--

Definition at line 54 of file Vec.h.

### 4.26.1.3 `vec_len`

```
#define vec_len(  
    vec ) ((vec)->length)
```

Returns the current length of the Vec written as a function-like macro

#### Parameters

<code>vec,a</code>	pointer to the vector we want to grab the len of.
--------------------	---

Definition at line 47 of file Vec.h.

## 4.26.2 Typedef Documentation

### 4.26.2.1 ptr\_dtor\_fn

```
typedef void(* ptr_dtor_fn) (ptr_t)
```

Definition at line 8 of file Vec.h.

### 4.26.2.2 ptr\_t

```
typedef void* ptr_t
```

Definition at line 7 of file Vec.h.

### 4.26.2.3 Vec

```
typedef struct vec_st Vec
```

## 4.26.3 Function Documentation

### 4.26.3.1 vec\_clear()

```
void vec_clear (
    Vec * self )
```

Erases all elements from the container. After this, the length of the vector is zero. Capacity of the vector is unchanged.

#### Parameters

<i>self</i>	a pointer to the vector we want to clear.
-------------	---

#### Precondition

Assumes self points to a valid vector.

**Postcondition**

The removed elements are destructed (cleaned up).

Definition at line 34 of file Vec.c.

```

34     {
35     if (self->ele_dtor_fn) {
36         for (int i = 0; i < self->length; i++) {
37             self->ele_dtor_fn(self->data[i]);
38         }
39     }
40
41     self->length = 0;
42 }
```

**4.26.3.2 vec\_destroy()**

```

void vec_destroy (
    Vec * self )
```

Destruct the vector. All elements are destructed and storage is deallocated. Must set capacity and length to zero. Data is set to NULL.

**Parameters**

<i>self</i>	a pointer to the vector we want to destruct.
-------------	--

**Precondition**

Assumes self points to a valid vector.

**Postcondition**

The removed elements are destructed (cleaned up) and data storage deallocated.

Definition at line 17 of file Vec.c.

```

17     {
18     if (self->ele_dtor_fn) {
19         for (int i = 0; i < self->length; i++) {
20             self->ele_dtor_fn(self->data[i]);
21         }
22     }
23     free(self->data);
24 }
```

**4.26.3.3 vec\_erase()**

```

void vec_erase (
    Vec * self,
    size_t index )
```

Erases an element at the specified valid location in the container

## Parameters

<i>self</i>	a pointer to the vector we want to erase from.
<i>index</i>	the index of the element we want to erase at. Elements after this index are "shifted" down one position.

## Precondition

Assumes *self* points to a valid vector. If the index is  $\geq$  *self*->length then this function will call perror

Definition at line 63 of file Vec.c.

```

63                                     {
64     if (index >= self->length) {
65         perror("vec_erase: index >= vec length");
66     }
67
68     if (self->ele_dtor_fn) {
69         self->ele_dtor_fn(self->data[index]);
70     }
71
72     for (unsigned int i = index; i < self->length - 1; i++) {
73         self->data[i] = self->data[i + 1];
74     }
75
76     self->length--;
77 }
```

## 4.26.3.4 vec\_erase\_no\_deletor()

```

void vec_erase_no_deletor (
    Vec * self,
    size_t index )
```

Erases an element at the specified location in the container, except it does not call the element deleter function on the element.

## Parameters

<i>self</i>	a pointer to the vector we want to erase from
<i>index</i>	the index of the element we want to erase at

Definition at line 79 of file Vec.c.

```

79                                     {
80     if (index >= self->length) {
81         perror("vec_erase: index >= vec length");
82     }
83
84     for (unsigned int i = index; i < self->length - 1; i++) {
85         self->data[i] = self->data[i + 1];
86     }
87
88     self->length--;
89 }
```

## 4.26.3.5 vec\_get()

```

ptr_t vec_get (
    Vec * self,
    size_t index )
```

Gets the specified element of the Vec

#### Parameters

<i>self</i>	a pointer to the vector who's element we want to get.
<i>index</i>	the index of the element to get.

#### Returns

the element at the specified index.

#### Precondition

Assumes self points to a valid vector. If the index is  $\geq$  self->length then this function will call perror

Definition at line 152 of file Vec.c.

```
152                                     {
153     if (index >= self->length) {
154         perror("vec_get: index greater than length");
155     }
156     return self->data[index];
157 }
```

#### 4.26.3.6 vec\_insert()

```
void vec_insert (
    Vec * self,
    size_t index,
    ptr_t new_ele )
```

Inserts an element at the specified location in the container

#### Parameters

<i>self</i>	a pointer to the vector we want to insert into.
<i>index</i>	the index of the element we want to insert at. Elements at this index and after it are "shifted" up one position. If index is equal to the length, then we insert at the end of the vector.
<i>new_ele</i>	the value we want to insert

#### Precondition

Assumes self points to a valid vector. If the index is  $>$  self->length then this function will call perror

#### Postcondition

If after the operation the new length is greater than the old capacity then a reallocation takes place and all elements are copied over. Capacity is doubled. Any pointers to elements prior to this reallocation are invalidated.

Definition at line 91 of file Vec.c.



```

91                                     {
92     if (index > self->length) {
93         perror("vec_insert: index > vec length");
94     }
95
96     if (index == self->length) { // Insertion at end = Adding at end
97         vec_push_back(self, new_ele);
98     } else { // Inserting not at the end
99         // Vector is full
100         if (self->length == self->capacity) {
101             vec_resize(self, self->capacity * 2);
102         }
103         // Insertion + Displacement
104         for (unsigned int i = self->length; i > index; i--) {
105             self->data[i] = self->data[i - 1];
106         }
107         self->data[index] = new_ele;
108
109         self->length++;
110     }
111 }

```

#### 4.26.3.7 vec\_new()

```

Vec vec_new (
    size_t initial_capacity,
    ptr_dtor_fn ele_dtor_fn )

```

Creates a new empty Vec(tor) with the specified initial\_capacity and specified function to clean up elements in the vector.

##### Parameters

<i>initial_capacity</i>	the initial capacity of the newly created vector, non negative
<i>ele_dtor_fn</i>	a function pointer to a function that takes in a ptr_t (a vector element) and cleans it up. This is commonly just free but custom functions can be passed in. NULL can also be passed in to specify that there is no cleanup function that needs to be called on each element.

##### Returns

a newly created vector with specified capacity, 0 length and the specified element destructor (cleanup) function.

##### Postcondition

if memory allocation fails, the function will perror

Definition at line 5 of file Vec.c.

```

5                                     {
6     Vec vector;
7     vector.capacity = initial_capacity;
8     vector.data = malloc(sizeof(void*) * initial_capacity);
9     if (vector.data == NULL) {
10         perror("vec_new: malloc failed");
11     }
12     vector.ele_dtor_fn = ele_dtor_fn;
13     vector.length = 0;
14     return vector;
15 }

```

#### 4.26.3.8 vec\_pop\_back()

```
bool vec_pop_back (
    Vec * self )
```

Removes and destroys the last element of the Vec

##### Parameters

<i>self</i>	a pointer to the vector we are popping.
-------------	---

##### Returns

true iff an element was removed.

##### Precondition

Assumes self points to a valid vector.

##### Postcondition

The capacity of self stays the same. The removed element is destructed (cleaned up) as specified by the `dtor_fn` provided in `vec_new`.

Definition at line 113 of file `Vec.c`.

```
113 {
114     if (self->length == 0) {
115         return false;
116     }
117     if (self->ele_dtor_fn) {
118         self->ele_dtor_fn(self->data[--self->length]);
119     } else {
120         self->length--;
121     }
122     return true;
123 }
```

#### 4.26.3.9 vec\_push\_back()

```
void vec_push_back (
    Vec * self,
    ptr_t new_ele )
```

Appends the given element to the end of the Vec

##### Parameters

<i>self</i>	a pointer to the vector we are pushing onto
<i>new_ele</i>	the value we want to add to the end of the container

**Precondition**

Assumes self points to a valid vector.

**Postcondition**

If a resize is needed and it fails, then this function will call perror

If after the operation the new length is greater than the old capacity then a reallocation takes place and all elements are copied over. Capacity is doubled. If initial capacity is zero, it is resized to capacity 1. Any pointers to elements prior to this reallocation are invalidated.

Definition at line 125 of file Vec.c.

```

125                                     {
126     if (self->capacity == self->length) {
127         if (self->capacity == 0) {
128             vec_resize(self, 1);
129         } else {
130             vec_resize(self, self->capacity * 2);
131         }
132     }
133
134     if (self->capacity == self->length) {
135         perror("vec_push_back: resize failed");
136     }
137
138     // The array is 0 indexed
139     self->data[self->length++] = new_ele;
140 }
```

**4.26.3.10 vec\_resize()**

```

void vec_resize (
    Vec * self,
    size_t new_capacity )
```

Resizes the container to a new specified capacity. Does nothing if new\_capacity <= self->length

**Parameters**

<i>self</i>	a pointer to the vector we want to resize.
<i>new_capacity</i>	the new capacity of the vector.

**Precondition**

Assumes self points to a valid vector.

**Postcondition**

If a resize takes place, then a reallocation takes place and all elements are copied over. Any pointers to elements prior to this reallocation are invalidated.

The removed elements are destructed (cleaned up).

Definition at line 44 of file Vec.c.

```

44                                     {
45     if (new_capacity * sizeof(void*) < new_capacity) {
46         perror("vec_resize: new capacity too large");
```

```

47  }
48  if (new_capacity > self->length) {
49      self->capacity = new_capacity;
50      ptr_t* new_data = malloc(sizeof(void*) * self->capacity);
51
52      // Copy over old elements
53      for (int i = 0; i < self->length; i++) {
54          new_data[i] = self->data[i];
55      }
56
57      free(self->data);
58
59      self->data = new_data;
60  }
61 }

```

#### 4.26.3.11 vec\_set()

```

void vec_set (
    Vec * self,
    size_t index,
    ptr_t new_ele )

```

Sets the specified element of the Vec to the specified value

##### Parameters

<i>self</i>	a pointer to the vector who's element we want to set.
<i>index</i>	the index of the element to set.
<i>new_ele</i>	the value we want to set the element at that index to

##### Returns

the element at the specified index.

##### Precondition

Assumes self points to a valid vector. If the index is  $\geq$  self->length then this function will call perror

Definition at line 142 of file Vec.c.

```

142                                     {
143     if (index >= self->length) {
144         perror("vec_set: idx >= len");
145     }
146     if (self->ele_dtor_fn) {
147         self->ele_dtor_fn(self->data[index]);
148     }
149     self->data[index] = new_ele;
150 }

```

## 4.27 SRC/pennfat.c File Reference

```

#include "shell/parser.h"
#include "shell/builtins.h"
#include "lib/pennos-errno.h"
#include "fs/fs_kfuncs.h"


```

```

#include "fs/fat_routines.h"
#include <signal.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>
#include <sys/mman.h>
#include <unistd.h>
#include <stdbool.h>

```

Include dependency graph for pennfat.c:



## Macros

- #define **PROMPT** "pennfat# "

## Functions

- int **main** (int argc, char \*argv[ ])

### 4.27.1 Macro Definition Documentation

#### 4.27.1.1 PROMPT

```
#define PROMPT "pennfat# "
```

Definition at line 16 of file pennfat.c.

### 4.27.2 Function Documentation

## 4.27.2.1 main()

```
int main (
    int argc,
    char * argv[] )
```

Definition at line 24 of file pennfat.c.

```
24     {
25         // register signal handlers
26         struct sigaction sa;
27         sa.sa_handler = signal_handler;
28         sigemptyset(&sa.sa_mask);
29         sa.sa_flags = SA_RESTART;
30
31         // set up handler for SIGINT (ctrl-c)
32         if (sigaction(SIGINT, &sa, NULL) == -1) {
33             P_ERRNO = P_ESIGNAL;
34             u_perror("Error setting up SIGINT handler");
35             return EXIT_FAILURE;
36         }
37
38         // set up handler for SIGTSTP (ctrl-z)
39         if (sigaction(SIGTSTP, &sa, NULL) == -1) {
40             P_ERRNO = P_ESIGNAL;
41             u_perror("Error setting up SIGTSTP handler");
42             return EXIT_FAILURE;
43         }
44
45         char input_buffer[1024];
46
47         while (true) {
48             // print prompt
49             if (k_write(STDOUT_FILENO, PROMPT, strlen(PROMPT)) < 0) {
50                 P_ERRNO = P_EWRITE;
51                 u_perror("prompt write error");
52                 break;
53             }
54
55             // read user input
56             int bytes_read = k_read(STDIN_FILENO, input_buffer, sizeof(input_buffer) - 1);
57
58             // check for EOF (ctrl-D)
59             if (bytes_read <= 0) {
60                 k_write(STDOUT_FILENO, "\n", 1);
61                 break;
62             }
63
64             // remove trailing newline if present
65             if (bytes_read > 0 && input_buffer[bytes_read - 1] == '\n') {
66                 input_buffer[bytes_read - 1] = '\0';
67             }
68
69             // parse command and check error
70             struct parsed_command *parsed_command = NULL;
71             int parse_result = parse_command(input_buffer, &parsed_command);
72             if (parse_result != 0) {
73                 if (parse_result == -1) {
74                     P_ERRNO = P_EINVAL;
75                     u_perror("Error parsing command");
76                 } else {
77                     print_parser_errcode(stderr, parse_result);
78                 }
79                 continue;
80             }
81
82             // skip empty commands
83             if (parsed_command->num_commands == 0) {
84                 free(parsed_command);
85                 continue;
86             }
87
88             // extract command and arguments
89             char **args = parsed_command->commands[0];
90
91             // execute command
92             if (strcmp(args[0], "mkfs") == 0) {
93                 if (args[1] == NULL || args[2] == NULL || args[3] == NULL) {
94                     P_ERRNO = P_EINVAL;
95                     u_perror("mkfs");
96                 } else {
97                     int blocks_in_fat = atoi(args[2]);
98                     int block_size = atoi(args[3]);
99                     if (mkfs(args[1], blocks_in_fat, block_size) != 0) {
100                         u_perror("mkfs");

```

```

101         }
102     }
103     } else if (strcmp(args[0], "mount") == 0) {
104         if (args[1] == NULL) {
105             P_ERRNO = P_EINVAL;
106             u_perror("mount");
107         } else {
108             if (mount(args[1]) != 0) {
109                 u_perror("mount");
110             }
111         }
112     } else if (strcmp(args[0], "unmount") == 0) {
113         if (unmount() != 0) {
114             u_perror("unmount");
115         }
116     } else if (strcmp(args[0], "ls") == 0) {
117         ls(args);
118     } else if (strcmp(args[0], "touch") == 0) {
119         touch(args);
120     } else if (strcmp(args[0], "cat") == 0) {
121         cat(args);
122     } else if (strcmp(args[0], "chmod") == 0) {
123         chmod(args);
124     } else if (strcmp(args[0], "mv") == 0) {
125         mv(args);
126     } else if (strcmp(args[0], "rm") == 0) {
127         rm(args);
128     } else if (strcmp(args[0], "cp") == 0) {
129         cp(args);
130     } else if (strcmp(args[0], "cmpctdir") == 0) { // extra credit
131         cmpctdir(args);
132     } else {
133         P_ERRNO = P_ECOMMAND;
134         u_perror("shell");
135     }
136     free(parsed_command);
137 }
138 return EXIT_SUCCESS;
139 }
140 }

```


## 4.28 SRC/pennos.c File Reference

```

#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>
#include "fs/fs_syscalls.h"
#include "kernel/kern_sys_calls.h"
#include "kernel/scheduler.h"
#include "shell/builtins.h"
#include "lib/pennos-errno.h"
#include "fs/fat_routines.h"

```

Include dependency graph for pennos.c:



## Functions

- int [main](#) (int argc, char \*argv[])

## Variables

- int `tick_counter`
- int `log_fd`

### 4.28.1 Function Documentation

#### 4.28.1.1 `main()`

```
int main (
    int argc,
    char * argv[] )
```

Definition at line 14 of file `pennos.c`.

```
14                                     {
15     // mount the filesystem
16     if (argc < 2) {
17         P_ERRNO = P_NEEDF;
18         u_perror("need a pennfat file to mount");
19         return -1;
20     } else {
21         if (mount(argv[1]) == -1) {
22             u_perror("mount failed");
23             return -1;
24         }
25     }
26
27     // get the log fd
28     if (argc >= 3) {
29         log_fd = open(argv[2], O_RDWR | O_CREAT | O_TRUNC, 0644);
30     } else {
31         log_fd = open("log/log", O_RDWR | O_CREAT | O_TRUNC, 0644);
32     }
33
34     // initialize scheduler architecture and init process
35     initialize_scheduler_queues();
36
37     pid_t init_pid = s_spawn_init();
38     if (init_pid == -1) {
39         P_ERRNO = P_INITFAIL;
40         u_perror("init spawn failed");
41         return -1;
42     }
43
44     scheduler();
45
46     // cleanup
47     s_cleanup_init_process();
48     free_scheduler_queues();
49     unmount();
50     close(log_fd);
51 }
```

### 4.28.2 Variable Documentation

#### 4.28.2.1 `log_fd`

```
int log_fd [extern]
```

Definition at line 36 of file `scheduler.c`.



## 4.28.2.2 tick\_counter


```
int tick_counter [extern]
```

Definition at line 35 of file scheduler.c.

## 4.29 SRC/shell/builtins.c File Reference

```
#include <stdio.h>
#include "builtins.h"
#include "../lib/pennos-errno.h"
#include "../fs/fs_syscalls.h"
#include <string.h>
```

Include dependency graph for builtins.c:



## Functions

- void [u\\_perror](#) (const char \*msg)

## 4.29.1 Function Documentation

## 4.29.1.1 u\_perror()

```
void u_perror (
    const char * msg )
```

Creates a user-level error message similar to perror.

## Parameters

<i>msg</i>	A string representing the error message from the shell.
------------	---

Definition at line 12 of file builtins.c.


```

12                                     {
13     char buffer[256];
14     const char *error_msg;
15
16     switch (P_ERRNO) {
17         case P_ENOENT:
18             error_msg = "file does not exist";
19             break;
20         case P_EBADF:
21             error_msg = "bad file descriptor";
22             break;
23         case P_EPERM:
24             error_msg = "operation not permitted";
25             break;
26         case P_EINVAL:
27             error_msg = "invalid arg";
28             break;
29         case P_EEXIST:
30             error_msg = "file already exists";
31             break;
32         case P_EBUSY:
33             error_msg = "file is busy or open";
34             break;
35         case P_EFULL:
36             error_msg = "no space left on device";
37             break;
38         case P_EINTR:
39             error_msg = "interrupted system call";
40             break;
41         case P_ENULL:
42             error_msg = "NULL returned unexpectedly";
43             break;
44         case P_EUNKNOWN:
45             error_msg = "unknown error";
46             break;
47         case P_EREAD:
48             error_msg = "interrupted read call";
49             break;
50         case P_ELSEEK:
51             error_msg = "interrupted lseek call";
52             break;
53         case P_EMAP:
54             error_msg = "interrupted mmap/munmap call";
55             break;
56         case P_EFUNC:
57             error_msg = "interrupted system call";
58             break;
59         case P_EOPEN:
60             error_msg = "interrupted open call";
61             break;
62         case P_EMALLOC:
63             error_msg = "error when trying to malloc";
64             break;
65         case P_EFS_NOT_MOUNTED:
66             error_msg = "file system not mounted yet";
67             break;
68         case P_ESIGNAL:
69             error_msg = "error with signal handling";
70             break;
71         case P_EWRITE:
72             error_msg = "interrupted write call";
73             break;
74         case P_ECLOSE:
75             error_msg = "interrupted close call";
76             break;
77         case P_EPARSE:
78             error_msg = "error when trying to parse a command";
79             break;
80         case P_ECOMMAND:
81             error_msg = "command not found";
82             break;
83         case P_NEEDF:
84             error_msg = "no file provided to mount";
85             break;
86         default:
87             error_msg = "Unknown error";
88             break;
89     }
90
91     snprintf(buffer, sizeof(buffer), "%s: %s\n", msg, error_msg);
92     s_write(STDERR_FILENO, buffer, strlen(buffer));
93 }

```

## 4.30 SRC/shell/builtins.h File Reference

This graph shows which files directly or indirectly include this file:



### Functions

- void [u\\_perror](#) (const char \*msg)

#### 4.30.1 Function Documentation

##### 4.30.1.1 u\_perror()

```
void u_perror (
    const char * msg )
```

Creates a user-level error message similar to perror.

#### Parameters

<i>msg</i>	A string representing the error message from the shell.
------------	---

Definition at line 12 of file builtins.c.

```

12      {
13      char buffer[256];
14      const char *error_msg;
15
16      switch (P_ERRNO) {
17      case P_ENOENT:
18          error_msg = "file does not exist";
19          break;
20      case P_EBADF:
21          error_msg = "bad file descriptor";
22          break;
23      case P_EPERM:
24          error_msg = "operation not permitted";
25          break;
26      case P_EINVAL:
27          error_msg = "invalid arg";
28          break;
29      case P_EEXIST:
30          error_msg = "file already exists";
31          break;
32      case P_EBUSY:
33          error_msg = "file is busy or open";
34          break;
35      case P_EFULL:
36          error_msg = "no space left on device";
37          break;
38      case P_EINTR:
39          error_msg = "interrupted system call";
```

```

40         break;
41     case P_ENULL:
42         error_msg = "NULL returned unexpectedly";
43         break;
44     case P_EUNKNOWN:
45         error_msg = "unknown error";
46         break;
47     case P_EREAD:
48         error_msg = "interrupted read call";
49         break;
50     case P_ELSEEK:
51         error_msg = "interrupted lseek call";
52         break;
53     case P_EMAP:
54         error_msg = "interrupted mmap/munmap call";
55         break;
56     case P_EFUNC:
57         error_msg = "interrupted system call";
58         break;
59     case P_EOPEN:
60         error_msg = "interrupted open call";
61         break;
62     case P_EMALLOC:
63         error_msg = "error when trying to malloc";
64         break;
65     case P_EFS_NOT_MOUNTED:
66         error_msg = "file system not mounted yet";
67         break;
68     case P_ESIGNAL:
69         error_msg = "error with signal handling";
70         break;
71     case P_EWRITE:
72         error_msg = "interrupted write call";
73         break;
74     case P_ECLOSE:
75         error_msg = "interrupted close call";
76         break;
77     case P_EPARSE:
78         error_msg = "error when trying to parse a command";
79         break;
80     case P_ECOMMAND:
81         error_msg = "command not found";
82         break;
83     case P_NEEDF:
84         error_msg = "no file provided to mount";
85         break;
86     default:
87         error_msg = "Unknown error";
88         break;
89 }
90
91 snprintf(buffer, sizeof(buffer), "%s: %s\n", msg, error_msg);
92 s_write(STDERR_FILENO, buffer, strlen(buffer));
93 }

```

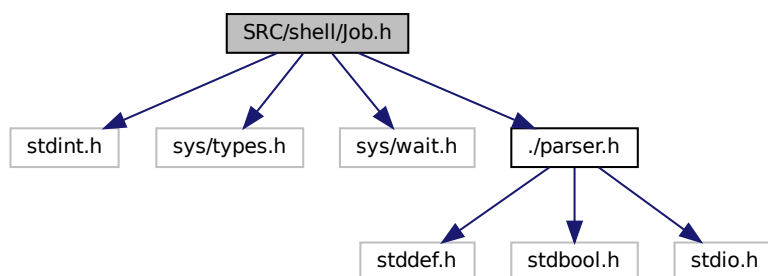
## 4.31 SRC/shell/Job.h File Reference

```


#include <stdint.h>
#include <sys/types.h>
#include <sys/wait.h>
#include "parser.h"

```

Include dependency graph for Job.h:



This graph shows which files directly or indirectly include this file:



## Classes

- struct [job\\_st](#)

## Typedefs

- typedef uint64\_t [jid\\_t](#)
- typedef struct [job\\_st](#) [job](#)

## Enumerations

- enum [job\\_state\\_t](#) { [RUNNING](#) , [STOPPED](#) , [FINISHED](#) }

### 4.31.1 Typedef Documentation

#### 4.31.1.1 `jid_t`

```
typedef uint64_t jid_t
```

Definition at line 10 of file Job.h.

#### 4.31.1.2 `job`

```
typedef struct job_st job
```

### 4.31.2 Enumeration Type Documentation

#### 4.31.2.1 `job_state_t`

```
enum job_state_t
```

##### Enumerator

RUNNING	
STOPPED	
FINISHED	

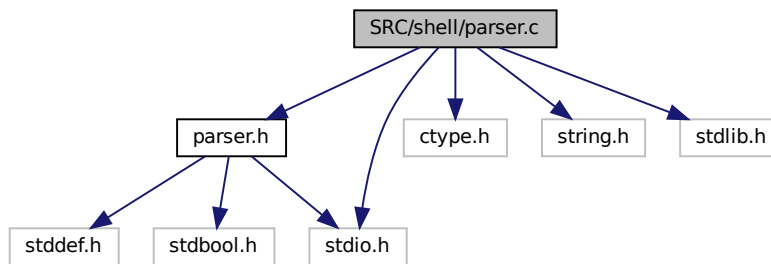
Definition at line 13 of file Job.h.

```
13 { RUNNING, STOPPED, FINISHED } job_state_t;
```

## 4.32 SRC/shell/parser.c File Reference

```
#include "parser.h"  
#include <ctype.h>  
#include <string.h>  
#include <stdlib.h>  
#include <stdio.h>
```

Include dependency graph for parser.c:



## Macros

- `#define JUMP_OUT(code) do {ret_code = code; goto PROCESS_ERROR;} while (0)`

## Functions

- `int parse_command (const char *const cmd_line, struct parsed_command **const result)`
- `void print_parsed_command (const struct parsed_command *const cmd)`
- `void print_parser_errcode (FILE *output, int err_code)`

### 4.32.1 Macro Definition Documentation

#### 4.32.1.1 JUMP\_OUT

```
#define JUMP_OUT(  
    code ) do {ret_code = code; goto PROCESS_ERROR;} while (0)
```

### 4.32.2 Function Documentation

#### 4.32.2.1 parse\_command()

```
int parse_command (
    const char * cmd_line,
    struct parsed_command ** result )
```

Arguments: cmd\_line: a null-terminated string that is the command line result: a non-null pointer to a struct `parsed_command` \*

Return value (int): an error code which can be, 0: parser finished succesfully -1: parser encountered a system call error 1-7: parser specific error, see error type above

This function will parse the given cmd\_line and store the parsed information into a struct `parsed_command`. The memory needed for the struct will be allocated by this function, and the pointer to the memory will be stored into the given \*result.

You can directly use the result in system calls. See demo for more information.

If the function returns a successful value (0), a struct `parsed_command` is guaranteed to be allocated and stored in the given \*result. It is the caller's responsibility to free the given pointer using `free(3)`.

Otherwise, no struct `parsed_command` is allocated and \*result is unchanged. If a system call error (-1) is returned, the caller can use `errno(3)` or `perror(3)` to gain more information about the error. layout of memory for struct `parsed_command` bool is\_background; bool is\_file\_append;

const char \*stdin\_file; const char \*stdout\_file;

size\_t num\_commands;

commands are pointers to arguments char \*\*commands[num\_commands];

below are hidden in memory \*\*

arguments are pointers to original\_string + num\_commands because all argv are null-terminated char \*arguments[total\_strings + num\_commands];

original\_string is a copy of the cmdline but with each token null-terminated char \*original\_string;

Definition at line 16 of file parser.c.

```
16 {
17 #define JUMP_OUT(code) do {ret_code = code; goto PROCESS_ERROR;} while (0)
18
19     int ret_code = -1;
20
21     const char *start = cmd_line;
22     const char *end = cmd_line + strlen(cmd_line);
23
24     for (const char *cur = start; cur < end; ++cur)
25         if (*cur == '#') {
26             // all subsequent characters following '#'
27             // shall be discarded as a comment.
28             end = cur;
29             break;
30         }
31
32     // trimming leading and trailing whitespaces
33     while (start < end && isspace(*start)) ++start;
34     while (start < end && isspace(end[-1])) --end;
35
36     struct parsed_command *pcmd = calloc(1, sizeof(struct parsed_command));
37     if (pcmd == NULL) return -1;
38     if (start == end) goto PROCESS_SUCCESS; // empty line, fast pass
39
40     // If a command is terminated by the control operator ampersand ( '&' ),
41     // the shell shall execute the command in background.
42     if (end[-1] == '&') {
43         pcmd->is_background = true;
44         --end;
```



```

45     }
46
47     // first pass, check token
48     int total_strings = 0; // number of total arguments
49     {
50         bool has_token_last = false, has_file_input = false, has_file_output = false;
51         const char *skipped;
52         for (const char *cur = start; cur < end; skip_space(&cur, end))
53             switch (cur[0]) {
54                 case '&':
55                     JUMP_OUT(UNEXPECTED_AMPERSAND); // does not expect anymore ampersand
56                 case '<':
57                     // if already had pipeline or had file input, error
58                     if (pcmd->num_commands > 0 || has_file_input) JUMP_OUT(UNEXPECTED_FILE_INPUT);
59
60                     ++cur; // skip '<'
61                     skip_space(&cur, end);
62
63                     // test if we indeed have a filename following '<'
64                     skipped = cur;
65                     skip_word(&skipped, end);
66                     if (skipped <= cur) JUMP_OUT(EXPECT_INPUT_FILENAME);
67
68                     // fast-forward to the end of the filename
69                     cur = skipped;
70                     has_file_input = true;
71                     break;
72                 case '>':
73                     // if already had file output, error
74                     if (has_file_output) JUMP_OUT(UNEXPECTED_FILE_OUTPUT);
75                     if (cur + 1 < end && cur[1] == '>') { // dealing with '»' append
76                         pcmd->is_file_append = true;
77                         ++cur;
78                     }
79
80                     ++cur; // skip '>'
81                     skip_space(&cur, end);
82
83                     // test filename, as the case above
84                     skipped = cur;
85                     skip_word(&skipped, end);
86                     if (skipped <= cur) JUMP_OUT(EXPECT_OUTPUT_FILENAME);
87
88                     // fast-forward to the end of the filename
89                     cur = skipped;
90                     has_file_output = true;
91                     break;
92                 case '|':
93                     // if already had file output but encounter a pipeline, it should
94                     // rather be a file output error instead of a pipeline one.
95                     if (has_file_output) JUMP_OUT(UNEXPECTED_FILE_OUTPUT);
96                     // if no tokens between two pipelines (or before the first one)
97                     // should throw a pipeline error
98                     if (!has_token_last) JUMP_OUT(UNEXPECTED_PIPELINE);
99                     has_token_last = false;
100                     ++pcmd->num_commands;
101                     ++cur; // skip '|'
102                     break;
103                 default:
104                     has_token_last = true;
105                     ++total_strings;
106                     skip_word(&cur, end); // skip that argument
107             }
108
109     if (total_strings == 0) {
110         // if there are no arguments but has ampersand or file input/output
111         // then we have an error
112         if (pcmd->is_background || has_file_input || has_file_output)
113             JUMP_OUT(EXPECT_COMMANDS);
114         // otherwise it's an empty line
115         goto PROCESS_SUCCESS;
116     }
117
118     // handle edge case where the command ends with a pipeline
119     // (not supporting line continuation)
120     if (!has_token_last) JUMP_OUT(UNEXPECTED_PIPELINE);
121 }
122 ++pcmd->num_commands;
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147     const size_t start_of_array = offsetof(struct parsed_command, commands) + pcmd->num_commands *
sizeof(char **);
148     const size_t start_of_str = start_of_array + (pcmd->num_commands + total_strings) * sizeof(char *);
149     const size_t slen = end - start;
150
151     char *const new_buf = realloc(pcmd, start_of_str + slen + 1);
152     if (new_buf == NULL) goto PROCESS_ERROR;

```

```

153     pcmd = (struct parsed_command *) new_buf;
154
155     // copy string to the new place
156     char *const new_start = memcpy(new_buf + start_of_str, start, slen);
157
158     // second pass, put stuff in
159     // no need to check for error anymore
160     size_t cur_cmd = 0;
161     char **argv_ptr = (char **) (new_buf + start_of_array);
162
163     pcmd->commands[cur_cmd] = argv_ptr;
164     for (const char *cur = start; cur < end; skip_space(&cur, end)) {
165         switch (cur[0]) {
166             case '<':
167                 ++cur;
168                 skip_space(&cur, end);
169                 // store input file name into 'stdin_file'
170                 pcmd->stdin_file = new_start + (cur - start);
171                 skip_word(&cur, end);
172                 // at end of the input file name
173                 new_start[cur - start] = '\\0';
174                 break;
175             case '>':
176                 if (pcmd->is_file_append) ++cur; // skip another '>'
177                 ++cur;
178                 skip_space(&cur, end);
179                 // store output file name into 'stdout_file'
180                 pcmd->stdout_file = new_start + (cur - start);
181                 skip_word(&cur, end);
182                 // at end of the output file name
183                 new_start[cur - start] = '\\0';
184                 break;
185             case '|':
186                 // null-terminate the current argv
187                 *(argv_ptr++) = NULL;
188                 // store the next argv head
189                 pcmd->commands[++cur_cmd] = argv_ptr;
190                 ++cur;
191                 break;
192             default:
193                 // at start of the argument string
194                 // store it into the arguments array
195                 *(argv_ptr++) = new_start + (cur - start);
196                 skip_word(&cur, end);
197                 // at end of the argument string
198                 new_start[cur - start] = '\\0';
199         }
200     }
201     // null-terminate the last argv
202     *argv_ptr = NULL;
203
204     PROCESS_SUCCESS:
205     *result = pcmd;
206     return 0;
207     PROCESS_ERROR:
208     free(pcmd);
209     return ret_code;
210 }

```

#### 4.32.2.2 print\_parsed\_command()

```

void print_parsed_command (
    const struct parsed_command *const cmd )

```

Definition at line 214 of file parser.c.

```

214     {
215         for (size_t i = 0; i < cmd->num_commands; ++i) {
216             for (char **arguments = cmd->commands[i]; *arguments != NULL; ++arguments)
217                 printf("%s ", *arguments);
218
219             if (i == 0 && cmd->stdin_file != NULL)
220                 printf("< %s ", cmd->stdin_file);
221
222             if (i == cmd->num_commands - 1) {
223                 if (cmd->stdout_file != NULL)
224                     printf(cmd->is_file_append ? "» %s " : "> %s ", cmd->stdout_file);
225             } else printf("| ");
226         }
227         puts(cmd->is_background ? "&" : "");
228     }

```

## 4.32.2.3 print\_parser\_errcode()


```
void print_parser_errcode (
    FILE * output,
    int err_code )
```

Definition at line 230 of file parser.c.

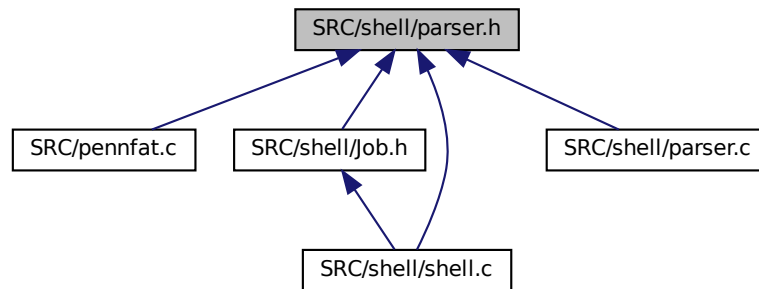
```
230                                     {
231     switch (err_code) {
232     case UNEXPECTED_FILE_INPUT:
233         fprintf(output, "UNEXPECTED INPUT REDIRECTION TO A FILE\n");
234         break;
235     case UNEXPECTED_FILE_OUTPUT:
236         fprintf(output, "UNEXPECTED OUTPUT REDIRECTION TO A FILE\n");
237         break;
238     case UNEXPECTED_PIPELINE:
239         fprintf(output, "UNEXPECTED PIPE\n");
240         break;
241     case UNEXPECTED_AMPERSAND:
242         fprintf(output, "UNEXPECTED AMPERESAND\n");
243         break;
244     case EXPECT_INPUT_FILENAME:
245         fprintf(output, "COULD NOT FINE FILENAME FOR INPUT REDIRECTION \\"<\\n");
246         break;
247     case EXPECT_OUTPUT_FILENAME:
248         fprintf(output, "COULD NOT FIND FILENAME FOR OUTPUT REDIRECTION \\"<\\n");
249         break;
250     case EXPECT_COMMANDS:
251         fprintf(output, "COULD NOT FIND ANY COMMANDS OR ARGS\n");
252         break;
253     default:
254         break;
255     }
256 }
```

## 4.33 SRC/shell/parser.h File Reference

```
#include <stddef.h>
#include <stdbool.h>
#include <stdio.h>
Include dependency graph for parser.h:
```



This graph shows which files directly or indirectly include this file:



## Classes

- struct [parsed\\_command](#)

## Macros

- `#define UNEXPECTED_FILE_INPUT 1`
- `#define UNEXPECTED_FILE_OUTPUT 2`
- `#define UNEXPECTED_PIPELINE 3`
- `#define UNEXPECTED_AMPERSAND 4`
- `#define EXPECT_INPUT_FILENAME 5`
- `#define EXPECT_OUTPUT_FILENAME 6`
- `#define EXPECT_COMMANDS 7`

## Functions

- int [parse\\_command](#) (const char \*cmd\_line, struct [parsed\\_command](#) \*\*result)
- void [print\\_parsed\\_command](#) (const struct [parsed\\_command](#) \*cmd)
- void [print\\_parser\\_errcode](#) (FILE \*output, int err\_code)

### 4.33.1 Macro Definition Documentation

#### 4.33.1.1 EXPECT\_COMMANDS

```
#define EXPECT_COMMANDS 7
```

Definition at line 30 of file parser.h.

#### 4.33.1.2 EXPECT\_INPUT\_FILENAME

```
#define EXPECT_INPUT_FILENAME 5
```

Definition at line 24 of file parser.h.

#### 4.33.1.3 EXPECT\_OUTPUT\_FILENAME

```
#define EXPECT_OUTPUT_FILENAME 6
```

Definition at line 27 of file parser.h.

#### 4.33.1.4 UNEXPECTED\_AMPERSAND

```
#define UNEXPECTED_AMPERSAND 4
```

Definition at line 21 of file parser.h.

#### 4.33.1.5 UNEXPECTED\_FILE\_INPUT

```
#define UNEXPECTED_FILE_INPUT 1
```

Definition at line 12 of file parser.h.

#### 4.33.1.6 UNEXPECTED\_FILE\_OUTPUT

```
#define UNEXPECTED_FILE_OUTPUT 2
```

Definition at line 15 of file parser.h.

#### 4.33.1.7 UNEXPECTED\_PIPELINE

```
#define UNEXPECTED_PIPELINE 3
```

Definition at line 18 of file parser.h.

## 4.33.2 Function Documentation

### 4.33.2.1 parse\_command()

```
int parse_command (
    const char * cmd_line,
    struct parsed_command ** result )
```

Arguments: cmd\_line: a null-terminated string that is the command line result: a non-null pointer to a struct `parsed_command` \*

Return value (int): an error code which can be, 0: parser finished succesfully -1: parser encountered a system call error 1-7: parser specific error, see error type above

This function will parse the given `cmd_line` and store the parsed information into a struct `parsed_command`. The memory needed for the struct will be allocated by this function, and the pointer to the memory will be stored into the given `*result`.

You can directly use the result in system calls. See demo for more information.

If the function returns a successful value (0), a struct `parsed_command` is guaranteed to be allocated and stored in the given `*result`. It is the caller's responsibility to free the given pointer using `free(3)`.

Otherwise, no struct `parsed_command` is allocated and `*result` is unchanged. If a system call error (-1) is returned, the caller can use `errno(3)` or `perror(3)` to gain more information about the error. layout of memory for struct `parsed_command` bool `is_background`; bool `is_file_append`;

const char \*stdin\_file; const char \*stdout\_file;

size\_t num\_commands;

commands are pointers to arguments char \*\*commands[num\_commands];

below are hidden in memory \*\*

arguments are pointers to original\_string + num\_commands because all argv are null-terminated char \*arguments[total\_strings + num\_commands];

original\_string is a copy of the cmdline but with each token null-terminated char \*original\_string;

Definition at line 16 of file parser.c.

```
16
17 #define JUMP_OUT(code) do {ret_code = code; goto PROCESS_ERROR;} while (0)
18
19     int ret_code = -1;
20
21     const char *start = cmd_line;
22     const char *end = cmd_line + strlen(cmd_line);
23
24     for (const char *cur = start; cur < end; ++cur)
25         if (*cur == '#') {
26             // all subsequent characters following '#'
27             // shall be discarded as a comment.
28             end = cur;
29             break;
30         }
31
32     // trimming leading and trailing whitespaces
33     while (start < end && isspace(*start)) ++start;
34     while (start < end && isspace(end[-1])) --end;
35
36     struct parsed_command *pcmd = calloc(1, sizeof(struct parsed_command));
```

```

37     if (pcmd == NULL) return -1;
38     if (start == end) goto PROCESS_SUCCESS; // empty line, fast pass
39
40     // If a command is terminated by the control operator ampersand ( '&' ),
41     // the shell shall execute the command in background.
42     if (end[-1] == '&') {
43         pcmd->is_background = true;
44         --end;
45     }
46
47     // first pass, check token
48     int total_strings = 0; // number of total arguments
49     {
50         bool has_token_last = false, has_file_input = false, has_file_output = false;
51         const char *skipped;
52         for (const char *cur = start; cur < end; skip_space(&cur, end))
53             switch (cur[0]) {
54                 case '&':
55                     JUMP_OUT(UNEXPECTED_AMPERSAND); // does not expect anymore ampersand
56                 case '<':
57                     // if already had pipeline or had file input, error
58                     if (pcmd->num_commands > 0 || has_file_input) JUMP_OUT(UNEXPECTED_FILE_INPUT);
59
60                     ++cur; // skip '<'
61                     skip_space(&cur, end);
62
63                     // test if we indeed have a filename following '<'
64                     skipped = cur;
65                     skip_word(&skipped, end);
66                     if (skipped <= cur) JUMP_OUT(EXPECT_INPUT_FILENAME);
67
68                     // fast-forward to the end of the filename
69                     cur = skipped;
70                     has_file_input = true;
71                     break;
72                 case '>':
73                     // if already had file output, error
74                     if (has_file_output) JUMP_OUT(UNEXPECTED_FILE_OUTPUT);
75                     if (cur + 1 < end && cur[1] == '>') { // dealing with '>' append
76                         pcmd->is_file_append = true;
77                         ++cur;
78                     }
79
80                     ++cur; // skip '>'
81                     skip_space(&cur, end);
82
83                     // test filename, as the case above
84                     skipped = cur;
85                     skip_word(&skipped, end);
86                     if (skipped <= cur) JUMP_OUT(EXPECT_OUTPUT_FILENAME);
87
88                     // fast-forward to the end of the filename
89                     cur = skipped;
90                     has_file_output = true;
91                     break;
92                 case '|':
93                     // if already had file output but encounter a pipeline, it should
94                     // rather be a file output error instead of a pipeline one.
95                     if (has_file_output) JUMP_OUT(UNEXPECTED_FILE_OUTPUT);
96                     // if no tokens between two pipelines (or before the first one)
97                     // should throw a pipeline error
98                     if (!has_token_last) JUMP_OUT(UNEXPECTED_PIPELINE);
99                     has_token_last = true;
100                    ++pcmd->num_commands;
101                    ++cur; // skip '|'
102                    break;
103                default:
104                    has_token_last = true;
105                    ++total_strings;
106                    skip_word(&cur, end); // skip that argument
107            }
108
109     if (total_strings == 0) {
110         // if there are no arguments but has ampersand or file input/output
111         // then we have an error
112         if (pcmd->is_background || has_file_input || has_file_output)
113             JUMP_OUT(EXPECT_COMMANDS);
114         // otherwise it's an empty line
115         goto PROCESS_SUCCESS;
116     }
117
118     // handle edge case where the command ends with a pipeline
119     // (not supporting line continuation)
120     if (!has_token_last) JUMP_OUT(UNEXPECTED_PIPELINE);
121 }
122 ++pcmd->num_commands;
123

```

```

146
147     const size_t start_of_array = offsetof(struct parsed_command, commands) + pcmd->num_commands *
148     sizeof(char **);
149     const size_t start_of_str = start_of_array + (pcmd->num_commands + total_strings) * sizeof(char *);
150     const size_t slen = end - start;
151
152     char *const new_buf = realloc(pcmd, start_of_str + slen + 1);
153     if (new_buf == NULL) goto PROCESS_ERROR;
154     pcmd = (struct parsed_command *) new_buf;
155
156     // copy string to the new place
157     char *const new_start = memcpy(new_buf + start_of_str, start, slen);
158
159     // second pass, put stuff in
160     // no need to check for error anymore
161     size_t cur_cmd = 0;
162     char **argv_ptr = (char **) (new_buf + start_of_array);
163
164     pcmd->commands[cur_cmd] = argv_ptr;
165     for (const char *cur = start; cur < end; skip_space(&cur, end)) {
166         switch (cur[0]) {
167             case '<':
168                 ++cur;
169                 skip_space(&cur, end);
170                 // store input file name into 'stdin_file'
171                 pcmd->stdin_file = new_start + (cur - start);
172                 skip_word(&cur, end);
173                 // at end of the input file name
174                 new_start[cur - start] = '\\0';
175                 break;
176             case '>':
177                 if (pcmd->is_file_append) ++cur; // skip another '>'
178                 ++cur;
179                 skip_space(&cur, end);
180                 // store output file name into 'stdout_file'
181                 pcmd->stdout_file = new_start + (cur - start);
182                 skip_word(&cur, end);
183                 // at end of the output file name
184                 new_start[cur - start] = '\\0';
185                 break;
186             case '|':
187                 // null-terminate the current argv
188                 *(argv_ptr++) = NULL;
189                 // store the next argv head
190                 pcmd->commands[++cur_cmd] = argv_ptr;
191                 ++cur;
192                 break;
193             default:
194                 // at start of the argument string
195                 // store it into the arguments array
196                 *(argv_ptr++) = new_start + (cur - start);
197                 skip_word(&cur, end);
198                 // at end of the argument string
199                 new_start[cur - start] = '\\0';
200         }
201     }
202     // null-terminate the last argv
203     *argv_ptr = NULL;
204
205     PROCESS_SUCCESS:
206     *result = pcmd;
207     return 0;
208     PROCESS_ERROR:
209     free(pcmd);
210     return ret_code;
211 }

```

#### 4.33.2.2 print\_parsed\_command()

```

void print_parsed_command (
    const struct parsed_command * cmd )

```

Definition at line 214 of file parser.c.

```

214     {
215     for (size_t i = 0; i < cmd->num_commands; ++i) {
216         for (char **arguments = cmd->commands[i]; *arguments != NULL; ++arguments)
217             printf("%s ", *arguments);
218     }

```



```

219         if (i == 0 && cmd->stdin_file != NULL)
220             printf("< %s ", cmd->stdin_file);
221
222         if (i == cmd->num_commands - 1) {
223             if (cmd->stdout_file != NULL)
224                 printf(cmd->is_file_append ? "» %s " : "> %s ", cmd->stdout_file);
225             } else printf("| ");
226         }
227         puts(cmd->is_background ? "& " : "");
228     }

```

#### 4.33.2.3 print\_parser\_errcode()

```

void print_parser_errcode (
    FILE * output,
    int err_code )

```

Definition at line 230 of file parser.c.

```

230     {
231         switch (err_code) {
232             case UNEXPECTED_FILE_INPUT:
233                 fprintf(output, "UNEXPECTED INPUT REDIRECTION TO A FILE\n");
234                 break;
235             case UNEXPECTED_FILE_OUTPUT:
236                 fprintf(output, "UNEXPECTED OUTPUT REDIRECTION TO A FILE\n");
237                 break;
238             case UNEXPECTED_PIPELINE:
239                 fprintf(output, "UNEXPECTED PIPE\n");
240                 break;
241             case UNEXPECTED_AMPERSAND:
242                 fprintf(output, "UNEXPECTED AMPERESAND\n");
243                 break;
244             case EXPECT_INPUT_FILENAME:
245                 fprintf(output, "COULD NOT FINE FILENAME FOR INPUT REDIRECTION \"><\n");
246                 break;
247             case EXPECT_OUTPUT_FILENAME:
248                 fprintf(output, "COULD NOT FIND FILENAME FOR OUTPUT REDIRECTION \"><\n");
249                 break;
250             case EXPECT_COMMANDS:
251                 fprintf(output, "COULD NOT FIND ANY COMMANDS OR ARGS\n");
252                 break;
253             default:
254                 break;
255         }
256     }

```

## 4.34 SRC/shell/shell.c File Reference

```

#include <string.h>
#include "../fs/fat_routines.h"
#include "../fs/fs_syscalls.h"
#include "../fs/fs_helpers.h"
#include "../kernel/kern_sys_calls.h"
#include "builtins.h"
#include "parser.h"
#include "shell_built_ins.h"
#include "stdlib.h"
#include "../kernel/stress.h"
#include <fcntl.h>
#include "../kernel/scheduler.h"
#include "../lib/Vec.h"
#include "Job.h"
#include "signal.h"
#include "lib/pennos-errno.h"

```



#### 4.34.1.1 MAX\_BUFFER\_SIZE

```
#define MAX_BUFFER_SIZE 4096
```

Definition at line 29 of file shell.c.

#### 4.34.1.2 MAX\_LINE\_BUFFER\_SIZE

```
#define MAX_LINE_BUFFER_SIZE 128
```

Definition at line 30 of file shell.c.

#### 4.34.1.3 PROMPT

```
#define PROMPT "$ "
```

Definition at line 26 of file shell.c.

### 4.34.2 Function Documentation

#### 4.34.2.1 execute\_command()

```
pid_t execute_command (  
    struct parsed_command * cmd )
```

Helper function to execute a parsed command from the shell. In particular, it spawns a child process to execute the command if the built-in should run as a separate process. Otherwise, it just calls the subroutine directly.

##### Parameters

<i>cmd</i>	the parsed command to execute, assumed non-null
------------	---

##### Returns

the created child id on successful spawn, 0 on successful subroutine call, -1 when nothing was called

Definition at line 257 of file shell.c.

```
257                                     {  
258  
259     // setup fds  
260     int input_fd = STDIN_FILENO; // standard fds  
261     int output_fd = STDOUT_FILENO;  
262
```

```

263 if (cmd->stdin_file != NULL) {
264     input_fd = s_open(cmd->stdin_file, F_READ);
265     if (input_fd < 0) {
266         input_fd = STDIN_FILENO; // reset to default
267     }
268 }
269
270 if (cmd->is_file_append) {
271     output_fd = s_open(cmd->stdout_file, F_APPEND);
272 } else {
273     output_fd = s_open(cmd->stdout_file, F_WRITE);
274 }
275 if (output_fd < 0) {
276     output_fd = STDOUT_FILENO; // reset to default
277 }
278
279 // check for independently scheduled processes
280 if (strcmp(cmd->commands[0][0], "cat") == 0) {
281     return s_spawn(u_cat, cmd->commands[0], input_fd, output_fd);
282 } else if (strcmp(cmd->commands[0][0], "sleep") == 0) {
283     return s_spawn(u_sleep, cmd->commands[0], input_fd, output_fd);
284 } else if (strcmp(cmd->commands[0][0], "busy") == 0) {
285     return s_spawn(u_busy, cmd->commands[0], input_fd, output_fd);
286 } else if (strcmp(cmd->commands[0][0], "echo") == 0) {
287     return s_spawn(u_echo, cmd->commands[0], input_fd, output_fd);
288 } else if (strcmp(cmd->commands[0][0], "ls") == 0) {
289     return s_spawn(u_ls, cmd->commands[0], input_fd, output_fd);
290 } else if (strcmp(cmd->commands[0][0], "touch") == 0) {
291     return s_spawn(u_touch, cmd->commands[0], input_fd, output_fd);
292 } else if (strcmp(cmd->commands[0][0], "mv") == 0) {
293     return s_spawn(u_mv, cmd->commands[0], input_fd, output_fd);
294 } else if (strcmp(cmd->commands[0][0], "cp") == 0) {
295     return s_spawn(u_cp, cmd->commands[0], input_fd, output_fd);
296 } else if (strcmp(cmd->commands[0][0], "rm") == 0) {
297     return s_spawn(u_rm, cmd->commands[0], input_fd, output_fd);
298 } else if (strcmp(cmd->commands[0][0], "chmod") == 0) {
299     return s_spawn(u_chmod, cmd->commands[0], input_fd, output_fd);
300 } else if (strcmp(cmd->commands[0][0], "ps") == 0) {
301     return s_spawn(u_ps, cmd->commands[0], input_fd, output_fd);
302 } else if (strcmp(cmd->commands[0][0], "kill") == 0) {
303     return s_spawn(u_kill, cmd->commands[0], input_fd, output_fd);
304 } else if (strcmp(cmd->commands[0][0], "zombify") == 0) {
305     return s_spawn(u_zombify, cmd->commands[0], input_fd, output_fd);
306 } else if (strcmp(cmd->commands[0][0], "orphanify") == 0) {
307     return s_spawn(u_orphanify, cmd->commands[0], input_fd, output_fd);
308 } else if (strcmp(cmd->commands[0][0], "hang") == 0) {
309     return s_spawn(hang, cmd->commands[0], input_fd, output_fd);
310 } else if (strcmp(cmd->commands[0][0], "nohang") == 0) {
311     return s_spawn(nohang, cmd->commands[0], input_fd, output_fd);
312 } else if (strcmp(cmd->commands[0][0], "recur") == 0) {
313     return s_spawn(recur, cmd->commands[0], input_fd, output_fd);
314 } else if (strcmp(cmd->commands[0][0], "crash") == 0) {
315     return s_spawn(crash, cmd->commands[0], input_fd, output_fd);
316 }
317
318 // check for sub-routines
319 if (strcmp(cmd->commands[0][0], "nice") == 0) {
320     u_nice(cmd->commands[0]);
321     return 0;
322 } else if (strcmp(cmd->commands[0][0], "nice_pid") == 0) {
323     u_nice_pid(cmd->commands[0]);
324     return 0;
325 } else if (strcmp(cmd->commands[0][0], "man") == 0) {
326     u_man(cmd->commands[0]);
327     return 0;
328 } else if (strcmp(cmd->commands[0][0], "bg") == 0) {
329     u_bg(cmd->commands[0]);
330     return 0;
331 } else if (strcmp(cmd->commands[0][0], "fg") == 0) {
332     u_fg(cmd->commands[0]);
333     return 0;
334 } else if (strcmp(cmd->commands[0][0], "jobs") == 0) {
335     u_jobs(cmd->commands[0]);
336     return 0;
337 } else if (strcmp(cmd->commands[0][0], "logout") == 0) {
338     u_logout(cmd->commands[0]);
339     return 0;
340 }
341
342 // otherwise, try to run command as a script
343 int script_fd_open = s_open(cmd->commands[0][0], F_READ);
344 if (script_fd_open < 0) { // if not a file, just move on
345     return -1;
346 }
347 if (has_executable_permission(script_fd_open) != 1) {
348     s_close(script_fd_open);
349     return -1;

```

```

350 } else {
351     script_fd = script_fd_open; // update global
352     input_fd_script = input_fd;
353     output_fd_script = output_fd;
354
355     char* script_argv[] = {cmd->commands[0][0], NULL};
356     pid_t wait_on = s_spawn(u_read_and_execute_script, script_argv, input_fd,
357                             output_fd);
358     int status;
359     s_waitpid(wait_on, &status, false); // wait for script to finish
360     script_fd = -1; // reset global
361     input_fd_script = STDIN_FILENO;
362     output_fd_script = STDOUT_FILENO;
363     s_close(script_fd_open);
364     return 0;
365 }
366
367 return -1; // no matches case
368 }

```

#### 4.34.2.2 fill\_buffer\_until\_full\_or\_newline()

```

void fill_buffer_until_full_or_newline (
    int fd,
    char * buffer )

```

Helper function that fills a buffer with characters read from a given file descriptor until the buffer is full (rare and impractical case), a newline is encountered, or EOF is reached.

##### Parameters

<i>fd</i>	the file descriptor to read from, assumed to be open
<i>buffer</i>	the buffer to fill with characters

Definition at line 113 of file shell.c.

```

113                                     {
114     int i = 0;
115     char currChar;
116     while (i < MAX_LINE_BUFFER_SIZE - 1) {
117         int bytes_read = s_read(fd, &currChar, 1);
118         if (bytes_read <= 0 || currChar == '\n') { // EOF or newline cases
119             break;
120         }
121         buffer[i] = currChar;
122         i++;
123     }
124     buffer[i] = '\0'; // Null-terminate the string, replaces \n
125 }

```

#### 4.34.2.3 free\_job\_ptr()

```

void free_job_ptr (
    void * ptr )

```

Definition at line 94 of file shell.c.

```

94                                     {
95     job* job_ptr = (job*)ptr;
96     free(job_ptr->pids);
97     free(job_ptr);
98 }

```

#### 4.34.2.4 setup\_terminal\_signal\_handlers()

```
void setup_terminal_signal_handlers (
    void )
```

Definition at line 73 of file shell.c.

```
73                                     {
74     struct sigaction sa_int = {0};
75     sa_int.sa_handler = shell_sigint_handler;
76     sigemptyset(&sa_int.sa_mask);
77     sa_int.sa_flags = SA_RESTART;
78     sigaction(SIGINT, &sa_int, NULL);
79
80     struct sigaction sa_stp = {0};
81     sa_stp.sa_handler = shell_sigstp_handler;
82     sigemptyset(&sa_stp.sa_mask);
83     sa_stp.sa_flags = SA_RESTART;
84     sigaction(SIGTSTP, &sa_stp, NULL);
85 }
```

#### 4.34.2.5 shell()

```
void* shell (
    void * input )
```

The main shell function that runs the shell. This is run via an `s_spawn` call from `init`'s process. It prompts the user for builtins and scripts to run.

##### Parameters

<i>input</i>	unused
--------------	--------

Definition at line 375 of file shell.c.

```
375                                     {
376
377     job_list = vec_new(0, free_job_ptr);
378
379     setup_terminal_signal_handlers();
380
381     while (true) {
382         int status;
383         pid_t child_pid;
384         while ((child_pid = s_waitpid(-1, &status, true)) > 0) {
385             // Child process has completed, no need to do anything special
386             // The s_waitpid function already handles cleanup
387         }
388
389         // prompt
390         if (s_write(STDOUT_FILENO, PROMPT, strlen(PROMPT)) < 0) {
391             u_perror("prompt write error");
392             break;
393         }
394
395         // parse user input
396         char buffer[MAX_BUFFER_SIZE];
397         ssize_t user_input = s_read(STDIN_FILENO, buffer, MAX_BUFFER_SIZE);
398         if (user_input < 0) {
399             u_perror("shell read error");
400             break;
401         } else if (user_input == 0) { // EOF case
402             s_shutdown_pennos();
403             break;
404         }
405
406         buffer[user_input] = '\0';
407         if (buffer[user_input - 1] == '\n') {
408             buffer[user_input - 1] = '\0';
409         }
410     }
```

```

411     struct parsed_command* cmd = NULL;
412     int cmd_parse_res = parse_command(buffer, &cmd);
413     if (cmd_parse_res != 0 || cmd == NULL) {
414         P_ERRNO = P_EPARSE;
415         u_perror("parse_command");
416         continue;
417     }
418
419     // handle the command
420     if (cmd->num_commands == 0) {
421         free(cmd);
422         continue;
423     }
424
425     child_pid = execute_command(cmd);
426     if (child_pid < 0) {
427         free(cmd);
428         continue;
429     } else if (child_pid == 0) {
430         free(cmd);
431         continue;
432     }
433
434     // If background, add the process to the job list.
435     if (cmd->is_background) {
436         // Create a new job entry.
437         job* new_job = malloc(sizeof(job));
438         if (new_job == NULL) {
439             perror("Error: mallocing new_job failed");
440             free(cmd);
441             continue;
442         }
443         new_job->id = next_job_id++;
444         new_job->pgid = child_pid; // For single commands, child's pid = pgid.
445         new_job->num_pids = 1;
446         new_job->pids = malloc(sizeof(pid_t));
447         if (new_job->pids == NULL) {
448             perror("Error: mallocing new_job->pids failed");
449             free(new_job);
450             free(cmd);
451             continue;
452         }
453         new_job->pids[0] = child_pid;
454         new_job->state = RUNNING;
455         new_job->cmd = cmd; // Retain command info; do not free here.
456         new_job->finished_count = 0;
457         vec_push_back(&job_list, new_job);
458
459         // Print job control information in the format: "[job_id] child_pid"
460         char msg[128];
461         snprintf(msg, sizeof(msg), "[%lu] %d\n", (unsigned long)new_job->id,
462                 child_pid);
463         s_write(STDOUT_FILENO, msg, strlen(msg));
464     } else {
465         // Foreground execution.
466         current_fg_pid = child_pid;
467         int status;
468         s_waitpid(child_pid, &status, false);
469         current_fg_pid = 2;
470     }
471 }
472
473 vec_destroy(&job_list);
474 s_exit();
475 return 0;
476 }

```

#### 4.34.2.6 shell\_sigint\_handler()

```

void shell_sigint_handler (
    int sig )

```

Definition at line 52 of file shell.c.

```

52     {
53         // If there's a foreground process, forward SIGINT (terminate) to it, as long
54         // as it's not the shell current_fg_pid will also never be 1 (INIT)
55         if (current_fg_pid != 2) {
56             s_kill(current_fg_pid, 2); // P_SIGTERM

```

```

57     }
58
59     s_write(STDOUT_FILENO, "\n", 1);
60 }

```

#### 4.34.2.7 shell\_sigstp\_handler()

```

void shell_sigstp_handler (
    int sig )

```

Definition at line 63 of file shell.c.

```

63     {
64         // If there's a foreground process, forward SIGTSTP (stop) to it
65         if (current_fg_pid != 2) {
66             s_kill(current_fg_pid, 0); // P_SIGSTOP
67         }
68
69         s_write(STDOUT_FILENO, "\n", 1);
70     }

```

#### 4.34.2.8 u\_execute\_command()

```

pid_t u_execute_command (
    struct parsed_command * cmd )

```

Helper function that will execute a given command so long as it's one of the built-ins. Notably, its output and input are determined by the spawning script process.

##### Parameters

<i>cmd</i>	the parsed command to try executing
------------	-------------------------------------

##### Returns

the pid of the process if one was spawned, 0 if a routine was run or -1 if not matches found

Definition at line 136 of file shell.c.

```

136     {
137
138         // check for independently scheduled processes
139         if (strcmp(cmd->commands[0][0], "cat") == 0) {
140             return s_spawn(u_cat, cmd->commands[0], input_fd_script, output_fd_script);
141         } else if (strcmp(cmd->commands[0][0], "sleep") == 0) {
142             return s_spawn(u_sleep, cmd->commands[0], input_fd_script, output_fd_script);
143         } else if (strcmp(cmd->commands[0][0], "busy") == 0) {
144             return s_spawn(u_busy, cmd->commands[0], input_fd_script, output_fd_script);
145         } else if (strcmp(cmd->commands[0][0], "echo") == 0) {
146             return s_spawn(u_echo, cmd->commands[0], input_fd_script, output_fd_script);
147         } else if (strcmp(cmd->commands[0][0], "ls") == 0) {
148             return s_spawn(u_ls, cmd->commands[0], input_fd_script, output_fd_script);
149         } else if (strcmp(cmd->commands[0][0], "touch") == 0) {
150             return s_spawn(u_touch, cmd->commands[0], input_fd_script, output_fd_script);
151         } else if (strcmp(cmd->commands[0][0], "mv") == 0) {
152             return s_spawn(u_mv, cmd->commands[0], input_fd_script, output_fd_script);
153         } else if (strcmp(cmd->commands[0][0], "cp") == 0) {
154             return s_spawn(u_cp, cmd->commands[0], input_fd_script, output_fd_script);
155         } else if (strcmp(cmd->commands[0][0], "rm") == 0) {
156             return s_spawn(u_rm, cmd->commands[0], input_fd_script, output_fd_script);
157         } else if (strcmp(cmd->commands[0][0], "chmod") == 0) {

```



```

158     return s_spawn(u_chmod, cmd->commands[0], input_fd_script, output_fd_script);
159 } else if (strcmp(cmd->commands[0][0], "ps") == 0) {
160     return s_spawn(u_ps, cmd->commands[0], input_fd_script, output_fd_script);
161 } else if (strcmp(cmd->commands[0][0], "kill") == 0) {
162     return s_spawn(u_kill, cmd->commands[0], input_fd_script, output_fd_script);
163 } else if (strcmp(cmd->commands[0][0], "zombify") == 0) {
164     return s_spawn(u_zombify, cmd->commands[0], input_fd_script, output_fd_script);
165 } else if (strcmp(cmd->commands[0][0], "orphanify") == 0) {
166     return s_spawn(u_orphanify, cmd->commands[0], input_fd_script, output_fd_script);
167 } else if (strcmp(cmd->commands[0][0], "hang") == 0) {
168     return s_spawn(hang, cmd->commands[0], input_fd_script, output_fd_script);
169 } else if (strcmp(cmd->commands[0][0], "nohang") == 0) {
170     return s_spawn(nohang, cmd->commands[0], input_fd_script, output_fd_script);
171 } else if (strcmp(cmd->commands[0][0], "recur") == 0) {
172     return s_spawn(recur, cmd->commands[0], input_fd_script, output_fd_script);
173 } else if (strcmp(cmd->commands[0][0], "crash") == 0) {
174     return s_spawn(crash, cmd->commands[0], input_fd_script, output_fd_script);
175 }
176
177 // check for sub-routines
178 if (strcmp(cmd->commands[0][0], "nice") == 0) {
179     u_nice(cmd->commands[0]);
180     return 0;
181 } else if (strcmp(cmd->commands[0][0], "nice_pid") == 0) {
182     u_nice_pid(cmd->commands[0]);
183     return 0;
184 } else if (strcmp(cmd->commands[0][0], "man") == 0) {
185     u_man(cmd->commands[0]);
186     return 0;
187 } else if (strcmp(cmd->commands[0][0], "bg") == 0) {
188     u_bg(cmd->commands[0]);
189     return 0;
190 } else if (strcmp(cmd->commands[0][0], "fg") == 0) {
191     u_fg(cmd->commands[0]);
192     return 0;
193 } else if (strcmp(cmd->commands[0][0], "jobs") == 0) {
194     u_jobs(cmd->commands[0]);
195     return 0;
196 } else if (strcmp(cmd->commands[0][0], "logout") == 0) {
197     u_logout(cmd->commands[0]);
198     return 0;
199 } else {
200     return -1; // no matches, no scripts now
201 }
202
203 return 0;
204 }

```

#### 4.34.2.9 u\_read\_and\_execute\_script()

```

void* u_read_and_execute_script (
    void * arg )

```

Helper function that reads a script file line by line, parses each line as a command, and executes it.

##### Parameters

<i>arg</i>	standard {function name, NULL} args
------------	-------------------------------------

##### Returns

NULL

Definition at line 214 of file shell.c.

```

214     {
215     // read the script line by line, parse each line, and execute the command
216     while (true) {
217         char buffer[MAX_LINE_BUFFER_SIZE];
218         fill_buffer_until_full_or_newline(script_fd, buffer);
219         if (buffer[0] == '\0') {

```

```

220     break; // EOF case
221 }
222
223 // parse the command
224 struct parsed_command* cmd = NULL;
225 int parse_result = parse_command(buffer, &cmd);
226 if (parse_result != 0 || cmd == NULL) {
227     P_ERRNO = P_EPARSE;
228     u_perror("parse_command");
229     free(cmd);
230 }
231
232 // execute the command
233 pid_t child_pid = u_execute_command(cmd);
234 if (child_pid > 0) { // if process was spawned, wait for it to finish
235     int status;
236     s_waitpid(child_pid, &status, false);
237 } else if (child_pid < 0) { // nothing spawning so safe to free cmd
238     free(cmd);
239 }
240 }
241
242 s_exit(); // exit the script
243 return NULL;
244 }

```

### 4.34.3 Variable Documentation

#### 4.34.3.1 current\_fg\_pid

pid\_t current\_fg\_pid [extern]

Definition at line 31 of file kern\_sys\_calls.c.

#### 4.34.3.2 input\_fd\_script

int input\_fd\_script = -1

Definition at line 42 of file shell.c.

#### 4.34.3.3 job\_list

Vec job\_list

Definition at line 37 of file shell.c.

#### 4.34.3.4 next\_job\_id

jid\_t next\_job\_id = 1

Definition at line 38 of file shell.c.

#### 4.34.3.5 output\_fd\_script

```
int output_fd_script = -1
```

Definition at line 43 of file shell.c.


#### 4.34.3.6 script\_fd

```
int script_fd = -1
```

Definition at line 41 of file shell.c.

## 4.35 SRC/shell/shell.h File Reference

This graph shows which files directly or indirectly include this file:



## Functions

- void \* [shell](#) (void \*input)

*The main shell function that runs the shell. This is run via an s\_spawn call from init's process. It prompts the user for builtins and scripts to run.*

### 4.35.1 Function Documentation

#### 4.35.1.1 shell()

```
void* shell (  
    void * input )
```

The main shell function that runs the shell. This is run via an s\_spawn call from init's process. It prompts the user for builtins and scripts to run.

## Parameters

<i>input</i>	unused
--------------	--------

Definition at line 375 of file shell.c.

```

375         {
376
377     job_list = vec_new(0, free_job_ptr);
378
379     setup_terminal_signal_handlers();
380
381     while (true) {
382         int status;
383         pid_t child_pid;
384         while ((child_pid = s_waitpid(-1, &status, true)) > 0) {
385             // Child process has completed, no need to do anything special
386             // The s_waitpid function already handles cleanup
387         }
388
389         // prompt
390         if (s_write(STDOUT_FILENO, PROMPT, strlen(PROMPT)) < 0) {
391             u_perror("prompt write error");
392             break;
393         }
394
395         // parse user input
396         char buffer[MAX_BUFFER_SIZE];
397         ssize_t user_input = s_read(STDIN_FILENO, buffer, MAX_BUFFER_SIZE);
398         if (user_input < 0) {
399             u_perror("shell read error");
400             break;
401         } else if (user_input == 0) { // EOF case
402             s_shutdown_pennos();
403             break;
404         }
405
406         buffer[user_input] = '\0';
407         if (buffer[user_input - 1] == '\n') {
408             buffer[user_input - 1] = '\0';
409         }
410
411         struct parsed_command* cmd = NULL;
412         int cmd_parse_res = parse_command(buffer, &cmd);
413         if (cmd_parse_res != 0 || cmd == NULL) {
414             P_ERRNO = P_EPARSE;
415             u_perror("parse_command");
416             continue;
417         }
418
419         // handle the command
420         if (cmd->num_commands == 0) {
421             free(cmd);
422             continue;
423         }
424
425         child_pid = execute_command(cmd);
426         if (child_pid < 0) {
427             free(cmd);
428             continue;
429         } else if (child_pid == 0) {
430             free(cmd);
431             continue;
432         }
433
434         // If background, add the process to the job list.
435         if (cmd->is_background) {
436             // Create a new job entry.
437             job* new_job = malloc(sizeof(job));
438             if (new_job == NULL) {
439                 perror("Error: mallocing new_job failed");
440                 free(cmd);
441                 continue;
442             }
443             new_job->id = next_job_id++;
444             new_job->pgid = child_pid; // For single commands, child's pid = pgid.
445             new_job->num_pids = 1;
446             new_job->pids = malloc(sizeof(pid_t));
447             if (new_job->pids == NULL) {
448                 perror("Error: mallocing new_job->pids failed");
449                 free(new_job);
450                 free(cmd);
451                 continue;
452             }

```

```

453     new_job->pids[0] = child_pid;
454     new_job->state = RUNNING;
455     new_job->cmd = cmd; // Retain command info; do not free here.
456     new_job->finished_count = 0;
457     vec_push_back(&job_list, new_job);
458
459     // Print job control information in the format: "[job_id] child_pid"
460     char msg[128];
461     snprintf(msg, sizeof(msg), "[%lu] %d\n", (unsigned long)new_job->id,
462             child_pid);
463     s_write(STDOUT_FILENO, msg, strlen(msg));
464 } else {
465     // Foreground execution.
466     current_fg_pid = child_pid;
467     int status;
468     s_waitpid(child_pid, &status, false);
469     current_fg_pid = 2;
470 }
471 }
472
473 vec_destroy(&job_list);
474 s_exit();
475 return 0;
476 }

```


## 4.36 SRC/shell/shell\_built\_ins.c File Reference

```

#include "shell_built_ins.h"
#include <stdbool.h>
#include <stddef.h>
#include <string.h>
#include <sys/types.h>
#include "../fs/fat_routines.h"
#include "../fs/fs_syscalls.h"
#include "../kernel/kern_sys_calls.h"
#include "../kernel/scheduler.h"
#include "../lib/Vec.h"
#include "../lib/spthread.h"
#include <errno.h>
#include <stdio.h>
#include <stdlib.h>

```

Include dependency graph for shell\_built\_ins.c:



## Functions

- void \* **u\_cat** (void \*arg)  
The usual *cat* program.
- void \* **u\_sleep** (void \*arg)  
Sleep for *n* seconds.
- void \* **u\_busy** (void \*arg)

- Busy wait indefinitely. It can only be interrupted via signals.*

  - void \* [u\\_echo](#) (void \*arg)  
*Echo back an input string.*
  - void \* [u\\_ls](#) (void \*arg)  
*Lists all files in the working directory. For extra credit, it should support relative and absolute file paths.*
  - void \* [u\\_chmod](#) (void \*arg)  
*Change permissions of a file. There's no need to error if a permission being added already exists, or if a permission being removed is already not granted.*
  - void \* [u\\_touch](#) (void \*arg)  
*For each file, create an empty file if it doesn't exist, else update its timestamp.*
  - void \* [u\\_mv](#) (void \*arg)  
*Rename a file. If the `dst_file` file already exists, overwrite it.*
  - void \* [u\\_cp](#) (void \*arg)
  - void \* [u\\_rm](#) (void \*arg)  
*Remove a list of files. Treat each file in the list as a separate transaction. (i.e. if removing file1 fails, still attempt to remove file2, file3, etc.)*
  - void \* [u\\_ps](#) (void \*arg)  
*List all processes on PennOS, displaying PID, PPID, priority, status, and command name.*
  - void \* [u\\_kill](#) (void \*arg)  
*Sends a specified signal to a list of processes. If a signal name is not specified, default to "term". Valid signals are -term, -stop, and -cont.*
  - void \* [u\\_nice](#) (void \*arg)  
*Spawn a new process for `command` and set its priority to `priority`.*
  - void \* [u\\_nice\\_pid](#) (void \*arg)  
*Adjust the priority level of an existing process.*
  - void \* [u\\_man](#) (void \*arg)  
*Lists all available commands.*
  - void \* [u\\_bg](#) (void \*arg)  
*Resumes the most recently stopped job in the background, or the job specified by `job_id`.*
  - void \* [u\\_fg](#) (void \*arg)  
*Brings the most recently stopped or background job to the foreground, or the job specified by `job_id`.*
  - void \* [u\\_jobs](#) (void \*arg)  
*Lists all jobs.*
  - void \* [u\\_logout](#) (void \*arg)  
*Exits the shell and shutdowns PennOS.*
  - void \* [zombie\\_child](#) (void \*arg)  
*Helper for zombify.*
  - void \* [u\\_zombify](#) (void \*arg)  
*Used to test zombifying functionality of your kernel.*
  - void \* [orphan\\_child](#) (void \*arg)  
*Helper for orphanify.*
  - void \* [u\\_orphanify](#) (void \*arg)  
*Used to test orphanifying functionality of your kernel.*

## Variables

- void (\*)(void \*) [get\\_associated\\_ufunc](#) (char \*func)  
*Helper function to get the associated "u-version" of a function given its standalone version. As a concrete example, if we pass in "cat", we will output "u\_cat".*

## 4.36.1 Function Documentation

### 4.36.1.1 orphan\_child()

```
void* orphan_child (
    void * arg )
```

Helper for orphanify.

Definition at line 314 of file shell\_built\_ins.c.

```
314     {
315     while (1)
316     ;
317     s_exit();
318 }
```

### 4.36.1.2 u\_bg()

```
void* u_bg (
    void * arg )
```

Resumes the most recently stopped job in the background, or the job specified by job\_id.

Example Usage: bg Example Usage: bg 2 (job\_id is 2)

Definition at line 271 of file shell\_built\_ins.c.

```
271     {
272     // TODO --> implement bg
273     return NULL;
274 }
```

### 4.36.1.3 u\_busy()

```
void* u_busy (
    void * arg )
```

Busy wait indefinitely. It can only be interrupted via signals.

Example Usage: busy

Definition at line 53 of file shell\_built\_ins.c.

```
53     {
54     while (1)
55     ;
56     s_exit();
57     return NULL;
58 }
```

#### 4.36.1.4 u\_cat()

```
void* u_cat (
    void * arg )
```

The usual `cat` program.

If `files` `arg` is provided, concatenate these files and print to stdout. If `files` `arg` is *not* provided, read from stdin and print back to stdout.

Example Usage: `cat f1 f2` (concatenates `f1` and `f2` and print to stdout) Example Usage: `cat f1 f2 < f3` (concatenates `f1` and `f2` and prints to stdout, ignores `f3`) Example Usage: `cat < f3` (concatenates `f3`, prints to stdout)

Definition at line 28 of file `shell_built_ins.c`.

```
28     {
29         cat(arg);
30         s_exit();
31         return NULL;
32     }
```

#### 4.36.1.5 u\_chmod()

```
void* u_chmod (
    void * arg )
```

Change permissions of a file. There's no need to error if a permission being added already exists, or if a permission being removed is already not granted.

Print appropriate error message if:

- `file` is not a file that exists
- `perms` is invalid

Example Usage: `chmod +x file` (adds executable permission to file) Example Usage: `chmod +rw file` (adds read + write permissions to file) Example Usage: `chmod -wx file` (removes write + executable permissions from file)

Definition at line 72 of file `shell_built_ins.c`.

```
72     {
73         chmod(arg);
74         s_exit();
75         return NULL;
76     }
```



#### 4.36.1.6 u\_cp()

```
void* u_cp (
    void * arg )
```

Copy a file. If the `dst_file` file already exists, overwrite it.

Print appropriate error message if:

- `src_file` is not a file that exists
- `src_file` does not have read permissions
- `dst_file` file already exists but does not have write permissions

Example Usage: `cp src_file dst_file`

Definition at line 90 of file `shell_built_ins.c`.

```
90     {
91         cp(arg);
92         s_exit();
93         return NULL;
94     }
```

#### 4.36.1.7 u\_echo()

```
void* u_echo (
    void * arg )
```

Echo back an input string.

Example Usage: `echo Hello World`

Definition at line 60 of file `shell_built_ins.c`.

```
60     {
61         s_echo(arg);
62         s_exit();
63         return NULL;
64     }
```

#### 4.36.1.8 u\_fg()

```
void* u_fg (
    void * arg )
```

Brings the most recently stopped or background job to the foreground, or the job specified by `job_id`.

Example Usage: `fg` Example Usage: `fg 2` (`job_id` is 2)

Definition at line 276 of file `shell_built_ins.c`.

```
276     {
277         // TODO --> implement fg
278         return NULL;
279     }
```

#### 4.36.1.9 u\_jobs()

```
void* u_jobs (
    void * arg )
```

Lists all jobs.

Example Usage: jobs

Definition at line 281 of file shell\_built\_ins.c.

```
281     {
282     // TODO --> implement jobs
283     return NULL;
284 }
```

#### 4.36.1.10 u\_kill()

```
void* u_kill (
    void * arg )
```

Sends a specified signal to a list of processes. If a signal name is not specified, default to "term". Valid signals are -term, -stop, and -cont.

Example Usage: kill 1 2 3 (sends term to processes 1, 2, and 3) Example Usage: kill -term 1 2 (sends term to processes 1 and 2) Example Usage: kill -stop 1 2 (sends stop to processes 1 and 2) Example Usage: kill -cont 1 (sends cont to process 1)

Definition at line 108 of file shell\_built\_ins.c.

```
108     {
109     char** argv = (char**)arg;
110     int sig = 2; // Default signal: term (2)
111     int start_index = 1; // Start after the "kill" command word.
112     char err_buf[128];
113
114     // Check if the first argument specifies a signal
115     if (argv[start_index] && argv[start_index][0] == '-') {
116         if (strcmp(argv[start_index], "-term") == 0) {
117             sig = 2;
118         } else if (strcmp(argv[start_index], "-stop") == 0) {
119             sig = 0;
120         } else if (strcmp(argv[start_index], "-cont") == 0) {
121             sig = 1;
122         } else {
123             // Construct error message
124             s_exit();
125             return NULL;
126         }
127         start_index++;
128     }
129
130     // Process each PID argument using strtol
131     for (int i = start_index; argv[i] != NULL; i++) {
132         char* endptr;
133         long pid_long = strtol(argv[i], &endptr, 10);
134         if (*endptr != '\0' || pid_long <= 0) {
135             snprintf(err_buf, 128, "Invalid PID: %s\n", argv[i]);
136             s_write(STDERR_FILENO, err_buf, strlen(err_buf));
137             continue;
138         }
139         pid_t pid = (pid_t)pid_long;
140         if (s_kill(pid, sig) < 0) {
141             snprintf(err_buf, 128, "b_kill error on PID %d\n", pid);
142             s_write(STDERR_FILENO, err_buf, strlen(err_buf));
143         }
144     }
145     s_exit();
146     return NULL;
147 }
```

#### 4.36.1.11 u\_logout()

```
void* u_logout (
    void * arg )
```

Exits the shell and shutdown PennOS.

Example Usage: logout

Definition at line 286 of file shell\_built\_ins.c.

```
286     {
287         s_shutdown_pennos();
288         return NULL;
289     }
```

#### 4.36.1.12 u\_ls()

```
void* u_ls (
    void * arg )
```

Lists all files in the working directory. For extra credit, it should support relative and absolute file paths.

Example Usage: ls (regular credit) Example Usage: ls ../../foo./bar/sample (only for EC)

Definition at line 66 of file shell\_built\_ins.c.

```
66     {
67         ls(arg);
68         s_exit();
69         return NULL;
70     }
```

#### 4.36.1.13 u\_man()

```
void* u_man (
    void * arg )
```

Lists all available commands.

Example Usage: man

Definition at line 233 of file shell\_built\_ins.c.

```
233     {
234         const char* man_string =
235             "cat f1 f2 ...      : concatenates provided files (if none, reads from "
236             "std in), and writes to std out\n"
237             "sleep n              : sleeps for n seconds\n"
238             "busy                  : busy waits indefinitely\n"
239             "echo str              : echoes back the input string str\n"
240             "ls                    : lists all files in the working directory\n"
241             "touch f1 f2 ...        : for each file, creates empty file if it doesn't "
242             "exist yet, otherwise updates its timestamp\n"
243             "mv f1 f2              : renames f1 to f2 (overwrites f2 if it exists)\n"
244             "cp f1 f2              : copies f1 to f2 (overwrites f2 if it exists)\n"
245             "rm f1 f2 ...          : removes the input list of files\n"
246             "chmod +_ f1           : changes f1 permissions to +_ specifications "
247             " (+x, +rw, etc)\n"
248             "ps                    : lists all processes on PennOS, displaying PID, "
249             "PPID, priority, status, and command name\n"
250             "kill (-_) pid1 pid2 : sends specified signal (term default) to list "
251             "of processes\n"
```

```

252     "nice n command      : spawns a new process for command and sets its "
253     "priority to n\n"
254     "nice_pid n pid      : adjusts the priority level of an existing "
255     "process to n\n"
256     "man                  : lists all available commands in PennOS\n"
257     "bg                   : resumes most recently stopped process in "
258     "background or the one specified by job_id\n"
259     "fg                   : brings most recently stopped or background job "
260     "to foreground or the one specifed by job_id\n"
261     "jobs                 : lists all jobs\n"
262     "logout               : exits the shell and shuts down PennOS\n"
263     "zombify              : creates a child process that becomes a zombie\n"
264     "orphanify            : creates a child process that becomes an "
265     "orphan\n";
266
267     s_write(STDOUT_FILENO, man_string, strlen(man_string));
268     return NULL;
269 }

```

#### 4.36.1.14 u\_mv()

```

void* u_mv (
    void * arg )

```

Rename a file. If the `dst_file` file already exists, overwrite it.

Print appropriate error message if:

- `src_file` is not a file that exists
- `src_file` does not have read permissions
- `dst_file` file already exists but does not have write permissions

Example Usage: `mv src_file dst_file`

Definition at line 84 of file `shell_built_ins.c`.

```

84     {
85     mv(arg);
86     s_exit();
87     return NULL;
88 }

```

#### 4.36.1.15 u\_nice()

```

void* u_nice (
    void * arg )

```

Spawn a new process for `command` and set its priority to `priority`.

Example Usage: `nice 2 cat f1 f2 f3` (spawns cat with priority 2)

Definition at line 194 of file `shell_built_ins.c`.

```

194     {
195     char* endptr;
196     errno = 0;
197     int new_priority = (int)strtol(((char**)arg)[1], &endptr, 10);
198     if (*endptr != '\0' || errno != 0 || new_priority > 2 ||
199         new_priority < 0) { // error catch
200         return NULL;
201     }

```

```

202
203 char* command = ((char**)arg)[2];
204 void* (*ufunc)(void*) = get_associated_ufunc(command);
205 if (ufunc == NULL) {
206     return NULL; // no matches, don't spawn
207 }
208
209 pid_t new_proc_pid = s_spawn(ufunc, &((char**)arg)[2], 0, 1);
210
211 if (new_proc_pid != -1) { // non-error case
212     s_nice(new_proc_pid, new_priority);
213 }
214
215 return NULL;
216 }

```

#### 4.36.1.16 u\_nice\_pid()

```

void* u_nice_pid (
    void * arg )

```

Adjust the priority level of an existing process.

Example Usage: nice\_pid 0 123 (sets priority 0 to PID 123)

Definition at line 218 of file shell\_built\_ins.c.

```

218 {
219     char* endptr;
220     errno = 0;
221     int new_priority = (int)strtol(((char**)arg)[1], &endptr, 10);
222     if (*endptr != '\0' || errno != 0) { // error catch
223         return NULL;
224     }
225     pid_t pid = (pid_t)strtol(((char**)arg)[2], &endptr, 10);
226     if (*endptr != '\0' || errno != 0) {
227         return NULL;
228     }
229     s_nice(pid, new_priority);
230     return NULL;
231 }

```

#### 4.36.1.17 u\_orphanify()

```

void* u_orphanify (
    void * arg )

```

Used to test orphanifying functionality of your kernel.

Example Usage: orphanify

Definition at line 320 of file shell\_built\_ins.c.

```

320 {
321     char* orphan_child_argv[] = {"orphan_child", NULL};
322     s_spawn(orphan_child, orphan_child_argv, STDIN_FILENO, STDOUT_FILENO);
323     s_exit();
324     return NULL;
325 }

```

**4.36.1.18 u\_ps()**

```
void* u_ps (
    void * arg )
```

List all processes on PennOS, displaying PID, PPID, priority, status, and command name.

Example Usage: ps

Definition at line 102 of file shell\_built\_ins.c.

```
102     {
103     s_ps(arg);
104     s_exit();
105     return NULL;
106 }
```

**4.36.1.19 u\_rm()**

```
void* u_rm (
    void * arg )
```

Remove a list of files. Treat each file in the list as a separate transaction. (i.e. if removing file1 fails, still attempt to remove file2, file3, etc.)

Print appropriate error message if:

- file is not a file that exists

Example Usage: rm f1 f2 f3 f4 f5

Definition at line 96 of file shell\_built\_ins.c.

```
96     {
97     rm(arg);
98     s_exit();
99     return NULL;
100 }
```

**4.36.1.20 u\_sleep()**

```
void* u_sleep (
    void * arg )
```

Sleep for n seconds.

Note that you'll have to convert the number of seconds to the correct number of ticks.

Example Usage: sleep 10

Definition at line 34 of file shell\_built\_ins.c.

```
34     {
35     char* endptr;
36     errno = 0;
37     if (((char**)arg)[1] == NULL) { // no arg case
38         s_exit();
39         return NULL;
40     }
41     int sleep_secs = (int)strtol(((char**)arg)[1], &endptr, 10);
42     if (*endptr != '\0' || errno != 0 || sleep_secs <= 0) {
43         s_exit();
44         return NULL;
45     }
46
47     int sleep_ticks = sleep_secs * 10;
48     s_sleep(sleep_ticks);
49     s_exit();
50     return NULL;
51 }
```

#### 4.36.1.21 u\_touch()

```
void* u_touch (
    void * arg )
```

For each file, create an empty file if it doesn't exist, else update its timestamp.

Example Usage: touch f1 f2 f3 f4 f5

Definition at line 78 of file shell\_built\_ins.c.

```
78     {
79     touch(arg);
80     s_exit();
81     return NULL;
82 }
```

#### 4.36.1.22 u\_zombify()

```
void* u_zombify (
    void * arg )
```

Used to test zombifying functionality of your kernel.

Example Usage: zombify

Definition at line 303 of file shell\_built\_ins.c.

```
303     {
304     char* zombie_child_argv[] = {"zombie_child", NULL};
305     s_spawn(zombie_child, zombie_child_argv, STDIN_FILENO, STDOUT_FILENO);
306     while (1)
307     ;
308     return NULL;
309 }
```

#### 4.36.1.23 zombie\_child()

```
void* zombie_child (
    void * arg )
```

Helper for zombify.

Definition at line 298 of file shell\_built\_ins.c.

```
298     {
299     s_exit();
300     return NULL;
301 }
```

### 4.36.2 Variable Documentation

#### 4.36.2.1 get\_associated\_ufunc

```
void*(*) (void*) get_associated_ufunc(char *func) (
    char * func )
```

Helper function to get the associated "u-version" of a function given its standalone version. As a concrete example, if we pass in "cat", we will output "u\_cat".

## Parameters

<i>func</i>	A string of the function name to get the associated ufunc for
-------------	---

## Returns

A ptr to the associated u-version function or NULL if no matches are found


Definition at line 164 of file shell\_built\_ins.c.

```

164                                     {
165     if (strcmp(func, "cat") == 0) {
166         return u_cat;
167     } else if (strcmp(func, "sleep") == 0) {
168         return u_sleep;
169     } else if (strcmp(func, "busy") == 0) {
170         return u_busy;
171     } else if (strcmp(func, "echo") == 0) {
172         return u_echo;
173     } else if (strcmp(func, "ls") == 0) {
174         return u_ls;
175     } else if (strcmp(func, "touch") == 0) {
176         return u_touch;
177     } else if (strcmp(func, "mv") == 0) {
178         return u_mv;
179     } else if (strcmp(func, "cp") == 0) {
180         return u_cp;
181     } else if (strcmp(func, "rm") == 0) {
182         return u_rm;
183     } else if (strcmp(func, "chmod") == 0) {
184         return u_chmod;
185     } else if (strcmp(func, "ps") == 0) {
186         return u_ps;
187     } else if (strcmp(func, "kill") == 0) {
188         return u_kill;
189     }
190
191     return NULL; // no matches case
192 }
```

## 4.37 SRC/shell/shell\_built\_ins.h File Reference

This graph shows which files directly or indirectly include this file:





## Functions

- void \* [u\\_cat](#) (void \*arg)  
*The usual `cat` program.*
- void \* [u\\_sleep](#) (void \*arg)  
*Sleep for `n` seconds.*
- void \* [u\\_busy](#) (void \*arg)  
*Busy wait indefinitely. It can only be interrupted via signals.*
- void \* [u\\_echo](#) (void \*arg)  
*Echo back an input string.*
- void \* [u\\_ls](#) (void \*arg)  
*Lists all files in the working directory. For extra credit, it should support relative and absolute file paths.*
- void \* [u\\_touch](#) (void \*arg)  
*For each file, create an empty file if it doesn't exist, else update its timestamp.*
- void \* [u\\_mv](#) (void \*arg)  
*Rename a file. If the `dst_file` file already exists, overwrite it.*
- void \* [u\\_cp](#) (void \*arg)
- void \* [u\\_rm](#) (void \*arg)  
*Remove a list of files. Treat each file in the list as a separate transaction. (i.e. if removing `file1` fails, still attempt to remove `file2`, `file3`, etc.)*
- void \* [u\\_chmod](#) (void \*arg)  
*Change permissions of a file. There's no need to error if a permission being added already exists, or if a permission being removed is already not granted.*
- void \* [u\\_ps](#) (void \*arg)  
*List all processes on PennOS, displaying PID, PPID, priority, status, and command name.*
- void \* [u\\_kill](#) (void \*arg)  
*Sends a specified signal to a list of processes. If a signal name is not specified, default to "term". Valid signals are -term, -stop, and -cont.*
- void \* [u\\_nice](#) (void \*arg)  
*Spawn a new process for `command` and set its priority to `priority`.*
- void \* [u\\_nice\\_pid](#) (void \*arg)  
*Adjust the priority level of an existing process.*
- void \* [u\\_man](#) (void \*arg)  
*Lists all available commands.*
- void \* [u\\_bg](#) (void \*arg)  
*Resumes the most recently stopped job in the background, or the job specified by `job_id`.*
- void \* [u\\_fg](#) (void \*arg)  
*Brings the most recently stopped or background job to the foreground, or the job specified by `job_id`.*
- void \* [u\\_jobs](#) (void \*arg)  
*Lists all jobs.*
- void \* [u\\_logout](#) (void \*arg)  
*Exits the shell and shutdowns PennOS.*
- void \* [u\\_zombify](#) (void \*arg)  
*Used to test zombifying functionality of your kernel.*
- void \* [u\\_orphanify](#) (void \*arg)  
*Used to test orphanifying functionality of your kernel.*

### 4.37.1 Function Documentation

#### 4.37.1.1 u\_bg()

```
void* u_bg (
    void * arg )
```

Resumes the most recently stopped job in the background, or the job specified by `job_id`.

Example Usage: `bg` Example Usage: `bg 2` (`job_id` is 2)

Definition at line 271 of file `shell_built_ins.c`.

```
271     {
272     // TODO --> implement bg
273     return NULL;
274 }
```

#### 4.37.1.2 u\_busy()

```
void* u_busy (
    void * arg )
```

Busy wait indefinitely. It can only be interrupted via signals.

Example Usage: `busy`

Definition at line 53 of file `shell_built_ins.c`.

```
53     {
54     while (1)
55     ;
56     s_exit();
57     return NULL;
58 }
```

#### 4.37.1.3 u\_cat()

```
void* u_cat (
    void * arg )
```

The usual `cat` program.

If `files arg` is provided, concatenate these files and print to `stdout` If `files arg` is *not* provided, read from `stdin` and print back to `stdout`

Example Usage: `cat f1 f2` (concatenates `f1` and `f2` and print to `stdout`) Example Usage: `cat f1 f2 < f3` (concatenates `f1` and `f2` and prints to `stdout`, ignores `f3`) Example Usage: `cat < f3` (concatenates `f3`, prints to `stdout`)

Definition at line 28 of file `shell_built_ins.c`.

```
28     {
29     cat(arg);
30     s_exit();
31     return NULL;
32 }
```

#### 4.37.1.4 u\_chmod()

```
void* u_chmod (
    void * arg )
```

Change permissions of a file. There's no need to error if a permission being added already exists, or if a permission being removed is already not granted.

Print appropriate error message if:

- `file` is not a file that exists
- `perms` is invalid

Example Usage: `chmod +x file` (adds executable permission to file) Example Usage: `chmod +rw file` (adds read + write permissions to file) Example Usage: `chmod -wx file` (removes write + executable permissions from file)

Definition at line 72 of file `shell_built_ins.c`.

```
72     {
73     chmod(arg);
74     s_exit();
75     return NULL;
76 }
```

#### 4.37.1.5 u\_cp()

```
void* u_cp (
    void * arg )
```

Copy a file. If the `dst_file` file already exists, overwrite it.

Print appropriate error message if:

- `src_file` is not a file that exists
- `src_file` does not have read permissions
- `dst_file` file already exists but does not have write permissions

Example Usage: `cp src_file dst_file`

Definition at line 90 of file `shell_built_ins.c`.

```
90     {
91     cp(arg);
92     s_exit();
93     return NULL;
94 }
```

#### 4.37.1.6 u\_echo()

```
void* u_echo (
    void * arg )
```

Echo back an input string.

Example Usage: echo Hello World

Definition at line 60 of file shell\_built\_ins.c.

```
60      {
61          s_echo(arg);
62          s_exit();
63          return NULL;
64      }
```

#### 4.37.1.7 u\_fg()

```
void* u_fg (
    void * arg )
```

Brings the most recently stopped or background job to the foreground, or the job specified by job\_id.

Example Usage: fg Example Usage: fg 2 (job\_id is 2)

Definition at line 276 of file shell\_built\_ins.c.

```
276      {
277          // TODO --> implement fg
278          return NULL;
279      }
```

#### 4.37.1.8 u\_jobs()

```
void* u_jobs (
    void * arg )
```

Lists all jobs.

Example Usage: jobs

Definition at line 281 of file shell\_built\_ins.c.

```
281      {
282          // TODO --> implement jobs
283          return NULL;
284      }
```

### 4.37.1.9 u\_kill()

```
void* u_kill (
    void * arg )
```

Sends a specified signal to a list of processes. If a signal name is not specified, default to "term". Valid signals are -term, -stop, and -cont.

Example Usage: kill 1 2 3 (sends term to processes 1, 2, and 3) Example Usage: kill -term 1 2 (sends term to processes 1 and 2) Example Usage: kill -stop 1 2 (sends stop to processes 1 and 2) Example Usage: kill -cont 1 (sends cont to process 1)

Definition at line 108 of file shell\_built\_ins.c.

```
108     {
109     char** argv = (char**)arg;
110     int sig = 2; // Default signal: term (2)
111     int start_index = 1; // Start after the "kill" command word.
112     char err_buf[128];
113
114     // Check if the first argument specifies a signal
115     if (argv[start_index] && argv[start_index][0] == '-') {
116         if (strcmp(argv[start_index], "-term") == 0) {
117             sig = 2;
118         } else if (strcmp(argv[start_index], "-stop") == 0) {
119             sig = 0;
120         } else if (strcmp(argv[start_index], "-cont") == 0) {
121             sig = 1;
122         } else {
123             // Construct error message
124             s_exit();
125             return NULL;
126         }
127         start_index++;
128     }
129
130     // Process each PID argument using strtol
131     for (int i = start_index; argv[i] != NULL; i++) {
132         char* endptr;
133         long pid_long = strtol(argv[i], &endptr, 10);
134         if (*endptr != '\0' || pid_long <= 0) {
135             snprintf(err_buf, 128, "Invalid PID: %s\n", argv[i]);
136             s_write(STDERR_FILENO, err_buf, strlen(err_buf));
137             continue;
138         }
139         pid_t pid = (pid_t)pid_long;
140         if (s_kill(pid, sig) < 0) {
141             snprintf(err_buf, 128, "b_kill error on PID %d\n", pid);
142             s_write(STDERR_FILENO, err_buf, strlen(err_buf));
143         }
144     }
145     s_exit();
146     return NULL;
147 }
```

### 4.37.1.10 u\_logout()

```
void* u_logout (
    void * arg )
```

Exits the shell and shutdowns PennOS.

Example Usage: logout

Definition at line 286 of file shell\_built\_ins.c.

```
286     {
287     s_shutdown_pennos();
288     return NULL;
289 }
```

#### 4.37.1.11 u\_ls()

```
void* u_ls (
    void * arg )
```

Lists all files in the working directory. For extra credit, it should support relative and absolute file paths.

Example Usage: ls (regular credit) Example Usage: ls ../../foo/./bar/sample (only for EC)

Definition at line 66 of file shell\_built\_ins.c.

```
66         {
67     ls(arg);
68     s_exit();
69     return NULL;
70 }
```

#### 4.37.1.12 u\_man()

```
void* u_man (
    void * arg )
```

Lists all available commands.

Example Usage: man

Definition at line 233 of file shell\_built\_ins.c.

```
233     {
234     const char* man_string =
235         "cat f1 f2 ...      : concatenates provided files (if none, reads from "
236         "std in), and writes to std out\n"
237         "sleep n             : sleeps for n seconds\n"
238         "busy                : busy waits indefinitely\n"
239         "echo str             : echoes back the input string str\n"
240         "ls                  : lists all files in the working directory\n"
241         "touch f1 f2 ...      : for each file, creates empty file if it doesn't "
242         "exist yet, otherwise updates its timestamp\n"
243         "mv f1 f2             : renames f1 to f2 (overwrites f2 if it exists)\n"
244         "cp f1 f2             : copies f1 to f2 (overwrites f2 if it exists)\n"
245         "rm f1 f2 ...         : removes the input list of files\n"
246         "chmod +_ f1          : changes f1 permissions to +_ specifications "
247         "(+x, +rw, etc)\n"
248         "ps                  : lists all processes on PennOS, displaying PID, "
249         "PPID, priority, status, and command name\n"
250         "kill (--) pid1 pid2 : sends specified signal (term default) to list "
251         "of processes\n"
252         "nice n command        : spawns a new process for command and sets its "
253         "priority to n\n"
254         "nice_pid n pid        : adjusts the priority level of an existing "
255         "process to n\n"
256         "man                  : lists all available commands in PennOS\n"
257         "bg                    : resumes most recently stopped process in "
258         "background or the one specified by job_id\n"
259         "fg                    : brings most recently stopped or background job "
260         "to foreground or the one specified by job_id\n"
261         "jobs                  : lists all jobs\n"
262         "logout                : exits the shell and shuts down PennOS\n"
263         "zombify                : creates a child process that becomes a zombie\n"
264         "orphanify             : creates a child process that becomes an "
265         "orphan\n";
266
267     s_write(STDOUT_FILENO, man_string, strlen(man_string));
268     return NULL;
269 }
```

#### 4.37.1.13 u\_mv()

```
void* u_mv (
    void * arg )
```

Rename a file. If the `dst_file` file already exists, overwrite it.

Print appropriate error message if:

- `src_file` is not a file that exists
- `src_file` does not have read permissions
- `dst_file` file already exists but does not have write permissions

Example Usage: `mv src_file dst_file`

Definition at line 84 of file `shell_built_ins.c`.

```
84         {
85     mv(arg);
86     s_exit();
87     return NULL;
88 }
```

#### 4.37.1.14 u\_nice()

```
void* u_nice (
    void * arg )
```

Spawn a new process for `command` and set its priority to `priority`.

Example Usage: `nice 2 cat f1 f2 f3` (spawns cat with priority 2)

Definition at line 194 of file `shell_built_ins.c`.

```
194         {
195     char* endptr;
196     errno = 0;
197     int new_priority = (int)strtol(((char**)arg)[1], &endptr, 10);
198     if (*endptr != '\0' || errno != 0 || new_priority > 2 ||
199         new_priority < 0) { // error catch
200         return NULL;
201     }
202
203     char* command = ((char**)arg)[2];
204     void* (*ufunc)(void*) = get_associated_ufunc(command);
205     if (ufunc == NULL) {
206         return NULL; // no matches, don't spawn
207     }
208
209     pid_t new_proc_pid = s_spawn(ufunc, &((char**)arg)[2], 0, 1);
210
211     if (new_proc_pid != -1) { // non-error case
212         s_nice(new_proc_pid, new_priority);
213     }
214
215     return NULL;
216 }
```

#### 4.37.1.15 u\_nice\_pid()

```
void* u_nice_pid (
    void * arg )
```

Adjust the priority level of an existing process.

Example Usage: nice\_pid 0 123 (sets priority 0 to PID 123)

Definition at line 218 of file shell\_built\_ins.c.

```
218     {
219     char* endptr;
220     errno = 0;
221     int new_priority = (int)strtol(((char**)arg)[1], &endptr, 10);
222     if (*endptr != '\0' || errno != 0) { // error catch
223         return NULL;
224     }
225     pid_t pid = (pid_t)strtol(((char**)arg)[2], &endptr, 10);
226     if (*endptr != '\0' || errno != 0) {
227         return NULL;
228     }
229     s_nice(pid, new_priority);
230     return NULL;
231 }
```

#### 4.37.1.16 u\_orphanify()

```
void* u_orphanify (
    void * arg )
```

Used to test orphanifying functionality of your kernel.

Example Usage: orphanify

Definition at line 320 of file shell\_built\_ins.c.

```
320     {
321     char* orphan_child_argv[] = {"orphan_child", NULL};
322     s_spawn(orphan_child, orphan_child_argv, STDIN_FILENO, STDOUT_FILENO);
323     s_exit();
324     return NULL;
325 }
```

#### 4.37.1.17 u\_ps()

```
void* u_ps (
    void * arg )
```

List all processes on PennOS, displaying PID, PPID, priority, status, and command name.

Example Usage: ps

Definition at line 102 of file shell\_built\_ins.c.

```
102     {
103     s_ps(arg);
104     s_exit();
105     return NULL;
106 }
```



#### 4.37.1.18 u\_rm()

```
void* u_rm (
    void * arg )
```

Remove a list of files. Treat each file in the list as a separate transaction. (i.e. if removing file1 fails, still attempt to remove file2, file3, etc.)

Print appropriate error message if:

- `file` is not a file that exists

Example Usage: `rm f1 f2 f3 f4 f5`

Definition at line 96 of file `shell_built_ins.c`.

```
96         {
97     rm(arg);
98     s_exit();
99     return NULL;
100 }
```

#### 4.37.1.19 u\_sleep()

```
void* u_sleep (
    void * arg )
```

Sleep for `n` seconds.

Note that you'll have to convert the number of seconds to the correct number of ticks.

Example Usage: `sleep 10`

Definition at line 34 of file `shell_built_ins.c`.

```
34     {
35     char* endptr;
36     errno = 0;
37     if (((char**)arg)[1] == NULL) { // no arg case
38         s_exit();
39         return NULL;
40     }
41     int sleep_secs = (int)strtol(((char**)arg)[1], &endptr, 10);
42     if (*endptr != '\0' || errno != 0 || sleep_secs <= 0) {
43         s_exit();
44         return NULL;
45     }
46
47     int sleep_ticks = sleep_secs * 10;
48     s_sleep(sleep_ticks);
49     s_exit();
50     return NULL;
51 }
```

#### 4.37.1.20 u\_touch()

```
void* u_touch (
    void * arg )
```

For each file, create an empty file if it doesn't exist, else update its timestamp.

Example Usage: touch f1 f2 f3 f4 f5

Definition at line 78 of file shell\_built\_ins.c.

```
78     {
79     touch(arg);
80     s_exit();
81     return NULL;
82 }
```

#### 4.37.1.21 u\_zombify()

```
void* u_zombify (
    void * arg )
```

Used to test zombifying functionality of your kernel.

Example Usage: zombify

Definition at line 303 of file shell\_built\_ins.c.

```
303     {
304     char* zombie_child_argv[] = {"zombie_child", NULL};
305     s_spawn(zombie_child, zombie_child_argv, STDIN_FILENO, STDOUT_FILENO);
306     while (1)
307     ;
308     return NULL;
309 }
```

# Index

- `_GNU_SOURCE`
  - `spthread.c`, [195](#)
- `ack`
  - `spthread_signal_args_st`, [18](#)
- `actual_arg`
  - `spthread_fwd_args_st`, [16](#)
- `actual_routine`
  - `spthread_fwd_args_st`, [16](#)
- `add_file_entry`
  - `fs_helpers.c`, [57](#)
  - `fs_helpers.h`, [72](#)
- `alarm_handler`
  - `scheduler.c`, [163](#)
  - `scheduler.h`, [173](#)
- `allocate_block`
  - `fs_helpers.c`, [58](#)
  - `fs_helpers.h`, [74](#)
- `block_size`
  - `fs_helpers.c`, [69](#)
  - `fs_helpers.h`, [87](#)
- `builtins.c`
  - `u_perror`, [231](#)
- `builtins.h`
  - `u_perror`, [233](#)
- `capacity`
  - `vec_st`, [20](#)
- `cat`
  - `fat_routines.c`, [24](#)
  - `fat_routines.h`, [41](#)
- `child_in_zombie_queue`
  - `scheduler.c`, [163](#)
  - `scheduler.h`, [174](#)
- `child_meta`
  - `spthread_fwd_args_st`, [16](#)
- `child_pcb`
  - `pcb_st`, [13](#)
- `child_with_changed_process_status`
  - `scheduler.c`, [164](#)
  - `scheduler.h`, [174](#)
- `chmod`
  - `fat_routines.c`, [26](#)
  - `fat_routines.h`, [43](#)
- `cmd`
  - `job_st`, [9](#)
- `cmd_str`
  - `pcb_st`, [13](#)
- `cmpctdir`
  - `fat_routines.c`, [28](#)
  - `fat_routines.h`, [45](#)
- `commands`
  - `parsed_command`, [11](#)
- `compact_directory`
  - `fs_helpers.c`, [59](#)
  - `fs_helpers.h`, [75](#)
- `CONT_BY_SIG`
  - `signal.h`, [183](#)
- `copy_host_to_pennfat`
  - `fs_helpers.c`, [61](#)
  - `fs_helpers.h`, [77](#)
- `copy_pennfat_to_host`
  - `fs_helpers.c`, [63](#)
  - `fs_helpers.h`, [79](#)
- `copy_source_to_dest`
  - `fs_helpers.c`, [64](#)
  - `fs_helpers.h`, [80](#)
- `cp`
  - `fat_routines.c`, [28](#)
  - `fat_routines.h`, [46](#)
- `crash`
  - `stress.c`, [186](#)
  - `stress.h`, [187](#)
- `create_pcb`
  - `kern_pcb.c`, [124](#)
  - `kern_pcb.h`, [129](#)
- `curr_priority_arr_index`
  - `scheduler.c`, [170](#)
- `current_fg_pid`
  - `fs_kfuncs.c`, [99](#)
  - `kern_sys_calls.c`, [142](#)
  - `shell.c`, [256](#)
- `current_pcb`
  - `kern_pcb.c`, [127](#)
  - `kern_sys_calls.c`, [142](#)
  - `scheduler.c`, [170](#)
- `current_running_pcb`
  - `fs_kfuncs.c`, [99](#)
  - `kern_pcb.c`, [127](#)
  - `kern_sys_calls.c`, [142](#)
  - `scheduler.c`, [170](#)
- `data`
  - `vec_st`, [21](#)
- `decrement_fd_ref_count`
  - `fs_helpers.c`, [65](#)
  - `fs_helpers.h`, [81](#)
- `delete_from_explicit_queue`
  - `kern_sys_calls.c`, [135](#)

- kern\_sys\_calls.h, 145
- delete\_from\_queue
  - kern\_sys\_calls.c, 135
  - kern\_sys\_calls.h, 146
- delete\_process\_from\_all\_queues
  - scheduler.c, 164
  - scheduler.h, 175
- delete\_process\_from\_all\_queues\_except\_current
  - scheduler.c, 164
  - scheduler.h, 175
- delete\_process\_from\_particular\_queue
  - scheduler.c, 165
  - scheduler.h, 176
- det\_priorities\_arr
  - scheduler.c, 171
- determine\_index\_in\_queue
  - kern\_sys\_calls.c, 135
  - kern\_sys\_calls.h, 146
- dir\_entry\_t, 5
  - firstBlock, 5
  - mtime, 5
  - name, 6
  - perm, 6
  - reserved, 6
  - size, 6
  - type, 6
- ele\_dtor\_fn
  - vec\_st, 21
- execute\_command
  - shell.c, 249
- EXITED\_NORMALLY
  - signal.h, 183
- EXPECT\_COMMANDS
  - parser.h, 242
- EXPECT\_INPUT\_FILENAME
  - parser.h, 242
- EXPECT\_OUTPUT\_FILENAME
  - parser.h, 243
- F\_APPEND
  - fat\_routines.h, 38
- F\_READ
  - fat\_routines.h, 38
- F\_WRITE
  - fat\_routines.h, 38
- fat
  - fs\_helpers.c, 69
  - fs\_helpers.h, 87
- FAT\_EOF
  - fat\_routines.h, 39
- FAT\_FREE
  - fat\_routines.h, 39
- fat\_routines.c
  - cat, 24
  - chmod, 26
  - cmpctdir, 28
  - cp, 28
  - ls, 29
  - mkfs, 30
  - mount, 31
  - mv, 32
  - rm, 34
  - touch, 35
  - unmount, 36
- fat\_routines.h
  - cat, 41
  - chmod, 43
  - cmpctdir, 45
  - cp, 46
  - F\_APPEND, 38
  - F\_READ, 38
  - F\_WRITE, 38
  - FAT\_EOF, 39
  - FAT\_FREE, 39
  - ls, 47
  - mkfs, 48
  - mount, 50
  - mv, 51
  - PERM\_EXEC, 39
  - PERM\_NONE, 39
  - PERM\_READ, 39
  - PERM\_READ\_EXEC, 39
  - PERM\_READ\_WRITE, 40
  - PERM\_READ\_WRITE\_EXEC, 40
  - PERM\_WRITE, 40
  - rm, 52
  - touch, 53
  - TYPE\_DIRECTORY, 40
  - TYPE\_REGULAR, 40
  - TYPE\_SYMLINK, 40
  - TYPE\_UNKNOWN, 41
  - unmount, 54
- fat\_size
  - fs\_helpers.c, 69
  - fs\_helpers.h, 87
- fd\_entry\_t, 7
  - filename, 7
  - first\_block, 7
  - in\_use, 7
  - mode, 7
  - position, 8
  - ref\_count, 8
  - size, 8
- fd\_table
  - fs\_helpers.c, 70
  - fs\_helpers.h, 87
  - pcb\_st, 13
- FILE\_DESCRIPTOR\_TABLE\_SIZE
  - kern\_pcb.h, 129
- filename
  - fd\_entry\_t, 7
- fill\_buffer\_until\_full\_or\_newline
  - shell.c, 251
- find\_file
  - fs\_helpers.c, 65
  - fs\_helpers.h, 82

- FINISHED
  - Job.h, 236
- finished\_count
  - job\_st, 9
- first\_block
  - fd\_entry\_t, 7
- firstBlock
  - dir\_entry\_t, 5
- format\_file\_info
  - fs\_kfuncs.c, 89
- free\_job\_ptr
  - shell.c, 251
- free\_pcb
  - kern\_pcb.c, 124
  - kern\_pcb.h, 130
- free\_scheduler\_queues
  - scheduler.c, 165
  - scheduler.h, 176
- fs\_fd
  - fs\_helpers.c, 70
  - fs\_helpers.h, 87
- fs\_helpers.c
  - add\_file\_entry, 57
  - allocate\_block, 58
  - block\_size, 69
  - compact\_directory, 59
  - copy\_host\_to\_pennfat, 61
  - copy\_pennfat\_to\_host, 63
  - copy\_source\_to\_dest, 64
  - decrement\_fd\_ref\_count, 65
  - fat, 69
  - fat\_size, 69
  - fd\_table, 70
  - find\_file, 65
  - fs\_fd, 70
  - get\_free\_fd, 66
  - has\_executable\_permission, 67
  - increment\_fd\_ref\_count, 67
  - init\_fd\_table, 68
  - is\_mounted, 70
  - mark\_entry\_as\_deleted, 68
  - MAX\_FDS, 70
  - num\_fat\_blocks, 70
- fs\_helpers.h
  - add\_file\_entry, 72
  - allocate\_block, 74
  - block\_size, 87
  - compact\_directory, 75
  - copy\_host\_to\_pennfat, 77
  - copy\_pennfat\_to\_host, 79
  - copy\_source\_to\_dest, 80
  - decrement\_fd\_ref\_count, 81
  - fat, 87
  - fat\_size, 87
  - fd\_table, 87
  - find\_file, 82
  - fs\_fd, 87
  - get\_free\_fd, 83
  - has\_executable\_permission, 84
  - increment\_fd\_ref\_count, 85
  - init\_fd\_table, 85
  - is\_mounted, 88
  - mark\_entry\_as\_deleted, 86
  - MAX\_FDS, 88
  - num\_fat\_blocks, 88
- fs\_kfuncs.c
  - current\_fg\_pid, 99
  - current\_running\_pcb, 99
  - format\_file\_info, 89
  - k\_close, 89
  - k\_ls, 90
  - k\_lseek, 91
  - k\_open, 92
  - k\_read, 94
  - k\_unlink, 95
  - k\_write, 96
- fs\_kfuncs.h
  - k\_close, 101
  - k\_ls, 102
  - k\_lseek, 103
  - k\_open, 105
  - k\_read, 107
  - k\_unlink, 109
  - k\_write, 110
  - SEEK\_CUR, 101
  - SEEK\_END, 101
  - SEEK\_SET, 101
- fs\_syscalls.c
  - s\_close, 114
  - s\_ls, 115
  - s\_lseek, 115
  - s\_open, 115
  - s\_read, 116
  - s\_unlink, 116
  - s\_write, 116
- fs\_syscalls.h
  - s\_close, 118
  - s\_ls, 119
  - s\_lseek, 119
  - s\_open, 120
  - s\_read, 121
  - s\_unlink, 122
  - s\_write, 122
  - STDERR\_FILENO, 118
  - STDIN\_FILENO, 118
  - STDOUT\_FILENO, 118
- generate\_next\_priority
  - scheduler.c, 165
  - scheduler.h, 176
- get\_associated\_ufunc
  - shell\_built\_ins.c, 269
- get\_free\_fd
  - fs\_helpers.c, 66
  - fs\_helpers.h, 83
- get\_next\_pcb
  - scheduler.c, 166

- scheduler.h, 177
- get\_pcb\_in\_queue
  - scheduler.c, 166
  - scheduler.h, 178
- handle\_signal
  - scheduler.c, 167
  - scheduler.h, 178
- hang
  - stress.c, 186
  - stress.h, 187
- has\_executable\_permission
  - fs\_helpers.c, 67
  - fs\_helpers.h, 84
- id
  - job\_st, 9
- in\_use
  - fd\_entry\_t, 7
- increment\_fd\_ref\_count
  - fs\_helpers.c, 67
  - fs\_helpers.h, 85
- init\_fd\_table
  - fs\_helpers.c, 68
  - fs\_helpers.h, 85
- init\_func
  - kern\_sys\_calls.c, 136
  - kern\_sys\_calls.h, 147
- initialize\_scheduler\_queues
  - scheduler.c, 167
  - scheduler.h, 179
- input\_fd
  - pcb\_st, 13
- input\_fd\_script
  - shell.c, 256
- is\_background
  - parsed\_command, 11
- is\_file\_append
  - parsed\_command, 11
- is\_mounted
  - fs\_helpers.c, 70
  - fs\_helpers.h, 88
- is\_sleeping
  - pcb\_st, 13
- jid\_t
  - Job.h, 235
- job
  - Job.h, 236
- Job.h
  - FINISHED, 236
  - jid\_t, 235
  - job, 236
  - job\_state\_t, 236
  - RUNNING, 236
  - STOPPED, 236
- job\_list
  - shell.c, 256
- job\_st, 8
  - cmd, 9
  - finished\_count, 9
  - id, 9
  - num\_pids, 9
  - pgid, 9
  - pids, 10
  - state, 10
- job\_state\_t
  - Job.h, 236
- JUMP\_OUT
  - parser.c, 237
- k\_close
  - fs\_kfuncs.c, 89
  - fs\_kfuncs.h, 101
- k\_ls
  - fs\_kfuncs.c, 90
  - fs\_kfuncs.h, 102
- k\_lseek
  - fs\_kfuncs.c, 91
  - fs\_kfuncs.h, 103
- k\_open
  - fs\_kfuncs.c, 92
  - fs\_kfuncs.h, 105
- k\_proc\_cleanup
  - kern\_pcb.c, 125
  - kern\_pcb.h, 131
- k\_proc\_create
  - kern\_pcb.c, 126
  - kern\_pcb.h, 132
- k\_read
  - fs\_kfuncs.c, 94
  - fs\_kfuncs.h, 107
- k\_unlink
  - fs\_kfuncs.c, 95
  - fs\_kfuncs.h, 109
- k\_write
  - fs\_kfuncs.c, 96
  - fs\_kfuncs.h, 110
- kern\_pcb.c
  - create\_pcb, 124
  - current\_pcb, 127
  - current\_running\_pcb, 127
  - free\_pcb, 124
  - k\_proc\_cleanup, 125
  - k\_proc\_create, 126
  - next\_pid, 127
  - remove\_child\_in\_parent, 126
- kern\_pcb.h
  - create\_pcb, 129
  - FILE\_DESCRIPTOR\_TABLE\_SIZE, 129
  - free\_pcb, 130
  - k\_proc\_cleanup, 131
  - k\_proc\_create, 132
  - pcb\_t, 129
  - remove\_child\_in\_parent, 133
- kern\_sys\_calls.c
  - current\_fg\_pid, 142
  - current\_pcb, 142

- current\_running\_pcb, 142
- delete\_from\_explicit\_queue, 135
- delete\_from\_queue, 135
- determine\_index\_in\_queue, 135
- init\_func, 136
- move\_pcb\_correct\_queue, 136
- one\_priority\_queue, 143
- s\_cleanup\_init\_process, 137
- s\_echo, 137
- s\_exit, 138
- s\_kill, 138
- s\_nice, 138
- s\_ps, 139
- s\_sleep, 139
- s\_spawn, 140
- s\_spawn\_init, 140
- s\_waitpid, 141
- sleep\_blocked\_queue, 143
- tick\_counter, 143
- two\_priority\_queue, 143
- zero\_priority\_queue, 143
- zombie\_queue, 143
- kern\_sys\_calls.h
  - delete\_from\_explicit\_queue, 145
  - delete\_from\_queue, 146
  - determine\_index\_in\_queue, 146
  - init\_func, 147
  - move\_pcb\_correct\_queue, 147
  - s\_cleanup\_init\_process, 148
  - s\_echo, 148
  - s\_exit, 149
  - s\_kill, 149
  - s\_nice, 151
  - s\_ps, 151
  - s\_sleep, 152
  - s\_spawn, 153
  - s\_spawn\_init, 153
  - s\_waitpid, 154
- length
  - vec\_st, 21
- log\_fd
  - logger.h, 161
  - pennos.c, 230
  - scheduler.c, 171
- log\_generic\_event
  - logger.c, 156
  - logger.h, 159
- log\_nice\_event
  - logger.c, 157
  - logger.h, 160
- log\_scheduling\_event
  - logger.c, 158
  - logger.h, 161
- logger.c
  - log\_generic\_event, 156
  - log\_nice\_event, 157
  - log\_scheduling\_event, 158
- logger.h
  - log\_fd, 161
  - log\_generic\_event, 159
  - log\_nice\_event, 160
  - log\_scheduling\_event, 161
  - tick\_counter, 161
- ls
  - fat\_routines.c, 29
  - fat\_routines.h, 47
- main
  - pennfat.c, 227
  - pennos.c, 230
- mark\_entry\_as\_deleted
  - fs\_helpers.c, 68
  - fs\_helpers.h, 86
- MAX\_BUFFER\_SIZE
  - shell.c, 248
- MAX\_FDS
  - fs\_helpers.c, 70
  - fs\_helpers.h, 88
- MAX\_LINE\_BUFFER\_SIZE
  - shell.c, 249
- meta
  - spthread\_st, 20
- meta\_mutex
  - spthread\_meta\_st, 17
- MILISEC\_IN\_NANO
  - spthread.c, 195
- mkfs
  - fat\_routines.c, 30
  - fat\_routines.h, 48
- mode
  - fd\_entry\_t, 7
- mount
  - fat\_routines.c, 31
  - fat\_routines.h, 50
- move\_pcb\_correct\_queue
  - kern\_sys\_calls.c, 136
  - kern\_sys\_calls.h, 147
- mtime
  - dir\_entry\_t, 5
- mv
  - fat\_routines.c, 32
  - fat\_routines.h, 51
- name
  - dir\_entry\_t, 6
- next\_job\_id
  - shell.c, 256
- next\_pid
  - kern\_pcb.c, 127
- nohang
  - stress.c, 186
  - stress.h, 188
- num\_commands
  - parsed\_command, 11
- num\_fat\_blocks
  - fs\_helpers.c, 70
  - fs\_helpers.h, 88

- num\_pids
  - job\_st, 9
- one\_priority\_queue
  - kern\_sys\_calls.c, 143
  - scheduler.c, 171
- orphan\_child
  - shell\_built\_ins.c, 261
- output\_fd
  - pcb\_st, 14
- output\_fd\_script
  - shell.c, 256
- P\_EBADF
  - pennos-errno.h, 190
- P\_EBUSY
  - pennos-errno.h, 190
- P\_ECLOSE
  - pennos-errno.h, 190
- P\_ECOMMAND
  - pennos-errno.h, 190
- P\_EEXIST
  - pennos-errno.h, 190
- P\_EFS\_NOT\_MOUNTED
  - pennos-errno.h, 190
- P\_EFULL
  - pennos-errno.h, 191
- P\_EFUNC
  - pennos-errno.h, 191
- P\_EINTR
  - pennos-errno.h, 191
- P\_EINVAL
  - pennos-errno.h, 191
- P\_ELSEEK
  - pennos-errno.h, 191
- P\_EMALLOC
  - pennos-errno.h, 191
- P\_EMAP
  - pennos-errno.h, 192
- P\_ENOENT
  - pennos-errno.h, 192
- P\_ENULL
  - pennos-errno.h, 192
- P\_EOPEN
  - pennos-errno.h, 192
- P\_EPARSE
  - pennos-errno.h, 192
- P\_EPERM
  - pennos-errno.h, 192
- P\_EREAD
  - pennos-errno.h, 193
- P\_ERRNO
  - pennos-errno.c, 189
  - pennos-errno.h, 194
- P\_ESIGNAL
  - pennos-errno.h, 193
- P\_EUNKNOWN
  - pennos-errno.h, 193
- P\_EWRITE
  - pennos-errno.h, 193
- P\_INITFAIL
  - pennos-errno.h, 193
- P\_NEEDF
  - pennos-errno.h, 193
- P\_SIGCONT
  - signal.h, 183
- P\_SIGSTOP
  - signal.h, 184
- P\_SIGTERM
  - signal.h, 184
- P\_WIFEXITED
  - signal.h, 184
- P\_WIFSIGNALED
  - signal.h, 184
- P\_WIFSTOPPED
  - signal.h, 184
- par\_pid
  - pcb\_st, 14
- parse\_command
  - parser.c, 237
  - parser.h, 244
- parsed\_command, 10
  - commands, 11
  - is\_background, 11
  - is\_file\_append, 11
  - num\_commands, 11
  - stdin\_file, 11
  - stdout\_file, 11
- parser.c
  - JUMP\_OUT, 237
  - parse\_command, 237
  - print\_parsed\_command, 240
  - print\_parser\_errcode, 240
- parser.h
  - EXPECT\_COMMANDS, 242
  - EXPECT\_INPUT\_FILENAME, 242
  - EXPECT\_OUTPUT\_FILENAME, 243
  - parse\_command, 244
  - print\_parsed\_command, 246
  - print\_parser\_errcode, 247
  - UNEXPECTED\_AMPERSAND, 243
  - UNEXPECTED\_FILE\_INPUT, 243
  - UNEXPECTED\_FILE\_OUTPUT, 243
  - UNEXPECTED\_PIPELINE, 243
- pcb\_st, 12
  - child\_pcbs, 13
  - cmd\_str, 13
  - fd\_table, 13
  - input\_fd, 13
  - is\_sleeping, 13
  - output\_fd, 14
  - par\_pid, 14
  - pid, 14
  - priority, 14
  - process\_state, 14
  - process\_status, 14
  - signals, 15



- thread\_handle, 15
- time\_to\_wake, 15
- pcb\_t
  - kern\_pcb.h, 129
- pennfat.c
  - main, 227
  - PROMPT, 227
- pennos-errno.c
  - P\_ERRNO, 189
- pennos-errno.h
  - P\_EBADF, 190
  - P\_EBUSY, 190
  - P\_ECLOSE, 190
  - P\_ECOMMAND, 190
  - P\_EEXIST, 190
  - P\_EFS\_NOT\_MOUNTED, 190
  - P\_EFULL, 191
  - P\_EFUNC, 191
  - P\_EINTR, 191
  - P\_EINVAL, 191
  - P\_ELSEEK, 191
  - P\_EMALLOC, 191
  - P\_EMAP, 192
  - P\_ENOENT, 192
  - P\_ENULL, 192
  - P\_EOPEN, 192
  - P\_EPARSE, 192
  - P\_EPERM, 192
  - P\_EREAD, 193
  - P\_ERRNO, 194
  - P\_ESIGNAL, 193
  - P\_EUNKNOWN, 193
  - P\_EWRITE, 193
  - P\_INITFAIL, 193
  - P\_NEEDF, 193
- pennos.c
  - log\_fd, 230
  - main, 230
  - tick\_counter, 230
- perm
  - dir\_entry\_t, 6
- PERM\_EXEC
  - fat\_routines.h, 39
- PERM\_NONE
  - fat\_routines.h, 39
- PERM\_READ
  - fat\_routines.h, 39
- PERM\_READ\_EXEC
  - fat\_routines.h, 39
- PERM\_READ\_WRITE
  - fat\_routines.h, 40
- PERM\_READ\_WRITE\_EXEC
  - fat\_routines.h, 40
- PERM\_WRITE
  - fat\_routines.h, 40
- pgid
  - job\_st, 9
- pid
  - pcb\_st, 14
- pids
  - job\_st, 10
- position
  - fd\_entry\_t, 8
- print\_parsed\_command
  - parser.c, 240
  - parser.h, 246
- print\_parser\_errcode
  - parser.c, 240
  - parser.h, 247
- priority
  - pcb\_st, 14
- process\_state
  - pcb\_st, 14
- process\_status
  - pcb\_st, 14
- PROMPT
  - pennfat.c, 227
  - shell.c, 249
- pthread\_fn
  - sphthead.c, 196
- ptr\_dtor\_fn
  - Vec.h, 219
- ptr\_t
  - Vec.h, 219
- put\_pcb\_into\_correct\_queue
  - scheduler.c, 168
  - scheduler.h, 180
- recur
  - stress.c, 186
  - stress.h, 188
- ref\_count
  - fd\_entry\_t, 8
- remove\_child\_in\_parent
  - kern\_pcb.c, 126
  - kern\_pcb.h, 133
- reserved
  - dir\_entry\_t, 6
- rm
  - fat\_routines.c, 34
  - fat\_routines.h, 52
- RUNNING
  - Job.h, 236
- s\_cleanup\_init\_process
  - kern\_sys\_calls.c, 137
  - kern\_sys\_calls.h, 148
- s\_close
  - fs\_syscalls.c, 114
  - fs\_syscalls.h, 118
- s\_echo
  - kern\_sys\_calls.c, 137
  - kern\_sys\_calls.h, 148
- s\_exit
  - kern\_sys\_calls.c, 138
  - kern\_sys\_calls.h, 149
- s\_kill

- kern\_sys\_calls.c, 138
- kern\_sys\_calls.h, 149
- s\_ls
  - fs\_syscalls.c, 115
  - fs\_syscalls.h, 119
- s\_lseek
  - fs\_syscalls.c, 115
  - fs\_syscalls.h, 119
- s\_nice
  - kern\_sys\_calls.c, 138
  - kern\_sys\_calls.h, 151
- s\_open
  - fs\_syscalls.c, 115
  - fs\_syscalls.h, 120
- s\_ps
  - kern\_sys\_calls.c, 139
  - kern\_sys\_calls.h, 151
- s\_read
  - fs\_syscalls.c, 116
  - fs\_syscalls.h, 121
- s\_shutdown\_pennos
  - scheduler.c, 168
  - scheduler.h, 180
- s\_sleep
  - kern\_sys\_calls.c, 139
  - kern\_sys\_calls.h, 152
- s\_spawn
  - kern\_sys\_calls.c, 140
  - kern\_sys\_calls.h, 153
- s\_spawn\_init
  - kern\_sys\_calls.c, 140
  - kern\_sys\_calls.h, 153
- s\_unlink
  - fs\_syscalls.c, 116
  - fs\_syscalls.h, 122
- s\_waitpid
  - kern\_sys\_calls.c, 141
  - kern\_sys\_calls.h, 154
- s\_write
  - fs\_syscalls.c, 116
  - fs\_syscalls.h, 122
- scheduler
  - scheduler.c, 168
  - scheduler.h, 180
- scheduler.c
  - alarm\_handler, 163
  - child\_in\_zombie\_queue, 163
  - child\_with\_changed\_process\_status, 164
  - curr\_priority\_arr\_index, 170
  - current\_pcb, 170
  - current\_running\_pcb, 170
  - delete\_process\_from\_all\_queues, 164
  - delete\_process\_from\_all\_queues\_except\_current, 164
  - delete\_process\_from\_particular\_queue, 165
  - det\_priorities\_arr, 171
  - free\_scheduler\_queues, 165
  - generate\_next\_priority, 165
  - get\_next\_pcb, 166
  - get\_pcb\_in\_queue, 166
  - handle\_signal, 167
  - initialize\_scheduler\_queues, 167
  - log\_fd, 171
  - one\_priority\_queue, 171
  - put\_pcb\_into\_correct\_queue, 168
  - s\_shutdown\_pennos, 168
  - scheduler, 168
  - sleep\_blocked\_queue, 171
  - tick\_counter, 171
  - two\_priority\_queue, 171
  - zero\_priority\_queue, 172
  - zombie\_queue, 172
- scheduler.h
  - alarm\_handler, 173
  - child\_in\_zombie\_queue, 174
  - child\_with\_changed\_process\_status, 174
  - delete\_process\_from\_all\_queues, 175
  - delete\_process\_from\_all\_queues\_except\_current, 175
  - delete\_process\_from\_particular\_queue, 176
  - free\_scheduler\_queues, 176
  - generate\_next\_priority, 176
  - get\_next\_pcb, 177
  - get\_pcb\_in\_queue, 178
  - handle\_signal, 178
  - initialize\_scheduler\_queues, 179
  - put\_pcb\_into\_correct\_queue, 180
  - s\_shutdown\_pennos, 180
  - scheduler, 180
- script\_fd
  - shell.c, 257
- SEEK\_CUR
  - fs\_kfuncs.h, 101
- SEEK\_END
  - fs\_kfuncs.h, 101
- SEEK\_SET
  - fs\_kfuncs.h, 101
- setup\_cond
  - spthread\_fwd\_args\_st, 16
- setup\_done
  - spthread\_fwd\_args\_st, 16
- setup\_mutex
  - spthread\_fwd\_args\_st, 17
- setup\_terminal\_signal\_handlers
  - shell.c, 251
- shell
  - shell.c, 252
  - shell.h, 257
- shell.c
  - current\_fg\_pid, 256
  - execute\_command, 249
  - fill\_buffer\_until\_full\_or\_newline, 251
  - free\_job\_ptr, 251
  - input\_fd\_script, 256
  - job\_list, 256
  - MAX\_BUFFER\_SIZE, 248

- MAX\_LINE\_BUFFER\_SIZE, 249
- next\_job\_id, 256
- output\_fd\_script, 256
- PROMPT, 249
- script\_fd, 257
- setup\_terminal\_signal\_handlers, 251
- shell, 252
- shell\_sigint\_handler, 253
- shell\_sigstp\_handler, 254
- u\_execute\_command, 254
- u\_read\_and\_execute\_script, 255
- shell.h
  - shell, 257
- shell\_built\_ins.c
  - get\_associated\_ufunc, 269
  - orphan\_child, 261
  - u\_bg, 261
  - u\_busy, 261
  - u\_cat, 261
  - u\_chmod, 262
  - u\_cp, 262
  - u\_echo, 263
  - u\_fg, 263
  - u\_jobs, 263
  - u\_kill, 264
  - u\_logout, 264
  - u\_ls, 265
  - u\_man, 265
  - u\_mv, 266
  - u\_nice, 266
  - u\_nice\_pid, 267
  - u\_orphanify, 267
  - u\_ps, 267
  - u\_rm, 268
  - u\_sleep, 268
  - u\_touch, 268
  - u\_zombify, 269
  - zombie\_child, 269
- shell\_built\_ins.h
  - u\_bg, 271
  - u\_busy, 272
  - u\_cat, 272
  - u\_chmod, 272
  - u\_cp, 273
  - u\_echo, 273
  - u\_fg, 274
  - u\_jobs, 274
  - u\_kill, 274
  - u\_logout, 275
  - u\_ls, 275
  - u\_man, 276
  - u\_mv, 276
  - u\_nice, 277
  - u\_nice\_pid, 277
  - u\_orphanify, 278
  - u\_ps, 278
  - u\_rm, 278
  - u\_sleep, 279
  - u\_touch, 279
  - u\_zombify, 280
- shell\_sigint\_handler
  - shell.c, 253
- shell\_sigstp\_handler
  - shell.c, 254
- shutup\_mutex
  - sphthread\_signal\_args\_st, 18
- signal
  - sphthread\_signal\_args\_st, 19
- signal.h
  - CONT\_BY\_SIG, 183
  - EXITED\_NORMALLY, 183
  - P\_SIGCONT, 183
  - P\_SIGSTOP, 184
  - P\_SIGTERM, 184
  - P\_WIFEXITED, 184
  - P\_WIFSIGNALED, 184
  - P\_WIFSTOPPED, 184
  - STOPPED\_BY\_SIG, 185
  - TERM\_BY\_SIG, 185
- signals
  - pcb\_st, 15
- SIGPTHD
  - sphthread.h, 203
- size
  - dir\_entry\_t, 6
  - fd\_entry\_t, 8
- sleep\_blocked\_queue
  - kern\_sys\_calls.c, 143
  - scheduler.c, 171
- sphthread.c
  - \_GNU\_SOURCE, 195
  - MILISEC\_IN\_NANO, 195
  - pthread\_fn, 196
  - sphthread\_cancel, 197
  - sphthread\_continue, 197
  - sphthread\_create, 198
  - sphthread\_disable\_interrupts\_self, 199
  - sphthread\_enable\_interrupts\_self, 199
  - sphthread\_equal, 199
  - sphthread\_exit, 200
  - sphthread\_fwd\_args, 196
  - sphthread\_join, 200
  - sphthread\_meta\_t, 197
  - SPTHREAD\_RUNNING\_STATE, 195
  - sphthread\_self, 200
  - SPTHREAD\_SIG\_CONTINUE, 196
  - SPTHREAD\_SIG\_SUSPEND, 196
  - sphthread\_signal\_args, 197
  - sphthread\_suspend, 200
  - sphthread\_suspend\_self, 201
  - SPTHREAD\_SUSPENDED\_STATE, 196
  - SPTHREAD\_TERMINATED\_STATE, 196
- sphthread.h
  - SIGPTHD, 203
  - sphthread\_cancel, 203
  - sphthread\_continue, 204

- spthread\_create, 204
- spthread\_disable\_interrupts\_self, 205
- spthread\_enable\_interrupts\_self, 206
- spthread\_equal, 206
- spthread\_exit, 206
- spthread\_join, 206
- spthread\_meta\_t, 203
- spthread\_self, 207
- spthread\_suspend, 207
- spthread\_suspend\_self, 208
- spthread\_t, 203
- spthread\_cancel
  - spthread.c, 197
  - spthread.h, 203
- spthread\_continue
  - spthread.c, 197
  - spthread.h, 204
- spthread\_create
  - spthread.c, 198
  - spthread.h, 204
- spthread\_disable\_interrupts\_self
  - spthread.c, 199
  - spthread.h, 205
- spthread\_enable\_interrupts\_self
  - spthread.c, 199
  - spthread.h, 206
- spthread\_equal
  - spthread.c, 199
  - spthread.h, 206
- spthread\_exit
  - spthread.c, 200
  - spthread.h, 206
- spthread\_fwd\_args
  - spthread.c, 196
- spthread\_fwd\_args\_st, 15
  - actual\_arg, 16
  - actual\_routine, 16
  - child\_meta, 16
  - setup\_cond, 16
  - setup\_done, 16
  - setup\_mutex, 17
- spthread\_join
  - spthread.c, 200
  - spthread.h, 206
- spthread\_meta\_st, 17
  - meta\_mutex, 17
  - state, 17
  - suspend\_set, 18
- spthread\_meta\_t
  - spthread.c, 197
  - spthread.h, 203
- SPTHREAD\_RUNNING\_STATE
  - spthread.c, 195
- spthread\_self
  - spthread.c, 200
  - spthread.h, 207
- SPTHREAD\_SIG\_CONTINUE
  - spthread.c, 196
- SPTHREAD\_SIG\_SUSPEND
  - spthread.c, 196
- spthread\_signal\_args
  - spthread.c, 197
- spthread\_signal\_args\_st, 18
  - ack, 18
  - shutup\_mutex, 18
  - signal, 19
- spthread\_st, 19
  - meta, 20
  - thread, 20
- spthread\_suspend
  - spthread.c, 200
  - spthread.h, 207
- spthread\_suspend\_self
  - spthread.c, 201
  - spthread.h, 208
- SPTHREAD\_SUSPENDED\_STATE
  - spthread.c, 196
- spthread\_t
  - spthread.h, 203
- SPTHREAD\_TERMINATED\_STATE
  - spthread.c, 196
- SRC/fs/fat\_routines.c, 23
- SRC/fs/fat\_routines.h, 36
- SRC/fs/fs\_helpers.c, 55
- SRC/fs/fs\_helpers.h, 71
- SRC/fs/fs\_kfuncs.c, 88
- SRC/fs/fs\_kfuncs.h, 100
- SRC/fs/fs\_syscalls.c, 113
- SRC/fs/fs\_syscalls.h, 117
- SRC/kernel/kern\_pcb.c, 123
- SRC/kernel/kern\_pcb.h, 128
- SRC/kernel/kern\_sys\_calls.c, 133
- SRC/kernel/kern\_sys\_calls.h, 144
- SRC/kernel/logger.c, 155
- SRC/kernel/logger.h, 158
- SRC/kernel/scheduler.c, 162
- SRC/kernel/scheduler.h, 172
- SRC/kernel/signal.c, 182
- SRC/kernel/signal.h, 183
- SRC/kernel/stress.c, 185
- SRC/kernel/stress.h, 187
- SRC/lib/pennos-errno.c, 188
- SRC/lib/pennos-errno.h, 189
- SRC/lib/spthread.c, 194
- SRC/lib/spthread.h, 202
- SRC/lib/Vec.c, 208
- SRC/lib/Vec.h, 216
- SRC/pennfat.c, 226
- SRC/pennos.c, 229
- SRC/shell/builtins.c, 231
- SRC/shell/builtins.h, 233
- SRC/shell/Job.h, 234
- SRC/shell/parser.c, 236
- SRC/shell/parser.h, 241
- SRC/shell/shell.c, 247
- SRC/shell/shell.h, 257

- SRC/shell/shell\_built\_ins.c, [259](#)
- SRC/shell/shell\_built\_ins.h, [270](#)
- state
  - job\_st, [10](#)
  - spthread\_meta\_st, [17](#)
- STDERR\_FILENO
  - fs\_syscalls.h, [118](#)
- stdin\_file
  - parsed\_command, [11](#)
- STDIN\_FILENO
  - fs\_syscalls.h, [118](#)
- stdout\_file
  - parsed\_command, [11](#)
- STDOUT\_FILENO
  - fs\_syscalls.h, [118](#)
- STOPPED
  - Job.h, [236](#)
- STOPPED\_BY\_SIG
  - signal.h, [185](#)
- stress.c
  - crash, [186](#)
  - hang, [186](#)
  - nohang, [186](#)
  - recur, [186](#)
- stress.h
  - crash, [187](#)
  - hang, [187](#)
  - nohang, [188](#)
  - recur, [188](#)
- suspend\_set
  - spthread\_meta\_st, [18](#)
- TERM\_BY\_SIG
  - signal.h, [185](#)
- thread
  - spthread\_st, [20](#)
- thread\_handle
  - pcb\_st, [15](#)
- tick\_counter
  - kern\_sys\_calls.c, [143](#)
  - logger.h, [161](#)
  - pennos.c, [230](#)
  - scheduler.c, [171](#)
- time\_to\_wake
  - pcb\_st, [15](#)
- touch
  - fat\_routines.c, [35](#)
  - fat\_routines.h, [53](#)
- two\_priority\_queue
  - kern\_sys\_calls.c, [143](#)
  - scheduler.c, [171](#)
- type
  - dir\_entry\_t, [6](#)
- TYPE\_DIRECTORY
  - fat\_routines.h, [40](#)
- TYPE\_REGULAR
  - fat\_routines.h, [40](#)
- TYPE\_SYMLINK
  - fat\_routines.h, [40](#)
- TYPE\_UNKNOWN
  - fat\_routines.h, [41](#)
- u\_bg
  - shell\_built\_ins.c, [261](#)
  - shell\_built\_ins.h, [271](#)
- u\_busy
  - shell\_built\_ins.c, [261](#)
  - shell\_built\_ins.h, [272](#)
- u\_cat
  - shell\_built\_ins.c, [261](#)
  - shell\_built\_ins.h, [272](#)
- u\_chmod
  - shell\_built\_ins.c, [262](#)
  - shell\_built\_ins.h, [272](#)
- u\_cp
  - shell\_built\_ins.c, [262](#)
  - shell\_built\_ins.h, [273](#)
- u\_echo
  - shell\_built\_ins.c, [263](#)
  - shell\_built\_ins.h, [273](#)
- u\_execute\_command
  - shell.c, [254](#)
- u\_fg
  - shell\_built\_ins.c, [263](#)
  - shell\_built\_ins.h, [274](#)
- u\_jobs
  - shell\_built\_ins.c, [263](#)
  - shell\_built\_ins.h, [274](#)
- u\_kill
  - shell\_built\_ins.c, [264](#)
  - shell\_built\_ins.h, [274](#)
- u\_logout
  - shell\_built\_ins.c, [264](#)
  - shell\_built\_ins.h, [275](#)
- u\_ls
  - shell\_built\_ins.c, [265](#)
  - shell\_built\_ins.h, [275](#)
- u\_man
  - shell\_built\_ins.c, [265](#)
  - shell\_built\_ins.h, [276](#)
- u\_mv
  - shell\_built\_ins.c, [266](#)
  - shell\_built\_ins.h, [276](#)
- u\_nice
  - shell\_built\_ins.c, [266](#)
  - shell\_built\_ins.h, [277](#)
- u\_nice\_pid
  - shell\_built\_ins.c, [267](#)
  - shell\_built\_ins.h, [277](#)
- u\_orphanify
  - shell\_built\_ins.c, [267](#)
  - shell\_built\_ins.h, [278](#)
- u\_perror
  - builtins.c, [231](#)
  - builtins.h, [233](#)
- u\_ps
  - shell\_built\_ins.c, [267](#)
  - shell\_built\_ins.h, [278](#)

- u\_read\_and\_execute\_script
  - shell.c, 255
- u\_rm
  - shell\_built\_ins.c, 268
  - shell\_built\_ins.h, 278
- u\_sleep
  - shell\_built\_ins.c, 268
  - shell\_built\_ins.h, 279
- u\_touch
  - shell\_built\_ins.c, 268
  - shell\_built\_ins.h, 279
- u\_zombify
  - shell\_built\_ins.c, 269
  - shell\_built\_ins.h, 280
- UNEXPECTED\_AMPERSAND
  - parser.h, 243
- UNEXPECTED\_FILE\_INPUT
  - parser.h, 243
- UNEXPECTED\_FILE\_OUTPUT
  - parser.h, 243
- UNEXPECTED\_PIPELINE
  - parser.h, 243
- unmount
  - fat\_routines.c, 36
  - fat\_routines.h, 54
- Vec
  - Vec.h, 219
- Vec.c
  - vec\_clear, 209
  - vec\_destroy, 210
  - vec\_erase, 210
  - vec\_erase\_no\_deletor, 211
  - vec\_get, 211
  - vec\_insert, 212
  - vec\_new, 213
  - vec\_pop\_back, 213
  - vec\_push\_back, 214
  - vec\_resize, 215
  - vec\_set, 216
- Vec.h
  - ptr\_dtor\_fn, 219
  - ptr\_t, 219
  - Vec, 219
  - vec\_capacity, 218
  - vec\_clear, 219
  - vec\_destroy, 220
  - vec\_erase, 220
  - vec\_erase\_no\_deletor, 221
  - vec\_get, 221
  - vec\_insert, 222
  - vec\_is\_empty, 218
  - vec\_len, 218
  - vec\_new, 223
  - vec\_pop\_back, 223
  - vec\_push\_back, 224
  - vec\_resize, 225
  - vec\_set, 226
- vec\_capacity
  - Vec.h, 218
- vec\_clear
  - Vec.c, 209
  - Vec.h, 219
- vec\_destroy
  - Vec.c, 210
  - Vec.h, 220
- vec\_erase
  - Vec.c, 210
  - Vec.h, 220
- vec\_erase\_no\_deletor
  - Vec.c, 211
  - Vec.h, 221
- vec\_get
  - Vec.c, 211
  - Vec.h, 221
- vec\_insert
  - Vec.c, 212
  - Vec.h, 222
- vec\_is\_empty
  - Vec.h, 218
- vec\_len
  - Vec.h, 218
- vec\_new
  - Vec.c, 213
  - Vec.h, 223
- vec\_pop\_back
  - Vec.c, 213
  - Vec.h, 223
- vec\_push\_back
  - Vec.c, 214
  - Vec.h, 224
- vec\_resize
  - Vec.c, 215
  - Vec.h, 225
- vec\_set
  - Vec.c, 216
  - Vec.h, 226
- vec\_st, 20
  - capacity, 20
  - data, 21
  - ele\_dtor\_fn, 21
  - length, 21
- zero\_priority\_queue
  - kern\_sys\_calls.c, 143
  - scheduler.c, 172
- zombie\_child
  - shell\_built\_ins.c, 269
- zombie\_queue
  - kern\_sys\_calls.c, 143
  - scheduler.c, 172